HARVARD BUSINESS REVIEW

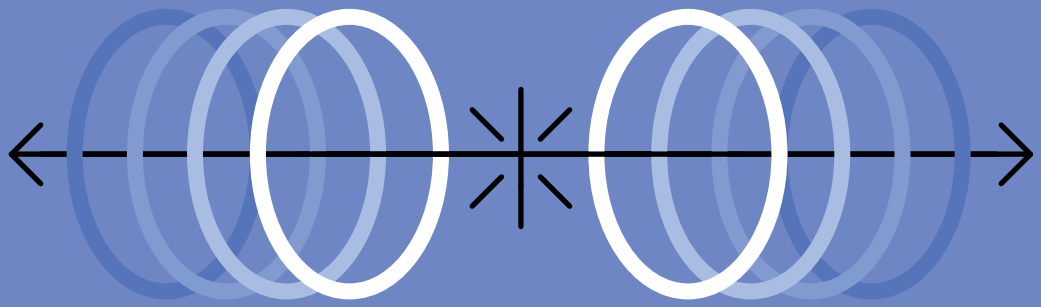

Driving Business Value With Agentic AI

HBR Analytic Services **WHITE PAPER**

Sponsored by

wipro

A New Era With Agentic AI: How Intelligent Systems Are Transforming Businesses

Imagine a world where artificial intelligence (AI) systems don't just help, they take charge. This is the future that agentic AI is making possible. Agentic AI is a smart AI system that can sense the environment, make decisions, and act without needing human intervention. This development represents more than a tech upgrade; it's a game changer for businesses everywhere.

Companies used to wonder whether they should invest in AI. Now they're racing to see how quickly they can adopt agentic AI. Why? Because these systems are not like the old ones, which needed constant human guidance. Agentic AI systems learn on their own, make decisions, and take actions based on context. This advancement means businesses can be faster, more agile, and scalable, adapting to real-time changes with ease.

Agentic AI combines advanced analytics, an awareness of the environment, and real-time decision-making capabilities. This integration moves businesses from just reacting to problems to anticipating and dynamically engaging with them. It's not about AI replacing humans; it's about boosting human creativity and strategic thinking with precise and scalable solutions.

But the benefits don't stop there. Agentic AI reduces the time it takes to bring products to market, improves customer experiences, and even helps create entirely new products and services. It sets the stage for resilient AI ecosystems that are self-sustaining and always improving.

This progression isn't just about new technology, it's a shift in mindset. Business executives need to rethink what's possible while they embed intelligence into every layer of their operations. This fosters resilience and encourages executives to consider what productivity, value creation, and competitive advantage mean in this new era.

At Wipro, we see this transformation happening up close. Companies are moving beyond experiments and starting to scale agentic AI across their operations. These systems don't just make things better—they transform them, driving faster innovation, stronger customer engagement, and real-time responsiveness.

This report from Harvard Business Review Analytic Services, sponsored by Wipro, takes a closer look at how forward-looking organizations are navigating their agentic AI journeys. It shares real-world examples, highlighting both the successes and the challenges. These insights can serve as a guide for those ready to start or speed up their own AI journeys.

Nagendra P. Bandaru
President, Global Technology Services
Wipro Ltd.

Agentic AI is here and it's moving fast. The organizations that prepare today with the right strategies and governance will thrive tomorrow. We invite you to explore this report, absorb its insights, and see how agentic AI can unlock your business's boldest ambitions—boosting autonomy and efficiency and driving lasting impact.

Driving Business Value With Agentic AI

Agentic AI is revolutionary, because it can interact with other forms of artificial intelligence (AI) within an organization and has the ability to make informed decisions and then initiate and complete tasks without the need for human intervention. Organizations are already using agentic AI, whether it's for automating back-office tasks so financial advisors have more time to help clients, enabling more efficient industrial processes, or adjusting drill paths to optimize extraction in an oil patch.

STILL OTHER ORGANIZATIONS are investigating use cases—and for good reason. Stamford, Conn.-based research firm Gartner Inc. predicts that by 2028, 33% of all enterprise software applications will include agentic AI, up from less than 1% in 2024. Further, Gartner believes that this shift will enable 15% of day-to-day work decisions to be made autonomously.[1]

As agentic AI becomes embedded in workflows, its implications stretch far beyond automation or productivity gains. Organizations may be entering a new era where human teams and AI agents operate in tandem, which will demand fresh approaches to processes, management, governance, and workforce planning.

Given such demands and the cultural changes required, agentic AI makes strong leadership critical for success, according to Zar Toolan, principal, wealth platforms, data and AI, at St. Louis-based financial services firm Edward D. Jones & Co. LP. "From a leadership standpoint, you will have humans doing work and you will have digital agents doing work," he says. "And that ushers in a whole new realm of change management and leadership capabilities."

Such capabilities are especially crucial because leveraging agentic AI at scale will necessitate rethinking training, skills, and the kinds of jobs people are required to do, according to Mohammed Rahim, group chief data officer (CDO) of the London-based global bank Standard Chartered plc. Organizational structures might change, too. "People will likely need

HIGHLIGHTS

Organizations may be entering **a new era where human teams and artificial intelligence (AI) agents operate in tandem**, demanding fresh approaches to processes, management, governance, and workforce planning.

While it's clear that agentic AI holds huge promise, **the technology is not a magic wand**—organizations very likely have work to do to before they can reap the benefits.

Preparing an organization and its people to adopt agentic AI can be challenging, as **leadership may face indifference, skepticism, or resistance.**

> **“You can’t just drop agentic AI on top of a slow, misaligned organization and expect miracles. You’ll just do the wrong things faster.”**
>
> Joe McKendrick, a business technology trends expert and coauthor of the *Harvard Business Review* article, “Overcoming the C-Suite’s Distrust of AI.”

to use different cognitive skills in their work if manual tasks are taken care of, and that requires some adjustment,” Rahim explains. “Organizations may have to provide training or change management to deal with those shifts on a large scale. Even the culture of an organization might change if the nature of work is different.”

While it’s clear that agentic AI holds huge promise, the technology is not a magic wand—organizations very likely have work to do before they can reap the benefits. “You can’t just drop agentic AI on top of a slow, misaligned organization and expect miracles,” says Joe McKendrick, a business technology trends expert and coauthor of the March 2022 *Harvard Business Review* article, “Overcoming the C-Suite’s Distrust of AI.” “You’ll just do the wrong things faster.”

This report looks at how organizations are either preparing for agentic AI or implementing it in an effort to drive business value. It highlights the challenges and successes they’ve encountered on those journeys. It also examines agentic AI’s implications for skills and new ways of working, as well as the governance structures and processes that may be needed to operationalize the technology.

Teaching Accountability

Organizations aren’t just jumping into the agentic AI pool. Standard Chartered, for example, is taking a cautious approach to agentic AI adoption. “We’re in the beginning stages of the journey,” says Rahim, the CDO. “We’re still designing what we want the future of AI to look like because it’s a sensitive and complex undertaking. Essentially, what we’d like to achieve is getting the agents or bots to interact with each other, across the businesses, to perform tasks with appropriate human oversight.”

To enable this vision, Standard Chartered is working closely with other stakeholders in the financial services ecosystem to test use cases and establish safety guidelines. “We’re working very closely with regulators, consultants, and our peer banking community to understand what can be done to ensure AI safety—for all of us,” says Rahim. “We’re also drawing on experts with industry and regulatory experience. We aim to learn from each other to make sure we get the best value from technology in the safest way.”

Standard Chartered tracks evolving country-specific requirements and helps shape AI regulation for the finance sector through active engagement with policymakers and central banks. “We’ve been working very closely with the Dubai International Financial Centre on safe adoption of AI and have recently signed an MOU [memorandum of understanding] on how we might collaborate to test agentic AI use cases,” Rahim explains.

Accordingly, Standard Chartered has started to engage with its employees on agentic AI and how to use it. Rahim cites “prompting,” the process of providing input to an AI model to generate a desired output, as an example of a critical skill needed to use agentic AI.

“AI literacy is a huge factor in today’s world. For example, prompting is one of the most important skills that humans need to learn,” Rahim asserts. “We’re conducting multiple mass-training sessions—called ‘promptathons’—to teach thousands of our employees how to prompt. They need to learn to think differently, to understand how a robot thinks, if they are going to ask an AI agent to perform a task.”

A central part of Standard Chartered’s efforts to prepare its workforce for agentic AI is teaching accountability. “If a plane is flown on autopilot, with the pilot monitoring from their chair in the cockpit, who is accountable for its safe landing?” Rahim asks. The answer is the pilot. “It’s the same with agentic AI,” he explains. “Individual accountability still exists, despite the circumstances or the level of technology involved.”

Taking a bigger-picture view, Rahim feels that a challenge worth solving centers on large-scale, long-term value. “A pain point that is difficult to address is the notion of adding value—not just as a business use case, but as a society. Collectively, we need to consider the type of society we are building in this fast-moving environment.”

Augmented Decision Making

Edward Jones, in another part of the financial services sector, is already using agentic AI to automate routine tasks and back-office functions so that financial advisors have more time to help clients. The firm is also piloting and rolling out more-advanced use cases. One is an "intelligent advice agent" that aims to achieve greater personalization to help financial advisors and branch teams serve their clients more effectively.

"We want to automate the ordinary and humanize the extraordinary," says Edward Jones' Toolan. "We're helping our teams with the things that might have been rote tasks so they can spend more time where their human judgment adds the most value."

Even if clients still want a human touch, they expect the firm to use AI to augment decision making. "Around 70% of our client base wants a human advisor, but what's interesting is that they want that advisor to be using AI in their work," he explains.

Such decision making is critical when it comes to managing client assets. "Our financial advisors need to personalize client portfolios to make sure that they're designed to help the client achieve the goals they want to achieve," Toolan notes. "We need to run the numbers on multiple factors such as market volatility, tax optimization, time horizon, and the client's risk appetite and personal circumstances to provide clients with choices. Agentic AI can assist with the analysis of that and provide the financial advisor with exactly what happened and why it happened."

Another use case involves AI agents that support Edward Jones branch teams on a day-to-day basis, in real time. "Soon we will have agentic service agents that sit digitally alongside our associates to help them to assist their clients—for instance, by answering questions and delivering the right information in real time at the moment of need," says Toolan. "Ultimately, agentic AI will deliver an enhanced experience for our clients and for our associates who help support our branch teams serving them."

> "A pain point that is difficult to address is the notion of adding value—not just as a business use case, but as a society," says Mohammed Rahim, group chief data officer of the global bank Standard Chartered plc.

Efforts are already bearing fruit. Toolan estimates that financial advisors are saving between one and two hours a week from using more generative AI and agentic AI. "I don't think of that productivity gain solely as hours saved," says Toolan. "I think of it as energy returned, freeing up time to add value in other ways, such as deeper conversations with clients."

The Need for Critical Thinking

Despite Edward Jones' agentic AI rollout, Toolan notes that the organization has a way to go before it would let the technology run completely without human intervention. "We do not have anything yet that is operating fully autonomously without a human in the loop," he asserts. "And I think that's really important. That would be more likely to occur at first in home-office processing and automation efforts. We want to learn and test the mechanics to make sure the responses we're getting are as good as they can be before we would want something to be fully autonomous. Even then, it would still have elements of human oversight."

A natural question might arise about agentic AI's decision-making capabilities and whether humans could be replaced as key decision makers. McKendrick, the business tech expert and author, doesn't think it will happen. "It's still being debated what impact agentic AI will have on the workforce, but I think you'll always need a human for

"We want to automate the ordinary and humanize the extraordinary," says Zar Toolan, principal, wealth platforms, data and AI, at Edward D. Jones & Co. LP.

decision making, especially for bigger-picture business-related decisions," he asserts. "Agentic AI might help you be more efficient and effective at delivering services, but it can't decide what services you need to provide or what's needed in the market." Critical thinking is an important human skill that will always be needed, regardless of how automated the world becomes—and agentic AI can augment skills and help people take on higher-level tasks, McKendrick stresses. But he also notes a possible unintended consequence of agentic AI's capabilities. "Organizations might need to pay attention to people becoming too reliant on AI's output, leading to a diminishing of human expertise," he says. "My fear is that too many people will rely on AI at the push of a button for answers, versus thinking things through. In fact, one CIO [chief information officer] told me they want critical thinkers on their staff, not people parroting what AI just told them."

Collaborative Approaches

For manufacturers that are using agentic AI or trying to figure out how to do so, disengagement and that reliance on a machine doing the thinking is a real threat. Preparing an organization and its people to adopt agentic AI can be challenging, as leadership may face indifference, skepticism, or resistance. As Visteon Corp., an automotive electronics supplier based in Van Buren Township, Mich., has discovered, clear communication across all levels of the organization can help get everyone on board.

A communicative approach to agentic AI was taken by Raj Singh, Visteon's vice president and global CIO. "The two most important tools are communication and collaboration," he asserts. "Not everyone needs to be on board—but if the majority understand the agentic AI vision and rally behind it, you'll get real value."

When Visteon started on its agentic AI journey, it held global town halls to introduce the topic and raise awareness about the technology and its possible applications. Out of an organization of more than 10,000 employees across 18 countries, around 1,800 people attended. "Town halls gave everyone, at all levels of technological maturity, an opportunity to see the art of the possible and understand the business value agentic AI could bring," says Singh.

The second part of Visteon's process was introducing a monthly company-wide AI newsletter to build momentum for change. "The newsletter details not only what we are planning to do with agentic AI but also what others in the industry are doing and how they are benefiting," Singh reports. "It illustrates real use cases in day-to-day roles and sparks the question 'If they can do that and gain efficiencies, why can't we?'"

But the agentic AI process doesn't end there. "The third element is hierarchical communication," explains Singh. "I present possible use cases and progress reports to senior department heads and the CEO every quarter. We broaden the discussion and encourage each person to take ideas back to their business areas and evaluate them. Ultimately, the ways in which we foster or adopt AI will become part of how performance is assessed."

Use cases that Visteon is investigating range from enabling faster and more proactive maintenance on production lines in its manufacturing facilities to making its software code regulation-compliant. "Today software is the heart of any automotive vehicle—from cockpit electronics and dashboards to battery management systems and power electronics for electric vehicles. All of this requires software, and we write millions of lines of code," says Singh. "Of course, in the automotive industry, safety is a priority, and our code

> **“The two most important tools are communication and collaboration. Not everyone needs to be on board—but if the majority understand the agentic AI vision and rally behind it, you’ll get real value.”**
>
> Raj Singh, vice president and global chief information officer, Visteon Corp.

has to adhere to Motor Industry Software Reliability Association standards. We are looking at ways in which agentic AI can assist in the compliance process, report on any issues, and help to correct any errors.”

In line with Visteon’s collaborative approach, each use case is vetted and input is sought from various parts of the business. “We’re cautiously optimistic about agentic AI—but not blindly so,” Singh explains. “Every use case goes through a vetting process: What value could it add? What investment would it need? How much efficiency will it deliver? Will it make us a better company?”

Every week, Singh convenes a panel, the AI Governing Council, which includes participants from across the organization, to discuss and assess use cases and decide whether or not to move forward with each of them. “The merit of each case and its business value is what drives those conversations,” he says.

McKendrick notes that a collaborative approach is crucial to agentic AI efforts because it can help reduce bias in AI and make the output of its analyses or decision making more useful to organizations and their customers. “There are instances of bias in all kinds of AI and machine learning because the data used to train the models can be limited and not representative of multiple groups or points of view,” he explains. “Software developers, data scientists, and IT people are doing a great job, but their world view might be limited. They may not have a 360-degree view of what’s going on in your business and with your employees who will be using agentic AI. They also may not know enough about your customers to come up with solutions that will benefit them.”

One way to remedy this situation is to include as many different kinds of people as possible in assessing use cases and in their design. Collaboration helps reduce bias in AI by drawing on a wider range of perspectives, skills, and cultural insights to design AI solutions that can meet the needs of a wider range of users. “Organizations could look at broadening perspectives by perhaps including an ethicist or a philosopher or people with humanities backgrounds in AI workshops,” McKendrick explains. “They may ask different questions, as they see the world in a different way. When you’re bringing AI into your organization, you need to have that rounded view. Everybody should be participating—it’s everybody’s project.”

Holy Grail of Automation

Other industrial organizations that are using agentic AI, like Wilmington, Del.-based Solenis LLC, a provider of water and hygiene solutions, have also taken a collaborative approach. One principle that spurred Solenis’ agentic AI efforts was to focus on the right use cases that would bring business value by solving real problems for the company and its clients—and collaboration guided execution. “Make sure functional and commercial teams are involved in the initial AI research stages,” says Charles D. Wallace, senior vice president and CIO at Solenis, “otherwise you might go down the path of investment without real outcomes. That’s not a business; that’s a research lab.”

At Solenis, agentic AI is being applied to optimize industrial processes in real time. Autonomous agents monitor chemical dosing across manufacturing plants, adjust formulas based on continuous sensor inputs, and even coordinate with other AI agents to monitor efficiency and cost. This process can be applied at Solenis’ facilities as well as at its customers’ manufacturing plants.

“Over time, our processes, technology, and machine learning capabilities have evolved to where we can look at customer data, assess it for how our chemicals are being used in their manufacturing processes, and make recommendations to optimize their processes and extend the life of their machines,” Wallace explains. “That kind of data analysis, automation, and autonomous action has real benefits for our customers, with significant impact.”

Wallace describes agentic AI as the holy grail of automation. “This is what we’ve been waiting for all this time. Automation at this level is what we’ve been talking about and now we can actually do it,” he says.

> **“I don’t think agentic AI will remove jobs from the market. I think it will elevate the jobs that are in the market.”**
>
> Charles D. Wallace, senior vice president and CIO, Solenis LLC

The company is now looking to adopt agentic AI in various nonindustrial business processes, including tax compliance and day-to-day transactions. “Enormous potential exists to automate decisions within order-to-cash processes and payment processes with minimal opportunity for risk to the business,” Wallace explains.

In fact, it is these types of processes that Wallace believes can weigh a business down and waste human potential. “There’s no sense in doing well that which you should not be doing at all,” he says. “It makes no sense having a human being sitting in front of a computer just to push a transaction through the system. Think about what other, more elevated, strategic tasks that employee could be doing instead.”

For this reason, Wallace believes that agentic AI does not pose a threat to the workforce. “I don’t think agentic AI will remove jobs from the market,” he asserts. “I think it will elevate the jobs that are in the market.”

The transition of human tasks and business processes to agentic AI requires trust in the technology, and that takes time. Shane Miller, vice president, IT, and CIO of Halliburton Co., an oil field services company, compares the process to onboarding a human employee. “One challenge is to verify that the AI agents are operational. We have to monitor them in real time and track their performance,” he says. “What if the pace of the agents is too slow, or performance does not meet expectations? Is that the agents’ fault? Is it a limitation of the system that [they’re] talking to, or is the technology platform the root of the issue? It can be quite difficult to find out what the cause of the problem is or even that there is a problem,” he adds. “But in the same way we work with humans, you have to take time to onboard, train, and correct AI agents to build trust in their capabilities.”

Halliburton, with dual headquarters in Houston and Dubai, uses agentic AI to make operations at oil and gas well sites safer. The company pioneers autonomous operations in several of its product service lines such as the use of real-time sensor data and agentic agents to guide drilling and fracturing operations. These agents make continuous decisions to adjust the drill path, avoid faults, optimize fracture placement, and maximize hydrocarbon extraction without the need for human input. “The result is far greater accuracy and production of more barrels, faster, with fewer errors,” says Miller. “The technology may reduce costs, but one of the most important values of automation is an increase in safety as a result of the reduction of workers on-site.”

In another use case, Halliburton’s proof-of-concept customer-facing AI agent integrates multiple systems. “If a customer wants to know if we’ve done a specific kind of work and the outcome of it, the agent will search all of our publicly available data, create a case study or a summary of a white paper, and set up a call with a company representative. This capability elevates our ability to collaborate and accelerates customer service,” says Miller.

Miller believes that agentic AI will prove most valuable in the unification of operational, sales, invoicing, tax, customer relationship management, and credit systems. “Integration is the ultimate goal,” he says. “Each of these platforms and agents has incremental value in and of itself. Once agents can assimilate all of that siloed data and see how it fits together, we will be able to make significant process changes and better decisions.”

Mitigating Risk

As agentic AI becomes more integrated across an organization, and more autonomous, its value to an organization will likely increase—but the risks could multiply, too. The technology’s ability to make decisions and act on them without human intervention could potentially lead to unforeseen consequences like errors, which could affect operations and customer experience.

Further, agentic AI’s speed means that it could make mistakes at scale. “That’s why we love it—it’s fast,” says Miller. “But it can go really fast in the wrong direction before we have a chance to correct it.”

The risk of errors requires organizations to set policies, guidelines, and guardrails to prevent incidents that could

affect governance and technological structures. "We realized that about 80% of what an agentic AI policy requires is already present for other applications for data security, certification, supplier due diligence, and such," says Miller. "Our approach was to keep the basics in place and add a separate policy for AI-specific issues to look for hallucinations, monitoring drift, or biases. Nevertheless, efforts in regard to policies and governance shouldn't be time-consuming. Company policies are not something we want to be adjusting every month," Miller adds. "That's why we've separated AI-specific requirements into a separate policy so we can iterate on those faster as the market demands."

Other organizations, too, are establishing protocols to guide their agentic AI usage. Edward Jones follows five guiding principles for responsible AI use, and these inform all decision making for agentic AI use cases and opportunities. "We hold the following five principles sacrosanct," says Toolan. "Our actions and our use of agentic AI must be human-centered, accountable, trustworthy, inclusive, and transparent."

Visteon has a process in place to ensure that there is always a human included in decision making, to limit risk. "We are somewhat cautious when it comes to accepting what agentic AI tells us. We trust the technology, but we will always verify its outputs," Singh, the CIO, explains. "We have rules of engagement listed as part of our AI framework stipulating that we will not make any decisions based solely on AI's recommendations."

Preparing an Agentic AI Journey

As more organizations either use agentic AI or contemplate doing so, more and more best practices have started to emerge. Some of the ones to keep in mind are:

Get started. "It's never too late to start, but it's always too late to wait," says Halliburton's Miller. "Start something—even if it's low risk or low reward. Just get your feet wet with agentic AI and build trust in the same way that you would with a new employee. Over time, you can give the agentic AI more responsibility and more power."

> "It's never too late to start, but it's always too late to wait. … Just get your feet wet with agentic AI and build trust in the same way that you would with a new employee," says Shane Miller, vice president, IT, and CIO at Halliburton Co.

Focus on outcomes. At Visteon, Singh stresses that any agentic AI applications should be business-focused and outcomes-driven. "If your agentic AI pilots or use cases are not driving business outcomes, then you might be wasting your time," he cautions.

Align AI and business strategy. "Understand your organization's strategy and make sure AI is amplifying it—not working in parallel to it, and not taking you away from it," says Rahim.

Understand your employees. "You need to listen to what your people are saying, identify their pain points, and assess what is really going on in the organization before embarking on agentic AI," says McKendrick.

Think about mindsets, skill sets, and tool sets. At Edward Jones, Toolan focuses on these three elements to help ensure that people are equipped and ready to embrace change. "Tech implementation is as much about the human psyche as it is about learning new skills and using new technology," he says.

Identify the right partners. Solenis' Wallace believes that achieving value from agentic AI depends on finding the right use cases and the right partners to help you implement them. "Leverage your third-party vendors and relationships with people outside of your organization to the best of your ability," he explains. "You'll get an astonishing amount of knowledge talking with providers, startups, and others."

Many organizations and executives are envisioning a future where agentic AI is far more integrated across their organization and can draw on greater sources of information for decision making.

Conclusion

Early agentic AI adopters are already seeing gains in productivity, quality, speed, and workplace safety. The organizations leading the way aren't waiting for perfection; they're starting small, testing use cases, and learning through controlled experimentation.

What they have in common is a shared perception that agentic AI is the way the world is moving, and what may be a novelty today will become the way business is done tomorrow.

Ultimately, many organizations and executives are envisioning a future where agentic AI is far more integrated across their organization and can draw on greater sources of information for decision making. Safeguards and data safety standards will need to be watertight before this level of integration becomes a reality.

There's another aspect to agentic AI that could empower organizations—and individuals—to marshal resources and achieve far more with less. "There's a concept floating around of a one-person billion-dollar company," says McKendrick. "This is unlikely to happen in a literal way, but it does demonstrate how agentic AI can provide the tools and processes to attain scale. It's theoretically possible to have a very small organization of entrepreneurial-minded people that can go out and achieve big things in the market."

This kind of mindset is applicable within larger organizations, too. As McKendrick notes, "Imagine what organizations might look like if they adopted an intrapreneurial approach and empowered their employees with agentic AI."

Endnotes

1 Coshow, Tom, “Intelligent Agents in AI Really Can Work Alone. Here’s How,” Gartner Inc., October 1, 2024. https://www.gartner.com/en/articles/intelligent-agent-in-ai.

HARVARD BUSINESS REVIEW

VISIT US ONLINE

hbr.org/hbr-analytic-services

Harvard Business Review Analytic Services is an independent commercial research unit within Harvard Business Review Group, conducting research and comparative analysis on important management challenges and emerging business opportunities. Seeking to provide business intelligence and peer-group insight, each report is published based on the findings of original quantitative and/or qualitative research and analysis. Quantitative surveys are conducted with the HBR Advisory Council, HBR's global research panel, and qualitative research is conducted with senior business executives and subject-matter experts from within and beyond the *Harvard Business Review* author community. Email us at hbranalyticservices@hbr.org.

CRE65030725