

Sandy Carter

With contributions from Nikhil Kassetty, Manas Talukdar, Yoshita Sharma, Olga Magnusson, Neeraj Madan, Priyanka Shrivastava, Bill Lesieur, and Sangeetha Rajkumar

AI Agents As Employees

AI GOVERNANCE | GEOPOLITICS OF TECHNOLOGY | AUTONOMOUS AGENTS



1.

Introduction

Artificial intelligence (AI) and automation are transforming the modern enterprise at an unprecedented pace. From customer service to logistics, financial analysis to product design, organizations are deploying AI to enhance decision-making, increase efficiency, and unlock new business models. Over the past decade, traditional AI methods like prediction, classification, clustering, and optimization, have delivered measurable improvements by analyzing data and supporting human tasks.

However, a new paradigm is emerging: AI is shifting from a tool that supports work to an entity that performs work.

This evolution is driven by the move from large language models (LLMs) that respond to prompts toward AI agents that drive action. An LLM is trained on vast amounts of text to understand and generate human-like language. Models like GPT-4 and Claude are generative and reactive, providing answers, content, or summaries upon request. AI agents go further. They are proactive, autonomous entities capable of initiating tasks, making decisions based on objectives, interacting with APIs and software systems, and collaborating with both humans and other agents.

Unlike classical AI, which is domain-specific and narrowly scoped, agents can be goal-driven, context-aware, and continuously learning participants in dynamic environments.

This paper explores a compelling frontier in AI adoption: the emergence of AI agents as legitimate “employees” within organizations. As these digital agents begin to take on roles traditionally held by human workers—executive assistants, financial analysts, or marketing strategists—they raise new questions about productivity, collaboration, governance, and the future of work.

The goal of this paper is to examine the architecture, applications, and implications of integrating AI agents into organizational structures. The central question is: *Can AI agents really be teammates and not just tools?*



2.

What Is an AI Agent?

An AI agent is a software-based autonomous system capable of perceiving its environment, reasoning about goals, and taking actions to achieve specific outcomes, often with minimal human intervention. Unlike traditional AI models that perform one-off predictions or static tasks, AI agents operate dynamically. They plan, adapt, and execute across multiple steps and interactions, behaving more like teammates than tools.

A defining characteristic of AI agents is autonomous decision-making. Instead of waiting for user input at each step, agents can proactively assess their context, determine appropriate actions, and carry out tasks toward a defined objective. For example, an AI agent can schedule meetings, summarize emails, answer customer questions, or organize notes without constant human directions. This marks a clear departure from traditional AI workflows, which depend on human orchestration.

Human orchestration refers to workflows where a person provides input, selects the model, runs the analysis, and interprets the result. In contrast, an AI agent can manage that entire pipeline across multiple systems.

To understand the evolution of agentic systems, it's useful to compare prompt chains and language agents. A prompt chain is a fixed sequence of instructions that guides an LLM through multiple steps, usually relying on hard-coded logic. While effective for narrow automation, prompt chains are rigid and fail when unexpected conditions arise. A language agent, by contrast, makes real-time decisions using environmental feedback, memory, and integrated tools. This adaptability enables it to operate in complex, shifting scenarios with minimal supervision.

Another critical trait of AI agents is goal-oriented behavior. Rather than executing a single instruction, agents are given objectives such as scheduling meetings, producing a competitive market report, adhering to an industry standard like Finra regulations, or drafting a customer response strategy. They then break these goals into sub-tasks, select the best course of action, and complete them. Modern AI agents also integrate external tools and models into their workflows. They can use APIs, query real-time databases, send emails, access calendars, and interact with enterprise software like Salesforce or SAP. Many are designed to dynamically select the most appropriate model for a given task, switching between LLMs depending on the context.