



Artificial intelligence impact assessment tool

The Digital Transformation Agency (DTA) provides this impact assessment tool and its supporting guidance to assist Australian Government agencies to assess their proposed use of AI. Agencies should not treat the tool as legal advice or as authorising proposed AI use. Agencies are responsible for any decisions relating to their use of AI and for seeking technical and legal advice as appropriate.

For questions and suggestions on the impact assessment tool and guidance, please email ai@dta.gov.au



Digital Transformation Agency



© Commonwealth of Australia (Digital Transformation Agency) 2025

With the exception of the Commonwealth Coat of Arms and where otherwise noted, this product is provided under a Creative Commons Attribution 4.0 International Licence.

(<http://creativecommons.org/licenses/by/4.0/legalcode>)

The Digital Transformation Agency (DTA) has tried to make the information in this product as accurate as possible. However, it does not guarantee that the information is totally accurate or complete. Therefore, you should not solely rely on this information when making a commercial decision.

The DTA is committed to providing web accessible content wherever possible. If you are having difficulties accessing this document, please email ai@dta.gov.au.

Document control

This document is maintained by the DTA and supports implementation of the Australian Government Policy for the responsible use of AI in government (the AI policy).

It will be updated periodically as the policy and technology evolve. For enquiries, please email ai@dta.gov.au.

To view the AI policy, and to check you have the most up-to-date versions of the AI impact assessment tool and guidance, please visit <https://digital.gov.au/ai/ai-in-government-policy>.

Version	Description	Date
v1.0	Published	01/12/2025

Contents

Introduction	4
1. Basic information.....	9
2. Purpose and expected benefits	14
3. Inherent risk assessment.....	16
4. Threshold assessment outcome	32
5. Fairness	34
6. Reliability and safety	36
7. Privacy protection and security	40
8. Transparency and explainability	42
9. Contestability.....	45
10. Human-centred values	47
11. Accountability	49
12. Use case review and next steps	50
Attachment A.....	53

Introduction

The artificial intelligence (AI) impact assessment tool is for Australian Government teams working on an AI use case. It helps teams identify, assess and manage AI use case impacts and risks against [Australia's AI Ethics Principles](#). Understanding and managing AI use case impacts and risks is critical for effective AI governance and to fulfilling the Australian Government's commitment to safe and responsible use of AI.

The impact assessment tool and its supporting guidance are intended to complement and strengthen – not duplicate – existing frameworks, legislation and practices that relate to government AI use. It does this by focusing on AI-specific impacts and risks that existing approaches may not fully address. It does not replace a comprehensive risk management plan, which captures all risks, treatments and ongoing monitoring measures.

Policy for the responsible use of AI in government

On 1 December 2025, the Digital Transformation Agency (DTA) published an updated [Policy for the responsible use of AI in government](#) (the AI policy). The policy update strengthens government's approach to safe and responsible AI through new measures on AI governance. It includes a new mandatory requirement for agencies to conduct an AI impact assessment for use cases identified as in scope of the AI policy.

The updated AI policy provides implementation timeframes for agencies to meet the new requirements. Agencies are required to implement the AI impact assessment requirement for in-scope use cases by 15 December 2026. While agencies may need this time to action the new AI policy requirements, including mandatory impact assessment, agencies should implement them sooner if practicable.

Refer to the AI policy for more information on the definition of 'AI use case' and the AI impact assessment requirements.

Assessing officers should familiarise themselves with the AI policy and Australia's AI Ethics Principles. Also consider other DTA resources designed to support government AI adoption, including:

- the [AI technical standard](#)
- AI procurement resources, including the [Guidance on AI procurement in government](#), [AI contract template](#) and [AI model clauses](#)
- guidance on the use of public generative AI tools [for agencies](#) and [for staff](#).

The DTA piloted a previous draft of this tool – known as the ‘Pilot AI assurance framework’ – with volunteer agencies from September to November 2024 and published the pilot draft in October 2024. The pilot findings that shaped this version of the tool and guidance are outlined in the [Pilot implementation report](#). The title has been updated to ‘AI impact assessment tool’ to better reflect its intended scope and purpose.

The DTA welcomes user feedback on the tool and supporting guidance.

Please send questions or suggestions to ai@dtg.gov.au.

AI impact assessment beyond the policy requirements

While the AI policy **requires** agencies to assess the impact of **in-scope** AI use cases, agencies are free to establish their own assessment requirements for AI use cases that are outside the scope of the AI policy.

Agencies may also use this tool to support other AI-related activities, such as procurements where suppliers use AI to provide goods and services. Refer to the [Guidance on AI procurement in government for further advice](#). Agencies should ensure specific internal requirements and expectations are clearly communicated to staff whose work involves AI.

For highly complex, novel or specific AI use cases, including AI-related activities which are outside the scope of the AI policy, agencies should also consider whether there are impacts which are not covered by this tool. These impacts must be assessed accordingly, and technical or legal advice sought where appropriate, as additional or different controls and mitigations may be required to address these impacts.

If you identify gaps in this tool, report these to the DTA. This feedback helps the DTA ensure that the AI policy and impact assessment tool remain fit for purpose and continue to evolve in response to emerging applications and risks.

Assessment roles and responsibilities

The tool and its supporting guidance are designed for Australian Government staff whose work involves AI. Each AI impact assessment must have an identified assessing officer and approving officer, and assessing officers should consult relevant experts for input. These roles are described below.

The updated AI policy also requires agencies to assign an **accountable use case owner** for each AI use case within the scope of the AI policy. This role is distinct from the assessment roles described above, although the accountable use case owner may also serve in one of those roles.

Assessing officer

An officer assigned to complete the assessment, coordinate the end-to-end process and serve as the contact point for any assessment queries. Depending on the specific use case and agency context, this officer may be a technical, data, governance or risk specialist, or they may be a policy or project officer from the business area that is implementing the AI use case in its operations.

Approving officer

An officer with appropriate authority to approve the AI use case assessment, including the inherent risk ratings. Like the assessing officer role above, the approving officer's specific role in the AI use case is not predetermined and will depend on the agency and use case context.

Expert contributors

Regardless of the assessing officer's role in the AI use case, they should seek peer review from colleagues and input from relevant experts as required, and document the experts consulted at section 1.10.

For some less complex, smaller scale AI use cases, the assessing officer may have all the information they need at their disposal to complete the assessment and may not require significant input from beyond their team. More complex use cases will likely require input from internal agency colleagues, including technical, data, risk management, policy and other domain area experts. If this expertise is not available in the agency, the assessing officer may need to seek external advice.

How to use this assessment tool

Check you have the latest version

The impact assessment tool will continue to evolve over time. To ensure you are using the latest version, please check <https://digital.gov.au/ai/ai-in-government-policy>.

Read the guidance

Before commencing this impact assessment tool and throughout the assessment process, you should [read the supporting guidance](#). The guidance mirrors the AI impact assessment tool's 12-section structure. For advice on completing a section in the tool, find the corresponding section number in the guidance.

Pre-assessment

Before assessing an AI use case, check if it is within the scope of the AI policy. The AI policy sets the criteria for in-scope use cases and the mandatory governance actions agencies must apply.

If your AI use case is out-of-scope of the AI policy, you may still find the impact assessment process useful. For example, impact assessment may support procurement, design and deployment decisions. Your agency may have specific requirements, in addition to the AI policy requirements, for assessing AI use case impacts. Check you are meeting any agency-specific requirements.

Start gathering use case information

You will need a broad range of information about the AI use case to complete the assessment. During the design phase the assessment is likely to be an iterative process, involving input from a range of experts. Regular updates may be needed as new information becomes available or design choices are refined.

Use case information you need for the assessment could include:

- the people the AI use case will affect, including demographic characteristics, needs or barriers they may face
- the potential impacts on individuals or groups, including direct and indirect impacts, their duration and reversibility, and how you intend to identify, assess and mitigate these
- the input data the AI system uses, including the type, source, collection method and security classification and how the input data will be stored and handled
- planned or existing security, monitoring, evaluation and quality assurance measures
- planned transparency measures to communicate information about the AI use case to individuals and groups
- how you will identify and consult relevant stakeholders
- how the AI system will record, log and explain any recommendations or decisions it makes, the extent of human oversight and validation of recommendations or decisions, and how any outcomes from the AI system will influence human decision-making
- who owns intellectual property rights in the inputs and outputs of the AI system, including copyright

- how your agency manages and delivers information technology services and solutions.

Threshold assessment: sections 1 to 4

First, complete sections 1 to 4 of the tool, which includes an assessment of inherent risks at section 3.

If all inherent risks are rated **low**, you can seek approving officer endorsement to conclude the impact assessment at section 4 and proceed with the use case, with appropriate plans for monitoring, evaluation and re-validation.

If any of the inherent risks at section 3 are rated **medium or high**, proceed to the full assessment and complete sections 5 to 12.

The AI policy also outlines additional requirements for use cases assessed as having an overall high inherent risk rating at the threshold assessment stage.

For any inherent risks rated medium or high, you may also consider seeking legal advice on whether the proposed use of AI is compliant with relevant laws and regulations before proceeding to the full assessment.

Full assessment: sections 5 to 12

The full impact assessment builds on the threshold risk assessment with more detailed analysis beyond the inherent risk level. It helps you examine specific potential harms, affected stakeholders, and contextual factors, to determine whether additional controls and mitigations are required.

Monitoring and evaluation

Regularly monitor and evaluate your AI use case throughout its lifecycle. If you identify a material change in the scope, usage or operation of the use case, you must formally re-validate your assessment, in line with AI policy requirements.

Re-validation

Check an approved use case assessment for accuracy and changes after deployment using re-validation. If re-validation results in assessment changes, the relevant officer or governance body must re-approve the changes. You may also specify other re-assessment intervals or triggers for your use case. This includes re-assessment to align with key project governance decision points.

1. Basic information

'AI use case' refers to the specific application of an AI system to achieve certain objectives or perform certain tasks. Refer to the AI policy for the definition of an AI use case and options to address situations where a single AI system encompasses multiple use cases.

1.1 AI use case profile

Name of AI use case

Enter your answer text here

Internal reference number or identifier

Enter your answer text here

Lead agency

Enter your answer text here

1.2 Establishing impact assessment responsibilities

Assessing officer

Officer responsible for completing this assessment, including coordinating input from relevant stakeholders, determining risk ratings and submitting the assessment for approval.

Assessing officer name

Enter your answer text here

Assessing officer position/team

Enter your answer text here

Assessing officer email

Enter your answer text here

Accountable use case owner(s)

If your use case is in scope of the AI policy, it must have a designated accountable use case owner. Refer to the AI policy and the Standard for accountability.

Name	Name	Position/team	Email

Approving officer

Approving officer name

Enter your answer text here

Approving officer position

Enter your answer text here

Approving officer email

Enter your answer text here

1.3 Additional roles and responsibilities

Record names and roles of people at your agency or third parties with responsibilities related to this AI use case. For example, responsibilities for developing the system, including external suppliers, monitoring the performance of the system or data governance.

Name	Role/title	Email	Responsibilities

1.4 AI use case description

In plain language, briefly explain how you are using or intend to use AI.

Enter your answer text here

1.5 In-scope use case

The following criteria will determine if the AI use case is in scope of the AI policy. If one of the below statements applies, your AI use case is in scope and must align with the requirements set out in the AI Policy. For details, refer to [Appendix C of the AI policy](#).

Select any criteria that place the use case within the scope of the AI policy.

- ☐ The use, misuse or failure of AI could lead to more than insignificant harm to individuals, communities, organisations, the environment or the collective rights of cultural groups including First Nations peoples.
- ☐ The use of AI will materially influence administrative decisions that affect individuals, communities, organisations, the environment or the collective rights of cultural groups including First Nations peoples.
- ☐ It is possible the public will directly interact with, or be significantly impacted by, the AI or its outputs without human review.
- ☐ The AI is designed to use personal or sensitive data, or security classified information.
- ☐ It is deemed an elevated risk AI use case as directed by the DTA.
- ☐ Not applicable.

1.6 Type of AI technology

Briefly explain the types of AI technology you are using or intend to use.

Enter your answer text here

1.7 Usage pattern

Select all that apply. Source: [Classification system for AI use](#).

- ☐ Decision-making and administrative action
- ☐ Analytics for insights
- ☐ Workplace productivity
- ☐ Image processing

Advisory note

Seek legal advice if AI automated decision-making is used for an administrative decision under an Act.

1.8 Administrative decisions

Only complete this section if you selected 'Decision-making and administrative action' in section 1.7 and if AI automated decision-making is used for an administrative decision under an Act.

Provide the legislative authority for automating such a decision.

Enter your answer text here

Advisory note

Express legislative authority is generally required to automate decision-making of an administrative decision under an Act. Legal advice should be obtained for any proposed use of AI in this context.

1.9 Domain

Select all that apply. Source: [Classification system for AI use](#)

- ☐ Service delivery
- ☐ Compliance and fraud detection
- ☐ Law enforcement, intelligence and security
- ☐ Policy and legal
- ☐ Scientific
- ☐ Corporate and enabling

1.10 Expert contributions

Record all experts you consulted during this impact assessment process, summarising their input and any resulting changes.

Name	Role and expertise	Date last consulted	Input and resulting changes

Advisory note

Agencies should consider engaging relevant internal or external expertise based on the complexity, novelty or potential impacts of the AI use case. This may include experts in ethics, law, discrimination, privacy, employment, intellectual property, technology, security, data, accessibility, automated decision-making, or the domain in which the AI system is being deployed.

1.11 Impact assessment review log

This impact assessment should be updated regularly throughout its lifecycle.

There are 2 types of review:

- a. **Pre-deployment review:** Continuous updates as part of the drafting process, before the assessment is approved. As new information emerges, design choices are refined and risks are reassessed, previous answers will be reviewed and updated.
- b. **Post-deployment re-validation:** Following the deployment of the use case, the AI policy requires agencies to:
 - regularly monitor and evaluate the use case to ensure it is operating as intended and that risks are effectively managed
 - re-validate the AI use case impact assessment by checking its accuracy and updating it when there is a material change in the scope, usage or operation of the use case.

Track impact assessment pre-deployment review and re-validation outcomes and summarise post-review changes.

Review type	Reasons for review	Review date	Post-review changes
☐ Pre-deployment			
☐ Re-validation			

2. Purpose and expected benefits

Under [*Australia's AI Ethics Principles*](#), the use of AI should have a clearly defined and beneficial purpose that is consistent with human, societal and environmental wellbeing.

2.1 Problem definition

Identify the problem you are trying to solve.

Enter your answer text here

2.2 AI use case purpose

Describe the purpose of your use of AI, focusing on how it will address the problem you have identified.

Enter your answer text here

2.3 Non-AI alternatives

Outline non-AI alternatives that could address this problem.

Enter your answer text here

2.4 Identifying stakeholders

Use the stakeholder mapping aid at **Attachment A** to identify stakeholder groups that may be affected by the AI use case. Record how they may be positively or negatively affected.

This will guide your consideration of expected benefits and potential risks. Stakeholder groups may include:

- end users of the AI outputs
- people who will be evaluated or monitored by the AI
- rights holders, for example copyright owners
- AI developers or engineers
- business or industry
- regulators
- agency staff and their unions, and other third-party personnel

- communities or other groups.

Stakeholder group	Briefly describe how they may be affected

2.5 Expected benefits

Consider the stakeholders identified in the previous question and briefly outline the expected benefits of the AI use case.

Consult the guidance for this section for resources to assist you.

What are the expected benefits of the AI use case?

Enter your answer text here

Advisory note

Your response should be supported by quantitative and/or qualitative analysis.

Qualitative analysis should consider whether there is an expected positive outcome and whether AI is a good fit to accomplish the relevant task, particularly compared to non-AI alternatives identified. Benefits may include gaining new insights or data.

3. Inherent risk assessment

This section identifies potential AI-specific risks that may arise from the AI use case. It is intended to complement existing risk management processes. It does not replace a comprehensive risk management plan capturing all risks, treatments and ongoing monitoring measures.

Base your impact assessment on the inherent risk, before it is treated with controls. Do not assess the residual risk, which would remain once controls are applied.

Consequence

Use the 5-tier consequence scale for consistent evaluation across impact areas. The consequence descriptions for sections 3.1 to 3.8 are summarised from the risk consequence assessment provided in the guidance. Refer to the table in the **guidance appendix** for further detail.

When assessing consequences, consider both intended and unintended consequences and outcomes. This includes evaluating the impact of system failure, the impact of misuse or malicious use, and of other deviations from expected use.

Likelihood

Use the risk likelihood scale in **Table 1** below to select a likelihood level for each of the impacts in sections 3.1 to 3.8.

Estimating likelihood for an AI use case can be harder due to limited historical data, changing behaviour and emerging challenges, such as the opaqueness of AI systems. If you are unsure, take a precautionary approach and rate the consequence as at least 'possible'. Seek advice from relevant experts – such as technical specialists, risk managers or domain experts – to inform or validate your assessment.

Risk rating

Use the risk matrix in **Table 2** below to determine your inherent risk rating for sections 3.1 to 3.8.

Provide a clear and concise rationale for each risk rating, explaining the relevant controls in place or planned. Ensure the assessment reflects the AI use case's scope, function, and current controls.

Table 1. Risk likelihood scale

Likelihood	Probability	<u>Description</u>
Almost certain	91% and above	Almost certain to eventuate within the foreseeable future
Likely	61-90%	Will probably eventuate within the foreseeable future
Possible	31-60%	May eventuate within the foreseeable future
Unlikely	5-30%	May eventuate at some time but is not likely to occur in the foreseeable future
Rare	Less than 5%	Will only eventuate in exceptional circumstances or as a result of a combination of unusual events

Table 2. Risk matrix

	Insignificant	Minor	Moderate	Major	Severe
Almost certain	Medium	Medium	High	High	High
Likely	Medium	Medium	Medium	High	High
Possible	Low	Medium	Medium	High	High
Unlikely	Low	Low	Medium	Medium	High
Rare	Low	Low	Low	Medium	Medium

3.1 Risk of reducing service accessibility and inclusion

This section asks you to assess the risk of your AI use making government services less accessible or inclusive for the public. Things to think about in assessing this risk include:

- Is the system being used for fundamental or critical government services to the public?
- Could any individuals or groups face new or increased barriers to accessing services or information?
- Could the system unintentionally exclude users with limited digital literacy or access to specific technologies?
- Could language barriers or cultural factors limit effective operation of the AI use case?
- Could the system produce answers which the public would misinterpret or find difficult to understand?

- Will the system allow users to seek verification of AI responses or review of AI outputs from a human? For example, will individuals be able to respond to misinformation produced by AI?

Why this matters

The use of AI can create barriers for some people if systems are not designed to be inclusive or accessible. This can range from minor inconvenience through to serious impacts.

Examples of minor inconvenience could include users not understanding how to interact with an AI system, or an AI system answering public enquiries with generic, inaccurate or confusing responses, requiring repeated attempts to resolve queries.

Serious impacts could include someone missing out on vital services because the AI has reached an incorrect decision without human input or oversight, or because the AI couldn't understand their language, disability or living situation. This could result in a breach of specific legislation relevant to the decision, if individuals are unable to access or understand the reasons for an AI-made decision or are unable to request a review of that decision.

What is the risk of your AI use making government services less accessible or inclusive for the public?

Consequence (select one)

- ☐ Insignificant – Minor glitch; no real barrier. Instantly resolved.
- ☐ Minor – Short-term, reversible barrier. Few users affected.
- ☐ Moderate – Noticeable access issues. Some groups impacted. Resolution requires effort.
- ☐ Major – Major service disruption. Widespread impact. Legal/public concern likely.
- ☐ Severe – Critical failure. Essential services inaccessible. Broad harm, urgent response.

Likelihood (select one)

- ☐ Almost certain
- ☐ Likely
- ☐ Possible
- ☐ Unlikely
- ☐ Rare

Risk rating (select one)

- ☐ Low
- ☐ Medium
- ☐ High

Rationale for risk rating including any existing controls

Enter your answer text here

3.2 Risk of unfair discrimination

This section asks you to assess the risk of your AI use unfairly discriminating against individuals, communities or groups. Things to think about in assessing this risk include:

- Could the system, through intended use, failure or misuse, replicate or amplify bias or discrimination? For example, from training data which is not diverse or representative, or embedded rules.
- Could certain groups be treated unfairly or inconsistently? This can happen if the system is trained on data reflecting existing inequalities or uses proxies correlating with race, class or gender and reproduces these biases in outputs. Proxies could include postcode, education level or country of birth.
- Is the AI system used as part of an assessment, eligibility, employment or resource allocation process? This includes making decisions with ethical or legal implications.
- Will users be able to identify bias in training data, and adjust the AI system accordingly? For example, users will gather more diverse sets of data if the AI system is reinforcing pre-existing patterns and therefore producing historical bias.

Why this matters

AI can produce discriminatory outcomes if:

- the parameters of the AI system are discriminatory
- the training data is unbalanced or contains hidden biases
- officials use AI outputs without proper scrutiny.

For example, if an AI system trained on biased hiring data screens job applications, it may disadvantage some demographics, causing issues from inconsistent shortlisting to excluding qualified candidates. This may not be compliant with Commonwealth anti-discrimination

legislation which prohibits discrimination on the basis of protected characteristics such as age, disability, race and sex.

What is the risk of your AI use unfairly discriminating against individuals, communities or groups?

Consequence (select one)

- ☐ Insignificant – Negligible discrimination; no harm. Issues caught and fixed early.
- ☐ Minor – Limited unfair treatment. Few individuals affected; quickly corrected.
- ☐ Moderate – Noticeable harm to some individuals/groups. Bias concerns raised; intervention needed.
- ☐ Major – Widespread harm to multiple communities. Significant impact; trust must be rebuilt.
- ☐ Severe – Systemic harm, especially to vulnerable groups. Public outrage; urgent reform required.

Likelihood (select one)

- ☐ Almost certain
- ☐ Likely
- ☐ Possible
- ☐ Unlikely
- ☐ Rare

Risk rating (select one)

- ☐ Low
- ☐ Medium
- ☐ High

Rationale for risk rating including any existing controls

Enter your answer text here

3.3 Risk of perpetuating stereotyping or demeaning representations

This section asks you to assess the risk of your AI use perpetuating stereotyping or demeaning representations of individuals, communities or groups. Things to think about in assessing this risk include:

- Could the training data cause the AI system to generalise the characteristics of individuals to a whole group? For example, not differentiating between individuals from racial or ethnic minorities.
- Could the system produce classifications or outputs that reinforce harmful stereotypes? For example, associating a certain role with a gender.
- Might outputs misrepresent individuals or communities, especially those historically marginalised? For example, the AI system processes and prioritises certain features of a particular race or ethnicity over others.
- Has the system been tested with diverse demographics, cultural contexts and perspectives?

Why this matters

AI systems can reinforce stereotypes or misrepresent people if they rely on overly broad categories or labels.

For example, if an AI tool used in social services classifies people based on postcode or income level, it might incorrectly assume certain behaviours or needs, leading to unfair treatment. This can result in minor impacts like poor communication, or more serious consequences such as inappropriate service responses. As noted above, this may not be compliant with Commonwealth anti-discrimination legislation.

Another example is if an AI tool publishes false information about an individual that causes significant harm to their reputation. This may bring the risk of defamatory action against an agency.

What is the risk of your AI use perpetuating stereotyping or demeaning representations of individuals, communities or groups?

Consequence (select one)

- ☐ Insignificant – Mild stereotyping, quickly fixed. No lasting harm.
- ☐ Minor – Isolated incidents. Few affected. Promptly resolved.
- ☐ Moderate – Noticeable public concern. Some groups harmed. Needs intervention.
- ☐ Major – Widespread harmful stereotypes. Public outcry. Trust damaged.
- ☐ Severe – Systemic harm across communities. Legal risk. Urgent reforms needed.

Likelihood (select one)

- ☐ Almost certain
- ☐ Likely
- ☐ Possible
- ☐ Unlikely
- ☐ Rare

Risk rating (select one)

- ☐ Low
- ☐ Medium
- ☐ High

Rationale for risk rating including any existing controls

Enter your answer text here

3.4 Risk of harm

This section asks you to assess the risk of your AI use harming individuals, communities, groups, organisations or the environment. Things to think about in assessing this risk include:

- Could this specific system unintentionally cause physical, psychological, social, economic, reputational or environmental harm through its intended or unintended operation or outputs?
- Might this AI use case deprioritise or misclassify certain individuals or groups, resulting in harm or exclusion within the context in which it is deployed?

- Does the system interact with physical environments or control machinery as part of its defined use case?
- Is the AI system intended for use in the management or operation of critical infrastructure or services?

Why this matters

AI systems may cause direct or indirect harm through biased outcomes, exclusion, misinformation, poor decision-making, or unintended consequences.

For example, an AI system used to analyse drone footage for environmental compliance in a specific region may misidentify land clearing activities, leading to enforcement action against land holders without appropriate human oversight.

Assessment of environmental impacts should focus on the environmental consequences of the specific AI use case, not general or systemic environmental impacts of AI technology.

What is the risk of your AI use harming individuals, communities, groups, organisations or the environment?

Consequence (select one)

- ☐ Insignificant – Minor glitch. No real harm. Easily managed.
- ☐ Minor – Isolated public/business/environmental impact. Low cost. Resolved quickly.
- ☐ Moderate – Noticeable harm to people, businesses, or ecosystems. Public concern raised.
- ☐ Major – Widespread harm. Financial losses, distress, or ecosystem damage.
- ☐ Severe – Critical, irreversible harm. Economic and environmental crisis. Urgent action needed.

Likelihood (select one)

- ☐ Almost certain
- ☐ Likely
- ☐ Possible
- ☐ Unlikely
- ☐ Rare

Risk rating (select one)

- ☐ Low
- ☐ Medium
- ☐ High

Rationale for risk rating including any existing controls

Enter your answer text here

3.5 Risk of AI use raising privacy concerns

This section asks you to assess the risk of your AI use raising privacy concerns. Things to think about in assessing this risk include:

- Will the AI use case involve the collection, storage or disclosure of [personal information](#)? Does this include personal information that is [sensitive](#)?
- Does the use case involve a new or changed way of handling personal information?
- Is the AI use case likely to have a significant impact on the privacy of individuals?

Under the [Australian Government Agencies Privacy Code](#), for all 'high privacy risk' projects agencies must conduct a [privacy impact assessment \(PIA\)](#). To determine whether a PIA is required for your AI use case, you should complete a privacy threshold assessment (PTA). A PTA will help you identify your use case's potential privacy impacts and screen for factors that point to a 'high privacy risk', which requires a PIA. Use the outcomes of your PTA and, if required, your PIA, to inform your risk ratings below.

Why this matters

Use of AI is likely to involve new or changed ways of handling personal information that are likely to have a significant impact on the privacy of individuals. AI technologies rely on large data sets that often include personal information, which can create new specific privacy risks, amplify existing risks and lead to serious harms.

AI systems:

- often repurpose existing datasets for new functions or activities, which alters the purpose for using or disclosing personal information
- may require the integration of data from multiple sources, creating new pathways of collection, storage, or disclosure

- frequently rely on large-scale or sensitive datasets increasing the likelihood of high privacy risk – including biometric or behavioural data
- can sometimes enable re-identification or generate inferences that amount to new personal information, even where data is de-identified
- are known to generate inaccurate or false results increasing the risk that personal information collected, used and disclosed by an agency is inaccurate
- can produce personal information about an identified or reasonably identifiable individual that is inferred, incorrect or artificially generated.

You may also wish to consider the Office of the Australian Information Commissioner

[Guidance on privacy and the use of commercially available AI products.](#)

What is the risk of your AI use raising privacy concerns?

Consequence (select one)

- ☐ Insignificant – Minor data mishandling. No sensitive info exposed. Trust intact.
- ☐ Minor – Small data breach. Few affected. Resolved quickly.
- ☐ Moderate – Noticeable breach of sensitive data. Some distress caused.
- ☐ Major – Large-scale misuse of private data. Public trust damaged.
- ☐ Severe – Widespread sensitive data exposure. Severe harm and lasting loss of trust.

Likelihood (select one)

- ☐ Almost certain
- ☐ Likely
- ☐ Possible
- ☐ Unlikely
- ☐ Rare

Risk rating (select one)

- ☐ Low
- ☐ Medium
- ☐ High

Rationale for risk rating including any existing controls

Enter your answer text here

3.6 Risk raising security concerns – data aspects

This section asks you to assess the risk of your AI use raising security concerns due to the sensitivity or security classification of the data or information the AI system uses. Things to think about in assessing this risk include:

- Would a breach of this data cause operational, reputational, or personal harm?
- Is any of the data subject to legislative, government classification, contractual or equitable confidentiality obligations, or protected status? This can include:
 - protected information
 - security classified information
 - commercial information
 - law enforcement information
 - national security information
 - confidential information
 - personal information in accordance with the *Privacy Act 1988* (Cth).
- Are all external data sources integrated with the AI system under appropriate governance?
- What measures are in place to confine access and handling of data and information? For example:
 - Will inputs and outputs leave the agency's systems, and if yes, how will the data be handled. For example, how will it be transferred and stored?
 - Will any third-party operating the system have access to inputs and outputs, and can they use this data for other purposes. For example, can the system ingest data or conduct data mining?
 - What records will the system keep of access and use of data?
- What measures are in place to respond to unauthorised access, use or disclosure of sensitive or security classified data, such as a data breach or incident response plan?

Why this matters

AI systems handling sensitive or security classified data could pose significant security risks if not properly protected. Unauthorised access, data breaches or leaks, or misuse may:

- breach contractual or statutory obligations

- disrupt critical operations or services
- result in financial losses
- compromise national security and public safety
- expose individuals or systems to harm
- erode confidence in government data safeguards.

Agencies should have particular regard to the way the AI system will handle inputs and outputs. For example, commercially available AI may use inputs for further training and development of the AI or allow third parties to access outputs generated by the AI, through channels such as an interface with search engines. This may constitute a breach of contractual obligations not to disclose confidential information, or statutory obligations not to disclose protected information, such as obligations under federal taxation or health legislation.

What is the risk of your AI use raising security concerns due to the sensitivity or security classification of the data or information the AI system uses?

Consequence (select one)

- ☐ Insignificant – Minor lapse. No data misuse. Promptly fixed.
- ☐ Minor – Small breach. Few records accessed. Quickly secured.
- ☐ Moderate – Moderate data compromise. Privacy concerns raised. Needs investigation.
- ☐ Major – Major breach of sensitive data. Public trust shaken. Urgent response needed.
- ☐ Severe – Massive breach. National security and privacy at risk. Emergency overhaul required.

Likelihood (select one)

- ☐ Almost certain
- ☐ Likely
- ☐ Possible
- ☐ Unlikely
- ☐ Rare

Risk rating (select one)

- ☐ Low
- ☐ Medium
- ☐ High

Rationale for risk rating including any existing controls

Enter your answer text here

3.7 Risk of raising security concerns – system aspects

This section asks you to assess the risk of your AI use raising security concerns due to the implementation, sourcing or characteristics of the AI system? Things to think about in assessing this risk include:

- Is the AI system developed externally, or does it use open source or pretrained models?
- Can the system's behaviour be explained, tested, or audited?
- What safeguards does the AI system have if it malfunctions or produces harmful outputs? For example, can the agency interrupt or shut the system down via a circuit-breaker?
- Will the agency monitor, and take steps to address, signs of anomalies, dysfunctions and unexpected performance?
- Are there known vulnerabilities or dependencies, or likely security risks?
- Will the AI system be deployed in a secure environment on the agency's premises, or will it be deployed through the cloud? If deploying through the cloud, are the servers located in Australia or overseas?
- Is the AI system using novel or experimental algorithms without established security benchmarks?

Why this matters

The way an AI system is built, sourced or configured can introduce security vulnerabilities. These can range from low-impact issues – such as a brief service disruption due to minor misconfiguration – to a high-impact situation where a third-party component enables a major breach of sensitive data or system compromise.

It can also result in a breach of an agency's obligation to comply with the:

- information and resource protection requirements under the [Protective Security Policy Framework](#)
- requirements to protect personal information from misuse, interference and loss specified in the [Australian Privacy Principles](#).

It may also result in data being subject to foreign laws where that data is taken offshore.

What is the risk of your AI use raising security concerns due to the implementation, sourcing or characteristics of the AI system?

Consequence (select one)

- ☐ Insignificant – Minor flaw (e.g. bug). No real security impact. Quickly fixed.
- ☐ Minor – Small vulnerability exploited. Limited breach. Resolved with updates.
- ☐ Moderate – System feature causes data leak or access issue. Contained but serious.
- ☐ Major – System feature causes data leak or access issue. Contained but serious.
- ☐ Severe – Critical vulnerability causes widespread breach. National/public safety at risk. Emergency response needed.

Likelihood (select one)

- ☐ Almost certain
- ☐ Likely
- ☐ Possible
- ☐ Unlikely
- ☐ Rare

Risk rating (select one)

- ☐ Low
- ☐ Medium
- ☐ High

Rationale for risk rating including any existing controls

Enter your answer text here

3.8 Risk to reputation or public confidence

This section asks you to assess the risk of your AI use posing a reputational risk or undermining public confidence in the government? Things to think about in assessing this risk include:

- Is there a risk that the use case will fail public expectations of government as an exemplar?
- Even if technically sound, outcomes that appear unfair, insensitive, or misaligned with community values may generate media scrutiny and or erode public confidence.

Why this matters

Even where an AI system is ethical, technically sound and fully compliant, there may still be reputational risk arising from how it is used. Public perception can be shaped by the context, purpose, or outcomes of the system.

What is the risk of your AI use posing a reputational risk or undermining public confidence in the government?

Consequence (select one)

- ☐ Insignificant – Minor, isolated issue; negligible impact on public trust; quickly addressed.
- ☐ Minor – Brief concern or media attention; small reputational dent; resolved with prompt action.
- ☐ Moderate – Public questions oversight; moderate trust impact; requires remedial response.
- ☐ Major – Widespread criticism; major reputational damage; trust rebuilding needed.
- ☐ Severe – Profound loss of public trust; seen as governance failure; long-term recovery and reform required.

Likelihood (select one)

- ☐ Almost certain
- ☐ Likely
- ☐ Possible
- ☐ Unlikely
- ☐ Rare

Risk rating (select one)

- ☐ Low
- ☐ Medium
- ☐ High

Rationale for risk rating including any existing controls

Enter your answer text here

3.9 Overall inherent risk rating

Determine the overall inherent risk rating for the AI use case based on the ratings selected in sections 3.1 to 3.8 **above**. Use the highest risk rating identified in the sections above as the overall inherent risk rating. This ensures the assessment reflects the most significant potential risk exposure, even if the other risks are rated lower.

This overall inherent risk rating is the risk level under standard operating conditions, assuming only existing baseline controls are in place, before any additional controls are applied. It does not consider proposed mitigation strategies or enhancements that are yet to be implemented.

Overall inherent risk (select one)

- ☐ Low
- ☐ Medium
- ☐ High

Rationale for overall inherent risk rating and explanation of relevant risk controls

Enter your answer text here

4. Threshold assessment outcome

4.1 Assessing officer recommendation

If the assessing officer is satisfied that all risks in the inherent risk assessment at section 3.9 are **low**, they may recommend that a full assessment is not needed and that the agency accept the low risk and proceed with the AI use case.

If the overall inherent risk rating recorded at section 3.9 is **medium** or **high**, then you must complete a full assessment.

Assessing officer recommendation

- ☐ A full assessment **is** necessary for this use case.
- ☐ A full assessment **is not** necessary for this use case.

Comments (optional)

Enter your answer text here

Assessing officer name

Enter your answer text here

Assessing officer position

Enter your answer text here

Date

Enter your answer text here

4.2 Approving officer review

Approving officer decision

- ☐ I have reviewed the recommendation, am satisfied by the supporting analysis and agree that a full assessment **is** necessary for this use case.
- ☐ I have reviewed the recommendation, am satisfied by the supporting analysis and agree that a full assessment **is not** necessary for this use case.

Comments (optional)

Enter your answer text here

Approving officer name

Enter your answer text here

Approving officer position

Enter your answer text here

Date

Enter your answer text here

5. Fairness

This section helps agencies ensure that their AI use case aligns with [Australia's AI Ethics Principle](#) of fairness. The intention is to ensure agencies can clearly define fairness in operational terms and demonstrate that outcomes can be assessed for equity and impartiality.

Under the AI Ethics Principles, AI systems should, throughout their lifecycle, be inclusive and accessible and should not involve or result in unfair discrimination against individuals, communities or groups.

5.1 Defining fairness

Where appropriate, you should consult relevant domain experts, affected parties and stakeholders to determine how to define and contextualise fairness for your use of AI. Consider inclusion and accessibility, as well as discrimination and bias. Consult the guidance for prompts and resources to assist you.

Do you have a clear definition of what constitutes a fair process and/or outcome in the context of your use of AI?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

5.2 Measuring fairness

Measuring fairness is an important step in identifying and mitigating fairness risks. A wide range of metrics are available to address various concepts of fairness. Consult the guidance for resources to assist you.

Do you have a way of measuring the fairness of system outcomes quantitatively or qualitatively?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

6. Reliability and safety

This section helps agencies ensure that their AI use case aligns with [Australia's AI Ethics Principle](#) of reliability and safety. Its intention is to ensure agencies consider reliability and safety in operational terms and demonstrate how these will be measured, monitored and acted on.

Under the AI Ethics Principles, AI systems should throughout their lifecycle reliably operate in accordance with their intended purpose.

6.1 Data suitability

Consider data quality and factors such as accuracy, timeliness, completeness, consistency, lineage, provenance and volume. Also ensure that all necessary data permissions and rights are in place.

If your AI system requires the input of data to operate, or you are training or evaluating an AI model, can you explain why the chosen data is suitable for your use case?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

6.2 Indigenous data

Consider if use of Indigenous data and AI outputs is consistent with the expectations of First Nations people, and the [Framework for Governance of Indigenous Data \(GID\)](#). See further advice in the guidance.

If your AI system uses Indigenous data, including where any outputs relate to Indigenous people, have you ensured that your AI use case is consistent with the Framework for Governance of Indigenous Data?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer*Enter your answer text here*

6.3 Suitability of procured AI system

This may include the procurement of a model, system, multiple systems or a class of systems from a third-party provider. It also covers situations where you are using open-source systems, application programming interfaces (APIs) or otherwise sourcing or adapting existing systems. Factors to consider are outlined in the guidance.

This includes ensuring that the AI is suitable for your intended use and that contracts used to procure the AI model or system protect against relevant risks. Ensure there are appropriate terms covering factors such as:

- testing
- support
- data governance
- privacy
- security
- intellectual property provisions
- warranties.

If you are procuring an AI model, can you explain its suitability for your use case?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer*Enter your answer text here*

6.4 Testing

Outline any areas of concern in results from testing. If the AI system has not yet undergone testing, outline elements to be considered in testing plan. For example, the model's accuracy.

Has the AI system been tested sufficiently and are you satisfied with its reliability and safety for the context of your use case?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

6.5 Pilot

If answering 'yes', explain what you have learned or hope to learn in relation to reliability and safety and, if applicable, outline how you adjusted the use of AI.

Have you conducted, or will you conduct, a pilot of your use case before deploying?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

6.6 Monitoring

Explain how you will monitor and evaluate performance.

Have you established a plan to monitor and evaluate the performance of your AI system?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

6.7 Preparedness to intervene or disengage

Refer to the guidance for resources to assist you in establishing appropriate processes.

Have you established clear processes for human intervention or safely disengaging the AI system where necessary? For example, if stakeholders raise valid concerns with insights or decisions or an unresolvable issue is identified.

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

6.8 Training of AI system operators

An operator of the system may include anyone who uses the AI to support their work, such as staff interpreting insights, responding to system alerts, or making decisions informed by the AI's outputs. With all automated systems, including AI systems, there is always the risk of overreliance on results. Operators should receive appropriate training to understand how to use the system responsibly, monitor and critically evaluate its outcomes.

Is there a process in place to ensure operators of the AI system are sufficiently skilled and trained?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

7. Privacy protection and security

This section helps agencies ensure their AI use case upholds [Australia's AI Ethics Principle](#) of privacy protection and security by ensuring that personal information is collected, used, stored and shared in a way that minimises risk and protects individuals' privacy in accordance with the [Australian Privacy Principles \(APPs\)](#).

Agencies must ensure their AI use case operates within an appropriate security framework, including the Protective Security Policy Framework (PSPF), with safeguards proportionate to the sensitivity of the data and potential harm.

Under the AI Ethics Principles, AI systems should throughout their lifecycle respect and uphold privacy rights and data protection, and ensure data security.

7.1 Minimise and protect personal information

See the guidance for advice on:

- collecting, using and disclosing personal information in accordance with the APPs
- privacy enhancing technologies.

You should also consider the Office of the Australian Information Commissioner [Guidance on AI and Privacy](#).

Are you satisfied that any collection, use or disclosure of personal information complies with the APPs?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

7.2 Privacy threshold and/or impact assessment

Record where you have stored your privacy threshold assessment and, if applicable, your privacy impact assessment are stored. For example, the document management system, case file, or other repository.

Enter your answer text here

7.3 Security risks

Engage with your agency's IT Security Adviser (ITSA) and consider the latest security guidance and strategies for AI use, such as the:

- Protective Security Policy Framework (PSPF)
- Information Security Manual (ISM)
- Australian Signals Directorate advice on [Engaging with AI](#).

What measures are in place to address security risks arising from the operation of the AI?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

8. Transparency and explainability

This section supports agencies to uphold [Australia's AI Ethics Principles](#) of transparency and explainability, which entail:

- consulting with affected stakeholders
- maintaining clear and accurate records
- disclosing when AI is used
- providing appropriate explanations of AI outputs that allow individuals to understand, question, or seek review where appropriate.

Under the AI Ethics Principles, there should be transparency and responsible disclosure so people can understand when they are being significantly impacted by AI and can find out when an AI system is engaging with them.

8.1 Consultation

Refer to the list of stakeholders identified in section 2.4. Seek out representatives with the appropriate skills, knowledge or experience to engage with AI ethics issues. Consult the guidance for prompts and resources to assist you.

Have you consulted stakeholders identified in section 2.4?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

8.2 Public visibility

When deciding to publish or not publish AI use information, consider the advice in the guidance on appropriate transparency mechanisms, information to include and factors to consider.

Will appropriate information be made publicly available, such as the scope and goals related to the use of AI?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

8.3 Maintain appropriate documentation and records

Comply with documentation and record-keeping requirements, including maintaining reliable records of decisions, testing and information and data assets used in an AI system. This is important to enable internal and external scrutiny, continuity of knowledge and accountability.

Have you put in place processes to maintain appropriate documentation and records throughout the lifecycle of the AI use case?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

8.4 Disclosing AI interactions and outputs

Consider members of the public, agency staff or other stakeholders who may interact with the AI system, or decision-makers who may rely on its outputs.

Will people be informed when they interact with the AI system or receive outputs generated by AI?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer, including, if applicable, how AI involvement will be disclosed.

Enter your answer text here

Will members of the public interacting with the system be provided with the ability to request a non-AI alternative?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

8.5 Offer appropriate explanations

You should be able to clearly explain how a government decision or outcome has been made or informed by AI. These explanations should be understandable to both technical and non-technical audiences. This is especially important if your AI system will materially influence administrative action or decision-making about individuals, groups, organisations or communities.

Will your AI system allow for appropriate explanation of the factors leading to AI-generated decisions, recommendations or insights?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

9. Contestability

This section supports agencies to uphold [Australia's AI Ethics Principle](#) of contestability. Agencies need to have clear processes for notifying individuals about AI use in administrative actions and for managing review or appeals in a timely manner.

Under the AI Ethics Principles, when an AI system significantly impacts a person, community, group or environment, there should be a timely process to allow people to challenge the use or outcomes of the AI system.

9.1 Notification of AI affecting rights

See the guidance for help interpreting 'administrative action', 'materially influenced' and 'legal or significant effect' as well as recommendations for notification content.

Will individuals, groups, organisations or communities be notified if an administrative action with a legal or significant effect on them was materially influenced by the AI system?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

9.2 Challenging administrative actions influenced by AI

Administrative law is the body of law that regulates government administrative action.

Access to review of government administrative action is a key component of access to justice. Consistent with best practice in administrative action, ensure that no person could lose a right, privilege or entitlement without access to a review process or an effective way to challenge an AI-generated or informed decision. The use of AI should not make it more difficult to access administrative law rights or place the agency in a position where it might contravene its administrative law duties.

Is there a timely and accessible process to challenge the administrative actions discussed at section 9.1?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

10. Human-centred values

Under [Australia's AI Ethics Principles](#), AI systems should throughout their lifecycle respect human rights, diversity and the autonomy of individuals.

This section supports agencies to uphold the AI Ethics Principle of human-centred values. It is intended to help agencies consider AI use case alignments with Australia's human rights obligations and to incorporate diverse perspectives. Incorporating diverse perspectives can help avoid biased or discriminatory outcomes that could undermine individual or community well-being.

10.1 Incorporating diversity

Consider how you have incorporated diverse perspectives through the lifecycle of your AI use case. For example, consider the choice of data, composition of development and deployment teams and the stakeholder and user groups to choose to consult.

Are you satisfied that you have incorporated diversity and people with appropriately diverse skills, experiences and backgrounds throughout the lifecycle of your AI use case?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

10.2 Human rights obligations

Complete this question after completing previous sections of the assessment. This approach will enable a more considered assessment of the human rights implications of your AI use case.

Have you consulted an appropriate source of advice or otherwise ensured that your AI use case and the use of data align with human rights obligations?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

11. Accountability

Under [Australia's AI Ethics Principles](#), AI systems should throughout their lifecycle be identifiable and accountable for the outcomes of the AI systems, and human oversight of AI systems should be enabled.

This section supports agencies to uphold the AI Ethics Principle of accountability. It is intended to help agencies consider its responsibility for the outcomes of the AI systems that they design, develop, deploy and operate.

11.1 Ensuring accountability during the lifecycle of the AI system

Consider the mechanisms you should put in place to ensure accountability for AI systems and their outcomes. This includes both before and after their design, development, deployment and operation.

Are there mechanisms in place to ensure responsibility and accountability by the agency and relevant individuals for AI systems and their outcomes?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

12. Use case review and next steps

12.1 Alignment with relevant legal frameworks

In addition to broader legislative obligations such as security, information, privacy, discrimination and human rights considered in sections 5 to 11, agencies must consider alignment with legislation and regulatory instruments specific to their context.

Have you identified and documented all agency specific legislation, or regulatory instruments that are relevant to the AI use case?

- ☐ Yes
- ☐ No
- ☐ Not applicable

Explain your answer

Enter your answer text here

12.2 Legal advice

Identifying and appropriately documenting the need for legal advice is an important part of managing risk and impacts. This section asks you to confirm if legal advice was sought during the assessment process and where that advice is recorded.

This information should not be disclosed to anyone other than those who need to know or access the information within the agency.

Have you identified the need for legal advice at any stage during the assessment process?

- ☐ Yes
- ☐ No
- ☐ Not applicable

If yes, where is the legal advice stored? For example, note the document management system, case file or other repository:

Enter your answer text here

12.3 Risk summary table

In the table below, list any inherent risks rated as **medium** or **high** in section 3 of the threshold assessment. Briefly explain the mitigations or controls that have been or will be applied and how these mitigations have influenced the residual risk rating.

Risk title	Risk treatments	Residual risk
<i>Enter risk title</i>	<i>Explain risk treatments</i>	☐ Low ☐ Medium ☐ High

12.4 Overall residual risk rating

Based on the assessment of residual risk in section 12.3, choose an overall residual risk rating for the AI use case.

Residual risk rating (select one)

- ☐ Low
- ☐ Medium
- ☐ High

Rationale for overall residual risk rating and explanation of relevant risk controls

Enter your answer text here

12.5 Internal governance body review

If your use case's inherent risk is rated **high** at section 3, you are required under the updated AI policy to apply specific actions, including submitting it to an appropriate internal governance body for review. Refer to the AI policy for further information on this requirement.

Your agency may also have specific requirements for referring AI use cases to an internal governance body, separate from the AI policy. Refer to your agency's AI-specific policies for more information.

Record the outcome of internal governance body review in the table below, including any recommendations and agreed next steps, if applicable.

Internal agency governance body recommendation

Enter your answer text here

Attachment A: Stakeholder mapping aid

Assessment section 2.4 asks you to identify stakeholder groups that may be affected by your AI use case, including its underlying AI system(s). The table below may assist with this exercise.

Please note the table has been provided as a discussion prompt and is not intended as a prescriptive or comprehensive list.

Stakeholder type	Identified use case stakeholders	Potential impacts (positive or negative)
End users People who will use the AI system and / or interpret its outputs.		
Evaluation or decision subjects People or groups who will be evaluated or monitored by the AI system (e.g. who the system is making predictions or recommendations about).		
Oversight team The person or team who is managing, operating, overseeing or controlling and monitoring the AI system throughout its lifecycle, including information managers.		
System owner or deployer The executive or executives responsible for authorising particular uses of an AI system.		
AI model or AI system engineers Those involved in AI model or system design, development and maintenance.		

Stakeholder type	Identified use case stakeholders	Potential impacts (positive or negative)
Rights holders Those who hold the rights to materials used by this AI use case (e.g. copyright owners or creators)		
Malicious actors Those who may intentionally misuse the AI system.		
Bystanders People who may be affected by the AI system without actively engaging with it, including those in the system's physical operating environment.		
Regulators Authorities responsible for creating, enforcing, or monitoring compliance with laws, regulations, and standards relevant to the AI system or its use.		
Other government agencies Agencies beyond your own agency that may be impacted by, collaborate on, or have policy or oversight responsibilities related to the AI system and its use.		
Civil society organisations Non-governmental organisations, advocacy groups, and community representatives concerned with ethical, social, or rights-based implications of the AI system and its use.		

Stakeholder type	Identified use case stakeholders	Potential impacts (positive or negative)
Communities or groups Communities or groups that are likely to be affected by the AI system or its use, including those that may be vulnerable or disproportionately impacted.		
Associated parties Third parties indirectly impacted by the AI system's evaluations or decisions, as well as those with an interest arising from their connection to other stakeholders—such as individuals, organisations, businesses, or industries.		
Staff and other personnel Personnel whose roles and workflows may be affected by your AI use case, including agency staff and their unions, as well as contractors and other third-party personnel.		
Intermediaries A facilitator or agent between 2 parties whose role may evolve with AI integration (e.g. tax agents).		



Guidance for the artificial intelligence impact assessment tool

The Digital Transformation Agency (DTA) provides the AI impact assessment tool and supporting guidance to assist Australian Government agencies to assess their proposed use of artificial intelligence (AI).

Agencies should not treat the tool or guidance as legal advice or as authorising proposed AI use.

Agencies are responsible for any decisions relating to their use of AI and for seeking technical and legal advice as appropriate.

For questions and suggestions on the impact assessment tool and guidance, please email ai@dta.gov.au



Digital Transformation Agency



© Commonwealth of Australia (Digital Transformation Agency) 2025

With the exception of the Commonwealth Coat of Arms and where otherwise noted, this product is provided under a Creative Commons Attribution 4.0 International Licence. (<http://creativecommons.org/licenses/by/4.0/legalcode>)

The Digital Transformation Agency (DTA) has tried to make the information in this product as accurate as possible. However, it does not guarantee that the information is totally accurate or complete. Therefore, you should not solely rely on this information when making a commercial decision.

The DTA is committed to providing web accessible content wherever possible. If you are having difficulties accessing this document, please email ai@dtg.gov.au.

Document control

This document is maintained by the DTA and provides practical advice for Australian Government staff using the AI impact assessment tool. The tool supports implementation of the Policy for the responsible use of AI in government (the AI policy).

It will be updated periodically as the policy and technology evolve. For enquiries, please email ai@dtg.gov.au.

To view the AI policy, and to check you have the most up-to-date versions of the AI impact assessment tool and guidance, please visit <https://digital.gov.au/ai/ai-in-government-policy>.

Version	Description	Date
v1.0	Published	01/12/2025



Contents

Introduction 4

1. Basic information 5

2. Purpose and expected benefits 10

3. Inherent risk assessment 12

4. Threshold assessment outcome 14

5. Fairness 15

6. Reliability and safety 18

7. Privacy protection and security 27

8. Transparency and explainability 30

9. Contestability 37

10. Human-centred values 40

11. Accountability 42

12 Use case review and next steps 43

Appendix: Risk consequence guidance table 45



Introduction

This document provides guidance for completing the Australian Government artificial intelligence (AI) impact assessment tool (the tool). Use this supporting guidance to understand, interpret and complete the tool.

The impact assessment tool is for Australian Government teams working on an AI use case. It helps teams identify, assess and manage AI use case impacts and risks against [Australia's AI Ethics Principles](#).

Agencies can use the tool to fulfill AI use case impact assessment requirements under the [Policy for the responsible use of AI in government](#) (the AI policy). Refer to the impact assessment tool document for assessment instructions and to the AI policy for key definitions and implementation requirements.

This guidance mirrors the AI impact assessment tool's 12-section structure. For advice on completing a section in the tool, find the corresponding section number in this guidance.

Assessing officers should familiarise themselves with the AI policy and the AI Ethics Principles. Also consider other Digital Transformation Agency (DTA) resources designed to support government AI adoption, including:

- the [AI technical standard](#)
- AI procurement resources, including the Guidance on [AI procurement in government](#), [AI contract template](#) and [AI model clauses](#)
- guidance on the use of public generative AI tools [for agencies](#) and [for staff](#).

The DTA welcomes user feedback on the tool and supporting guidance. Please send questions or suggestions to ai@dta.gov.au.

1. Basic information

1.1 AI use case profile

This section is intended to record basic information about the AI use case.

Name of AI use case

Choose a clear, simple name that accurately conveys the nature of the use case.

Internal reference number or identifier

Assign a unique reference number or other identifier for your assessment. This is intended to assist with internal record keeping and engagement with the DTA.

Lead agency

The agency with primary responsibility for the AI use case. Where 2 or more agencies are jointly leading, nominate one as the contact point for assessment.

1.2 Establishing impact assessment responsibilities

Assessing officer

An officer assigned to complete the assessment, coordinate the end-to-end process and serve as the contact point for any assessment queries. Depending on the use case and agency context, they may be a technical, data, governance or risk specialist, or a policy or project officer from the business area implementing the AI use case in its operations.

Accountable use case owner(s)

This role is described in the AI policy and the Standard for accountability.

Approving officer

This should be an officer with appropriate authority to approve the AI use case assessment, including the inherent risk ratings. Similar to the assessing officer role above, the approving officer's specific role in the AI use case will depend on the agency and use case context.

1.3 Additional roles and responsibilities

Clear roles and responsibilities are essential for ensuring accountability in the development and use of AI systems. In this section, you are asked to identify any additional individual officers that may have responsibilities related to your AI use case and the underlying AI system(s). Consider the roles already outlined – such as assessing officers or approving officers – and consider other positions that contribute to the AI system's lifecycle or oversight.

This is not intended to create new requirements for specific roles under the AI policy or this impact assessment. It is intended to help agencies record relevant roles and responsibilities, maintain transparency and facilitate accountability during AI use case implementation. For example, you could identify the person(s) responsible for:

- the decision to use the AI system and the scope of the AI system
- designing, developing and maintaining the system, such as key personnel of third-party suppliers
- applying interpreting the AI system's outputs, including decisions or actions based on those outputs
- controlling the AI system, with authority to start, stop, or deactivate the system under normal operating conditions
- monitoring and maintaining performance and safety, meeting quality standards and detecting errors, biases, and unintended consequences
- disengaging or stopping the system, if immediate intervention is required to prevent or stop harm
- the governance of the data used for operating, training or validating the AI system

You should consider distributing these roles among multiple officers where feasible, to avoid excessive concentration of responsibilities in a single individual, while ensuring responsible officers are appropriately skilled and senior.

1.4 AI use case description

Briefly explain how you are using or intending to use AI. This should be an 'elevator pitch' that gives the reader a clear idea of the kind of AI use intended, without going into unnecessary technical detail, which is captured in your other project documentation. Use simple, clear language, avoiding technical jargon where possible. You may wish to include:

- a high level description of the problem the AI use case is trying to solve
- the way AI will be used
- the outcome it is intended to achieve.

1.5 In-scope use case

Record whether your AI use case is in scope of the [Policy for the responsible use of AI in government](#) (the AI policy). [Appendix C](#) of the AI policy specifies the criteria to determine if an AI use case is in scope.

At a minimum, an AI use case is in scope of the AI policy if any of the following apply:

- The use, misuse or failure of AI could lead to more than insignificant harm to individuals, communities, organisations, the environment or the collective rights of cultural groups including First Nations peoples.
- The use of AI will materially influence administrative decisions that affect individuals, communities, organisations, the environment or the collective rights of cultural groups including First Nations peoples.
- It is possible the public will directly interact with, or be significantly impacted by, the AI or its outputs without human review.
- The AI is designed to use personal or sensitive data¹ or security classified information².
- It is deemed an elevated risk AI use case as directed by the DTA.

Agencies may wish to apply the AI policy to AI use cases that do not meet the above criteria. This includes use cases with specific characteristics or factors unique to an agency's operating environment that may benefit from applying an impact assessment and governance actions.

The AI policy has been designed to exclude incidental and lower risk uses of AI that do not meet the criteria. Incidental uses of AI may include off-the-shelf software with AI features such as grammar checks and internet searches with AI functionality. The AI policy recognises that incidental usage of AI will grow over time and focuses on uses that require additional oversight and governance.

In assessing whether a use case is in scope, agencies should also carefully consider AI use in the following areas:

- recruitment and other employment-related decision making
- automated decision making of discretionary decisions
- administration of justice and democratic processes
- law enforcement, profiling individuals, and border control
- health
- education
- critical infrastructure.

While use cases in these areas are not automatically high-risk, they are more likely to involve risks that require careful attention through an impact assessment.

For information on how the policy applies when doing early-stage experimentation, refer to [Appendix C](#) of the AI policy.

¹ As defined by the Privacy Act 1988 (Cth).

² As defined by the Australian Government Protective Security Policy Framework

If your use case is within scope, record the AI policy criteria that apply using the checklist. If your use case meets multiple criteria, tick each one. If you are unsure, it is best practice to select the criteria that most closely reflect your use case.

The criteria are designed to help identify uses of AI that require additional oversight and governance. This provides a clearer picture of the types of uses across government.

Refer to the AI policy for further detail on mandatory use case governance actions. Consult your accountable use case owner and your agency's AI accountable official for agency-specific guidance on fulfilling the mandatory AI policy actions and any internal agency requirements in addition to the mandatory actions.

Note you can also apply the AI policy and impact assessment tool to use cases that do not meet the criteria. In this case, you can select 'not applicable' for this question.

1.6 Type of AI technology

Briefly explain what type of AI technology you are using or intend to use. For example, supervised or unsupervised learning, computer vision, natural language processing, generative AI.

This may require a more technical answer than the use case description. Aim to be clear and concise with your answer and use terms that a reasonably informed person with experience in the AI field would understand.

1.7 Usage pattern

Select the AI system usage pattern or patterns that apply to your use case. For usage pattern definitions, refer to the [Classification systems for AI use](#).

1.8 Administrative decisions

Only complete this section if you selected 'Decision-making and administrative action' in assessment section 1.7 and if AI automated decision-making is used for an administrative decision under an Act.

Express legislative authority is generally required to automate decision-making of an administrative decision under an Act. Legal advice should be obtained for any proposed use of AI in this context.

Agencies using automated decision-making should review the [Commonwealth Ombudsman's Better Practice Guide on Automated Decision Making](#).

Agencies should generally consider:

- any legislation or framework which requires a particular decision-maker such as a minister to make a decision – for example, the minister may not have discharged their duty if they rely on the AI outputs without proper validation

- an official's duty of care and diligence under the *Public Governance, Performance and Accountability Act 2013* (Cth) – for example, an official fails to validate AI outputs
- administrative law requirements of legality and procedural fairness – for example, how an automated decision can be challenged.

1.9 Domain

Select the AI system domain or domains that apply to your use case. For domain definitions, refer to the [Classification systems for AI use](#).

1.10 Expert contributions

List any expert consultation undertaken during the assessment process, including the nature of their expertise, the specific contributions they made, and how their input informed the assessment process. While such consultation is not mandatory, agencies should consider engaging relevant internal or external expertise based on the complexity, novelty or potential impacts of the AI system.

Agencies should:

- Identify the areas where expert input is most needed, such as ethical considerations, legal risk and compliance, and technological challenges.
- Engage with a diverse range of experts to ensure a comprehensive assessment.
- Document the consultation process thoroughly, including the date of consultation, the experts' names and affiliations, and their key recommendations.
- Summarise how the expert input was integrated into the assessment and any resulting changes to the AI system or its deployment.
- Review and update the record of expert consultations periodically to ensure it remains relevant and accurate.

1.11 Impact assessment review log

As new information becomes available or design choices are refined, you should reassess all identified risks and consider whether previous responses still reflect the current state of the project. When data sources, functionality, user groups or other project elements change, revise previous answers to maintain clear and accurate records of the risk profile.

The AI policy specifies requirements for monitoring, evaluating and re-validating use cases following deployment. For details, refer to the AI policy.

2. Purpose and expected benefits

2.1 Problem definition

Describe the problem that you are trying to solve.

For example, the problem might be that your agency receives a high volume of public submissions, and that this volume makes it difficult to engage with the detail of issues raised in submissions in a timely manner.

Do not describe how you plan to fix the problem or how AI will be used.

Though 'problem' implies a negative framing, the problem may be that your agency is not able to take full advantage of an opportunity to do things in a better or more efficient way.

2.2 AI use case purpose

Clearly and concisely describe the purpose of your use of AI, focusing on how it will address the problem you described at section 2.1.

Your answer may read as a positive restatement of the problem and how it will be addressed.

For example, the purpose may be to enable you to process public submissions more efficiently and effectively and engage with the issues that they raise in more depth.

2.3 Non-AI alternatives

Briefly outline non-AI alternatives that could address the problem you described at section 2.1.

Non-AI alternatives may have advantages over solutions involving AI. For example, they may be cheaper, safer or more reliable.

Considering these alternatives will help clarify the benefits and drawbacks of using AI and help your agency make a more informed decision about whether to proceed with an AI based solution.

2.4 Identifying stakeholders

Conduct a mapping exercise to identify the individuals or groups who may be affected by the AI use case. Consider holding a workshop or brainstorm with a diverse team to identify the different direct and indirect stakeholders of your AI use case.

The stakeholder mapping aid attached to the impact assessment tool may help generate discussion on the types of stakeholder groups to consider. Please note the table has been provided as a prompt to aid discussion and is not intended as a prescriptive or comprehensive list.

2.5 Expected benefits

This section requires you to explain the expected benefits of the AI use case, considering the stakeholders identified in the previous question. The AI Ethics Principles specify that throughout their lifecycle, AI systems should benefit individuals, society and the environment.

This analysis should be supported by specific metrics or qualitative analysis. Metrics should be quantifiable measures of positive outcomes that can be measured after the AI is deployed to assess the value of using AI. Any qualitative analysis should consider whether there is an expected positive outcome and whether AI is a good fit to accomplish the relevant task, particularly when compared to the non-AI alternatives you identified previously. Benefits may include gaining new insights or data.

Consider consulting the following resources for further advice:

- [DTA Benefits management policy](#)
- [Australian Government Guide to Policy Impact Analysis.](#)

3. Inherent risk assessment

To complete the inherent risk assessment, follow these steps.

Sections 3.1 to 3.8

Definitions

Inherent risk: reflects the level of risk that exists before any additional or new controls are applied. This is the risk level under standard operating conditions, assuming only existing baseline or standard controls are in place.

Residual risk: reflects the level of risk that remains after new or additional treatments, controls or safeguards have been implemented.

Determine risk likelihood and consequence

For each risk category listed in the assessment table, determine the likelihood and consequence of the risk occurring for your AI use case. The likelihood descriptors are provided in **Table 1** of the impact assessment tool, and consequence descriptors are in the **appendix** to this guidance.

The inherent risk assessment should reflect the intended scope and function of the AI use case. In conducting your assessment, you should be clear on:

- key factors contributing to the likelihood and consequence of the risk
- any assumptions or uncertainties affecting your risk assessment.

Determine inherent risk rating

Use the risk matrix provided in **Table 2** of the impact assessment tool to determine the risk rating for each category.

Provide explanations

Provide clear and concise explanations for each risk rating.

Key considerations

When completing the inherent risk assessment, keep the following in mind:

- Try to be objective and honest in your assessment of risks. Underestimating risks at this stage could lead to inadequate risk management.
- Determining risk ratings can be challenging. Seek guidance from others to assist you, especially subject matter experts and those experienced in safe and responsible AI risk assessments.
- Consider the perspectives of stakeholders, including those identified at section 2.4, in assessing the likelihood and consequence of risks.
- Consider the perspectives of marginalised groups, including First Nations people, especially in relation to risks relating to discrimination and stereotyping. You may not have the background or life experience to fully appreciate these risks.
- Consider both intended and unintended consequences and outcomes. This includes evaluating the impact of system failure, the impact of misuse or malicious use and other deviations from expected use.
- Where there is uncertainty or disagreement about the appropriate inherent risk rating, choose the higher rating.
- Document key assumptions or evidence used in determining the inherent risk rating, as this will help explain the rationale for your assessment to reviewers.
- Consider the expected benefits of the AI use case before deciding whether to proceed based on significant but mitigable risks.

3.9 Overall inherent risk rating

In this section, you are required to determine the threshold risk rating for the AI use case based on the ratings selected in the sections above. The highest risk rating identified in any earlier sections must be used as the overall risk rating.

Resources

CSIRO Data61 [Responsible AI Pattern Catalogue](#)

Massachusetts Institute of Technology (MIT) [AI Risk Repository](#)

United States National Institute of Standards and Technology (NIST) [AI Risk Management Framework \(AI RMF\)](#) and [AI RMF Playbook](#)

4. Threshold assessment outcome

4.1 Assessing officer recommendation

Once completed, if the assessing officer is satisfied all risks are low, they may recommend that a full assessment is not required and that the approving officer accept the low risks and endorse the use case.

If one or more risks are medium or higher, the assessing officer must either:

- complete a full assessment
- amend the scope, or function to a point where the threshold assessment results in a low risk rating
- decide to not accept the risk and not proceed with the AI use case.

4.2 Approving officer review

Once the assessing officer has made their recommendation, the approving officer must:

- review the recommendation
- confirm whether they are satisfied by the supporting analysis
- agree that a full assessment is or is not necessary for the use case.

5. Fairness

5.1 Defining fairness

Fairness is a core principle in the design and use of AI systems, but it is a complex and contextual concept. Australia's AI Ethics Principles state that AI systems should be inclusive and accessible and should not involve or result in unfair discrimination. However, there are different and sometimes conflicting definitions of fairness and people may disagree on what is fair.

For example, there is a distinction between:

- individual fairness – treating individuals similarly
- group fairness – similar outcomes across different demographic groups.

Different approaches to fairness involve different trade-offs and value judgments. The most appropriate fairness approach will depend on the specific context and objectives of your AI use case.

When defining fairness for your AI use case, you should be aware that AI models are typically trained on broad sets of data that may contain bias. Bias can arise in data where it is incomplete, unrepresentative or reflects societal prejudices. AI models may reproduce biases present in the training data, which can lead to misleading or unfair outputs, insights or recommendations. This may disproportionately impact some groups, such as First Nations people, people with disability, LGBTIQ+ communities and culturally and linguistically diverse communities. For example, an AI tool used to screen job applicants might systematically disadvantage people from certain backgrounds if trained on hiring data that reflects past discrimination.

When defining fairness for your AI use case, consider the inclusivity and accessibility of the AI. AI can lead to unfairness if it creates barriers for individuals or groups who wish to access government services. For example, an AI chatbot designed to provide social security information may produce unfair outcomes because it is more difficult for vulnerable or underrepresented groups to access the digital technologies required to access the chatbot.

When defining fairness for your AI use case, it is recommended that you:

- consult relevant domain experts, affected parties and stakeholders (such as those you have identified at assessment section 2.4) to help you understand the trade-offs and value judgements that may be involved
- document your definition of fairness in your response to assessment section 5.1, including how you have balanced competing priorities and why you believe it to be appropriate to your use case
- be transparent about your fairness definition and be open to revisiting it based on stakeholder feedback and real-world outcomes.

You should also ensure that your definition of fairness complies with anti-discrimination laws. In Australia, it is unlawful to discriminate on the basis of a number of protected attributes including age, disability, race, sex, intersex status, gender identity and sexual orientation, in certain areas of public life including education and employment. Australia's federal anti-discrimination laws are contained in the following legislation:

- *Age Discrimination Act 2004*
- *Disability Discrimination Act 1992*
- *Racial Discrimination Act 1975*
- *Sex Discrimination Act 1984*.

Where the AI will produce information or be involved in decision-making, you should also ensure that your definition of fairness reflects the administrative law principle of procedural fairness, which requires that decision-making is transparent and challengeable.

5.2 Measuring fairness

You may be able to use a combination of quantitative and qualitative approaches to measuring fairness. Quantitative fairness metrics can allow you to compare outcomes across different groups and assess this against fairness criteria. Qualitative assessments, such as stakeholder engagement and expert review, can provide additional context and surface issues that metrics alone might miss.

Quantifying fairness

The specific quantitative metrics you use to measure fairness will depend on the definition of fairness you have adopted for your use case. When selecting fairness metrics, you should:

- choose metrics that align with your fairness definition, recognising the trade-offs between different fairness criteria and other objectives like accuracy
- confirm if you have appropriate data to assess those metrics, including compliance with the [Australian Privacy Principles](#) where personal or sensitive information is being collected and used
- set clear and measurable acceptance criteria (see guidance for section 6.4)
- establish a plan for monitoring these metrics (see section 6.6) and processes for remediation, intervention or safely disengaging the AI system if those thresholds are not met

For examples of commonly used fairness metrics, see the [Fairness Assessor Metrics Pattern](#) from the CSIRO's Data61 unit.

Qualitatively assessing fairness

Consider some of these qualitative approaches, which may be useful to overcome data limitations and to surface issues that metrics may overlook.

Stakeholder engagement

Consult affected communities, stakeholders and domain experts to understand their perspectives and identify potential issues

User testing and feedback

Test your AI system with diverse users and solicit their feedback on the fairness and appropriateness of the system's outputs. Seek out the perspectives of marginalised groups and groups that may be impacted by the AI system.

Expert review

Engage experts, such as AI ethicists or accessibility and inclusivity specialists, to review the fairness of your system's outputs and the overall approach to fairness. Identify potential gaps or unintended consequences.

Resources

For advice on bias measurement and minimisation techniques, see the National AI Centre's report on [Implementing Australia's AI Ethics Principles](#).

CSIRO Data61's [Responsible AI Pattern Catalogue](#) includes a [Fairness Assessor Metrics Pattern](#).

Consider resources on fairness in AI in the OECD [Catalogue of Tools & Metrics for Trustworthy AI](#).

6. Reliability and safety

6.1 Data suitability

The data used to operate, train and validate your AI system has a significant impact on its performance, fairness and safety. In your answer, explain why the chosen data is suitable for your use case. Some relevant considerations are outlined below.

When choosing between datasets, consider whether the data can be separated by marginalised groups, particularly by Indigenous status identifiers. If the data is Indigenous data, see section 6.2 below, which references the [Framework for Governance of Indigenous Data](#).

Agencies should also refer to the Australian Public Service (APS) Data Ethics Framework for guidance on managing and using data and analytics ethically in government, including where AI is used in analytics. The framework is underpinned by 3 key principles: trust, respect and integrity. It provides advice on implementation across different major use cases and agency operations and encourages agencies to assess potential risks and benefits, consider fairness and inclusivity, and engage with stakeholders where appropriate. Visit the Department of Finance website to access the [APS Data Ethics Framework](#).

Data quality should be assessed prior to use in AI systems. Agencies should select applicable metrics to determine a data set's quality and identify any remediation required before using it for training or validation in AI systems. Relevant metrics to consider include diversity, relevance, accuracy, completeness, timeliness, validity and lack of duplication. One method to ensure good quality data is to set minimum thresholds appropriate to specific use cases, such as through acceptance criteria discussed below at section 6.4. An example of a specific framework for determining data quality in statistical uses is the [ABS Data Quality Framework](#).

Where third party material or data is being used to operate, train or validate an AI system, it is important to protect the rights of intellectual property holders. If the AI may use, modify or otherwise handle material in which intellectual property exists, agencies should confirm that both the following are true:

- the AI provider holds the necessary intellectual property rights in the AI output material
- the agency holds the necessary intellectual property rights in the input material.

The AI may otherwise infringe third party intellectual property rights.

Agencies should also confirm that the AI system has safeguards in place to prevent the unauthorised use or disclosure of confidential information.

Where data used to operate, train and validate the AI system includes personal information, agencies should confirm that collection, use and disclosure is in accordance with the Australian Privacy Principles (APPs) under the *Privacy Act 1988* (Cth) (see section 7 below).

The relevance of the data used in training the AI model may influence the output and may not be relevant to the use case and Australian context. Consider whether the model is likely to make accurate or reliable predictions concerning matters relating to Australian subject matter if it has been trained on, for example, US-centric data.

You should also consider data provenance, lineage and volume – as outlined below:

Data provenance

Involves keeping records of the data collected, processed and stored by the AI system and creating an audit trail to assign custody and trace accountability for issues. It provides assurance of the chain of custody and its reliability, insofar as origins of the data are documented.

Data lineage

Involves documenting data origins and flows to enable stakeholders to better understand how datasets are constructed and processed. This fosters transparency and trust in AI systems.

Data volume

Consider the volume of data you need to support the operation, training and validation of your AI system.

6.2 Indigenous data

Describe how any components of your AI system have used or will use Indigenous data, or where any outputs relate to First Nations individuals, communities or groups.

All Australian Public Service (APS) agencies are required to implement the Framework for Governance of Indigenous Data. This framework adopts the definition of 'Indigenous data' as provided by Maïam nayri Wingara Indigenous Data Sovereignty Collective:

Information or knowledge, in any format or medium, which is about and may affect Indigenous peoples both collectively and individually.

If the data used to operate, train or validate your AI system, or any outputs from your AI system, meet this definition of Indigenous data, refer to the Framework for Governance of Indigenous Data for guidance on applying the framework.

The framework is based on the principles of:

- respect for cultural heritage
- informed consent
- privacy, including collective or group privacy
- trust.

The Framework for Governance of Indigenous Data is also informed by 2 complementary data governance frameworks:

- FAIR Guiding Principles (Findable, Accessible, Interoperable, Reusable) – providing technical standards for scientific data management and stewardship
- CARE Principles (Collective Benefit, Authority to Control, Responsibility, Ethics) – which focuses on Indigenous data governance, reflecting the crucial role of data in self-determination.

Relevant practices to consider in this context include:

- Checking if datasets used to train the AI included diverse and representative samples of cultural expression, artifacts, languages and practices. This supports the AI system being able to recognise and appropriately respond to a greater range of cultural contexts in a less biased manner.
- Describing any mechanisms in place for engaging with Indigenous individuals, communities or group representatives and collecting and incorporating their feedback on the AI system's performance, especially regarding cultural aspects.
- Describing processes to review documentation and protocols that ensure the project has incorporated the GID principles. Look for evidence of meaningful engagement with and input from suitably qualified and experienced Indigenous individuals, communities and groups. Assess if the system includes features or options that allow Indigenous stakeholders to control how their data is used and represented and describe how benefits of the project to First Nations Peoples, to which the data relate, have been considered.

Also consider the use of Indigenous data in the context of the United Nations Declaration on the Rights of Indigenous Peoples and apply the concept of 'free, prior and informed consent' in relation to the use of Indigenous data in AI systems.

6.3 Suitability of procured AI model

If you are procuring an AI model or system from a third-party provider, your procurement process should consider if the provider has appropriate data management including data quality and data provenance in relation to the model. This will help you to identify whether the AI model is fit for the context and purpose of your AI use case.

This may include:

- governance
- data sourcing
- privacy
- security
- intellectual property
- cybersecurity practices.

There are many other considerations you should take into account when selecting a procured AI model and contracting with a supplier. The following considerations may be relevant to your use case:

- Determine if data will be hosted overseas and if it could be subject to foreign laws. Consider the potential for foreign ownership, control, or influence (FOCI) and refer to the Department of Home Affairs FOCI Risk Assessment Guidance.
- Determine if processes and practices are in place to address risks along the supplier's supply chain, such as risks relating to FOCI, security, transparency and business practices). Agencies should refer to Australian Signals Directorate guidance on cyber supply chain risk management.
- Assess whether the AI model meets the functional requirements for your use case.
- Determine how the model was evaluated, including the test data and benchmarks used.
- Determine how versioning for the AI model is handled.
- Consider the support the supplier provides for users and procurers.
- Review provisions regarding potential liability issues and clarify accountability between your agency and the provider if the product fails.
- Establish security precautions, such as handover or destruction of agency data upon termination or expiry of the procurement contract, and identify any residual risks and mitigation measures.
- Confirm what controls the agency has if the AI system malfunctions, produces harmful outputs or behaves in an unintended way.
- Review any guarantees that data handling and management across the entire lifecycle of the data meet internal agency and legislative requirements.
- Review any warranties the supplier will provide, such as suitability of the AI system for the intended use, absence of defects, and development with reasonable care and skill.
- Ensure the supplier has a contractual obligation to comply with relevant legislation and frameworks, including privacy, discrimination, AI Ethics Principles, the AI policy and the Protective Security Policy Framework (PSPF).
- Clarify any supplier responsibilities for training, monitoring and validation of the AI system.
- Clarify ownership of intellectual property rights in relation to the AI model, inputs, outputs and other materials such as user manuals or technical documentation.
- Review measures taken to prevent or reduce hallucinations, unwanted bias and model drift. For example, evaluation of training data for harm or bias and adjustments made to compensate.
- Assess whether the level of human oversight, transparency, explainability and interpretability of the model is sufficient for your use case.
- Specify the kinds of records the supplier will provide to the agency, such as records of how agency data is used by the AI system.
- Determine the computing and storage capacity requirements for operating the model on premises.
- Assess the capability needed to maintain the AI model and whether this can be done in-house or require external sourcing.

- If using a platform as a service (PaaS) to run and support your AI system or AI model, consider risks associated with outsourcing.
- Evaluate whether the AI system could be designed or influenced to promote certain products or services, and how such behaviour could be detected and addressed. For example, if the supplier accepts advertising or sponsorship to give prominence to products or services.

Consider also how your agency will support transparency across the AI supply chain, for example, by notifying the developer of issues encountered in using the model or system.

Refer to the DTA's AI procurement resources including the:

- Guidance on AI procurement in government
- [AI contract template](#)
- Digital Sourcing ClauseBank [AI model clauses](#).

6.4 Testing

Testing is a key element for assuring the responsible and safe use of AI models – for both models developed in-house and externally procured – and in turn, of AI systems. Rigorous testing helps validate that the system performs as intended across diverse scenarios. Thorough and effective testing helps identify problems before deployment.

Testing AI systems against test datasets can reveal biases or possible unintended consequences or issues before real-world deployment. Testing on data that is limited or skewed can fail to reveal shortcomings.

Consider establishing clear and measurable acceptance criteria for the AI system that, if met, would be expected to control harms that are relevant in the context of your AI use case. Acceptance criteria should be specific, objective and verifiable. They are meant to specify the conditions under which a potential harm is adequately controlled.

Consider developing a test plan for the acceptance criteria to outline the proposed testing methods, tools and metrics. Documenting results through a test report will assist with demonstrating accountability and transparency. A test report could include the following:

- a summary of the testing objectives, methods and metrics used
- results for each test case
- an analysis of the root causes of any identified issues or failures
- recommendations for remediation or improvement, and whether the improvements should be done before deployment or as a future release.

In your explanation, outline any areas of concern in results from testing. If the AI system has not yet undergone testing, outline elements to be considered in testing plans.

Model accuracy

As an example, model accuracy is a key metric for evaluating the performance of an AI system. Accuracy should be considered in the specific context of the AI use case, as the consequences of errors or inaccuracies can vary significantly depending on the domain and application. This can include:

- unfairness – for example, where a decision has been made based on inaccurate data
- breach of individual rights – for example, where information produced by AI is defamatory
- non-compliance with legislation – for example, presenting false or misleading information in breach of Australian Consumer Law, or acting in a discriminatory manner in breach of anti-discrimination laws.

Some of the factors that can influence AI model output accuracy and reliability include:

- choice of AI model or model architecture
- quality, accuracy and representativeness of training data
- presence of bias in the training data or AI model
- robustness to noise, outliers and edge cases
- ability of the AI model to generalise to new data
- potential for errors or ‘hallucinations’ in outputs
- environmental factors (such as lighting conditions for computer vision systems)
- adversarial attacks (such as malicious actors manipulating input data to affect outputs)
- stability and consistency of performance over time
- whether AI model allows for sponsorship or advertising to give prominence to certain outputs.

Ways to assess and validate the accuracy of your model for your AI use case include:

- quantitative metrics
- qualitative analysis – such as manual review of output, error analysis, and user feedback
- domain-specific benchmarks or performance standards
- comparison to human performance or alternative models.

It is important to set accuracy targets that are appropriate for the risk and context of the use case.

For high stakes decisions, you should aim for a very high level of accuracy and have clear processes for handling uncertain or borderline cases.

6.5 Pilot

Conducting a pilot study is a valuable way to assess the real-world performance and impact of your AI use before full deployment. A well-designed pilot can surface issues related to reliability, safety, fairness and usability that may not be apparent in a controlled development environment.

If you are planning a pilot, your explanation should provide a brief overview of the pilot's:

- scope and duration
- objectives and key results (OKRs)
- key performance indicators (KPIs)
- participant selection and consent process
- risk mitigation strategies.

If you have already completed a pilot, reflect on the key findings and lessons learned, including by:

- assessing how the pilot outcomes compared to your expectations.
- identifying any issues or surprises that emerged during the pilot.
- documenting how you adapted your AI use case based on the pilot results.

If you are not planning to conduct a pilot, explain why not. Consider whether the scale, risk or novelty of your use case warrants a pilot phase. Discuss alternative approaches you are taking to validate the performance of your AI use case and gather user feedback prior to full deployment.

6.6 Monitoring

Monitoring is key to maintaining the reliability and safety of AI systems over time. It enables active rather than passive oversight and governance, and ensures the agency has ongoing accountability for the AI-assisted performance and decision-making processes.

Your monitoring plan should be tailored to the specific risks and requirements of your use case. In your explanation, describe your approach to monitoring any measurable acceptance criteria (as discussed above at section 6.4) and other relevant metrics such as performance metrics or anomaly detection. In your plan, include your proposed monitoring intervals for your use case. The AI policy requires agencies to establish a clear process to address AI incidents aligned to their ICT management approach. Incident remediation must be overseen by an appropriate governance body or senior executive and should be undertaken in line with any other legal obligations.

Periodically evaluate your monitoring and evaluation mechanisms to ensure they remain effective and aligned with evolving conditions throughout the lifecycle of your AI use case. Examples of events that could influence your monitoring plan are system upgrades, error reports, changes in input data, performance deviation or feedback from stakeholders.

Monitoring can help identify issues that can impact the safety and reliability of your AI system, such as:

- concept drift – changes in the relationship between input data and the feature being predicted
- data drift – changes in input data patterns compared to the data used to train the model.

Vendors offer monitoring tools that may be worth considering for your use case. For more information on continuous monitoring, refer to the NAIC's [Implementing Australia's AI Ethics Principles report](#).

6.7 Preparedness to intervene or disengage

Relevant stakeholders, including those who operate, use or interact with the AI system, those who monitor AI system performance, and affected stakeholders identified at section 2.4, should have the ability to raise concerns about insights or decisions assisted by the AI system.

Agencies must develop clear pathways for staff or other relevant stakeholders to report AI safety concerns, including AI incidents. Agencies should also document and take appropriate steps in relation to any interventions that occur to ensure consistency and fairness.

In addition, agencies should be prepared to quickly and safely disengage an AI system when an unresolvable issue is identified. This could include a data breach, unauthorised access or system compromise. Consider such scenarios in business continuity, data breach and security response plans.

Techniques to avoid overreliance on AI system outputs

Agencies should consider the following techniques to avoid overreliance on AI system outputs.

Three techniques to consider at the **system design stage**:

- Build in transparency about system limitations, by incorporating prompts to remind users to critically analyse outputs. These could include explanations of outputs, hallucination reminders, reference source checking and accuracy scores.
- Build in 2-way feedback pathways by prompting users to assess the quality of the AI system's outputs and provide feedback. Similarly, provide feedback to users on their interactions with the systems, such as feedback on ineffective prompts or alerts when the user has accepted a risky decision.
- Build in steps that require human decision-making, for example by designing the AI system to provide options to choose from rather than accept a single outcome, prompting users to engage with and evaluate AI outputs.

At the **evaluation stage**, focus on validating whether the system supports human judgement as intended. Engage directly with users to understand their experience, encourage them to assess outputs critically and suggest improvements. Review user behaviour, feedback loops and decision-making patterns and prompts to confirm that safeguards against overreliance are effective. Use these insights to refine system design, guidance and training materials.

6.8 Training of AI system operators

AI system operators play a crucial role in ensuring the responsible and effective use of AI. They must have the necessary skills, knowledge and judgment to understand the system's capabilities and limitations, how to appropriately use the system, interpret its outputs and make informed decisions based on those outputs.

In your answer, describe the process for ensuring AI system operators are adequately trained and skilled. This may include:

Initial training

Consider what training operators receive before being allowed to use the AI system. Does this training cover technical aspects of the system, as well as ethical and legal considerations?

As a baseline, you may expect that operators:

- understand the limitations of the AI system
- are able to monitor the AI system, so that anomalies, errors and unexpected performance can be detected and addressed
- are aware of the possible tendency of relying, or over-relying, on AI outputs
- are able to correctly interpret AI outputs, taking into account the particular characteristics of the system
- are able to decide when to disregard, override or reverse the AI outputs.

Ongoing training

This includes processes for continuous learning and skill development, and for keeping officers up to date with changes or updates to the AI system.

Evaluation

This can include skills and knowledge assessment, certification or qualification requirements for operators.

Support

Ensure resources and support are available to operators if they have questions or encounter issue. Consider whether this needs to be tailored to the specific needs and risks of your AI system or proposed use case or whether general AI training requirements are sufficient.

7. Privacy protection and security

7.1 Minimise and protect personal information

Compliance with the Australian Privacy Principles

Agencies should consider how the AI use case will comply with the Australian Privacy Principles (APPs) in Schedule 1 to the *Privacy Act 1988* (Cth). The APPs apply to personal information inputted into an AI system, as well as the output generated or inferred by an AI system that contains personal information. Under the APPs:

- **APP 1:** Agencies must implement practices, procedures and systems to ensure compliance with APPs. Agencies must also have a clearly expressed and up-to-date privacy policy. This can include establishing clear processes for verification of AI outputs containing personal information, and adding transparent information about its use of AI in its privacy policy.
- **APP 3:** AI inputs or outputs generated or inferred by AI, which contain personal information, must be reasonably necessary for, or directly related to, the agency's functions or activities. Additionally, if the AI input or output comprises sensitive personal information, the individual must consent unless another exception applies. Collection must occur by fair and lawful means.
- **APP 5:** Agencies should notify individuals of AI-related purposes for which their personal information is being collected and any proposed use of AI to generate outputs which contain personal information.
- **APP 6:** Agencies may only input an individual's personal information into an AI system, or use or disclose AI outputs which contain personal information, for the primary purpose for which the agency collected the information, unless they have consent or another exception applies – for example, if the agency can establish a related secondary use would be reasonably expected by the individual.
- **APP 10:** Agencies must take reasonable steps to ensure personal information collected, used and disclosed by the AI system is accurate, up-to-date, complete and relevant.
- **APP 11:** Agencies must take reasonable steps to protect personal information from misuse, interference and loss, as well as unauthorised access, modification or disclosure.

For more information, refer to the [APP guidelines](#) and the Office of the Australian Information Commissioner (OAIC) [Guidance on privacy and the use of commercially available AI products](#).

Also consider your agency's internal privacy policy and resources and consult your agency's privacy officer.

Privacy enhancing technologies

Your agency may want or need to use privacy enhancing technologies to assist in de-identifying personal information under the APPs or as a risk mitigation/trust building approach. Where the risk of re-identification is very low, de-identified information will no longer comprise personal information and agencies can use the information in ways that the Privacy Act would normally restrict.

Consider the Office of the Australian Information Commissioner's (OAIC) detailed guidance on [De-identification and the Privacy Act](#). The OAIC has also jointly developed a resource with CSIRO Data61 on [De-identification Decision-Making Framework](#).

7.2 Privacy Threshold and/or Impact Assessment

The [Australian Government Agencies Privacy Code](#) (the Privacy Code) requires Australian Government agencies subject to the *Privacy Act 1988* to conduct a privacy impact assessment (PIA) for all 'high privacy risk projects'. A project may be a high privacy risk if the agency reasonably considers that the project involves new or changed ways of handling personal information that are likely to have a significant impact on the privacy of individuals.

To determine whether a PIA is required, you should complete a privacy threshold assessment (PTA). A PTA will help you identify your use case's potential privacy impacts and screen for factors that point to a 'high privacy risk project' requiring a PIA under the Code.

Agencies should conduct a PTA and, if required, a PIA at an early stage of AI use case development or procurement– for example, after identifying the minimum viable product. This will enable the agency to fully consider whether to proceed with the AI use case or to change the approach if the PIA identifies significant negative privacy impacts. It may be appropriate to conduct a PTA and, if required, a PIA earlier than your AI impact assessment using this tool.

If you have not completed a PTA or PIA, explain how you considered potential privacy impacts – for example, if you have determined the AI use case will not involve personal information. Privacy assessments should consider if relevant individuals have provided informed consent, where required, to the collection, use and disclosure of their personal information in the AI system's training or operation, or as an output for making inferences. Also consider any consent obtained has been recorded, including a description of processes used to obtain the consent.

For more information, refer to the Office of the Australian Information Commissioner (OAIC) [advice for Australian Government agencies on when to conduct a privacy impact assessment](#). You can also consult your agency's privacy officer and internal privacy policy and resources.

If your AI system has used or will use Indigenous data, you should also consider whether principles of collective or group privacy of First Nations people are relevant and refer to the Framework for Governance of Indigenous Data (see section 6.2).

7.3 Security risks

Agencies should consider the digital and cyber security risks associated with operation of the AI. Agencies may wish to refer to the frameworks and guidance noted below in considering what measures the AI will have in place to address security risks.

The Protective Security Policy Framework (PSPF) applies to non-corporate Commonwealth entities subject to the *Public Governance, Performance and Accountability Act 2013* (PGPA Act). Agencies should refer to the PSPF to understand security requirements relevant to AI technologies. These include managing procurement risks, incorporating and enforcing security terms in contracts, addressing FOCI risks, protecting classified information, and ensuring systems are authorised in accordance with the Information Security Manual (ISM).

You should engage with your agency's ITSA early in the AI use case development and assessment process to ensure it meets all PSPF and ISM requirements.

Agencies should implement security measures to align with Australian Signals Directorate (ASD) guidance on [AI data security](#). This outlines data security risks in the development, testing and deployment of AI, and sets out best practices for securing AI data across stages of the AI lifecycle to address these risks.

Agencies should ensure appropriate procedures are in place to address a data breach or security incident. This may include processes to mitigate the immediate consequences of a data breach or security incident and to ensure any actual or potential ongoing loss to the agency is minimised.

For further mitigation considerations for organisations to consider refer to ASD's guidance on [Engaging with AI](#). It is highly recommended that your agency engages with and implements the mitigation considerations in the guidance. This includes:

- enforcing multi-factor authentication or privileged access for AI systems
- managing backups of the AI system and training data
- ensuring the AI system is secure-by-design, including across its supply chain
- conducting periodic health checks on the AI system.

Agencies should also consider the requirements outlined in the Department of Home Affairs [PSPF Policy Advisory on OFFICIAL Information Use with Generative AI](#). These include only providing access to certain generative AI products that meet hosting and other security criteria and ensuring staff have relevant training.

8. Transparency and explainability

8.1 Consultation

You should consult with a diverse range of internal and external stakeholders at every stage of your AI use case development and deployment to help identify potential biases, privacy concerns, and other ethical and legal issues present in your AI use case. This process can also help foster transparency, accountability, and trust with your stakeholders and can help improve their understanding of the technology's benefits and limitations. Refer to the stakeholders you identified in section 2.4.

If your project has the potential to significantly impact First Nations individuals, communities or groups, it is critical that you meaningfully consult with relevant community representatives.

Consultation resources

APS Framework for Engagement and Participation

Sets principles and standards that underpin effective APS engagement with citizens, community and business and includes practical guidance on engagement methods.

Best practice consultation guidance note

This resource from the Office of Impact Analysis details the Australian Government consultation principles outlined in the Guide to Policy Impact Analysis.

Principles for engagement in projects concerning Aboriginal and Torres Strait Islander peoples

This resource from the Australian Institute of Aboriginal and Torres Strait Islander Studies (AIATSIS) provides non-Indigenous policy makers and service designers with the foundational principles for meaningfully engaging with Aboriginal and Torres Strait Islander peoples on projects that impact their communities.

8.2 Public visibility

Where appropriate, you should consider options to make the scope and goals of your AI use case publicly available. For instance, consider including this information on the relevant program page on your agency website or through other official communications. This information could include:

- use case purpose
- overview of model and application, including how the AI will use data to provide relevant outputs
- benefits

- risks and mitigations
- training data sources
- contact information for public enquiries.

All agencies in scope of the AI policy are required to publish an AI transparency statement. Your agency's AI accountable official is responsible for ensuring your agency's transparency statement complies with the AI policy. More information on this requirement is contained in the AI policy and associated [Standard for transparency statements](#). Consult your agency's AI accountable official for specific advice on your use case.

Furthermore, to comply with APP 1 and APP 5, agencies should consider updating their privacy policies with information about their use of AI. For example, to advise that personal information may be disclosed to AI system developers or owners.

Considerations for publishing

In some circumstances it may not be appropriate to publish detailed information about your AI use case. When deciding whether to publish this information you should balance the public benefits of AI transparency with the potential risks as well as compatibility with any legal requirements around publication.

For example, you may choose to limit the information you publish, or not publish any information at all, if the use case is still in the experimentation phase, or if publishing may:

- have negative implications for national security
- have negative implications for law enforcement or criminal intelligence activities
- significantly increase the risk of fraud or non-compliance
- significantly increase the risk of cybersecurity threats
- jeopardise commercial competitiveness – for example, revealing trade secrets or commercially valuable information
- breach confidentiality obligations held by the agency under a contract
- breach statutory secrecy provisions.

8.3 Maintain appropriate documentation and records

Agencies should comply with legislation, policies and standards for maintaining reliable and auditable records of decisions, testing, and the information and data assets used in an AI system. This will enable internal and external scrutiny, continuity of knowledge and accountability. For example, when responding to information requests under the *Freedom of Information Act 1982* (Cth). This will also support transparency across the AI supply chain. For example, this documentation may be useful to any downstream users of AI models or systems developed by your agency.

Agencies should document AI technologies they are using to perform government functions as well as essential information about AI models, their versions, creators and owners. In addition, artefacts used and produced by AI – such as prompts, inputs and raw outputs – may constitute Commonwealth records under the *Archives Act 1983* and may need to be kept for certain periods of time identified in records authorities issued by the National Archives of Australia (NAA). Such Commonwealth records must not be destroyed, disposed of, transferred, damaged or altered except in limited circumstances listed in the Archives Act.

To identify their legal obligations, business areas implementing AI in agencies may want to consult with their information and records management teams. The NAA can also provide advice on how to manage data and records produced by different AI use cases.

Refer to NAA advice on:

- principles and expectations for the creation and management of government business information, contained in the [Information Management Standard for Australian Government](#)
- [Information management for records created using AI technologies](#).

AI documentation types

Where suitable, you should consider creating the following forms of documentation for any AI system you build. If you are procuring an AI system from an external provider, it may be appropriate to request these documents as part of your tender process.

System factsheet/model card

A system factsheet (sometimes called a model card) is a short document designed to provide an overview of an AI system to non-technical audiences (such as users, members of the public, procurers, and auditors). These factsheets usually include information about the AI system's purpose, intended use, limitations, training data, and performance against key metrics.

Datasheets

Datasheets are documents completed by dataset creators to provide an overview of the data used to train and evaluate an AI system. Datasheets provide key information about the dataset including its contents, data owners, composition, intended uses, sensitivities, provenance, labelling and representativeness.

System decision registries

System decision registries record key decisions made during the development and deployment of an AI system. These registries contain information about what decisions were made, when they were made, who made them and why they were made (the decision rationale).

Reliability and safety documentation

It is also best practice to maintain documentation on testing, piloting and monitoring and evaluation of your AI system and use case, in line with the practices outlined in section 6.

For more on AI documentation, see [Implementing Australia's AI Ethics Principles](#).

8.4 Disclosing AI interactions and outputs

You should design your use case to inform people that they are interacting with an AI system or are being exposed to content that has been generated by AI. This includes disclosing AI interactions and outputs to internal agency staff and decision-makers, as well as external parties such as members of the public engaging with government.

When to disclose use of AI

You should ensure that you disclose when a user is **directly interacting with an AI system**, especially:

- when AI plays a significant role in critical decision-making processes
- when AI has potential to influence opinions, beliefs or perceptions
- where there is a legal requirement regarding AI disclosure (for example, updated privacy policies under APP 1 and APP 5)
- where AI is used to generate recommendations for content, products or services.

You should ensure that you disclose when someone is being exposed to **AI-generated content** including where:

- any of the content has not been through a contextually appropriate degree of fact checking and editorial review by a human with the appropriate skills, knowledge or experience in the relevant subject matter
- the content purports to portray real people, places or events or could be misinterpreted that way

- the intended audience for the content would reasonably expect disclosure
- there is a legal requirement regarding AI disclosure (for example, updated privacy policies under APP 1 and APP 5).

Exercise judgment and consider the level of disclosure that the intended audience would expect, including where AI-generated content has been through rigorous fact-checking and editorial review. Err on the side of greater disclosure – norms around appropriate disclosure will continue to develop as AI-generated content becomes more ubiquitous.

Mechanisms for disclosure of AI interactions:

When designing or procuring an AI system, you should consider the most appropriate mechanism(s) for disclosing AI interactions. Some examples are outlined below:

Verbal or written disclosures

Verbal or written disclosures are statements that are heard by or shown to users to inform that they are interacting with (or will be interacting with) an AI system.

For example, disclaimers/warnings, specific clauses in privacy policy and/or terms of use, content labels, visible watermarks, by-lines, physical signage, communication campaigns.

Behavioural disclosures

Behavioural disclosure refers to the use of stylistic indicators that help users to identify that they are engaging with AI-generated content. These indicators should generally be used in combination with other forms of disclosure.

For example, using clearly synthetic voices or formal, structured language, robotic avatars.

Technical disclosures

Technical disclosures are machine-readable identifiers for AI-generated content.

For example, inclusion in metadata, technical watermarks, cryptographic signatures.

Agencies should consider using AI systems that use industry-standard provenance technologies, such as those aligned with the standard developed by the [Coalition for Content Provenance \(C2PA\)](#).

Ability to request a non-AI alternative

In certain contexts, it may be best practice not to provide a non-AI alternative, particularly where the AI system is low-risk, improves service delivery without affecting rights or entitlements, and where alternate pathways would create unnecessary cost, complexity, or delay. However, in other situations, offering the ability to request a non-AI alternative can be important.

8.5 Offer appropriate explanations

Explainability refers to accurately and effectively conveying an AI system's decision process to a stakeholder, even if they don't fully understand the specifics of how the model works. Explainability facilitates procedural fairness, transparency, independent expert scrutiny and access to justice by ensuring that agencies have the material that is required to provide affected individuals with evidence that forms the basis of a decision when needed. To interpret the AI's output and offer an explanation to relevant stakeholders, you should consider whether the agency can access:

- the inputs from the agency
- the logic behind an individual output
- the model that the AI System uses and the sources of data for the model
- information on which features of the AI contributed to the output
- automatic records of events which allow for traceability of the AI's functioning
- any risk management measures in place which would allow the agency to understand risks and adjust use of the AI accordingly (for example, technical limitations such as error rates of an AI model).

You should be able to clearly explain how a government decision or outcome has been made or informed by AI to a range of technical and non-technical audiences. You should also be aware of any requirements in legislation to provide reasons for decisions, both generally and in relation to the particular class of decisions that you are seeking to make using AI.

Explanations may apply globally (how a model broadly works) or locally (why the model has come to a specific decision). You should determine which is more appropriate for your audience.

Principles for providing effective explanations

Contrastive

Outline why the AI system output one outcome instead of another outcome.

Selective

Focus on the most-relevant factors contributing to the AI system's decision process.

Consistent with the audience's understanding

Align with the audience's level of technical (or non-technical) background.

Generalisation to similar cases

Generalise to similar cases to help the audience predict what the AI system will do.

Tools for explaining non-interpretable models

Providing explanations is relatively straightforward for interpretable models with low complexity and clear parameters. However, in practice, most AI systems have low interpretability and require effective post-hoc explanations that balance accuracy and simplicity. Among other matters, you should also consider defining appropriate timeframes for providing explanations in the context of your use case.

When developing explanations, consider the range of available approaches based on your model type and use case.

- For traditional machine learning models, feature importance methods and visualisation techniques can help explain individual predictions or overall model behaviour.
- For neural networks and deep learning systems, specialised interpretation methods have been developed that analyse network activations, attention patterns, and gradients.
- Large language models and foundation models require distinct approaches, including prompt-based explanations and emergent interpretability techniques.
- Model-agnostic methods offer flexibility across different architectures, while example-based approaches use counterfactuals and contrastive examples to make predictions more understandable.

Advice on appropriate explanations is available in the National AI Centre's [Implementing Australia's AI Ethics Principles report](#).

Other reputable resources for explainability tools include open-source libraries maintained by academic institutions and research communities and documentation from major cloud platform providers. When selecting tools, prioritise those with active maintenance, clear documentation, and validation through published research.

However, explainable AI algorithms are not the only way to improve system explainability. Human-centred design can also play an important part, including:

- developing effective explanation interfaces tailored to different stakeholder audiences
- determining appropriate levels of detail for various contexts
- ensuring explanations are actionable and meaningful for decision-makers

9. Contestability

9.1 Notification of AI affecting rights

You should notify individuals, groups, communities or businesses when an administrative action materially influenced by an AI system has a legal or significant effect on them. This promotes transparency and access to justice, by ensuring individuals can understand how government uses AI to perform actions that affect them and have the opportunity to seek review of that decision.

This notification should state that the action was materially influenced by an AI system and include information on available review rights and how the individual can challenge the action. The notification should be clear, up-to-date, concise and understandable, and should not be complex, lengthy, legalistic or vague. It may be appropriate to provide notification prior to the action being taken or at the same time that the action occurs (for example, an applicant may be asked to acknowledge that AI will be used to a stated extent to assess their application).

An action producing a **'legal effect'** is when an individual, group, community or business's legal status or rights are affected, and includes an effect on the:

- provision of rights or benefits granted by legislation or common law
- imposition of penalties or orders (civil or criminal), and
- contractual rights.

An action producing a **'significant effect'** is when an individual, group, community or business's circumstances, behaviours, interests or choices are affected, and includes an effect on the provision of:

- critical government services or support, such as housing, insurance, education enrolment, criminal justice, employment opportunities and health, disability or aged care services
- basic necessities, such as food and water.

An action may be considered to have been **'materially influenced'** by an AI system if:

- the action was automated by an AI system, with little to no human oversight
- a component of the action was automated by an AI system, with little to no human oversight – for example, a computer performs the first 2 limbs of an action, with the final limb made by a human
- the AI system is likely to influence actions that are performed – for example, the AI system output recommended a decision to a human for consideration or provided substantive analysis to inform a decision.

‘Administrative action’ is any of the following:

- making, or refusing or failing to make, a decision
- exercising, or refusing or failing to exercise, a power
- performing, or refusing or failing to perform, a function or duty.

Advisory note

This guidance is designed to supplement, not replace, existing administrative law requirements pertaining to notification of administrative decisions. The Attorney-General's Department is leading work to develop a consistent legislative framework for automated decision-making (ADM), as part of the government's response to recommendation 17.1 of the [Robodebt Royal Commission Report](#).

9.2 Challenging administrative actions influenced by AI

Individuals, groups, communities or businesses should be provided with a timely opportunity to challenge an administrative action that has a legal or significant effect on them when the action was materially influenced by an AI system. This is an important administrative law principle. It also promotes accountability and improves the quality and consistency of government decisions.

Administrative actions may be subject to both merits review and judicial review.

Merits review

Considers whether a decision made was the correct or preferable one in the circumstances, and may include internal review conducted by the agency or external review by the Administrative Review Tribunal.

Where an action can be challenged via internal review (as permitted by relevant legislation), you should consider what processes are in place to allow for internal review of an action materially influenced by AI, for example, by another or more senior officer in the agency.

Judicial review

Examines whether an action was lawful (for example, whether the decision maker had the power to make a decision or whether a legal error has occurred in making a decision), and is limited to actions which affect an individual's liberties, vested rights or legitimate expectations.

You should ensure review rights that ordinarily apply to human-made decisions or actions are not impacted or limited because an AI system has been used.

Notifications discussed at section 9.1 should include information about available review mechanisms so that people can make informed decisions about disputing administrative actions.

Ensure a person within your agency is able to answer questions in a court or tribunal about an administrative action taken by an AI system if that matter is ultimately challenged. Review mechanisms also impact on the obligation to provide reasons. For example, the *Administrative Decisions (Judicial Review) Act 1977* gives applicants a right to request reasons for administrative decisions.

10. Human-centred values

10.1 Incorporating diversity

Diversity of perspective promotes inclusivity, mitigates biases, supports critical thinking, mitigates the risk of non-compliance with anti-discrimination laws and should be incorporated in all AI system lifecycle stages.

AI systems require input from stakeholders from a variety of backgrounds, including different ethnicities, genders, ages, abilities and socio-economic statuses. This also includes people with diverse professional backgrounds, such as ethicists, social scientists and domain experts relevant to the AI application. Determining which stakeholders and user groups to consult, which data to use, and the optimal team composition will depend on your AI system.

Failing to adequately incorporate diversity into relevant AI lifecycle stages can have unintended negative consequences, as illustrated in a number of real-world examples:

- AI systems ineffective at predicting recidivism outcomes for defendants of colour and underestimating the health needs of patients from marginalised racial and ethnic backgrounds.
- AI job recruitment systems unfairly affecting employment outcomes.
- Algorithms used to prioritise patients for high-risk care management programs were less likely to refer black patients than white patients with the same level of health.
- An AI system designed to detect cancers had shown biases towards lighter skin tones stemming from an oversight in collecting a more diverse set of skin tone images, potentially delaying life-saving treatments.

Resources, including approaches, templates and methods to ensure sufficient diversity and inclusion of your AI system, are described in the NAIC's [Implementing Australia's AI Ethics Principles report](#).

10.2 Human rights obligations

You should consult an appropriate source of advice or otherwise ensure that your AI use case and use of data align with human rights obligations. If you have not done so, explain your reasoning.

It is recommended that you complete this question after you have completed the previous sections of the assessment. This will provide more complete information to enable an assessment of the human rights implications of your AI use case.

In Australia, it is unlawful to discriminate on the basis of a number of protected attributes including age, disability, race, sex, intersex status, gender identity and sexual orientation, in certain areas of public life

including education and employment. Australia's federal anti-discrimination laws are contained in the following legislation.

- *Age Discrimination Act 2004*
- *Disability Discrimination Act 1992*
- *Racial Discrimination Act 1975*
- *Sex Discrimination Act 1984.*

Human rights are defined in the *Human Rights (Parliamentary Scrutiny) Act 2011* as the rights and freedoms contained in the 7 core international human rights treaties to which Australia is a party, namely the:

- International Covenant on Civil and Political Rights (ICCPR)
- International Covenant on Economic, Social and Cultural Rights (ICESCR)
- International Convention on the Elimination of All Forms of Racial Discrimination (CERD)
- Convention on the Elimination of All Forms of Discrimination against Women (CEDAW)
- Convention against Torture and Other Cruel, Inhuman or Degrading Treatment or Punishment (CAT)
- Convention on the Rights of the Child (CRC)
- Convention on the Rights of Persons with Disabilities (CRPD).

In addition to other rights referred to in this guidance, human rights you may consider as part of your assessment of the AI use case include:

- a right to privacy – for example, where AI is being used for tracking and surveillance)
- freedom of expression and information – for example, where AI is used to moderate a forum and therefore possibly suppress legitimate forms of expression
- human agency – for example, where AI makes an automated decision on an individual's behalf.

11. Accountability

11.1 Ensuring accountability during the life cycle of the AI system

Agencies should consider putting mechanisms in place during the life cycle of the AI system to ensure that the agency itself, or the relevant decision-maker, remains responsible and accountable for a government decision which involves the use of AI. Such mechanisms should clearly define how ultimate responsibility for the decision is retained, even when AI is used to analyse data or generate recommended outcomes.

Accountability should be considered at all stages of the AI system lifecycle. Some of the relevant considerations for different stages are outlined below, including:

During the design and development phase

- How the AI will be constructed in a way that is consistent with the scope of a decision-maker's discretion and any legislative framework which confers authority on a decision-maker
- How the AI will be designed in a way that ensures that the decision-maker takes into account any matters which it is required to consider as part of decision-making
- Whether the decision-maker will have the ability to override or disregard decisions made by AI—for example, where its outputs are based on biased data.

During the deployment and operation phase

- What information the decision-maker needs to have oversight of the AI (for example, information on the capacities and limitations of the AI, and the process that the AI will use to reach a conclusion)
- What processes are in place to ensure that, where appropriate, final discretion or judgement lies with the decision-maker (for example, the decision-maker analyses the information provided by the AI before making a decision)
- What records will be kept of the decision-maker's reasoning at any decision points which require discretion and judgment. This contributes to the contestability of decisions in accordance with section 9.

Where the scope of the use case changes or is developed during the life of the AI system

- What assessment, review and acceptance-testing processes are to be applied to the changes to the AI system to ensure the above.

12 Use case review and next steps

12.1 Alignment with relevant legal frameworks

This question looks to confirm that you have identified and documented any agency specific legislation, regulations, or binding policy instruments that are relevant to your AI use case.

When completing this section:

- review your agency's legislative and regulatory frameworks. Identify any provisions that may be affected by, or place restrictions on, the design, operation, or outputs of the AI system
- if there is any uncertainty, engage your agency's legal area early, and maintain legal professional privilege where appropriate.

12.2 Legal advice

This section asks whether your agency has sought or obtained legal advice in relation to the AI use case. If you answer 'yes', you should summarise the nature of the legal issue without including the content of the advice. This information should not be disclosed to anyone other than those who need to know or access the information within the agency.

Note that including the actual content of legal advice in this tool may result in waiver of legal professional privilege, meaning the advice could be legally required to be disclosed to others. To avoid unintended waiver, only summarise the subject matter of the advice—for example, 'privacy compliance' or 'intellectual property risks' rather than reproducing or paraphrasing the advice itself.

12.3 Risk summary table

To complete the risk summary table:

- list any risks assessed as **medium** or **high** at the inherent risk assessment stage in section 3
- summarise any mitigations or controls that have been or will be applied
- explain how these mitigations have influenced the residual risk rating

12.4 Overall residual risk rating

To complete this section, choose an overall residual risk rating for the AI use case. Refer to your response to section 12.3.

12.5 Internal governance body review

If your use case's inherent risk is rated as **high** at section 3, you are required under the AI policy to apply specific actions, including creating or reusing a governance body for the purpose of governing high-risk AI. You may document the outcome of the governance body review here, including any recommendations and agreed next steps.

Appendix: Risk consequence guidance table

This table is designed to help you select the appropriate consequence level for the risk questions in sections 3.1 to 3.8. Examples are illustrative, not exhaustive.

Risk	Insignificant	Minor	Moderate	Major	Severe
3.1 Negatively affecting public accessibility or inclusivity of government services	Insignificant compromises to accessibility or inclusivity of services. Minor technical issues causing brief inconvenience but no actual barriers to access or inclusion. Issues rapidly resolved with minimal impact on user experience.	Limited, reversible compromises to accessibility or inclusivity of services. Some people experience difficulties accessing services due to technical issues or design oversights. Barriers are short-term and addressed once identified, with additional support provided to people affected.	Many compromises are made to the accessibility or inclusivity of services. Considerable access challenges for a modest number of users. Resolving access issues requires substantial effort and resources. Certain groups may be disproportionately impacted. Affected users experience frustration and delays in receiving services.	Extensive compromises are made to the accessibility or inclusivity of services, may include some essential services. Ongoing delays that require external technical assistance to resolve. Widespread inconvenience, frustration, public distress and potential legal implications. Vulnerable user groups disproportionately impacted.	Widespread irreversible ongoing compromises are made to the accessibility or inclusivity of services, including some essential services. Majority of users, especially vulnerable groups affected. Essential services inaccessible for extended periods, causing significant public distress, legal implications, and a loss of trust in government efficiency. Comprehensive and immediate actions are urgently needed to rectify the situation.
3.2 Unfair discrimination against individuals, communities or groups	Negligible instances of discrimination, with virtually no discernible effect on individuals, communities, or groups. Issues are proactively identified and rapidly addressed before causing harm.	Limited instances of unfair discrimination occur, affecting a small number of individuals. Relatively isolated cases, and corrective measures minimise their impact.	Moderate levels of discrimination leading to noticeable harm to certain individuals, communities, or groups. These incidents raise bias and fairness concerns and require targeted interventions.	Significant discrimination results in major, tangible harm to individuals and multiple communities or groups. Rebuilding trust requires substantial reforms and remediation efforts.	Pervasive and systemic discrimination causes severe harm across a broad spectrum of the population, particularly marginalised and vulnerable groups. Public outrage, potential legal action, and a profound loss of trust in government. Immediate, sweeping reforms and accountability measures are required.

Risk	Insignificant	Minor	Moderate	Major	Severe
3.3 Perpetuating stereotyping or demeaning representations of individuals, communities or groups	Inadvertently reinforce mild stereotypes, but these instances are quickly identified and rectified with no lasting harm or public concern.	Isolated cases of stereotyping, affecting limited members of community with some noticing and raising concerns. Prompt action mitigates the issue, preventing broader impact.	Moderate stereotyping by AI systems leads to noticeable public discomfort and criticism. Disproportionally affecting certain communities or groups. Requires targeted corrective measures to address and prevent recurrence.	Significant and widespread reinforcement of harmful stereotypes and demeaning representations. Causes public outcry and damages the relationship between communities and government entities. Urgent, comprehensive strategies are needed to rectify these representations and restore trust.	Pervasive and damaging stereotyping severely harms multiple communities, leading to widespread distress. Potential legal consequences, and a profound breach of trust in government use of technology. Requires immediate, sweeping actions to address the harm, including system overhauls and public apologies.
3.4 Harm to individuals, communities, groups, organisations or the environment	Inconsequential glitches with no real harm to the public, business operations or ecosystems. Easily managed through routine measures.	Isolated incidents mildly affecting the public. Slight inconveniences or disruptions to businesses, leading to manageable financial costs. Limited manageable environmental disturbances affecting local ecosystems or resource consumption.	Noticeable negative effects on the public. Businesses face operational challenges or financial losses, affecting their competitiveness. Obvious environmental degradation, including pollution or habitat disruption, prompting public concern.	Significant public harm causing distress and potentially lasting damage. Significant harm to a wide range of businesses, resulting in substantial financial losses, layoffs, and long-term reputational damage. Compromises ecosystem wellbeing causing substantial pollution, loss of biodiversity, and resource depletion.	Widespread, profound harm and severe distress affecting broad segments of the public. Profound damage across the business sector, leading to bankruptcies, major job losses, and a lasting negative impact on the economy. Comprehensive environmental destruction, leading to critical loss of biodiversity, irreversible ecosystem damage, and severe resource scarcity.

Risk	Insignificant	Minor	Moderate	Major	Severe
3.5 Raising privacy concerns	Insignificant data handling errors occur without compromising sensitive information. Incidents are quickly rectified, maintaining public trust in data security.	Isolated exposure of limited sensitive data affects a small group of individuals. Swift actions taken to secure the data and prevent further incidents.	Breach of moderate amounts of sensitive data, leading to privacy concerns among the affected populace. Some individuals experience inconvenience and distress.	Serious misuse of sensitive private data affects a large segment of the population, leading to widespread privacy violations and a loss of public trust. Comprehensive measures are urgently required to secure data and address the privacy breaches.	Significant potential to expose sensitive information of a vast number of individuals, causing severe harm, identity-theft risks; use of sensitive personal information in a way that is likely to draw public criticism with limited ability for individuals to choose how their information is used. Significant potential to harm trust in government information handling with potential for lasting consequences.
3.6 Raising security concerns due to the sensitivity or security classification of the data being used by an AI system	Inconsequential security lapses occur without actual misuse of sensitive data. Quickly identified and corrected with no real harm done. These types of incidents may serve as prompts for reviewing security protocols.	A limited security breach involves unauthorised access to protected data affecting a small number of records with minimal impact. Immediate actions secure the breach, and affected individuals are notified and supported. Incident is catalyst for review of security protocols.	Security incident leads to the compromise of a moderate volume of sensitive data, raising concerns over data protection and privacy. The breach necessitates a thorough investigation, enhanced security measures.	A significant security breach results in extensive unauthorised access to sensitive or protected data, causing considerable concern and distress among the public. Urgent security upgrades and support measures for impacted individuals are implemented, to restore security and trust.	A massive security breach exposes a vast amount of sensitive and protected data, leading to severe implications for national security, public safety, and individual privacy. This incident triggers an emergency response, including legal actions, a major overhaul of security systems, and long-term support for those affected.

Risk	Insignificant	Minor	Moderate	Major	Severe
3.7 Raising security concerns due to implementation, sourcing or characteristics of the AI system	Inconsequential security concerns arise due to characteristics of the AI system, such as software bugs, which are promptly identified and fixed with no adverse effects on overall security. These issues may serve as lessons, leading to slight improvements in the system's security framework.	Certain characteristics of the AI system lead to vulnerabilities that are exploited in a limited manner, causing minor security breaches. Immediate remediation measures are taken, and the system is updated to prevent similar issues.	A moderate security risk is realised when intrinsic features of the AI system allow for unintended access or data leaks. Incident affects a noticeable but contained component of the AI system. Prompts a comprehensive security review of the AI system and the implementation of more robust safeguards.	Significant security flaws in the AI system's design result in major breaches, compromising a large amount of data and severely affecting system integrity. Incident leads to an urgent overhaul of security measures and protocols, alongside efforts to mitigate the damage.	Critical security vulnerabilities inherent to the AI system lead to widespread breaches, exposing vast quantities of sensitive data and jeopardising national security or public safety. The incident results in severe consequences, necessitating emergency responses, extensive system redesigns, and long-term efforts to recover from the breach and prevent recurrence.
3.8 Posing a reputational risk or undermining public confidence in the government	Isolated reputational issues arise, quickly addressed and explained. Causes negligible damage to public trust in government capabilities.	Small-scale AI mishaps lead to brief public concern, slightly denting the government's reputation. Prompt clarification and corrective measures minimize long-term impact on public confidence Seen by the government as poor management.	Misapplications result in moderate public dissatisfaction and questioning of government oversight. Requires remedial actions to mend trust and address concerns. Seen by government and opposition as failed management.	Widespread public scepticism and criticism, majorly affecting the government's image. Requires substantial efforts to rebuild public confidence through transparency, accountability, and improvement of AI governance. High profile negative stories, seen by government and opposition as significant failed management.	Severe misuse or failure of AI systems leads to profound public distrust and criticism. Significantly undermining confidence in government effectiveness and integrity. Requires comprehensive, long-term strategies for rehabilitation of public trust, including systemic changes and ongoing engagement. Seen by government and opposition as catastrophic failure of management. Minister expresses loss of confidence or trust in agency.