

```
In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/dock
# For example, here's several helpful packages to load
```

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list a

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/) that gets
# You can also write temporary files to /kaggle/temp/, but they won't be saved ou
```

```
In [3]: df=pd.read_csv("/kaggle/input/bank-customer-churn-prediction/Churn_Modelling.csv")
```

```
In [6]: df.shape
```

```
Out[6]: (10000, 14)
```

```
In [7]: df.sample(5)
```

```
Out[7]:
```

| | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender | Age | Tenure | Balance |
|------|-----------|------------|------------|-------------|-----------|--------|-----|--------|-----------|
| 5613 | 5614 | 15689412 | Christie | 604 | France | Female | 32 | 7 | 127849.38 |
| 5278 | 5279 | 15799300 | Kao | 510 | Germany | Male | 31 | 0 | 113688.63 |
| 5864 | 5865 | 15803840 | Forbes | 729 | France | Female | 32 | 9 | 0.00 |
| 4532 | 4533 | 15739194 | Manfrin | 548 | Spain | Male | 38 | 0 | 178056.54 |
| 6615 | 6616 | 15792934 | Carruthers | 661 | France | Male | 26 | 8 | 0.00 |

```
In [8]: df.drop(columns=["RowNumber", "CustomerId", "Surname"], axis=1, inplace=True)
df.sample(5)
```

```
Out[8]:
```

| | CreditScore | Geography | Gender | Age | Tenure | Balance | NumOfProducts | HasCrCard | IsActive |
|------|-------------|-----------|--------|-----|--------|-----------|---------------|-----------|----------|
| 3750 | 629 | France | Male | 39 | 2 | 129669.32 | 2 | 1 | |
| 9879 | 486 | Germany | Male | 62 | 9 | 118356.89 | 2 | 1 | |
| 1891 | 584 | France | Female | 37 | 1 | 0.00 | 2 | 1 | |
| 9732 | 724 | Spain | Male | 39 | 3 | 0.00 | 2 | 0 | |
| 1174 | 705 | Spain | Female | 40 | 5 | 203715.15 | 1 | 1 | |

```
In [9]: df.duplicated().sum()
```

```
Out[9]: 0
```

```
In [10]: df.isna().sum()
```

```
Out[10]: CreditScore      0
Geography      0
Gender         0
Age            0
Tenure         0
Balance        0
NumOfProducts  0
HasCrCard      0
IsActiveMember 0
EstimatedSalary 0
Exited         0
dtype: int64
```

```
In [11]: df.describe()
```

```
Out[11]:
```

| | CreditScore | Age | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember |
|-------|--------------|--------------|--------------|---------------|---------------|--------------|----------------|
| count | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 |
| mean | 650.528800 | 38.921800 | 5.012800 | 76485.889288 | 1.530200 | 0.70550 | 0.70550 |
| std | 96.653299 | 10.487806 | 2.892174 | 62397.405202 | 0.581654 | 0.45584 | 0.45584 |
| min | 350.000000 | 18.000000 | 0.000000 | 0.000000 | 1.000000 | 0.00000 | 0.00000 |
| 25% | 584.000000 | 32.000000 | 3.000000 | 0.000000 | 1.000000 | 0.00000 | 0.00000 |
| 50% | 652.000000 | 37.000000 | 5.000000 | 97198.540000 | 1.000000 | 1.00000 | 1.00000 |
| 75% | 718.000000 | 44.000000 | 7.000000 | 127644.240000 | 2.000000 | 1.00000 | 1.00000 |
| max | 850.000000 | 92.000000 | 10.000000 | 250898.090000 | 4.000000 | 1.00000 | 1.00000 |

```
In [12]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   CreditScore            10000 non-null  int64  
1   Geography              10000 non-null  object  
2   Gender                 10000 non-null  object  
3   Age                   10000 non-null  int64  
4   Tenure                 10000 non-null  int64  
5   Balance                10000 non-null  float64 
6   NumOfProducts          10000 non-null  int64  
7   HasCrCard              10000 non-null  int64  
8   IsActiveMember         10000 non-null  int64  
9   EstimatedSalary        10000 non-null  float64 
10  Exited                 10000 non-null  int64  
dtypes: float64(2), int64(7), object(2)
memory usage: 859.5+ KB
```

```
In [13]: df["Exited"].value_counts()
```

```
Out[13]: Exited
0      7963
1      2037
Name: count, dtype: int64
```

```
In [14]: df['Geography'].value_counts()
```

```
Out[14]: Geography
France      5014
Germany     2509
Spain       2477
Name: count, dtype: int64
```

```
In [15]: df.Gender.value_counts()
```

```
Out[15]: Gender
Male        5457
Female      4543
Name: count, dtype: int64
```


```
In [16]: df.columns
```

```
Out[16]: Index(['CreditScore', 'Geography', 'Gender', 'Age', 'Tenure', 'Balance',
               'NumOfProducts', 'HasCrCard', 'IsActiveMember', 'EstimatedSalary',
               'Exited'],
              dtype='object')
```

```
In [17]: df=pd.get_dummies(df, ["Geography","Gender"], drop_first=True, dtype='int')
df.sample(5)
```

```
Out[17]:
```

| | CreditScore | Age | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary |
|------|-------------|-----|--------|-----------|---------------|-----------|----------------|-----------------|
| 3473 | 682 | 42 | 0 | 0.00 | 1 | 0 | 1 | 919 |
| 1377 | 768 | 44 | 6 | 60603.40 | 1 | 1 | 1 | 1780 |
| 7989 | 645 | 39 | 8 | 0.00 | 2 | 0 | 0 | 968 |
| 1488 | 596 | 30 | 6 | 121345.88 | 4 | 1 | 0 | 419 |
| 1021 | 485 | 32 | 6 | 102238.01 | 2 | 1 | 1 | 1940 |



```
In [18]: from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler, StandardScaler
```

```
In [19]: # def mask_outliers_iqr(df, columns): # Not required
#         df_out = df.copy()
#         for col in columns:
#             if not np.issubdtype(df_out[col].dtype, np.number):
#                 continue # Skip non-numeric columns

#         Q1 = df_out[col].quantile(0.25)
#         Q3 = df_out[col].quantile(0.75)
#         IQR = Q3 - Q1
#         lower_bound = Q1 - 1.5 * IQR
#         upper_bound = Q3 + 1.5 * IQR

#         # Mask outliers with NaN
#         df_out[col] = df_out[col].mask((df_out[col] < lower_bound) | (df_out[col] > upper_bound))


#         return df_out
```

```
In [80]: X,y= df.drop(["Exited"], axis=1), df[["Exited"]]
```

```
In [81]: X.head()
```

```
Out[81]:
```

| | CreditScore | Age | Tenure | Balance | NumOfProducts | HasCrCard | IsActiveMember | EstimatedSalary |
|---|-------------|-----|--------|-----------|---------------|-----------|----------------|-----------------|
| 0 | 619 | 42 | 2 | 0.00 | 1 | 1 | 1 | 101348.8 |
| 1 | 608 | 41 | 1 | 83807.86 | 1 | 0 | 1 | 112542.5 |
| 2 | 502 | 42 | 8 | 159660.80 | 3 | 1 | 0 | 113931.5 |
| 3 | 699 | 39 | 1 | 0.00 | 2 | 0 | 0 | 93826.6 |
| 4 | 850 | 43 | 2 | 125510.82 | 1 | 1 | 1 | 79084.7 |



```
In [82]: col=['CreditScore',  
            'Age',  
            'NumOfProducts',  
            ]
```

```
In [83]: # X=mask_outliers_iqr(X, col)
```

```
In [84]: import seaborn as sns, matplotlib.pyplot as plt
```

```
In [85]: x_train, X_test, y_train, Y_test = train_test_split(X, y , stratify=y, test_size=
```

```
In [86]: scaller=StandardScaler()
```

```
In [87]: x_train_scaled= scaller.fit_transform(x_train)
```

```
In [88]: X_test_scaled=scaller.transform(X_test)
```

```
In [90]: import tensorflow
```

```
from tensorflow import keras
```

```
In [91]: from tensorflow.keras import Sequential
```

```
In [92]: from tensorflow.keras.layers import Dense, Dropout  
from tensorflow.keras.regularizers import l2
```

```
In [93]: df.shape[1]
```

```
Out[93]: 12
```

```
In [94]: model = Sequential()
model.add(Dense(units=32,activation="relu", kernel_regularizer=l2(0.01), input_dim=100))
model.add(Dense(units=12,activation="relu"))
model.add(Dense(units=12,activation="relu", kernel_regularizer=l2(0.01)))
model.add(Dense(units=12,activation="relu", kernel_regularizer=l2(0.01)))
model.add(Dropout(0.2))
model.add(Dense(units=6,activation="relu",kernel_regularizer=l2(0.01)))
model.add(Dense(units=3,activation="relu", kernel_regularizer=l2(0.01)))
model.add(Dropout(0.2))
model.add(Dense(units=1,activation = 'sigmoid'))
```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [95]: model.summary()
```

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|---------------------|--------------|---------|
| dense_7 (Dense) | (None, 32) | 3210 |
| dense_8 (Dense) | (None, 12) | 1212 |
| dense_9 (Dense) | (None, 12) | 1212 |
| dense_10 (Dense) | (None, 12) | 1212 |
| dropout_2 (Dropout) | (None, 12) | 0 |
| dense_11 (Dense) | (None, 6) | 612 |
| dense_12 (Dense) | (None, 3) | 312 |
| dropout_3 (Dropout) | (None, 3) | 0 |
| dense_13 (Dense) | (None, 1) | 13 |



Total params: 1,195 (4.67 KB)

Trainable params: 1,195 (4.67 KB)

Non-trainable params: 0 (0.00 B)

```
In [96]: print(np.unique(y_train))
y_train.info()
```

```
[0 1]
<class 'pandas.core.frame.DataFrame'>
Index: 8000 entries, 2515 to 2304
Data columns (total 1 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Exited  8000 non-null         int64
dtypes: int64(1)
memory usage: 125.0 KB
```

```
In [97]: from tensorflow.keras.metrics import Recall, Precision, Accuracy
```

```
model.compile(
    loss='binary_crossentropy',
    optimizer='adam',
    metrics=[Recall(), Precision(), Accuracy()]
)
```

```
In [99]: from tensorflow.keras.callbacks import EarlyStopping
```

```
early_stop = EarlyStopping(
    monitor='val_recall',
    mode='max',
    patience=10,
    restore_best_weights=True,
    min_delta=0.0001
)
```

```
In [100]: history=model.fit(x_train_scaled, y_train, epochs=100, validation_split=.2, \
                           callbacks=[early_stop],class_weight={0: 1, 1: 10})
```

```
Epoch 1/100
200/200 ————— 6s 9ms/step - accuracy: 0.0000e+00 - loss: 2.1716
- precision_1: 0.2024 - recall_1: 0.9935 - val_accuracy: 0.0000e+00 - val_loss: 1.2140 - val_precision_1: 0.2050 - val_recall_1: 1.0000
```

```
Epoch 2/100
57/200 ————— 0s 3ms/step - accuracy: 0.0000e+00 - loss: 1.8722
- precision_1: 0.2013 - recall_1: 1.0000
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/callbacks/early_stopping.py:
153: UserWarning: Early stopping conditioned on metric `val_recall` which is not available. Available metrics are: accuracy,loss,precision_1,recall_1,val_accuracy,val_loss,val_precision_1,val_recall_1
current = self.get_monitor_value(logs)
```

```
200/200 ————— 1s 4ms/step - accuracy: 0.0000e+00 - loss: 1.8375
- precision_1: 0.2051 - recall_1: 1.0000 - val_accuracy: 0.0000e+00 - val_loss: 1.1105 - val_precision_1: 0.2050 - val_recall_1: 1.0000
```

```
Epoch 3/100
200/200 ————— 1s 4ms/step - accuracy: 0.0000e+00 - loss: 1.6502
- precision_1: 0.2011 - recall_1: 1.0000 - val_accuracy: 0.0000e+00 - val_loss:
```

```
In [101]: w,b=model.layers[0].get_weights()
b
```

```
Out[101]: array([-0.18660632,  0.04815373, -0.24157415,  0.00117472,  0.02556655,
                0.17147207, -0.09948713, -0.3484123 , -0.32381824,  0.08960334,
               -0.08077571, -0.03213134,  0.05704772, -0.01498128,  0.0861932 ,
               -0.5179489 , -0.1961962 ,  0.36160317, -0.0240429 , -0.05631261,
                0.09280947,  0.09838948,  0.00956321,  0.01863383, -0.0962879 ,
                0.00358349,  0.02820406,  0.28084746, -0.0657841 , -0.00603615,
                0.01156903,  0.07417452], dtype=float32)
```

```
In [102]: test_pred_proba=model.predict(X_test_scaled)
test_pred=np.where(test_pred_proba>=0.5, 1,0)
```

63/63 ————— 0s 5ms/step

```
In [103]: from sklearn.metrics import accuracy_score
```

```
In [104]: from sklearn.metrics import classification_report, confusion_matrix, accuracy_sco
```

```
In [105]: Y_test.value_counts()
```

```
Out[105]: Exited
0          1593
1           407
Name: count, dtype: int64
```

```
In [106]: accuracy_score( Y_test, test_pred)
```

```
Out[106]: 0.7145
```

```
In [107]: train_pred_proba=model.predict(x_train_scaled)
train_pred=np.where(train_pred_proba>=0.5, 1,0)
```

250/250 ————— 0s 1ms/step

```
In [108]: print(classification_report(y_train, train_pred))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.96 | 0.69 | 0.80 | 6370 |
| 1 | 0.42 | 0.88 | 0.57 | 1630 |
| accuracy | | | 0.73 | 8000 |
| macro avg | 0.69 | 0.79 | 0.69 | 8000 |
| weighted avg | 0.85 | 0.73 | 0.76 | 8000 |

```
In [109]: print(classification_report(Y_test, test_pred))
```


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.94 | 0.68 | 0.79 | 1593 |
| 1 | 0.40 | 0.83 | 0.54 | 407 |
| accuracy | | | 0.71 | 2000 |
| macro avg | 0.67 | 0.76 | 0.67 | 2000 |
| weighted avg | 0.83 | 0.71 | 0.74 | 2000 |

```
In [110]: type(history.history)
```

```
Out[110]: dict
```

```
In [111]: import matplotlib.pyplot as plt
```

```
In [112]: plt.plot(history.history["loss"], label="Train Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.title("Training vs Validation Loss")
plt.legend()
plt.grid(True)
plt.show()
```




```
In [113]: list(history.history.keys())
```


```
Out[113]: ['accuracy',
'loss',
'precision_1',
'recall_1',
'val_accuracy',
'val_loss',
'val_precision_1',
'val_recall_1']
```



```
In [114]: plt.plot(history.history["recall_1"], label="Train Recall")
plt.plot(history.history["val_recall_1"], label="Validation Recall")
plt.xlabel("Epochs")
plt.ylabel("Recall")
plt.title("Training vs Validation Loss")
plt.legend()
plt.grid(True)
plt.show()
```



```
In [116]: plt.plot(history.history["accuracy"], label="Train Recall")
plt.plot(history.history["val_accuracy"], label="Validation Recall")
plt.xlabel("Epochs")
plt.ylabel("Recall")
plt.title("Training vs Validation Loss")
plt.legend()
plt.grid(True)
plt.show()
```



Make the App

```
In [4]: import ipywidgets as widgets
from IPython.display import display
```

```
In [5]: # Create widgets for each feature
credit_score = widgets.IntSlider(description='CreditScore', min=300, max=900, value=400)
age = widgets.IntSlider(description='Age', min=18, max=100, value=40)
tenure = widgets.IntSlider(description='Tenure', min=0, max=20, value=5)
balance = widgets.FloatText(description='Balance', value=50000.0)
products = widgets.Dropdown(description='Products', options=[1,2,3,4], value=1)
has_card = widgets.ToggleButtons(description='HasCrCard', options=[0,1], value=1)
is_active = widgets.ToggleButtons(description='IsActive', options=[0,1], value=1)
salary = widgets.FloatText(description='Salary', value=100000.0)

gender = widgets.Dropdown(description='Gender', options=['Male', 'Female'], value='Male')
geography = widgets.Dropdown(description='Geography', options=['Germany', 'Spain', 'France'], value='Germany')

predict_btn = widgets.Button(description='Predict Churn', button_style='success')
output = widgets.Output()
```

```
In [6]: def make_prediction(b):
    data = {
        'CreditScore': credit_score.value,
        'Age': age.value,
        'Tenure': tenure.value,
        'Balance': balance.value,
        'NumOfProducts': products.value,
        'HasCrCard': has_card.value,
        'IsActiveMember': is_active.value,
        'EstimatedSalary': salary.value,
        'Geography_Germany': 1 if geography.value == 'Germany' else 0,
        'Geography_Spain': 1 if geography.value == 'Spain' else 0,
        'Gender_Male': 1 if gender.value == 'Male' else 0
    }

    input_df = pd.DataFrame([data])
    input_scaled = scaler.transform(input_df) # Use the same scaler
    proba = model.predict(input_scaled)[0][0]
    prediction = 1 if proba >= 0.5 else 0

    with output:
        output.clear_output()
        print(f"🔍 Predicted Probability: {proba:.2f}")
        print(f"🚨 Prediction:", "Will Churn (Exited = 1)" if prediction == 1 else "Will not Churn (Exited = 0)")
```

```
In [7]: predict_btn.on_click(make_prediction)
```

```
In [9]: form_items = widgets.VBox([
        credit_score, age, tenure, balance, products,
        has_card, is_active, salary, geography, gender,
        predict_btn, output
    ])
display(form_items)
```

| | | |
|---|-------------------------------------|-----|
| CreditScore | <input type="radio"/> | 600 |
| Age | <input type="radio"/> | 40 |
| Tenure | <input type="radio"/> | 5 |
| Balance | <input type="text" value="50000"/> | |
| Products | <input type="text" value="1"/> | |
| HasCrCard | | |
| <input type="radio"/> 0 <input type="radio"/> 1 | | |
| IsActive | | |
| <input type="radio"/> 0 <input type="radio"/> 1 | | |
| Salary | <input type="text" value="100000"/> | |
| Geography | <input type="text" value="France"/> | |
| Gender | <input type="text" value="Male"/> | |
| Predict Churn | | |

```
In [ ]:
```