

tuzvfmyfn

January 30, 2025

```
[1]: import keras

from keras.datasets import cifar10
from tensorflow.keras.preprocessing.image import ImageDataGenerator

from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten
from keras.layers import Conv2D, MaxPooling2D

import os

import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report
import itertools

%matplotlib inline

#Loading required header files
```

```
[2]: # Load the CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Print dataset shapes
print("x_train shape:", x_train.shape)
print("y_train shape:", y_train.shape)
print("x_test shape:", x_test.shape)
print("y_test shape:", y_test.shape)

# Print data types
print("x_train dtype:", x_train.dtype)
print("y_train dtype:", y_train.dtype)
```

```
x_train shape: (50000, 32, 32, 3)
y_train shape: (50000, 1)
x_test shape: (10000, 32, 32, 3)
y_test shape: (10000, 1)
```

```
x_train dtype: uint8
y_train dtype: uint8
```


```
[3]: num_classes = 10

data_augmentation = False
```

```
[4]: #checking distribution
fig, axs = plt.subplots(1,2,figsize=(15,5))
sns.countplot(x=y_train.ravel(), ax = axs[0],legend=False)
axs[0].set_title("Distribution of training data")
axs[0].set_xlabel("Classes")

sns.countplot(x=y_test.ravel(), ax = axs[1],legend=False)
axs[1].set_title("Distribution of testing data")
axs[1].set_xlabel("Classes")

plt.tight_layout()
plt.show()
```



```
[5]: x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255 #since picxcel max value is 255 dividing by it will normalize
           ↳ the values
x_test /= 255
#one hot encoding the target variable
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
[6]: model = Sequential()

model.add(Conv2D(filters = 32,kernel_size = (3,3), padding =
           ↳ 'same',activation="relu", input_shape=x_train.shape[1:]))
model.add(Conv2D(filters = 32,kernel_size = (3,3),activation="relu"))
```

```

model.add(MaxPooling2D(pool_size =(2,2)))
model.add(Dropout(.25))

model.add(Conv2D(filters = 64, kernel_size = (3,3), padding =_
↳ 'same',activation="relu"))
model.add(Conv2D(filters = 64, kernel_size = (3,3), activation="relu"))
model.add(MaxPooling2D(pool_size = (2,2)))
model.add(Dropout(0.25))

model.add(Flatten())
model.add(Dense(512,activation="relu"))

model.add(Dropout(0.5))

model.add(Dense(num_classes,activation="softmax" ))

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
conv2d_1 (Conv2D)	(None, 30, 30, 32)	9248
max_pooling2d (MaxPooling2D)	(None, 15, 15, 32)	0
dropout (Dropout)	(None, 15, 15, 32)	0
conv2d_2 (Conv2D)	(None, 15, 15, 64)	18496
conv2d_3 (Conv2D)	(None, 13, 13, 64)	36928
max_pooling2d_1 (MaxPooling2D)	(None, 6, 6, 64)	0
dropout_1 (Dropout)	(None, 6, 6, 64)	0
flatten (Flatten)	(None, 2304)	0
dense (Dense)	(None, 512)	1180160
dropout_2 (Dropout)	(None, 512)	0

dense_1 (Dense) (None, 10) 5130

```
=====
Total params: 1,250,858
Trainable params: 1,250,858
Non-trainable params: 0
-----
```

```
[7]: opt = keras.optimizers.RMSprop(learning_rate = 0.0001, decay = 0.000001)

model.compile(loss = "categorical_crossentropy", optimizer = opt,
              metrics = ['accuracy'])
```

```
[8]: # Set parameters
batch_size = 64 # Adjusted batch size for better CPU performance
epochs = 100

# Check if using data augmentation
if not data_augmentation:
    print("Not using data augmentation.")
    history = model.fit(x_train, y_train, batch_size=batch_size, epochs=epochs,
                        validation_data=(x_test, y_test), workers=4,
                        use_multiprocessing=True) #using multiple core for better CPU performance
else:
    print("Using realtime data augmentation.")

    # Initialize ImageDataGenerator with augmentation settings
    datagen = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1,
    use_multiprocessing=True)

    # Fit the generator to the training data (not necessary if augmentation
    parameters are fixed)
    datagen.fit(x_train)

    # Train the model using the augmented data
    history = model.fit(datagen.flow(x_train, y_train, batch_size=batch_size,
    use_multiprocessing=True),
                        validation_data=(x_test, y_test),
    workers=4, use_multiprocessing=True))
```

Not using data augmentation.

Epoch 1/100

782/782 [=====] - 133s 168ms/step - loss: 1.8923 -
accuracy: 0.3073 - val_loss: 1.6324 - val_accuracy: 0.4078

Epoch 2/100

782/782 [=====] - 142s 182ms/step - loss: 1.6026 -
accuracy: 0.4171 - val_loss: 1.4816 - val_accuracy: 0.4726

Epoch 3/100
782/782 [=====] - 111s 143ms/step - loss: 1.4646 - accuracy: 0.4676 - val_loss: 1.3915 - val_accuracy: 0.4984

Epoch 4/100
782/782 [=====] - 112s 143ms/step - loss: 1.3590 - accuracy: 0.5117 - val_loss: 1.2822 - val_accuracy: 0.5385

Epoch 5/100
782/782 [=====] - 113s 145ms/step - loss: 1.2894 - accuracy: 0.5394 - val_loss: 1.2219 - val_accuracy: 0.5706

Epoch 6/100
782/782 [=====] - 119s 152ms/step - loss: 1.2244 - accuracy: 0.5632 - val_loss: 1.2124 - val_accuracy: 0.5665

Epoch 7/100
782/782 [=====] - 115s 147ms/step - loss: 1.1764 - accuracy: 0.5850 - val_loss: 1.1837 - val_accuracy: 0.5925

Epoch 8/100
782/782 [=====] - 118s 150ms/step - loss: 1.1267 - accuracy: 0.6010 - val_loss: 1.0791 - val_accuracy: 0.6171

Epoch 9/100
782/782 [=====] - 151s 193ms/step - loss: 1.0889 - accuracy: 0.6187 - val_loss: 1.0161 - val_accuracy: 0.6453

Epoch 10/100
782/782 [=====] - 146s 185ms/step - loss: 1.0463 - accuracy: 0.6304 - val_loss: 0.9871 - val_accuracy: 0.6540

Epoch 11/100
782/782 [=====] - 121s 155ms/step - loss: 1.0157 - accuracy: 0.6444 - val_loss: 0.9883 - val_accuracy: 0.6589

Epoch 12/100
782/782 [=====] - 128s 163ms/step - loss: 0.9821 - accuracy: 0.6553 - val_loss: 0.9578 - val_accuracy: 0.6666

Epoch 13/100
782/782 [=====] - 130s 167ms/step - loss: 0.9512 - accuracy: 0.6675 - val_loss: 1.0064 - val_accuracy: 0.6521

Epoch 14/100
782/782 [=====] - 126s 161ms/step - loss: 0.9266 - accuracy: 0.6757 - val_loss: 0.8914 - val_accuracy: 0.6899

Epoch 15/100
782/782 [=====] - 139s 177ms/step - loss: 0.9084 - accuracy: 0.6835 - val_loss: 0.8507 - val_accuracy: 0.6986

Epoch 16/100
782/782 [=====] - 130s 166ms/step - loss: 0.8881 - accuracy: 0.6889 - val_loss: 0.8638 - val_accuracy: 0.7006

Epoch 17/100
782/782 [=====] - 125s 160ms/step - loss: 0.8635 - accuracy: 0.6989 - val_loss: 0.8252 - val_accuracy: 0.7119

Epoch 18/100
782/782 [=====] - 124s 159ms/step - loss: 0.8453 - accuracy: 0.7060 - val_loss: 0.9144 - val_accuracy: 0.6839

Epoch 19/100
782/782 [=====] - 122s 156ms/step - loss: 0.8278 - accuracy: 0.7124 - val_loss: 0.9245 - val_accuracy: 0.6828

Epoch 20/100
782/782 [=====] - 133s 170ms/step - loss: 0.8122 - accuracy: 0.7161 - val_loss: 0.8040 - val_accuracy: 0.7212

Epoch 21/100
782/782 [=====] - 129s 165ms/step - loss: 0.7994 - accuracy: 0.7219 - val_loss: 0.7843 - val_accuracy: 0.7290

Epoch 22/100
782/782 [=====] - 158s 202ms/step - loss: 0.7800 - accuracy: 0.7284 - val_loss: 0.8516 - val_accuracy: 0.7137

Epoch 23/100
782/782 [=====] - 1752s 2s/step - loss: 0.7687 - accuracy: 0.7337 - val_loss: 0.7754 - val_accuracy: 0.7316

Epoch 24/100
782/782 [=====] - 94s 120ms/step - loss: 0.7591 - accuracy: 0.7360 - val_loss: 0.7600 - val_accuracy: 0.7404

Epoch 25/100
782/782 [=====] - 98s 125ms/step - loss: 0.7503 - accuracy: 0.7404 - val_loss: 0.7904 - val_accuracy: 0.7313

Epoch 26/100
782/782 [=====] - 101s 129ms/step - loss: 0.7349 - accuracy: 0.7455 - val_loss: 0.8020 - val_accuracy: 0.7269

Epoch 27/100
782/782 [=====] - 105s 134ms/step - loss: 0.7233 - accuracy: 0.7502 - val_loss: 0.7555 - val_accuracy: 0.7397

Epoch 28/100
782/782 [=====] - 111s 142ms/step - loss: 0.7185 - accuracy: 0.7530 - val_loss: 0.7126 - val_accuracy: 0.7569

Epoch 29/100
782/782 [=====] - 1524s 2s/step - loss: 0.7124 - accuracy: 0.7542 - val_loss: 0.7214 - val_accuracy: 0.7553

Epoch 30/100
782/782 [=====] - 115s 147ms/step - loss: 0.7023 - accuracy: 0.7584 - val_loss: 0.7222 - val_accuracy: 0.7520

Epoch 31/100
782/782 [=====] - 115s 147ms/step - loss: 0.6953 - accuracy: 0.7584 - val_loss: 0.7486 - val_accuracy: 0.7467

Epoch 32/100
782/782 [=====] - 107s 137ms/step - loss: 0.6868 - accuracy: 0.7635 - val_loss: 0.7302 - val_accuracy: 0.7577

Epoch 33/100
782/782 [=====] - 117s 149ms/step - loss: 0.6791 - accuracy: 0.7667 - val_loss: 0.7155 - val_accuracy: 0.7598

Epoch 34/100
782/782 [=====] - 109s 139ms/step - loss: 0.6780 - accuracy: 0.7671 - val_loss: 0.6986 - val_accuracy: 0.7634

Epoch 35/100
782/782 [=====] - 113s 144ms/step - loss: 0.6686 - accuracy: 0.7680 - val_loss: 0.7142 - val_accuracy: 0.7598
Epoch 36/100
782/782 [=====] - 115s 147ms/step - loss: 0.6610 - accuracy: 0.7733 - val_loss: 0.6881 - val_accuracy: 0.7696
Epoch 37/100
782/782 [=====] - 117s 149ms/step - loss: 0.6652 - accuracy: 0.7739 - val_loss: 0.6897 - val_accuracy: 0.7674
Epoch 38/100
782/782 [=====] - 124s 158ms/step - loss: 0.6538 - accuracy: 0.7752 - val_loss: 0.6639 - val_accuracy: 0.7718
Epoch 39/100
782/782 [=====] - 119s 152ms/step - loss: 0.6504 - accuracy: 0.7778 - val_loss: 0.6775 - val_accuracy: 0.7707
Epoch 40/100
782/782 [=====] - 120s 153ms/step - loss: 0.6417 - accuracy: 0.7808 - val_loss: 0.7556 - val_accuracy: 0.7543
Epoch 41/100
782/782 [=====] - 118s 152ms/step - loss: 0.6367 - accuracy: 0.7808 - val_loss: 0.6912 - val_accuracy: 0.7699
Epoch 42/100
782/782 [=====] - 122s 156ms/step - loss: 0.6410 - accuracy: 0.7836 - val_loss: 0.6804 - val_accuracy: 0.7712
Epoch 43/100
782/782 [=====] - 121s 155ms/step - loss: 0.6323 - accuracy: 0.7844 - val_loss: 0.6685 - val_accuracy: 0.7755
Epoch 44/100
782/782 [=====] - 125s 160ms/step - loss: 0.6298 - accuracy: 0.7836 - val_loss: 0.6438 - val_accuracy: 0.7830
Epoch 45/100
782/782 [=====] - 128s 164ms/step - loss: 0.6234 - accuracy: 0.7884 - val_loss: 0.6949 - val_accuracy: 0.7659
Epoch 46/100
782/782 [=====] - 122s 157ms/step - loss: 0.6212 - accuracy: 0.7891 - val_loss: 0.6553 - val_accuracy: 0.7802
Epoch 47/100
782/782 [=====] - 124s 158ms/step - loss: 0.6223 - accuracy: 0.7886 - val_loss: 0.6710 - val_accuracy: 0.7740
Epoch 48/100
782/782 [=====] - 122s 156ms/step - loss: 0.6125 - accuracy: 0.7928 - val_loss: 0.6352 - val_accuracy: 0.7879
Epoch 49/100
782/782 [=====] - 125s 159ms/step - loss: 0.6119 - accuracy: 0.7936 - val_loss: 0.6695 - val_accuracy: 0.7786
Epoch 50/100
782/782 [=====] - 125s 160ms/step - loss: 0.6077 - accuracy: 0.7933 - val_loss: 0.6418 - val_accuracy: 0.7834

Epoch 51/100
782/782 [=====] - 123s 158ms/step - loss: 0.6073 - accuracy: 0.7937 - val_loss: 0.6556 - val_accuracy: 0.7751

Epoch 52/100
782/782 [=====] - 124s 158ms/step - loss: 0.6042 - accuracy: 0.7962 - val_loss: 0.6415 - val_accuracy: 0.7842

Epoch 53/100
782/782 [=====] - 128s 164ms/step - loss: 0.5959 - accuracy: 0.8011 - val_loss: 0.6485 - val_accuracy: 0.7837

Epoch 54/100
782/782 [=====] - 128s 164ms/step - loss: 0.5984 - accuracy: 0.7981 - val_loss: 0.6875 - val_accuracy: 0.7788

Epoch 55/100
782/782 [=====] - 126s 161ms/step - loss: 0.5920 - accuracy: 0.7979 - val_loss: 0.6484 - val_accuracy: 0.7819

Epoch 56/100
782/782 [=====] - 125s 160ms/step - loss: 0.5955 - accuracy: 0.8002 - val_loss: 0.6262 - val_accuracy: 0.7932

Epoch 57/100
782/782 [=====] - 126s 162ms/step - loss: 0.5883 - accuracy: 0.8014 - val_loss: 0.6695 - val_accuracy: 0.7776

Epoch 58/100
782/782 [=====] - 125s 160ms/step - loss: 0.5888 - accuracy: 0.7994 - val_loss: 0.6349 - val_accuracy: 0.7873

Epoch 59/100
782/782 [=====] - 128s 163ms/step - loss: 0.5824 - accuracy: 0.8012 - val_loss: 0.6337 - val_accuracy: 0.7873

Epoch 60/100
782/782 [=====] - 123s 158ms/step - loss: 0.5820 - accuracy: 0.8027 - val_loss: 0.6200 - val_accuracy: 0.7947

Epoch 61/100
782/782 [=====] - 128s 163ms/step - loss: 0.5785 - accuracy: 0.8058 - val_loss: 0.6174 - val_accuracy: 0.7956

Epoch 62/100
782/782 [=====] - 126s 161ms/step - loss: 0.5815 - accuracy: 0.8047 - val_loss: 0.6571 - val_accuracy: 0.7816

Epoch 63/100
782/782 [=====] - 124s 159ms/step - loss: 0.5764 - accuracy: 0.8064 - val_loss: 0.6478 - val_accuracy: 0.7885

Epoch 64/100
782/782 [=====] - 128s 164ms/step - loss: 0.5726 - accuracy: 0.8061 - val_loss: 0.6116 - val_accuracy: 0.7966

Epoch 65/100
782/782 [=====] - 126s 162ms/step - loss: 0.5794 - accuracy: 0.8062 - val_loss: 0.6548 - val_accuracy: 0.7893

Epoch 66/100
782/782 [=====] - 124s 159ms/step - loss: 0.5694 - accuracy: 0.8083 - val_loss: 0.6390 - val_accuracy: 0.7920

Epoch 67/100
782/782 [=====] - 128s 163ms/step - loss: 0.5735 - accuracy: 0.8063 - val_loss: 0.6546 - val_accuracy: 0.7868
Epoch 68/100
782/782 [=====] - 129s 165ms/step - loss: 0.5664 - accuracy: 0.8090 - val_loss: 0.6011 - val_accuracy: 0.8007
Epoch 69/100
782/782 [=====] - 128s 164ms/step - loss: 0.5612 - accuracy: 0.8104 - val_loss: 0.6131 - val_accuracy: 0.7969
Epoch 70/100
782/782 [=====] - 130s 166ms/step - loss: 0.5639 - accuracy: 0.8111 - val_loss: 0.6695 - val_accuracy: 0.7912
Epoch 71/100
782/782 [=====] - 132s 169ms/step - loss: 0.5618 - accuracy: 0.8105 - val_loss: 0.6283 - val_accuracy: 0.7933
Epoch 72/100
782/782 [=====] - 133s 170ms/step - loss: 0.5608 - accuracy: 0.8123 - val_loss: 0.6180 - val_accuracy: 0.7967
Epoch 73/100
782/782 [=====] - 143s 183ms/step - loss: 0.5604 - accuracy: 0.8153 - val_loss: 0.6558 - val_accuracy: 0.7921
Epoch 74/100
782/782 [=====] - 197s 252ms/step - loss: 0.5562 - accuracy: 0.8144 - val_loss: 0.6027 - val_accuracy: 0.7988
Epoch 75/100
782/782 [=====] - 138s 177ms/step - loss: 0.5539 - accuracy: 0.8134 - val_loss: 0.6244 - val_accuracy: 0.7922
Epoch 76/100
782/782 [=====] - 130s 167ms/step - loss: 0.5595 - accuracy: 0.8126 - val_loss: 0.6237 - val_accuracy: 0.7990
Epoch 77/100
782/782 [=====] - 138s 176ms/step - loss: 0.5556 - accuracy: 0.8139 - val_loss: 0.6069 - val_accuracy: 0.8019
Epoch 78/100
782/782 [=====] - 169s 216ms/step - loss: 0.5578 - accuracy: 0.8133 - val_loss: 0.6235 - val_accuracy: 0.8005
Epoch 79/100
782/782 [=====] - 162s 207ms/step - loss: 0.5539 - accuracy: 0.8154 - val_loss: 0.6022 - val_accuracy: 0.8039
Epoch 80/100
782/782 [=====] - 133s 170ms/step - loss: 0.5442 - accuracy: 0.8169 - val_loss: 0.6299 - val_accuracy: 0.7953
Epoch 81/100
782/782 [=====] - 131s 168ms/step - loss: 0.5490 - accuracy: 0.8173 - val_loss: 0.6212 - val_accuracy: 0.7989
Epoch 82/100
782/782 [=====] - 128s 163ms/step - loss: 0.5448 - accuracy: 0.8181 - val_loss: 0.6493 - val_accuracy: 0.7912

Epoch 83/100
782/782 [=====] - 125s 160ms/step - loss: 0.5471 - accuracy: 0.8167 - val_loss: 0.6570 - val_accuracy: 0.7883

Epoch 84/100
782/782 [=====] - 128s 164ms/step - loss: 0.5419 - accuracy: 0.8194 - val_loss: 0.6107 - val_accuracy: 0.7997

Epoch 85/100
782/782 [=====] - 130s 166ms/step - loss: 0.5430 - accuracy: 0.8207 - val_loss: 0.6182 - val_accuracy: 0.7947

Epoch 86/100
782/782 [=====] - 127s 163ms/step - loss: 0.5414 - accuracy: 0.8193 - val_loss: 0.6128 - val_accuracy: 0.7987

Epoch 87/100
782/782 [=====] - 138s 177ms/step - loss: 0.5458 - accuracy: 0.8175 - val_loss: 0.6226 - val_accuracy: 0.7950

Epoch 88/100
782/782 [=====] - 156s 200ms/step - loss: 0.5362 - accuracy: 0.8204 - val_loss: 0.6178 - val_accuracy: 0.8049

Epoch 89/100
782/782 [=====] - 174s 223ms/step - loss: 0.5351 - accuracy: 0.8218 - val_loss: 0.6945 - val_accuracy: 0.7826

Epoch 90/100
782/782 [=====] - 4318s 6s/step - loss: 0.5391 - accuracy: 0.8205 - val_loss: 0.6419 - val_accuracy: 0.7913

Epoch 91/100
782/782 [=====] - 98s 126ms/step - loss: 0.5372 - accuracy: 0.8219 - val_loss: 0.6340 - val_accuracy: 0.7936

Epoch 92/100
782/782 [=====] - 111s 143ms/step - loss: 0.5322 - accuracy: 0.8201 - val_loss: 0.6407 - val_accuracy: 0.7983

Epoch 93/100
782/782 [=====] - 114s 146ms/step - loss: 0.5335 - accuracy: 0.8208 - val_loss: 0.6139 - val_accuracy: 0.8006

Epoch 94/100
782/782 [=====] - 111s 143ms/step - loss: 0.5344 - accuracy: 0.8221 - val_loss: 0.6296 - val_accuracy: 0.7975

Epoch 95/100
782/782 [=====] - 4057s 5s/step - loss: 0.5358 - accuracy: 0.8233 - val_loss: 0.6082 - val_accuracy: 0.8043

Epoch 96/100
782/782 [=====] - 97s 124ms/step - loss: 0.5347 - accuracy: 0.8227 - val_loss: 0.6165 - val_accuracy: 0.8021

Epoch 97/100
782/782 [=====] - 97s 123ms/step - loss: 0.5324 - accuracy: 0.8236 - val_loss: 0.6302 - val_accuracy: 0.8033

Epoch 98/100
782/782 [=====] - 98s 125ms/step - loss: 0.5314 - accuracy: 0.8219 - val_loss: 0.6272 - val_accuracy: 0.8012


```
Epoch 99/100
782/782 [=====] - 107s 136ms/step - loss: 0.5326 -
accuracy: 0.8242 - val_loss: 0.5953 - val_accuracy: 0.8077
Epoch 100/100
782/782 [=====] - 109s 139ms/step - loss: 0.5300 -
accuracy: 0.8224 - val_loss: 0.6039 - val_accuracy: 0.8034
```

```
[9]: def plotmodelhistory(history):
    fig,axs = plt.subplots(1,2, figsize = (15,5))
    axs[0].plot(history.history['accuracy'])
    axs[0].plot(history.history['val_accuracy'])
    axs[0].set_title('Model Accuracy')
    axs[0].set_ylabel('Accuracy')
    axs[0].set_xlabel('Epoch')
    axs[0].legend(['train','validate'], loc = 'upper left')

    axs[1].plot(history.history['loss'])
    axs[1].plot(history.history['val_loss'])
    axs[1].set_title('Model loss')
    axs[1].set_ylabel('Loss')
    axs[1].set_xlabel('Epoch')
    axs[1].legend(['Train','Validate'], loc = 'upper left')

    plt.tight_layout()
    plt.show()
```

```
[10]: plotmodelhistory(history)
```



```
[11]: scores = model.evaluate(x_test, y_test, verbose =1)
print('Test loss:', scores[0])
print('Test accuracy:', scores[1])

pred = model.predict(x_test)
```

```

313/313 [=====] - 7s 22ms/step - loss: 0.6039 -
accuracy: 0.8034
Test loss: 0.6038806438446045
Test accuracy: 0.8033999800682068
313/313 [=====] - 9s 26ms/step

```

```

[13]: Y_pred_classes = np.argmax(pred, axis = 1)
      Y_true = np.argmax(y_test, axis = 1)

      print(classification_report(Y_true,Y_pred_classes))

```

	precision	recall	f1-score	support
0	0.82	0.85	0.84	1000
1	0.90	0.89	0.90	1000
2	0.79	0.63	0.70	1000
3	0.66	0.65	0.66	1000
4	0.70	0.85	0.77	1000
5	0.83	0.61	0.71	1000
6	0.76	0.92	0.83	1000
7	0.84	0.86	0.85	1000
8	0.86	0.91	0.89	1000
9	0.90	0.85	0.87	1000
accuracy				0.80 10000
macro avg	0.81	0.80	0.80	10000
weighted avg	0.81	0.80	0.80	10000

```


[21]: from sklearn.metrics import confusion_matrix, classification_report
      import seaborn as sns
      import matplotlib.pyplot as plt
      class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog',
        ↪ 'horse', 'ship', 'truck']

      # Generate confusion matrix
      conf_matrix = confusion_matrix(Y_true, Y_pred_classes)

      # Plot confusion matrix
      plt.figure(figsize=(10, 7))
      sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues',
        ↪ xticklabels=class_names, yticklabels=class_names)
      plt.xlabel('Predicted')
      plt.ylabel('True')
      plt.title('Confusion Matrix')

```

```
plt.show()
```



```
[22]: test_loss, test_acc = model.evaluate(x_test, y_test, verbose=2)
print(f'Test accuracy: {test_acc:.4f}')
print(f'Test loss: {test_loss:.4f}')
```

```
313/313 - 6s - loss: 0.6039 - accuracy: 0.8034 - 6s/epoch - 19ms/step
Test accuracy: 0.8034
Test loss: 0.6039
```