



AI & Partners

Amsterdam - London - Singapore



EU AI Act

Fundamental Rights Impact
Assessments (FRIAs):

An Overview

October 2025

Authors

Sean Musch, Michael Borrelli, Charles Kerrigan, Jessica
Chappelow, Elena Knezovic, Susanne Chishti, Andrey Prozorov



AI & Partners

Amsterdam - London - Singapore

AI & Partners defends and extends the digital rights of users at risk around the world. By combining direct technical support, comprehensive policy engagement, global advocacy, grassroots professional services, regulatory interventions, and participating in industry groups such as AI Commons, we fight for fundamental rights in the artificial intelligence age.

This report was prepared by Sean Donald John Musch and Michael Charles Borrelli. For more information visit <https://www.ai-and-partners.com/>.

Contact: Michael Charles Borrelli | Director | m.borrelli@ai-and-partners.com.

This report is an AI & Partners publication.





Who Are We

AI That You Can Trust

Why Us?

Stay on the right side of history. At AI & Partners, we believe AI should unlock potential—not cause harm. We've seen the fear and fallout when teams lose control of AI, but also the trust and innovation that follow when it's handled responsibly. That's why we exist: to help you build AI you can trust and stand behind—for the long run.

80%

of AI systems
are unknown

What Do We Do?


We enable safe AI usage—for your organization and your clients. Unknown AI adoption leads to confusion, risk, and reputational damage. We help you take control with tools to identify, monitor, and govern all AI systems—so you're not reacting to AI, you're leading it.



How Do We Do It?


Do you know what AI systems you have? Identify all known and unknown AI systems (algorithms, LLMs, prompts, and models) from all internal and external AI vendors, automated by generating your inventory. Overall, 80% of AI inventory is unknown to our clients.

How do you guarantee ongoing safe AI use? Continuously monitor deployed AI systems for performance drift, anomalies or failures, real-world impacts, and emerging risks (e.g. data poisoning). Any malfunction of an AI system has severe implications for organisations (e.g. inability to assess online misinformation that leads to widespread public mistrust), so monitoring becomes a matter of urgency.



AI Discovery & AI Inventory

Automatically detect all AI systems, including models, algorithms, and prompts, and maintain a live, always-updated register for full visibility and compliance.



Responsible AI

Embed fairness, transparency, and control into every stage of AI use—aligning with the EU AI Act and building 'Trustworthy-by-Design'.



Model Monitoring

Continuously track your AI models after deployment to detect drift, bias, or failure—so you stay in control and prevent harm before it happens.



Contents

Who Are We	2
Introduction	4
Key questions being asked about Fundamental Rights Impact Assessments	5
1. What is a Fundamental Rights Impact Assessment (FRIA)?	6
2. Who must carry out a FRIA?	6
3. When must a FRIA be completed?	6
4. What should a FRIA include?	6
5. How does a FRIA differ from a Data Protection Impact Assessment (DPIA)?	6
6. What happens if the FRIA reveals serious risks?	7
7. Do FRIAs need to be made public?	7
8. What role does the AI provider play in the FRIA?	7
9. What tools or templates are available to help with a FRIA?	7
10. What are the consequences of not performing a FRIA?	7
Understanding Fundamental Rights Impact Assessments	8
1. Obligation to Perform a Fundamental Rights Impact Assessment (FRIA)	9
2. Core Elements of the FRIA	10
3. Timing and Updates to the FRIA	11
4. Notification to Market Surveillance Authorities	12
5. Integration with Data Protection Impact Assessments (DPIAs)	13
6. Template and Support from the AI Office	14
Mapping Fundamental Rights Impact Assessments to AI Ecosystem Actors	15
Calls to Action	0
Conclusion	2
Appendix A: FRIA Implementation & Templates	3
Purpose	3
A. Fundamental Rights — Structured Checklist	3
B. FRIA vs DPIA — Comparison Table & Integration Guidance	4
C. FRIA Risks vs General AI Risks — Side-by-Side	5
D. Clear Triggers for Conducting or Updating a FRIA	5
E. FRIA Template — Full fields & guidance	6
Appendix B: Expert Vignette	8
Elena Knezović	8
References	9





Introduction

Executive Summary - As the EU Artificial Intelligence Act (AI Act) approaches full implementation, Fundamental Rights Impact Assessments (FRIAs) stand as essential tools for safeguarding human dignity, freedoms, and equality in high-risk AI deployments. This report outlines FRIA obligations under Article 27, including who must conduct them (primarily deployers of high-risk systems in public or sensitive contexts), when (pre-deployment with updates for changes), and what they must include (six core elements focusing on risks, oversight, and mitigation). Key takeaways emphasize collaboration among deployers, providers, and stakeholders; integration with DPIAs; and practical steps for compliance. By embedding rights-based evaluations, FRIAs promote trustworthy AI, reduce risks, and foster innovation aligned with EU values. For a detailed breakdown, refer to the Contents on Page 4.

As the EU Artificial Intelligence Act (AI Act) moves toward implementation, **Fundamental Rights Impact Assessments (FRIAs)** have emerged as a core safeguard for ensuring that high-risk AI systems respect and protect human dignity, freedoms, and equality. These structured assessments require deployers to identify, document, and—where necessary—disclose the potential impacts their systems may have on fundamental rights, with a particular focus on vulnerable individuals and sensitive use cases.

This report serves as a **practical blueprint** for public and private deployers, providers, regulators, and civil society actors seeking to design, conduct, and integrate FRIAs into their AI governance practices. Drawing on the EU AI Act—specifically **Article 27** on FRIA obligations—it distils compliance requirements, ethical considerations, and operational guidance for effective implementation. From defining scope and affected groups to embedding mitigation strategies and notification protocols, the report identifies opportunities to align FRIA practice with broader EU rights-protection frameworks.

The FRIA obligation represents a decisive step in the EU’s strategy to balance AI innovation with **accountability and public trust**. It ensures that rights-based risk assessment is not an afterthought but a pre-deployment requirement, complemented by ongoing review as AI systems evolve. This process not only protects individuals from discriminatory or harmful outcomes but also helps organizations anticipate legal exposure and reputational risks before market entry.

In mapping the **roles, responsibilities, and interactions** among deployers, providers, and affected stakeholders, the document promotes a collaborative, evidence-based approach to rights-impact evaluation. It highlights how deployers can accelerate compliance, how providers can support with technical documentation, and how users and advocacy groups can contribute to more transparent and trustworthy AI systems.

Whether you are a municipality deploying biometric verification in public services, a healthcare provider implementing diagnostic AI, or a technology firm supplying high-risk models, this report provides a **structured reference for conducting FRIAs**—laying the foundation for lawful, ethical, and rights-aligned AI deployment in the EU.

Best regards,

Sean Musch

Founder/CEO

AI & Partners





Key questions
being
asked about
Fundamental Rights
Impact Assessments



1. What is a Fundamental Rights Impact Assessment (FRIA)?

A Fundamental Rights Impact Assessment (FRIA) is a formal evaluation conducted before deploying high-risk AI systems. It identifies potential impacts on people's fundamental rights, including risks like discrimination, exclusion, surveillance, or denial of access to services. The assessment involves analyzing how and when the AI will be used, who may be affected, how humans will oversee it, and what happens if something goes wrong. It's a key safeguard introduced by the EU AI Act to ensure that the use of advanced AI respects human dignity, rights, and freedoms — particularly in sensitive public-sector or high-risk service contexts. For example, a municipality using AI for biometric verification in welfare services might assess risks of excluding vulnerable migrants due to algorithmic biases.

2. Who must carry out a FRIA?

FRAs are required for deployers of high-risk AI systems in public bodies and private entities providing public services. Additionally, it applies to deployers of specific high-risk systems listed in Annex III of the EU AI Act. While providers (the developers or suppliers of AI systems) are not directly responsible for completing the FRIA, they must supply essential information to deployers. This ensures that the deployer can perform a well-informed assessment. The obligation applies regardless of whether the deployer is in the public or private sector, as long as the AI system falls into one of the specified high-risk categories.

3. When must a FRIA be completed?

The FRIA must be completed before the first use of the high-risk AI system. This pre-deployment requirement ensures that potential risks are proactively identified and addressed. However, the obligation doesn't end there. If anything changes — such as the context of use, the people affected, or the risks identified — the deployer must update the FRIA to reflect the new circumstances. This dynamic approach acknowledges that AI systems can evolve or behave differently in real-world conditions, and continuous evaluation is essential to protecting rights throughout the system's lifecycle, not just at the point of launch.

4. What should a FRIA include?

A FRIA must include six key elements: (1) how the AI system fits into the deployer's processes, (2) how often and how long the system will be used, (3) who may be affected, (4) the specific risks to those individuals or groups, (5) human oversight mechanisms, and (6) actions to take if risks occur. This structure ensures a comprehensive look at ethical, legal, and social risks. It's not enough to describe the AI's technical features — the focus must be on impact: who gets hurt, how, and what will be done about it if something goes wrong.

5. How does a FRIA differ from a Data Protection Impact Assessment (DPIA)?

A DPIA, required under GDPR, focuses on risks related to personal data processing — such as data minimization, consent, or profiling. A FRIA, however, covers broader human rights risks, such as discrimination, freedom of expression, or access to public services. If an organization has already done a DPIA, the FRIA should complement it by addressing rights issues not captured by GDPR alone. This integration ensures both data privacy and other human rights are considered when deploying AI. In essence, the FRIA adds a layer of ethical scrutiny that goes beyond data and looks at societal impact more broadly.





6. What happens if the FRIA reveals serious risks?

If a FRIA reveals serious risks to fundamental rights, the deployer must put in place mitigation measures — such as enhanced human oversight, access to redress mechanisms, or limiting the scope of the AI's use. In some cases, they may decide not to use the system at all. These measures should be clearly documented in the FRIA and communicated to the appropriate authorities. If the risks cannot be adequately mitigated, the system may not comply with the EU AI Act, and its deployment could be blocked or penalized. The goal is proactive prevention, not retroactive correction of harm.

7. Do FRIAs need to be made public?

Currently, the EU AI Act does not require FRIAs to be made public, but the results must be reported to the market surveillance authority (unless exempted). However, deployers may choose to publish a summary of the assessment to demonstrate transparency and accountability, especially when public trust is critical. In practice, voluntary disclosure can help build credibility and open constructive dialogue with affected communities or watchdog groups. Some sectors, like healthcare or education, may even benefit from co-creating assessments with community input. Still, sensitive security or operational details can remain confidential as long as the rights assessment is thorough.

8. What role does the AI provider play in the FRIA?

While the legal obligation to complete the FRIA lies with the deployer, providers play a critical supporting role. They must supply detailed information about the AI system — including its intended purpose, known risks, limitations, and human oversight instructions — as required under Article 13. Without this input, deployers cannot properly assess how the system will affect fundamental rights in real-world use. Good providers will go further, offering templates, guidance, and even post-market monitoring support. This cooperative relationship between provider and deployer is key to achieving meaningful, accurate, and actionable FRIAs that live up to their protective purpose.

9. What tools or templates are available to help with a FRIA?

The AI Office is tasked with developing a standardized template and possibly an automated tool to assist deployers in completing the FRIA. This template will help ensure consistent assessments across different sectors and simplify the process, especially for smaller organizations. It's expected to cover each required FRIA element with structured questions and prompts. Using the official template not only supports compliance but also helps deployers think critically about risk. Over time, this may evolve into a centralized tool with guidance, examples, or best practices. Until then, many organizations may choose to adapt existing risk or ethics frameworks to get started.

10. What are the consequences of not performing a FRIA?

Failure to perform a required FRIA — or submitting an incomplete or misleading one — could lead to regulatory penalties under the EU AI Act. More importantly, it exposes people to real harm if AI systems are used irresponsibly, especially in high-risk sectors. Organizations that skip this step may face legal action, reputational damage, or even be ordered to cease using the AI system. Given the high stakes involved, especially for vulnerable populations, the FRIA is not just a box-ticking exercise — it's a vital tool for ensuring ethical, lawful, and rights-respecting use of artificial intelligence in critical domains.





Understanding Fundamental Rights Impact Assessments



1. Obligation to Perform a Fundamental Rights Impact Assessment (FRIA)



What is this?

- This is a legal requirement for public bodies and private entities providing public services to evaluate how high-risk AI systems might affect individuals' fundamental rights before deploying them.
- It's a pre-deployment analysis focused on ethical and legal impacts.

How is it relevant?

- It ensures that AI is implemented responsibly in sensitive environments like healthcare, education, or law enforcement, where misuse or errors could infringe on basic human rights.

Why does it matter?

- Without this obligation, AI could be deployed without safeguards, risking discrimination, privacy violations, and loss of accountability — particularly affecting vulnerable communities.





2. Core Elements of the FRIA



What is this?

- These are the specific content requirements of the FRIA, including use context, frequency, affected groups, risks, human oversight mechanisms, and mitigation plans.
- They form the structural backbone of the assessment.

How is it relevant?

- They guide deployers in systematically identifying and evaluating potential harms, ensuring all necessary angles — from operational detail to ethical oversight — are covered before the AI system is used.

Why does it matter?

- These elements ensure that rights risks are not overlooked.
- For example, in a healthcare AI diagnostic tool, affected groups might include elderly patients; risks could involve biased outcomes leading to misdiagnosis; and mitigation might include regular bias audits.



3. Timing and Updates to the FRIA



What is this?

- This provision mandates when the FRIA must be completed — prior to first use — and requires updates if the context of AI use changes or new risks emerge during operation.

How is it relevant?

- It enforces a dynamic responsibility for deployers, ensuring the FRIA isn't a one-off task but a living document that reflects evolving realities and risks of AI deployment.

Why does it matter?

- Technology and society evolve. If FRIAs aren't updated, systems could cause harm no longer aligned with original assumptions, leaving users unprotected..
- For instance, if an educational AI tool's usage expands to include new student demographics, an update might reveal emerging privacy risks not initially foreseen.



4. Notification to Market Surveillance Authorities



What is this?

- This is the requirement for deployers to report the results of their FRIA to the relevant authorities, using an official template — unless legally exempt.
- It provides external oversight of high-risk AI systems.

How is it relevant?

- It makes the assessment transparent and verifiable by public bodies, allowing regulators to track systemic risks, ensure compliance, and intervene if necessary.

Why does it matter?

- Without this notification, oversight bodies can't monitor how AI is used in high-risk areas.
- Reporting ensures public accountability and builds trust in how institutions use AI.



5. Integration with Data Protection Impact Assessments (DPIAs)



What is this?

- This rule allows the FRIA to complement existing DPIAs (under GDPR), expanding from data privacy to broader human rights concerns.
- It encourages synergy rather than redundancy in compliance efforts.

How is it relevant?

- Many organizations already conduct DPIAs.
- Adding FRIA as a complement ensures that both data privacy and wider human rights are covered without duplicating effort.

Why does it matter?

- AI systems can harm rights beyond data privacy — such as freedom of expression or protection from bias.
- Integrating these perspectives ensures more holistic, meaningful protections for individuals.



6. Template and Support from the AI Office



What is this?

- This is a standardized questionnaire and/or digital tool provided by the AI Office to guide and support deployers in completing the FRIA accurately, efficiently, and consistently.

How is it relevant?


- It reduces uncertainty and administrative burden for deployers, especially smaller organizations, and ensures assessments meet regulatory expectations and are comparable across sectors.

Why does it matter?

- Templates help prevent omissions and encourage good practice.
- They also standardize reporting for authorities, making oversight easier and helping spot systemic issues or patterns in AI-related rights risks.



Mapping Fundamental Rights Impact Assessments to AI Ecosystem Actors



To promote inclusivity and collaboration, deployers are encouraged to co-create FRIAs with civil society, affected users, or advocacy groups. This involves consultations to identify overlooked risks, ensuring assessments are transparent and representative. Such engagement aligns with the EU AI Act's emphasis on accountability and can enhance public trust.

Section	Description	Deployers	Providers	Users	Challenge
Obligation to Perform a Fundamental Rights Impact Assessment (FRIA)	Before using a high-risk AI system, certain public and private deployers must conduct a detailed assessment of how the system could affect fundamental rights. This ensures responsible AI deployment, especially for systems impacting vulnerable individuals or key public functions.	Deployers must perform the FRIA before using a high-risk AI system. They are directly responsible for identifying fundamental rights risks in their operational context. This includes detailing how the system integrates into their processes and services to ensure lawful, fair, and safe deployment.	While the legal obligation to conduct the FRIA falls on the deployer, providers must supply sufficient information (as per Article 13) about the AI system's intended use, risks, and limitations to enable deployers to complete the FRIA accurately and responsibly.	Users are generally not responsible for completing the FRIA, but their experiences and feedback can highlight risks not foreseen during the assessment. They may also rely on assurances that the deployer has properly assessed the system's fundamental rights impact.	Ensuring timely access to provider data; balancing assessment depth with resource constraints for smaller deployers.
Core Elements of the FRIA	The assessment must include: how the AI system will be used, when and how often, who it may affect, specific risks identified, human oversight plans, and risk mitigation steps. This structured approach ensures a comprehensive understanding of potential impacts	Deployers must describe the AI system's role in their operations, its usage patterns, the affected individuals, and the associated risks. They must also detail human oversight strategies and define corrective measures if risks materialize, ensuring they actively manage any rights-related harm.	Providers must ensure that deployers have access to comprehensive, clear information regarding the AI system's functionalities, intended use, potential limitations, and any known risks. This input is essential for deployers to fill in critical parts of the FRIA.	Users are indirectly protected by the FRIA's core elements. Knowing that risk identification, human oversight, and mitigation plans are required reassures users that their rights will be considered in the AI system's operation and any potential harm will be managed.	Identifying all affected groups accurately; addressing evolving risks in dynamic environments.



	before deployment.				
Timing and Updates to the FRIA	The FRIA must be completed before the AI system's first use. If circumstances change (e.g., new risks emerge or usage patterns shift), deployers are required to update the assessment to keep it current and relevant throughout the AI system's lifecycle.	Deployers must complete the FRIA before first use and remain vigilant during the system's operation. If usage circumstances change — such as new user groups, different deployment environments, or unexpected risk materialization — they must promptly update the FRIA to reflect these changes.	Providers should monitor post-market performance of their systems and inform deployers of any new risks or changes that could impact the deployers' obligation to update the FRIA. This cooperation ensures ongoing protection of fundamental rights throughout the system's lifecycle.	Users might notice when AI behavior shifts or unexpected impacts occur. Their feedback could trigger a deployer's obligation to review and update the FRIA, helping maintain high standards of rights protection during the system's continued use.	Detecting subtle changes in real-time; coordinating updates across teams.
Notification to Market Surveillance Authorities	After completing the FRIA, deployers must submit the results using a standardized template to the designated market surveillance authority. Some cases may be exempt from this reporting requirement under specific conditions outlined separately in the legislation.	Deployers must submit the FRIA results to the market surveillance authority, using the AI Office's standardized template. This notification ensures public oversight and regulatory accountability for high-risk AI deployments, unless an exemption applies under specific circumstances.	Providers are not required to submit the FRIA, but they must support deployers by providing detailed system documentation that will likely be attached to or referenced in the notification to authorities, ensuring accurate compliance reporting.	Users have no direct responsibility regarding notifications, but the public availability of some information from FRIA notifications could empower users to better understand their rights, scrutinize AI deployments, or raise complaints if rights appear at risk.	Maintaining confidentiality of sensitive details; handling exemptions consistently.



Integration with Data Protection Impact Assessments (DPIAs)	If a deployer has already conducted a Data Protection Impact Assessment under GDPR or related regulations, the FRIA should build on that work. The FRIA complements existing assessments by adding a broader human rights dimension beyond just data protection concerns.	Deployers who already conducted a Data Protection Impact Assessment (DPIA) under GDPR must now ensure that the FRIA complements and extends it. Where DPIAs focus on personal data, FRIAs must add a broader assessment of impacts on other fundamental rights.	Providers must clarify to deployers how their AI system aligns with privacy and human rights obligations. When designing AI systems, they should ensure their documentation supports both DPIA and FRIA processes for easier integration by deployers.	Users benefit from the integration of DPIA and FRIA because it ensures their privacy and broader fundamental rights — such as non-discrimination, freedom of expression, and dignity — are simultaneously and systematically protected when AI systems are deployed.	Avoiding redundancy in assessments; aligning with evolving GDPR-AI Act intersections.
Template and Support from the AI Office	The AI Office will provide a standardized questionnaire and automated tools to help deployers complete the FRIA more easily and uniformly. This support aims to reduce administrative burden and ensure that assessments are thorough, consistent, and accessible to a range of deployers.	Deployers will use the standardized FRIA template developed by the AI Office, simplifying compliance and ensuring consistency in how risks are assessed and reported. This practical tool helps deployers focus on substantive rights issues rather than administrative complexity.	Providers should familiarize themselves with the AI Office template to better anticipate deployers' needs. This allows providers to prepare technical documentation and support materials that align with the structure and expectations of the FRIA process.	Users indirectly benefit from the template's existence because it promotes transparency and consistency across AI deployments. A uniform FRIA process increases the likelihood that users' fundamental rights are properly considered before and during system operation.	Adapting templates for sector-specific needs; ensuring accessibility for non-experts.



Calls to Action



Establish a Cross-Functional Team for FRIA Compliance

Create a dedicated team involving legal, IT, ethics, and service operations to identify high-risk AI systems and oversee FRIA obligations. Integrate this requirement into internal governance protocols to ensure assessments are conducted before deployment and reflect human rights considerations.

- **KPI for Success:** Track completion rates of FRIAs pre-deployment and audit findings related to team involvement (e.g., aim for 100% cross-functional review).



Develop a Standardized Internal FRIA Checklist

Design a reusable FRIA template that captures the six core elements: context, duration, affected groups, risks, oversight, and mitigation. Collaborate with stakeholders through structured interviews or review sessions to ensure thorough and consistent documentation of risks across all high-risk AI systems.

- **KPI for Success:** Measure checklist usage adoption (e.g., 90% of high-risk systems assessed via the template) and reductions in identified risks post-mitigation.



Implement a Trigger-Based FRIA Update Process

Introduce internal protocols that automatically flag when updates to a FRIA are needed—e.g., system upgrades, expanded use cases, or changes in affected user groups. Assign responsibility to a designated compliance lead to review changes and initiate reassessment where necessary.

- **KPI for Success:** Monitor update frequency (e.g., quarterly reviews) and incident rates (e.g., reduced bias incidents by 20% through timely updates).



Build Reporting Mechanisms to Notify Authorities

Create internal workflows and documentation standards to support notification to market surveillance authorities. Use the forthcoming AI Office template to ensure consistency. Store completed FRIAs securely and designate team members responsible for liaising with regulators in case of audits, inspections, or updates.

- **KPI for Success:** Achieve 100% on-time notifications and track regulatory feedback or compliance scores.





Conclusion

The introduction of **Fundamental Rights Impact Assessments (FRIAs)** under the EU AI Act represents a transformative shift in how high-risk AI systems are evaluated, deployed, and governed—moving toward a more **rights-centered, transparent, and risk-based compliance model**. By requiring structured, pre-deployment analysis of potential impacts on fundamental rights, FRIAs translate the AI Act's legal obligations into practical safeguards, reducing uncertainty for deployers while reinforcing user protections and regulatory accountability.

Throughout this report, we have highlighted how FRIAs serve as a **critical convergence point** for deployers, providers, regulators, and affected communities—facilitating early identification of risks, informed oversight, and the development of AI systems that are both lawful and explainable. When executed with diligence and inclusivity, FRIAs can accelerate alignment with the AI Act's requirements for transparency, human oversight, and post-market monitoring, while also fostering trust in AI across the EU.

However, successful implementation of FRIAs depends on more than meeting the minimum legal requirements. It demands **operational readiness, stakeholder engagement, and a commitment to continuous review**.

Organizations must adapt governance models, documentation practices, and risk management processes to reflect the AI Act's expectations, particularly for high-risk scenarios and contexts involving vulnerable individuals.

Smaller deployers and first-time practitioners may face steeper learning curves, making early planning, training, and collaboration with regulators and civil society essential.

FRIAs not only strengthen EU compliance but also align with international standards, such as UNESCO's Recommendation on the Ethics of Artificial Intelligence and OECD AI Principles, which emphasize human rights, transparency, and inclusivity. This positions the EU AI Act as a potential global benchmark, encouraging harmonized approaches to rights-based AI governance worldwide.

Despite these challenges, forward-looking deployers are already showing that well-designed FRIAs can **streamline compliance, enhance system quality, and mitigate reputational and legal risks**. As AI becomes increasingly embedded across sectors—from healthcare and education to public administration and civic infrastructure—FRIAs provide a powerful tool for ensuring that innovation advances **without compromising fundamental rights**. To conclude, robust and transparent FRIA practices could not only strengthen compliance under the EU AI Act but also set a **global benchmark for rights-based AI governance**.





Appendix A: FRIA Implementation & Templates

Purpose

This Appendix collects the requested amendments and practical tools for deployers, providers and compliance teams. It is intended to be attached to the FRIA report as a single, self-contained reference that: (1) turns narrative guidance into actionable checklists and templates, (2) clarifies interactions with DPIAs, (3) distinguishes rights risks from technical risks, (4) provides clear triggers for when to run a FRIA, and (5) offers a ready-to-use FRIA template and one-page checklist.

A. Fundamental Rights — Structured Checklist

Use this checklist when scoping and running a FRIA. Mark each item as **Yes / No / N/A** and add brief evidence notes.



- **Right to non-discrimination and equality** — Could the system produce or amplify disparate outcomes by race, gender, age, disability, nationality, or other protected characteristic?
- **Right to dignity and freedom from degrading treatment** — Could the system create intrusive surveillance, demeaning profiling, or humiliating outcomes?
- **Right to privacy and data protection** — Does the system entail collection, use or combination of personal data? Are safeguards proportionate and documented?
- **Right to freedom of expression and assembly** — Could the system censor, downrank, or chill lawful speech or lawful assembly?
- **Right to access essential services** — Could the system deny, delay, or degrade access to healthcare, education, welfare, justice or other essential public services?
- **Right to due process, contestability and effective remedy** — Can affected persons understand, challenge or seek remedy for decisions materially affecting them?
- **Rights of vulnerable groups** — Will children, elderly, migrants, persons with disabilities, or other vulnerable groups be affected differently or more severely?
- **Economic and social rights (where relevant)** — Could deployment worsen access to employment, housing, social support or other socio-economic entitlements?
- **Freedom from arbitrary or discriminatory automated decision-making** — Is there a risk of unsupported automated exclusion or automated profiling without human oversight?



- **Security and safety of persons** — Could errors lead to bodily harm, life-threatening outcomes, or dangerous physical states?

Checklist guidance: for any item flagged “Yes”, the FRIA must document: (a) nature of risk, (b) affected groups, (c) likelihood and severity, (d) mitigation(s) already planned, (e) residual risk, (f) monitoring plan and (g) whether notification to authorities is required.

B. FRIA vs DPIA — Comparison Table & Integration Guidance

Dimension	FRIA (Fundamental Rights Impact Assessment)	DPIA (Data Protection Impact Assessment)
Legal basis	EU AI Act (Article 27)	GDPR (Article 35)
Primary focus	Impacts on fundamental rights (discrimination, access to services, expression, dignity)	Risks to personal data and privacy (confidentiality, integrity, lawful basis)
Typical triggers	Deployment of high-risk AI systems (Annex III) or use affecting vulnerable groups	Personal data processing likely to result in high risk to rights/freedoms of natural persons
Responsible party	Deployer of the high-risk AI system	Data controller (or jointly with processor)
Required when	Mandatory for high-risk AI deployers; Mandatory where processing is likely to updates when context changes	Mandatory where processing is likely to result in high risk to individuals under GDPR
Overlap	Rights that stem from data processing (privacy) are covered, but FRIA goes further	Privacy-specific mitigations feed into FRIA; DPIA outputs are inputs to FRIA if personal data used
Outcome	Rights-oriented mitigation, possible restriction or halting of deployment	Privacy mitigations, possible limitation or halt to processing

Integration rule (explicit): When an AI system processes personal data, the deployer **must** integrate DPIA results into the FRIA. That means: attach the DPIA as an annex (or link), map DPIA mitigations to FRIA mitigation measures, and ensure any residual DPIA-level risk is carried through into the FRIA’s residual risk assessment and monitoring plan.

When to run which assessment (practical guidance):

- Run a **DPIA** when new personal data processing is introduced or a materially new use of personal data occurs.
- Run a **FRIA** when the system is high-risk under Annex III, affects vulnerable groups, or has broader societal impacts beyond privacy.
- **If both apply**, complete the DPIA first (or in parallel), then explicitly reference and incorporate DPIA findings into the FRIA. The FRIA must still separately address non-data rights (e.g., discrimination, access).



C. FRIA Risks vs General AI Risks — Side-by-Side

General (Technical / Organisational) AI Risks

- Model robustness failures (drift, poor generalisation)
- Data integrity and poisoning
- Cybersecurity vulnerabilities and supply-chain risks
- Explainability and traceability gaps
- Performance variability across environments
- Operational continuity and scalability issues

FRIA (Fundamental Rights) Risks

- Discrimination and unequal treatment
- Exclusion from essential services
- Unlawful surveillance and dignity harms
- Chilling effects on free expression or assembly
- Denial of due process or routes for redress
- Disproportionate harm to vulnerable groups

Practical note: Organisational risk management should cover the general AI risks; FRIA covers the rights-centred consequences of those technical failures. For example, model drift (technical risk) can lead to discriminatory outputs (FRIA risk) — both require mitigation, but the FRIA frames the mitigation in terms of rights protection and remedies.

D. Clear Triggers for Conducting or Updating a FRIA

Conduct a FRIA (or update an existing FRIA) when any of the following apply:

1. **System classification:** the AI system is high-risk as per Annex III of the AI Act.
2. **New deployment context:** the system is deployed in a new geographic, cultural, or sectoral context that alters affected populations.
3. **New affected groups:** new categories of people (including vulnerable groups) are brought into scope.
4. **Functional change or upgrade:** material changes in algorithms, training data, model scale, or decision logic.
5. **Integration of personal data:** adding or changing personal data inputs that were not covered in prior assessments.
6. **Evidence of harm or incidents:** complaint, audit, incident report, or monitored output indicating potential rights harm (e.g., bias metrics exceed thresholds).
7. **Regulatory or policy change:** new guidance from AI Office or national authority that affects risk thresholds.





8. **Change in interconnections:** new system integrations, API linkages, or third-party model dependencies.

Timing & minimum cadence: FRIA must be completed before first use; thereafter, schedule periodic reviews (e.g., quarterly for high-impact deployments, annually otherwise) and immediate updates on any trigger.

E. FRIA Template — Full fields & guidance

Below is a pragmatic template to paste into compliance tools or documentation systems. Fields marked **(Required)** should always be completed for high-risk systems.

FRIA — Template Fields

1. **Document metadata** *(Required)*
 - FRIA title
 - Unique FRIA ID
 - Deployer name and contact
 - Provider/vendor name and contact
 - System name and version
 - Date completed and version history
2. **Executive summary** *(Required)*
 - 2-3 sentence overview of purpose, main risks identified, and recommended Go/No-Go or conditional approval.
3. **System description & scope** *(Required)*
 - Functional description, intended purpose, primary outputs/decisions, human-facing interfaces, deployment environment.
4. **Process integration** *(Required)*
 - How the system fits into existing business processes or public services; decision points and escalation pathways.
5. **Usage profile** *(Required)*
 - Frequency, duration, user roles (operators, affected persons), expected volumes, geographic scope.
6. **Data map & DPIA linkage** *(Required where personal data used)*
 - Data types, data flows, retention, lawful basis, reference to DPIA (attach or link), pseudonymisation/ anonymisation status.
7. **Affected individuals & stakeholder map** *(Required)*
 - Primary and secondary affected groups, vulnerable subgroups, estimated population sizes, stakeholder engagement plan.
8. **Fundamental rights checklist & mapping** *(Required)*





- For each right in Section A, describe whether affected (Y/N), how, and evidence.

9. Risk assessment (per identified right) (Required)

- For each identified risk: description, severity (High/Medium/Low), likelihood (High/Medium/Low), justification, and quantification where possible.

10. Mitigation measures (Required)

- Technical (filters, bias mitigation, differential privacy), organisational (training, governance), legal/contractual (procurement clauses), and procedural (human oversight, escalation).
- Map each mitigation to the risk(s) it addresses and include implementation owner and timeline.

11. Human oversight & contestability (Required)

- Human-in-the-loop points, onboarding/ escalation procedures, how individuals can contest decisions, and redress mechanisms.

12. Monitoring, metrics & thresholds (Required)

- Define KPIs, fairness metrics, performance metrics, drift detection thresholds, incident definitions, and who monitors them.

13. Residual risk & decision (Required)

- Residual risks after mitigation, acceptability decision (Accept / Accept with Conditions / Reject), signoff (names, roles, date).

14. Notification & publication (Required)

- Whether the FRIA will be notified to market surveillance authorities (yes/no), attachments included, and whether a public summary will be published.

15. Review & update log (Required)

- Triggers for update, review cadence, latest update summary.

16. Annexes

- DPIA (if applicable), technical documentation from provider (Article 13 materials), test reports, stakeholder consultation notes, relevant procurement contracts.


Practical tips for completion:

- Use quantifiable metrics where possible (e.g., demographic disparity ratios, false positive rates by group).
- Keep the executive summary short but definitive — reviewers and authorities often read this first.
- Attach supporting evidence (test datasets, code-level descriptions) in annexes





Appendix B: Expert Vignette



Elena Knezović

*AI Visualisation
and Storytelling*



“

In an era where AI touches every layer of society, trust becomes the true currency of innovation. Regulation alone can't create it, design, intention, and accountability must work together to make AI genuinely human-centered.

”

Biography

Elena Knezović, is a Conceptual Brand Visionary for founders and brands who want more than design they want depth, emotion and distinction. With clients like Laura Mercier and other forward-thinking names in fashion, beauty, and lifestyle, she helps brands transform abstract ideas into immersive visual stories that connect.



AI & Partners

Amsterdam - London - Singapore



References

Artificial Intelligence Act, (2025), 'Article 27: Fundamental Rights Impact Assessment for High-Risk AI Systems', accessible at <https://artificialintelligenceact.eu/article/27/> (last accessed 7th May 2025)

European Parliament and The Council of the European Union, (2024), 2024/1689 Regulation (EU) 2024/1689 of the European Parliament and of The Council of 13 June 2024 laying down harmonised rules on artificial intelligence and amending Regulations (EC) No 300/2008, (EU) No 167/2013, (EU) No 168/2013, (EU) 2018/858, (EU) 2018/1139 and (EU) 2019/2144 and Directives 2014/90/EU, (EU) 2016/797 and (EU) 2020/1828 (Artificial Intelligence Act), accessible at https://eur-lex.europa.eu/legal-content/EN/TXT/PDF/?uri=OJ:L_202401689 (last accessed 7th May 2025)

OECD, (2019), OECD Principles on Artificial Intelligence, accessible at <https://oecd.ai/en/ai-principles> (last accessed August 09, 2025).

UNESCO, (2021), Recommendation on the Ethics of Artificial Intelligence, accessible at <https://unesdoc.unesco.org/ark:/48223/pf0000380455> (last accessed August 09, 2025).





Important notice

This document has been prepared by AI & Partners B.V. for the sole purpose of enabling the parties to whom it is addressed to evaluate the capabilities of AI & Partners B.V. to supply the proposed services.

Other than as stated below, this document and its contents are confidential and prepared solely for your information, and may not be reproduced, redistributed or passed on to any other person in whole or in part. If this document contains details of an arrangement that could result in a tax or National Insurance saving, no such conditions of confidentiality apply to the details of that arrangement (for example, for the purpose of discussion with tax authorities). No other party is entitled to rely on this document for any purpose whatsoever and we accept no liability to any other party who is shown or obtains access to this document.

This document is not an offer and is not intended to be contractually binding. Should this proposal be acceptable to you, and following the conclusion of our internal acceptance procedures, we would be pleased to discuss terms and conditions with you prior to our appointment. Images used throughout the document have either been produced in-house or sourced from publicly available sources (see **References** for details).

AI & Partners B.V. is the Dutch headquarters of AI & Partners, a global professional services firm. Please see <https://www.ai-and-partners.com/> to learn more about us.

© 2025 AI & Partners B.V. All rights reserved.

Designed and produced by AI & Partners B.V.