



Agentic AI advantage: Unlocking next-level value

KPMG. Make the Difference.

kpmg.com/agenticAI



At KPMG, the future is agent-centric

AI has reached another inflection point. Just as organizations begin to scale generative AI, a more transformative shift is already underway: the rise of AI agents.

Agents aren't just the next wave — they're a step-change. These systems will act independently, pursue goals and operate across complex workflows. That shift changes everything: how work gets done, how decisions are made and how value is delivered.

This paper breaks down what agentic AI really means: how to define it, what it enables and where to begin. The perspectives are grounded in the real world — drawn from the work we're doing with clients and in our own business.

The time to prepare is now.

Unlike cloud, which took nearly two decades to reshape enterprise infrastructure, this shift is playing out in months. Leading organizations aren't waiting — they're already investing, scaling and positioning themselves to capture the benefits of AI agents at speed.

This isn't just a tech conversation — it's a leadership one. Agents will demand tighter alignment between business and technology leaders than anything that's come before. Success will require a shared vision, coordinated execution and a clear plan for embedding intelligence at scale.

That includes preparing your people. The organizations that will lead will likely be those that help their workforces to adapt by training teams to work with agents, clarifying evolving roles and leading the cultural shift that this level of transformation demands.

And as agents gain more autonomy and influence, trust becomes more critical. A scalable strategy isn't only about what agents can do — it's about ensuring that they operate reliably, transparently and within guardrails that build confidence.

With all the hype, it's easy to lose focus. Agents can be a critical part of AI strategy, but they aren't the strategy. The leaders pulling ahead of the pack are moving fast and with clarity — staying anchored on the "no regrets" moves: aligning leadership, modernizing tech and data, upskilling their people and embedding trust at every layer.

The shift won't be easy. But it can be defining. The decisions that leaders make today will likely reshape industries and redefine what it takes to secure a competitive advantage.

As we continue our AI journey, we'll continue to share what we learn — and welcome your ideas and perspectives as we collectively navigate what's next.



Steve Chase

Global Head of AI and Digital Innovation, KPMG International and Vice Chair – Artificial Intelligence & Digital Innovation KPMG in the US

Contents

04 Welcome to the age of agentic AI

05 The power behind the promise of AI

06 The size of the prize could be game-changing

07 How much value could be at stake?

08 Four ways AI agents can unlock value

10 Staying ahead of the regulatory curve

11 The spectrum of agents is evolving to meet critical needs

12 From experimentation to implementation — Choosing the right agent for the task

18 Navigating the agentic AI obstacle course

28 How to accelerate your agentic-AI journey

30 How KPMG can help

Welcome to the age of agentic AI

As organizations begin to scale the game-changing capabilities of generative AI (Gen AI), an even more advanced capability is emerging: the rise of AI agents. While still in its early days, agentic AI is evolving quickly and demonstrating a transformative potential that no corporate executive or board can afford to ignore.

In simple terms, **Gen AI** builds “digital assistants” that help employees free up time to be more productive, effective and creative. Digital assistants *augment* their human colleagues by providing information and guidance to help them take better decisions and act faster.

Agentic AI is the next leap forward — going beyond *assisting* humans by fully *automating* what humans do and acting on their behalf to deliver the intended business outcomes. By integrating Gen AI tools, data sources and insights, AI agents are able to unlock significantly greater business value.

They’re already capable of taking on human tasks that were once considered too complex to

automate, and they may be an even better option for some activities previously handled by earlier AI tools that are now being reimaged.


The dawn of human/AI collaboration

Agents’ ability to automate complex tasks is growing rapidly — currently doubling every three to seven months. In coding, their capabilities have already grown from handling tasks that took humans seconds to complete to automating hours-long tasks. If this exponential growth continues, agents could be four times more capable within six months and sixteen times more powerful in a year, enabling them to move from automating human *tasks* to automating human *roles*, creating an entirely new talent pool of “digital co-workers” or “teammates”.


Agents are also increasingly capable of communicating and collaborating with each other, further amplifying their business utility. By taking on a series of complex tasks as a group, “swarms”¹ of agents with different skills may eventually automate entire business processes from end to end because once agents are automated, they can autonomously re-engineer and continuously improve these processes.

The number of digital workers may start out small but has the potential to scale quickly, flexing up or down with demand. By taking on routine tasks, digital colleagues free human workers to focus on higher-value, complex and stimulating work — while learning to collaborate seamlessly with their digital counterparts in the shared pursuit of breakthrough performance.

This multi-layered approach can unlock operational performance while also creating opportunities to reimagine entire business models. However, realizing this value will require careful management of risks to ensure it is achieved safely and securely.



of organizations are
using or exploring
agentic AI, with



already
extensively active
in this space.

¹ National Library of Medicine, “The road forward with swarm systems,” 30 January 2025

Source: Intelligent enterprises, KPMG International, 2025

The power behind the promise of AI

Generative AI unlocked insights at scale. Agentic AI is taking the next step — turning insights into intelligent, trusted actions.

AI agents bring the agency that AI systems need.

They blend advanced reasoning from LLMs with planning, orchestration, knowledge, data mining, tool calling (including MCP and connectors), guardrails and governance to achieve organizational goals through meaningful action.

They make real-time decisions, adapt to new situations and learn from their interactions and

reinforced feedback — all under appropriate levels of human oversight both in the loop and on the loop.


The range of real-world applications is vast, from task-oriented agents that extract and compare data or documents against standards, to fully autonomous agents capable of transforming entire value chains.

Organizational goals are the primary input, guiding agents to interpret requirements and formulate actionable strategies by ingesting both structured and unstructured data.

The tools and approaches they use are selected and adapted dynamically, based on defined policies, frameworks, success criteria and external conditions.

KPMG International is pleased to present this collaborative report built on pivotal research findings by KPMG in the US. Their innovative approach to industry trend analysis forms the foundation of this report, now expanded to a global context and reflecting our shared commitment to delivering timely, relevant, and actionable insights worldwide.

We encourage readers who are interested in the original report from KPMG in the US to explore it [here](#).



The size of the prize could be game-changing

While the scenarios may sound futuristic, KPMG professionals believe that most companies will eventually transform into hybrid organizations where both humans and AI agents collaborate seamlessly — driving efficiencies, transforming operations, growing revenue, reducing costs and gaining a lasting competitive edge.

The IDC, for example, projects that Gen AI alone will add nearly US\$10 trillion to the global GDP over the next 10 years.² The value potential for agentic AI promises to be exponentially higher, based on KPMG proprietary research of AI's on more than 17 million companies across major global markets.³ KPMG estimates that agentic AI alone will be key to unlocking a staggering US\$3T in corporate productivity improvements.⁴

It's also anticipated that agentic AI has the potential to unlock at least five percent of EBITDA in labor productivity for the average Fortune 1000 company — a potential that is echoed by many industry analysts. A recent study from IDC, for example, suggests that organizations adopting agentic AI can achieve an 18 percent improvement in employee productivity and satisfaction.⁵

And according to Erik Brynjolfsson, Professor and Senior Fellow at the Stanford Institute for Human-Centered AI, “The next 10 years will be the biggest business transformation in all of history.” He suggests, “As the applications of AI become more widespread, we will see the economy-wide productivity growth rate at least double.”⁶ Given the size of the prize, this paper explores how business leaders can navigate the next phase of the AI journey and the value at stake for organizations that get it right.

It introduces the various types of agents that are emerging in the KPMG TACO framework™ — Taskers, Automators, Collaborators and Orchestrators — to optimize agentic AI and outlines the challenges ahead, including **four key barriers** and **practical actions** that can be taken *today* to drive real value in the world of agentic AI *tomorrow*.

But, while agentic AI unlocks powerful capabilities, it also shifts how organizations must think about control, responsibility and trust. For enterprise leaders, that means moving early and deliberately to ensure the use of AI is safe, legal and aligned with their organization's values. They must also address the urgency of staying ahead of the regulatory, risk and compliance curve that will define where these agents can (and cannot) be applied, and how they should be evaluated.

KPMG estimates that agentic AI alone will be key to unlocking a staggering US\$3T in corporate productivity improvements.⁴

² IDC, “The Business Opportunity of AI” infographic, November 2023

³ [Quantifying the Gen AI Opportunity](#), KPMG in the US, 2025.

⁴ Ibid

⁵ IDC, “The Global Impact of Artificial Intelligence on the Economy and Jobs: AI will Steer 3.5% of GDP in 2030,” August 2024

⁶ KPMG International, “Intelligent Industries, A blueprint for creating value through AI-driven transformation,” 2025

How much value could be at stake?

AI is evolving so rapidly in terms of both advanced capabilities and lower costs that any predictions of future value must be considered to be speculative. In the last 18 months alone, usage costs for leading AI models with equivalent performance have dropped from US\$20.00 to US\$0.07 per million tokens — a 280-fold improvement.⁷ The exponential pace of development for both generative and agentic AI will likely continue for the foreseeable future, with performance promising to improve at a faster rate than silicon chips did according to Moore's Law.

In terms of market spend on agentic AI, estimates from leading analysts are bullish and growing fast. In addition to the US\$50B in market value that many analysts have predicted, returns on agentic-AI investments are also promising, according to an IDC study. Their research finds that for every US\$1 a company invests in AI, it can realize an average of US\$3.50 in return, with five percent of organizations worldwide realizing an average return of US\$8.⁸

The KPMG AI Value Assessment model has taken a first step in attempting to quantify the addressable opportunity for agentic AI by industry, company, function and roles. It is too early to make reliable predictions, and this is a conservative approach that looks only at the impact on labor productivity. Assuming that agentic AI will be the key to automating high-complexity tasks in the model, it equates to US\$3T in corporate productivity improvements and a 5.4 percent EBITDA improvement for the average company. (See sidebar on the right)

⁷ Stanford University, Human Centered Artificial Intelligence (HAI), "AI Index 2025: State of AI in 10 charts," 7 April 2025

⁸ IDC, "The Business Opportunity of AI" infographic, November 2023

The KPMG AI value assessment



In 2024, KPMG launched a massive study with the aim of developing the world's largest database to help quantify the potential impact of Gen AI on labor productivity for enterprises across the world.



Leveraging KPMG's patent pending AI Value Assessment model, the equivalent of an organizational digital twin was built for more than 17 million companies using more than 3 billion data points.



Due to the availability of financial information, 7,000 public companies were identified in the total sample, which collectively employ more than 72 million workers in 2,000 different roles, with average revenue of \$7.5B.



The complexity of tasks exposed to AI were graded into low, medium or high, using internal AI models, based, for example, on the degree of reasoning and interactions required.



Of these, it was determined that agentic AI is best positioned to unlock value in high-complexity tasks.

Source: KPMG in the US, "The Agentic AI advantage: Unlocking the next level of AI value," June 2025

Four ways AI agents can unlock value

1 Expanding the boundaries of automation

Agentic AI is redefining what processes can be reimaged and automated, enabling organizations to offload a wider variety of labor-intensive tasks. Knowledge workers' time is freed up to pursue higher-value strategic priorities. And like their human counterparts, agents can continuously improve their skills by leveraging advanced algorithms and machine learning to adapt, reinforce and learn. Ultimately, the seamless collaboration between digital and human workers will unlock significant organizational value in responding to changing market demands and opportunities in near-real time. Further, when there is agent-to-agent communication and coordination, execution can occur instantly, delivering immediate outcomes across the business.

2 24/7/365 availability

Because agents are designed to be independent and available outside normal working hours, they can automate more complex tasks and processes, while also expanding the number of daily work hours. They can allow organizations to monitor regulatory changes in real time, infer and extract mandatory obligations and compare and update compliance documentation in GRC systems across multiple jurisdictions.

When several agents are working simultaneously, they can also accomplish tasks much faster than a human team, while also finding ways to collaborate across workstreams. In manufacturing, for example, agents can manage production lines by optimizing schedules and workflows. They can perform predictive maintenance by analyzing equipment data to predict failures, and schedule timely repairs that reduce downtime significantly. When supply disruptions occur, procurement agents can find alternatives and production agents can reconfigure schedules automatically.

3 Building organizational resilience through continuous optimization

Agents automate routine workflow tasks while allowing human talent to focus on high-impact decisions where their judgment adds the greatest value. In healthcare, for example, pre-authorization agents evaluate medical-necessity criteria against patient records and payer requirements, thereby accelerating approvals, ensuring regulatory compliance and reducing the administrative burden. Further, when multiple specialized agents work together as a coordinated swarm, they fundamentally transform business processes — creating seamless handoffs, making real-time adjustments, optimizing resources and making autonomous decisions within the required guardrails.

This orchestrated intelligence personalizes customer interactions, ensures continuous compliance and drives adaptive workflows. In financial services, for example, AI agents streamline month-end close processes by identifying anomalies and preparing journal entries, significantly reducing the close cycle time. Simultaneously, vouching agents automatically cross-reference invoices with contracts, purchase orders and receipts, flagging discrepancies without the need for manual review.

4 Turning knowledge into action

Although Gen AI can significantly enhance productivity, the potential value from agentic AI can be much greater due to its ability to convert knowledge into coordinated action. As organizations deploy multiple agents, they need to feed them the right knowledge to drive the most effective outcomes. This requires systematically capturing and structuring the tacit expertise that typically resides in employees' minds, making it available for the agents' use.

For example, in retail store performance audits, audit agents continuously monitor store operational data, customer feedback and sales metrics to flag compliance issues and performance anomalies across locations — while simultaneously generating corrective action recommendations for regional managers.

Case study: AI agent fully automates retail returns, driving faster customer service and scale

KPMG in Canada developed and deployed an AI-powered agent that fully automates the in-store returns process for a large, multinational retailer. Built using the Salesforce Agentforce platform capabilities, the solution delivers faster customer service while unlocking immediate operational value.

The AI agent eliminates the need for human intervention, reducing the return resolution time for customers from several days to just minutes. Beyond the immediate gains, it has provided a strategic roadmap to address the company's people, process and technology considerations for integrating AI agents into core customer service operations at scale. It also includes a full rollout plan for implementing retail return agents across its North American customer contact centers.

Staying ahead of the regulatory curve

Governments and industry bodies are racing to write laws about how AI agents can operate, especially around privacy, decision-making and fairness. Enterprise leaders need to watch these developments closely and prepare to adapt quickly. That means compliance readiness must start well before the rules are formalized to anticipate and prepare for the regulatory standards as they evolve.

While the AI regulations are in development, enforcement expectations are already taking shape. Waiting for them to be finalized is a risk in itself.

While the regulations are emerging, there is a window of opportunity for enterprise leaders to start preparing by using the standards and frameworks that are already in place.

Those organizations that start building responsible AI governance now will be better positioned to earn stakeholder trust, avoid future compliance shocks and scale AI safely and strategically when the regulations are enacted.

The **EU AI Act**, **NIST Risk Management Framework (RMF)** and **ISO 42001**, for example, are all practical foundations for helping to build responsible, scalable AI agents by:

- Aligning agentic strategies with the EU AI Act — even for organizations that operate outside the EU, the Act will be setting the tone for regulation globally.
- Building risk-aware systems by applying the NIST RMF to agent development, especially if the agents are expected to operate autonomously.
- Using ISO 42001 as a common baseline to set up organizational structures and controls.

Taking these early actions goes well beyond risk mitigation, with the potential for achieving a competitive advantage. Proactively building your AI strategy based on emerging regulatory standards can also help to reduce future exposure, demonstrate due diligence and avoid costly remediation later.

The spectrum of agents is evolving to meet critical needs


Determining where to use specific agent types can be complex. At KPMG, we've developed a structured classification system called the KPMG TACO framework to help organizations make sense of the evolving landscape and determine how best to leverage agentic AI.

Across categories, there is consistency in the core computational components such as foundation and reasoning models, including LLMs, knowledge representation systems, contextual processing, memory architectures and tool integration frameworks such as model context protocol (MCP).

However, there are key distinctions emerging in terms of the increasing complexity of goals that these agents can fulfill as well as the corresponding planning and orchestration capabilities they can execute.

This progressive framework provides organizations with a methodical approach to matching their business requirements with the appropriate level of agent sophistication. As the types and complexity of agents are evolving rapidly, organizations should also prepare for how agents will collaborate and communicate with each other. Industry leaders are already aligning around this eventuality.

Finally, having a centralized "AI factory" where new agentic capabilities can be conceived and developed will be critical, particularly for organizations that want to build their own customized agents instead of buying off-the-shelf products.



From experimentation to implementation — Choosing the right agent for the task

There are several different types of AI agents, and determining which types are best suited to the tasks can be complex. KPMG has developed a structured classification system called the KPMG TACO framework that helps organizations make sense of the evolving agentic AI landscape and how to leverage its full capabilities.

KPMG TACO framework overview

	Taskers	Automators	Collaborators	Orchestrators
Overall complexity	Low	Low-to-medium	Medium-to-high	High
Primary use	Execute singular goals with one or many tasks	Automate goals that span multiple systems and process areas	Human/AI partnerships, with AI acting as a teammate for problem solving	Multi-agent coordination across ecosystems in pursuit of complex goals
Planning capabilities	Simple: Planning with prompt-based execution, including chaining and gated logic	Basic: Several sub-goals, deterministic logic with decision points, and adaptive prompts	Intermediate: Adaptive and contextual planning with human collaboration	Sophisticated: Multi-agent coordination with contingency planning and resource optimization
Value propositions	Free up workforce for higher-value tasks. Efficiency, innovation, and new ways of working.	Coordinate multiple goals across systems. Reimagine workflows and processes. Cycle time reduction. Improved customer experience.	Enhance workforce creativity and innovation. Create unlimited human multiplier effect. Improved employee experience.	Drive new revenue streams and models. Unleash economic transformation.
Required knowledge and tools, including MCP needs	- Explicit goal instructions - Standardized API connectors for specific actions - MCP: Minimal complexity with direct tool invocation	- Complex, cross-application and procedural goal instructions w/deep knowledge across systems - Enterprise API connectors w/authentication handling - MCP: Moderate complexity w/sequential tool operations	- Availability of curated data and knowledge corpus with domain expertise - Goal expression w/constraints and preferences - Several advanced connectors w/error handling - MCP: Sophisticated/high complexity for handling tool interdependencies and state management	- Comprehensive knowledge and goal expression both outside of and within the enterprise - Knowledge of complex policies; complex connector ecosystem w/dynamic discovery - MCP: Advanced orchestration w/parallel tool execution and complex state reconciliation

Taskers

Taskers require guidance on *what* needs to be achieved, rather than *how* to achieve it. They execute well-defined singular goals that feature repetitive tasks and minimal complexity, such as reviewing invoices or screening for compliance risks. They rely on foundational data that can be used across multiple applications and business processes. The “human in the loop” supplies detailed instructions outlining the steps and the actions the agent must execute. Looking to the future, KPMG in Canada has developed a risk management agent that can continuously monitor and assess organizational and external data for risk mitigation in client organizations.

Example Tasker applications

- **Customer service:** Categorize incoming support tickets based on keywords and assign to the appropriate team.
- **Vendor due diligence:** Screen vendor names against global sanction lists and flag potential issues.
- **Invoice validation:** Extract and verify data against compliance rules.
- **Legal affairs:** Convert past legal rulings into a standardized format.
- **Administrative:** Review employee expense reports for policy compliance and supporting documentation.

Automators

Automators deconstruct more complex goals and manage tasks that span multi-system workflows that can comprise applications, systems and functional areas. They draw from “tacit knowledge,” which is contained in workers’ minds based on their past experience with the organization. Automators execute end-to-end processes, orchestrate related tasks, manage dependencies between tasks and dynamically engage tools, APIs or MCP servers, as needed. A prime value area for Automators is customer service. According to Gartner, by 2029 agentic AI will autonomously resolve 80 percent of common customer service issues without human intervention.⁹

Example Automator applications

- **Supply chain:** Streamline workflows by integrating order and transportation management systems, as well as supplier APIs.
- **Healthcare:** Automate claim adjudication by integrating hospital billing systems, payer systems and regulatory databases.
- **Finance:** Accelerate month-end close by managing accruals, reconciliations and journal entries across multiple systems.
- **E2E order-to-cash:** Automate invoice generation, payment reconciliation and reminders.
- **Customer onboarding:** Adapt dynamically to missing info or compliance checks.

⁹ Gartner, “Gartner Predicts Agentic AI Will Autonomously Resolve 80% of Common Customer Service Issues Without Human Intervention by 2029,” 5 March 2025

Collaborators

Collaborators function as adaptive AI teammates, working interactively with humans to achieve multi-dimensional goals, while responding to user feedback and evolving contexts. These agents integrate multiple AI capabilities (machine learning, NLP, speech, computer vision) and can work with both structured and unstructured data to provide expertise-driven support and augmentation. Unlike simpler agents, collaborators maintain contextual awareness and long-term memory across extended complex interactions, enabling them to function as true digital colleagues rather than passive tools.

One example is Ema, a California-based AI agent startup offering agentic AI employees and autonomous agents in areas such as customer support and sales. Another example is Sierra AI's Customer Service Resolution Agent, which accesses multiple data sources to resolve complex issues.

Example Collaborator applications

- **Product development:** Synthesize user feedback, suggesting feature prioritization and generating release plans with timelines and resource estimates.
- **Marketing:** Brainstorm campaign ideas with marketing teams, refine ad copy based on style preferences and automatically schedule social media posts.
- **Financial planning and analysis:** Work alongside analysts to develop scenario models, evaluate assumptions and provide insights based on historical data patterns, while adapting to the specific analytical approach preferred by each team member.

Orchestrators

Orchestrators are an advanced class of agentic system designed to coordinate multiple agents, tools and interdependent workflows at scale. These agents operate as intelligent control towers, dynamically selecting from a broad ecosystem of specialized agents and services to achieve complex, cross-functional objectives. They rely on explicit knowledge, which is more objective, quantifiable, and often technical. Operating in real time, orchestrators enable dynamic resource optimization, multi-agent choreography, task delegation, inter-agent communication (leveraging protocols such as the agent-to-agent (A2A) protocol) and the reconciliation of dependencies across systems. Preparing for their adoption requires addressing the technology stack and codifying robust Trusted AI protocols.

As the most sophisticated of current AI agents, orchestrators have the potential to truly unlock new business models, foster cross-industry ecosystems and stimulate economic transformation at scale.

Example Orchestrator applications

- **Finance:** Manage inter-company reconciliation processes by orchestrating agents for currency conversion, account matching and exception reporting.
- **Regulatory compliance:** Orchestrate multi-jurisdictional compliance monitoring by assigning agents to review local regulations, ensure adherence and escalate violations for manual review.
- **Multi-entity procurement optimization:** Orchestrate sourcing, negotiation and compliance agents across global subsidiaries — within and outside of the enterprise.

AI case studies

Agentic AI “taskers” power up customer service and unlock engineering capacity at a leading Australia energy company

A major energy provider implemented a pioneering agentic AI solution developed by KPMG Australia to streamline customers’ connections to the energy grid. Previously, engineers manually handled applications for services such as solar panel or EV charger installations — time-consuming tasks that diverted their focus from higher-impact initiatives.

Now, autonomous AI “taskers” are processing these applications in the background, guided by human-in-the-loop oversight and trusted AI principles. This has accelerated approvals and issue resolution, improving customer satisfaction and turnaround times. Critically, it has also unlocked engineering capacity, allowing specialist teams to focus on complex, high-value projects such as building new data centers and industrial facilities.

Online trading company streamlines procure-to-pay process with a multi-agent “Automator”

An online trading company turned to KPMG in the US and agentic-AI platform company Ema to create an automated procure-to-pay system. Ema’s product is a multi-agent system that can connect to enterprise tools to fulfill complex, multi-step tasks using both structured and unstructured data. In this case, KPMG and Ema created a streamlined accrual solution for their joint client in which Ema pulled Coupa data that included invoice purchase orders and goods received, identifying anomalies that could then be escalated to a human for reconciliation. The result was the company will now be able to complete an accrual process in a few hours, instead of the 10 days of effort historically required by three full-time employees.

A leading telecom uses AI to navigate complex global regulations with speed and precision

As regulatory demands grow in complexity and volume, organizations are under pressure to respond faster and more effectively. Faced with thousands of regulatory obligations, risks and controls across 20+ countries, a major telecom provider turned to KPMG Australia to address this challenge. Partnering with Microsoft, KPMG Australia developed a powerful AI-powered compliance solution that leveraged KPMG’s proven frameworks and extensive risk, controls and compliance experience.

The AI agent is able to transform dense regulatory text into clear, actionable insights, enabling the organization’s teams to quickly interpret new regulations, map them to internal controls and assess the controls’ effectiveness — reducing a process that once took months to just days.

Government agency clears a five-month payments backlog with agentic AI document-processing

A large government entity responsible for disbursing wage-subsidy payments was drowning in a five-month backlog of ~25,000 reimbursement claims. Each was accompanied by complex, variably-formatted documents that required rule-based checks, calculations and human judgment. KPMG in Bahrain deployed an agentic AI solution on the agency’s existing OutSystems platform, leveraging Azure OpenAI to orchestrate specialized document-understanding agents that extract data, apply numerous business rules and escalate only true exceptions for review. Within four weeks, the agents validated 95 percent of the outstanding claims, cut average processing time by five times and restored timely cash-flow support to more than 40,000 enterprises nationwide, all while embedding data protection controls for sustainable scale.

Different agent types — the same core components

Across the agent types, the core computational components remain consistent, including foundation and reasoning models like LLMs, knowledge representation systems, contextual processing engines, memory architectures and tool integration frameworks such as the Model Context Protocol (MCP). (See sidebar on the right)

However, key distinctions emerge in terms of the increasing complexity of the goals that these agents can fulfill and the corresponding planning and orchestration capabilities they can execute.

This progressive framework provides organizations with a methodical approach to matching their business requirements with the appropriate level of agent sophistication. As agent capabilities advance, organizations should prepare for complex interactions between agents, including autonomous coordination and communication — an emerging frontier already drawing attention from industry leaders.

Finally, having a centralized “AI factory” where new agentic capabilities can be conceived and developed will be critical, particularly for organizations that want to build their own customized agents instead of buying off-the-shelf products. (See sidebar on page 15)

Common standards for cross-agent communication and data access

Currently, what an AI model “knows” is fed into the system during an initial process called pre-training, which typically occurs once or during periodic updates. In these cases, the models are limited to data available up to a cutoff date. Their functionality typically comprises only user prompts, a running conversation history and limited external information introduced during interactions.

As agents develop rapidly, parallel efforts are enabling them to communicate across organizational boundaries while continuously accessing external information sources, tools and services. Though Anthropic’s Model Context Protocol (MCP) is a work in progress, it has established standards and a platform that removes the need for custom integrations, paving the way for more capable and context-aware AI applications.

Core components of MCP include:



MCP client

An AI app or tool that allows MCP servers to access external data and functionality via a standardized protocol.



MCP server

Programs that use MCP protocols to leverage specific capabilities from data sources, systems of records and agents.



MCP transport

A communication layer that handles the exchange of messages between MCP clients and MCP servers.

Ultimately, this protocol is designed to allow organizations to deploy smaller, more targeted and efficient AI systems that interact relatively seamlessly with other agents and with external data — eliminating the need for tailored fine-tuning.


Source: MCP: The new “USB-C for AI” that’s bringing fierce rivals together, Ars Technica, 1 April 2025

KPMG Workbench: A federated work platform for enterprise AI, including agentic systems

KPMG Workbench is a federated platform to discover, build, evaluate, manage and operate a growing ecosystem of AI-enabled apps, knowledge assistants and intelligent agents. It's designed to work alongside KPMG professionals and client delivery teams in building AI apps, knowledge assistants and agents that encapsulate firmwide knowledge, operate securely within country-specific regulatory and data boundaries and integrate securely into diverse client environments. Components, including agents built on the Workbench platform, are modular, reusable and instantly deployable with the ability to adapt to specific client contexts while also maintaining consistency and governance at scale.

Co-developed with Microsoft, KPMG Workbench supports global replication and deployment, allowing any certified assistant or agent to be made available across the world. These agents can be customized for each organization's needs. By embedding knowledge, automating complex workflows and enabling human-agent collaboration, Workbench accelerates how KPMG professionals deliver insights and services that are dynamic, responsive and future-ready.

Navigating the agentic AI obstacle course



Agentic AI appears to be advancing rapidly — potentially outpacing the disruption caused by generative AI. While still in its early stages, its impact may be felt faster than many expect. Companies that move first stand to gain a significant edge, while others may find themselves struggling to catch up. The gap between the initial hype and real-world performance is often underestimated and may be shorter than imagined within 12 to 36 months.

To improve their odds, organizations should act now to begin the journey, envision where agentic AI may be headed and prepare to overcome the four key barriers we see today.



In the words of renowned physicist Niels Bohr, “Prediction is very difficult, especially if it’s about the future.” Agentic AI will likely force and enable organizations to be more strategically agile. It will be necessary to review and challenge corporate strategies more frequently, especially those that may have been formed years ago and reviewed only annually.

In many industries, barriers to entry are falling as labor becomes digital and knowledge becomes free. For example, new digital moats¹⁰ are appearing with access to large, high-quality datasets and affordable power. To compete in a market reshaped by AI, many will need to rethink their entire businesses from the ground up — including the customers they serve, the value they provide and how they compete.

Review your enterprise strategy.

Instead of the annual off-site strategy meeting, a more frequent review cycle will likely be needed to cope with the level of agentic AI uncertainty and speed of change. Use scenario building to inform your choices and include at least three components in the approach: quantify the value opportunity at stake (for the current state business); assess the risk of others disrupting you; and consider opportunities to disrupt others. Companies may start with an analysis of opportunities at their own organization but should also consider the impact of all their customers, suppliers and competitors who may also be adopting agentic AI.

Shape your agent strategy.

In conjunction with the broader enterprise strategy, and in parallel with the Gen AI adoption journey, companies need to determine the level of urgency to start building an agentic workforce. Decide whether you want to be an early adopter, fast follower or active monitor. Then, pinpoint the highest-value areas for initial deployments within the organization — by function, workflow or types of tasks. These analyses require quantifying the incremental value that agents can unlock, assessing the incremental risks and understanding agent readiness from both a technology and people perspective.

Evolve your partner ecosystem.

As agents move AI into the core operations of the enterprise, reliance on technology providers may increase. Your choice of partners, and your importance to them, could materially impact your speed to market.¹¹ Partnerships with leading AI-centric companies across both Big Tech and startups can be critical — and they’re already in high demand. Don’t overlook new and smaller potential partners that can bring a competitive edge within a specific function, industry or area of use.

¹⁰ A “digital moat” is a key competitive advantage that sets a company apart from its competitors.

¹¹ [Accelerate growth and innovation with partner ecosystems, KPMG in the US, 2025.](#)



Workforce

By taking the time now to envision the impact of agentic AI on the organization, companies can better manage the necessary shifts in roles, skills and workforce structure. This foresight will help organizations cope with the magnitude of change and increase the chances of gaining a competitive edge. (See sidebar on page 22)

Codify the ways you work.

Human activities can be divided into structured work, which follows standardized processes with clear rules, and tacit work that involves applying judgment, creativity and experience. Structured work offers near-term potential for agents and can be aided by existing software tools for process mining and documentation. Tacit work is harder to automate, but agentic AI may have a larger impact over time, with several start-ups emerging to capture tacit work. The more work that can be codified now, the faster and greater the impact of agentic AI will be. In a recent KPMG AI Pulse Survey, 78 percent of respondents said they would use agentic AI to analyze complex data sets requiring tacit work, while 66 percent plan to employ agents for routine administrative tasks.¹² Establishing systematic methods to capture and codify subject matter experts' tacit knowledge to enable agents' effectiveness will likely become operational necessities and will likely lead to competitive advantage.

Change the way you change.

For many companies, adoption of AI in the flow of work is proving more challenging than anticipated. While agentic AI relies less on humans changing their day-to-day behavior, similar challenges will eventually apply to the initial transfer of work and interactions between digital and human co-workers. Traditional top-down change management approaches, relying on upskilling, training, performance metrics and accountability, have not yet proven effective for AI. Companies should instead explore behavioral change approaches that feature role modeling, peer-to-peer learning and psychological safety as key ingredients in a cocktail of simultaneous interventions designed for their unique cultures.

Imagine a fluid hybrid organization.

As agents join the digital workforce, they will become part of the organizational structure alongside humans. This will require new ways of reporting, goal setting and performance management for agents and for humans. Adding to the challenge, the number of digital workers in the future could conceivably change by hundreds or thousands within seconds. For humans, their roles, reporting lines and performance metrics may also need to change significantly and more frequently. An apt analogy would be conducting a major merger integration every year. Companies that currently rely on offshoring may soon be able to reshore operations (at a fraction of the cost) by fully automating workflows with agentic AI. In some cases, it may even be feasible to launch entirely new business units or startups with minimal or no human involvement from the start.

¹² KPMG AI Pulse Survey Q1, KPMG in the US, April 2025.



Governance/trust

AI agents inherently require more stringent controls and an even greater commitment to trusted AI principles since they operate with relative autonomy. Errors that do occur can also be amplified at a greater scale with AI agents that may propagate and repeat unfavorable actions before human intervention.

Many enterprises are already recognizing the need to proactively measure how AI agents are performing — not just whether they work, but whether they work safely, fairly and reliably. Without clear metrics and oversight, agentic AI introduces invisible risks and potential missed opportunities in terms of efficiency, reliability, alignment with human intent and organizational policies as well as ethical and security risks.

Without this more robust governance, organizations will be unable to realize the full value from agentic AI and will not be able to safely move as fast as competitors that do.

Elevate security and privacy.

In addition to standard security considerations for AI, such as transparency, explainability, data privacy and compliance, organizations should conduct regular stress testing, bias detection and fail-safe mechanisms to prevent unintended consequences from agentic AI. Imagine the planner within an agent coming up with inconsistent plans every time it executes or, even worse, the wrong tools being used for taking action. Human errors are typically contained incidents, but a single error in agentic-AI programming or decision-making logic can lead to widespread consequences, potentially affecting numerous simultaneous outcomes before the issue is even detected.

Avoid potential ethical violations.

Throughout any AI implementation, but particularly those involving agentic AI, ethical considerations must be top of mind. Ethics shouldn't be treated as an afterthought or added at the end; they must be embedded throughout the entire development lifecycle. Establish clear ethical guidelines and governance protocols from the outset to proactively address risks around bias, fairness, and accountability.

Put humans in and on the loop.

While generative AI benefits from having a human in the loop between the advice generated by AI and the action taken, this is not the case for the more autonomous agents. KPMG's philosophy is that agents are going significantly augment human beings so humans can focus on higher-value work. We are not talking about agents replacing humans. From an ethical perspective, it will be important make sure that these agents are working under human guidance with 'humans in the loop' and 'on the loop'.

That kind of oversight is essential to ensure that agents don't do unnecessary things. It also helps to ensure they adhere to regulatory compliance and control risks.

Technology, data and security

With the extensive use of data by AI agents on the horizon, companies need to ensure that their proprietary data is accessible and of sufficient quality for agents to ingest. As the use of agents expands, organizations must also ensure that their agents can interact with each other effectively and have digital identities to manage them as co-workers within company HR systems.

Start building your agentic supply chain.

To source agents, organizations currently have three primary options: build a tailored custom solution, buy a pre-built agent, or partner with an external vendor. Each approach has its own benefits and considerations. The choice should align with specific organizational needs, resources and strategic goals, as well as the complexity of the agents required. Many will find it advantageous to create a multi-faceted agent supply chain with a mix of more than one of these options that can be applied across different parts of the organization over time. [See sidebar on page 24.]

A crucial question that's emerging, however, is "who owns the data"? Agents will soon be doing more work on behalf of users, including orchestrating the goals of users across SaaS applications. It may be the natural way businesses will want to interact with SaaS applications.

Ensure a robust data foundation.

Organizations sit on a wealth of proprietary data that they are not able to access effectively today. To succeed in the agentic era, a modern, cloud-enabled data platform and robust data foundation will be required, with high-quality data that is accurate, accessible, unbiased and enriched with context. Establish robust data governance protocols to ensure that data is collected, stored and used ethically and securely and in compliance with regulatory requirements and data privacy standards. The MCP plays a critical role in enabling agents to dynamically access and inject relevant data into their reasoning processes. Finally, organizations should enforce enterprise-grade data governance to secure data integrity, enable safe access, protect sensitive information through DLP controls, mitigate bias, ensure data quality and maintain compliance with privacy regulations such as GDPR, HIPAA and CCPA.

Enable consistent agent-to-agent interoperability.

Agent-to-agent (A2A) communication must move beyond custom integrations or brittle handoffs. We recommend adopting the emerging protocol developed by Google, which provides a standardized framework for secure, explainable and auditable agent interactions. A2A supports message passing, intent negotiation and structured collaboration between agents — enabling use cases such as multi-agent task orchestration, agent chaining and dynamic role assignment. Adopting standards for agent interoperability ensures that enterprise agents remain modular, composable and future-proof, while avoiding vendor lock-in and duplicated effort.

Strengthen agent identity and security.

Each agent must have a unique identity, with scoped permissions, credentials and continuous oversight. Equally essential is runtime isolation, ensuring agents operate in segmented environments where data exposure, tool misuse or anomalous behavior can be detected and contained. Every agent action must be auditable, revocable and attributable. Organizations must also be proactive in mitigating emerging security threats, such as memory poisoning, agent hijacking and identity spoofing.¹³ Remedies include implementing execution logging, identity attestation, runtime containment and real-time monitoring to ensure agent behavior remains secure, explainable and compliant.

Build in a “kill switch”

A kill switch is essential for agentic AI risk governance. It provides organizations with the ability to halt an agent’s actions immediately if it behaves unexpectedly, fails to comply with an organizational policy or introduces potential harm.

In fast-moving systems, real-time control is a business imperative. If there is a fundamental threat, it means “bringing the agent down” temporarily and finding alternative mechanisms to handle the work. It requires continuous monitoring of the agents, the ability to evaluate them, quantify the evaluations and identify scenarios where it may be necessary to “stop the elevator” and activate backup plans.

Remedies include implementing execution logging, identity attestation, runtime containment and real-time monitoring to ensure agent behavior remains secure, explainable and compliant — mechanisms best delivered through an agent control system designed for enterprise-scale governance.

¹³ Top 10 for LLM Applications and Agentic Systems, The Open Web Application Security Project (OWASP) Generative AI, LLM Security Guidance and Initiatives.

What about the workforce? Driving adoption and engagement

The adoption of Gen AI has been a gradual process, often hindered by workforce resistance, skepticism and a “swivel chair” approach, since workers have to pivot from their day-to-day workflow to use Gen AI. The introduction of agentic AI is changing this dynamic. By embedding AI agents into the unique workflows and tasks of different roles, organizations can realize much greater adoption, acceptance and utilization.

But even given agentic AI’s relative ease of use, some employees may be reticent to working side by side with AI agents. Nearly a quarter of organizations have reported that employee resistance to change and reluctance to use AI tools was a key challenge with AI.¹⁴ Therefore, it is necessary to establish a robust change management program that considers the impact of AI agents on team dynamics and morale, and on a culture of collaboration and innovation. Further, it is important to communicate that integration of AI agents into the workforce is not about replacing human jobs; it is designed to enhance the capabilities of the existing workforce.

As AI agents become embedded into the daily workflows of professionals and are combined with the actions they take, organizations will be positioned to realize greater workforce specialization and retention, higher trust in outcomes and enhanced innovation. Agents also promise to drive a 30 percent increase in workforce efficiency and a 25 percent reduction in operational costs by 2027, according to a Harvard Business Review report.¹⁵

Organizations that adopt both Gen AI and AI agents while concurrently dissecting and reconstructing workforce roles will be more likely to realize significant productivity and capacity gains. A recent study surveying more than 100,000 workers from 11 Gen AI-connected occupations found that AI tools can reduce working times by 50 percent for one-third of employees’ job tasks.¹⁶ According to a recent KPMG AI Pulse Survey, the functions that organizational leaders believe would benefit the most from agentic AI include information technology (76 percent), operations (74 percent), risk and compliance (56 percent), finance (39 percent) and marketing (35 percent).¹⁷

The introduction of agentic AI is changing this dynamic. By embedding AI agents into the unique workflows and tasks of different roles, organizations can realize much greater adoption, acceptance and utilization.

¹⁴ KPMG AI Pulse Survey Q1, KPMG in the US, April 2025

¹⁵ Harvard Business Review, “What is agentic AI, and How Will It Change Work?,” 12 December 2024

¹⁶ Ibid.

¹⁷ KPMG AI Pulse Survey Q1, KPMG in the US, April 2025.

These are some examples: of how humans and AI agents can work together:

- In finance, teams no longer need to consciously engage with a Gen AI system during month-end close processes. Instead, agents can automatically identify unusual transactions, prepare journal entries, and alert controllers when human judgment is needed.
- In the consumer and retail sector, AI agents can serve as virtual shopping assistants, providing hyper-personalized customer interactions. This allows human employees to focus on more strategic tasks, such as developing new products and services.
- In the life sciences, AI agents can collaborate with human researchers, leading to faster breakthrough treatments — particularly for rare diseases and hard-to-treat cancers — that might not be possible with human effort alone.

An enhanced workforce, empowered by agentic AI, can drive significant advancements and innovation, positioning organizations at the forefront of their industries. However, many employees continue to be concerned about the impact that AI may have. Will it take their job away? Can they be reassured that it will not — though it is possible that they could be replaced by the person who knows how to work and engage with AI?

It will be important for organizations to reinforce these key messages: *“Agents are here to augment what you do. They’re here to help you become more innovative in what you do. It isn’t just about productivity — it’s also about making you more innovative, helping to serve your customers better, helping you launch products and services in the market better and avoiding compliance issues. Overall, it’s going to make you and your work better as a whole.”*

In the consumer and retail sector, AI agents can serve as virtual shopping assistants, providing hyper-personalized customer interactions. This allows human employees to focus on more strategic tasks, such as developing new products and services.

How to acquire agents: Build, buy or partner?

As C-suite executives consider the most economically viable way to acquire AI agents, it’s important to consider the pros and cons of all the options.

One school of thought suggests that organizations should keep buying agents and not invest in building their own. However, if all agentic AI solutions look the same, what will differentiate your organization from all the others?

What is your strategy?

In reality, there is a likely to be a balance where an organization will buy some core, horizontal, essential agentic capabilities — such as a sales coach that will help in preparing for meetings. These solutions likely won’t provide your enterprise with a competitive advantage, though it may be possible to configure or customize them with additional data to make them uniquely fit for your organization and needs.

On the other hand, if you’re a bank, a healthcare institution or operating in other highly regulated industries you may need a unique agent that understands and is able to respond to highly specific customer enquiries and complaints, Not routing customer concerns quickly and to the right department has the potential to violate regulatory issues, jeopardize trust and damage the reputation of your organization.

In this scenario, an “agent out of the box” is not possible, especially if your organization is managing large numbers of customer contact types. The data that’s required and the way you handle them will be unique to your processes. And, in that case, you’re likely going to ‘build’ your own agents to reflect your unique value proposition.

Here are some of the considerations for helping to choose the best option for your organization.

Build vs. buy — Considerations

Core	Skills and expertise	TCO	Cost implications
• IP differentiation • Customization flexibility • Infra requirements	• Technical proficiency • Delivery velocity • Upkeep day-2+ • Keeping up with advancements	• Upfront vs. hidden costs • Getting full visibility	• Subscription fatigue • Hidden costs • Adoption waste
UX	Data privacy	Vendor	Risk and compliance
• Change and adoption • Brand consistency	• Data privacy and protection • Indemnifications • Used for training? • Usage rights and data lock-in	• Stability and longevity • Lock-in • Support and partnership	• IP protection • Model transparency • Contract term changes

As C-suite executives, and particularly the CFO, consider the most economically viable way to acquire AI agents, they will want to consider the pros and cons of each option. Here's a detailed look at each one.

Build: Customization and control

Building agentic AI agents in-house allows organizations to tailor agents to their specific needs and integrate them seamlessly with existing systems and processes. The organization will also own the intellectual property output by the agent. On the other hand, due to the complexity of this process, organizations will need an internal and/or external team with deep expertise in AI, data science and software development to guide the process. For example, a large financial services firm decided to build its own AI tasker to automate invoice processing. By developing the agent in-house, it was customized to handle specific financial regulations and internal processes relevant to their organization.

Buy: Rapid execution and efficiency

According to 67 percent of respondents to a recent KPMG AI Pulse Survey¹⁸, buying existing agentic-AI capabilities is the quickest and most straightforward option and is the preferred method of acquiring agents.

On the one hand, this option is the most expeditious, as it involves acquiring pre-built AI agents from established vendors, which can be deployed immediately to address specific business needs. It can also be more cost-effective than building from scratch and has the added benefit of offering access to well-tested and proven AI solutions.

However, this option is not without its challenges: Off-the-shelf AI agents have limited customization options and, over time, they may become obsolete for certain customers. There are also potential challenges in integrating these tools with existing technology stacks, as well as the potential that the agent won't have data security capabilities that are on par with the organization's own standards.

Partner: Access to advanced expertise

Partnering with external vendors combines the benefits of both building and buying. This approach involves collaborating with AI experts to develop and implement AI agents, leveraging their expertise and infrastructure, while sharing risk and cost. Challenges include less control over the development and customization of AI agents compared to building in-house, as well as the potential for higher costs.

As with any technology partnership, third-party risks need to be top of mind and built into the service level agreement. Protections need to be in place to protect the company if a partner pivots on the development plan originally agreed upon, if a system fails or if incorrect data is fed into the system. The latter is particularly critical as plugging highly sensitive information into an AI agent without guardrails in place could lead to a highly damaging data breach.

Finally, in evaluating the avenue to take, these are some key questions to consider: Does the organization have the in-house expertise and resources to build AI agents? Are pre-built solutions sufficient, or is there a need for highly tailored agents? Do agents need to be deployed quickly? What types of agents align best with long-term strategic goals? Is there an operating model in place to upkeep and maintain these agents?

¹⁸ KPMG AI Pulse Survey Q1, April 2025.

How to accelerate your agentic-AI journey

At the time of this paper's writing, many companies in the US have taken the first steps on their agentic journeys by starting to explore or pilot agents. However, the majority (88 percent) have not yet actively deployed any agents in their operations. Regardless of where you are on your agentic-AI journey, what matters most is how rapidly, effectively and securely you are able to move forward.

To be successful, companies need a clear vision for how to incorporate agents into their overall AI transformation journey, including a solid rationale for where and when to start and an approach for how to scale across key functions over time. Most companies would do well to accelerate their data transformation efforts and evolve their Trusted AI governance practices, given higher agentic risks.

Talent leadership will be busy preparing the organization for the arrival of digital co-workers, such as putting new performance metrics in place for both digital and human workers.

To accelerate an agentic-AI program, leaders should consider these five next steps:

1. Articulate the vision

Start with a clear vision and strategy for agentic-AI integration. The vision should articulate how AI agents will transform the business, desired outcomes, key business goals to pursue and how to start supplying AI agents. Identify areas where agents can address pain points and add the most value, such as routine tasks, complex workflows or customer interactions.

2. Start agentic pilots

Executives starting an agentic-AI journey should consider where in the organization they can provide the most value and, therefore, how they should approach pilot projects. Organizations venturing into AI-enhanced operations can attempt one or more of the following three approaches:

- **Focus on "hot spots":** Target areas where the organization has previously conducted successful Gen AI pilots, especially those with high user engagement. Those will likely have the highest potential for value creation, as well as the most mature AI infrastructure and skilled personnel.
- **Go deep into a function:** Embed AI agents extensively within a specific, highly repetitive, low-risk function, such as finance. Use this pilot as a scalable model to apply agentic AI to other functions across the organization.
- **Broad utilization:** Identify a critical cross-functional value stream that is central to business operations and implement AI agents throughout this end-to-end process.



3. Prepare to scale agents within key functions and/or workflows

Consider whether there are a few AI agent platforms that can be piloted quickly to achieve early wins that can help to build executive buy-in and momentum. Use tools like the KPMG TACO framework to identify the types of AI agents that can best address your needs and determine which departments and processes are most appropriate for agentic-AI applications. Eventually, develop a plan to scale to an enterprise-wide strategy with multiple agent types working in tandem.

4. Evolve your trusted AI governance and playbook for agents

It is critical that organizations expand their existing AI governance programs to account for agentic AI with mechanisms in place to stay ahead of legal and regulatory considerations. Consider creating a living catalog of all AI agents to track their purposes, dependencies and performance metrics for effective scaling and governance. Develop robust operational processes to manage and maintain AI agents, including monitoring systems, performance metrics and continuous improvement methods.

5. Implement Trusted AI evaluations specific to AI agents

As organizations align their AI governance practices with increasing regulatory conventions and leading standards, a rigorous evaluation and reporting approach for AI agents will be essential. AI system cards document how technology is used in an organization, acting as single, authoritative sources for the results of AI system evaluations that detail intended use, data considerations, AI components and limitations. Further, before launching an agentic-AI system into production, conduct thorough testing, which increasingly includes “purple teams” that run rigorous offensive testing and defensive evaluation. Inputs from these processes ultimately inform a trust score, which reflects an AI system’s adherence to the organization’s AI principles.

6. Establish agentic talent performance metrics

Put metrics in place to evaluate the impact of pilot projects, as well as mechanisms for gathering feedback and learnings from stakeholders during the pilot phase. Continuously measure and optimize the performance of AI agents. Develop clear performance metrics to evaluate their effectiveness, including error rates, processing times and customer satisfaction. Use data and feedback to continuously improve performance, including refining algorithms, updating training data and implementing new features. It should be noted that AI agents implement telemetry to monitor their own SLA compliance, system health, tool performance and anomalies.

“While new types of AI agents are being introduced at a rapid pace, it’s important to remember that this is only the beginning of the beginning. By laying the groundwork today for the advent of the age of agentic AI, organizations will be well-positioned to revolutionize not only their operations, but also their entire businesses.”

Swami Chandrasekaran

Global Head of AI & Data Labs
KPMG International
Principal, KPMG in the US

How KPMG can help

Accelerating responsible value with AI

AI transformation is not just about deploying new technology. By optimizing the collaboration between what humans and agents can do today, you can unlock the full potential of what will be possible tomorrow and help create lasting value together.

The KPMG AI transformation journey is a structured, end-to-end approach that helps organizations unlock value from AI while managing risks and enabling trust. It's designed to guide enterprises through every stage of AI transformation — from opportunity identification, to strategic alignment and scalable implementation — while embedding trusted AI practices throughout.

Turn AI into your competitive edge — from strategy to execution, with a people-centric approach grounded in trust. KPMG offers a portfolio of services tailored to help you overcome key barriers to agentic AI adoption, with trust and responsible use at every stage of the journey.

KPMG AI jumpstart — Starting the AI journey

Accelerate your path to value from proof-of-concept to launch to adoption at scale

- Identify agentic AI use cases and accelerate the path to value with a replicable proof-of-concept factory approach.
- Safely experiment with AI agents and scale adoption of tested use cases across your organization.

KPMG AI strategy — Creating a path to value

Help shape your agentic AI strategy and business case with an actionable roadmap.

- Shape your agentic AI strategy and execution plan tailored to your unique starting point and business needs.
- Build a credible business case with quantified returns/metrics to help drive investments and release product funding.

KPMG AI Trust — Enabling trust

AI Trust to help you design, build and deploy agentic AI solutions in a responsible and ethical way through:

- Trusted and robust AI governance frameworks, operating models, policies and procedures.
- Responsible AI security and privacy strategies, processes and tools to help prevent, detect, respond to, and recover from cyber intrusions.

KPMG AI Workforce — Aligning people and change

Unlock the full potential of agentic AI for, and with, your people during the journey

- Augment your workforce with AI agents, enabling more strategic work and accelerating efficiencies.
- Reshape your workforce and define agentic AI governance to help realize the full value of your investments.

KPMG AI Technology — Building data foundations for AI

High-quality data, advanced technology and trust come together to help to ensure that AI and other emerging technologies can deliver actionable, accurate and reliable outcomes that create real value.

- Design and develop scalable, flexible cloud-native architecture to support agentic AI models and applications.
- Conduct AI model evaluations and selection using a robust framework to assess the performance accuracy of large language models (LLMs) on various tasks.

Find out more: kpmg.com/AIservices

Authors



Steve Chase

Global Head of AI and Digital Innovation
KPMG International
schase@kpmg.com

Stephen (Steve) Chase was appointed KPMG International's Global Head of AI and Digital Innovation on August 1, 2025, leading AI-enabled client solutions and internal transformation. He also serves as Vice Chair of AI and Digital Innovation and sits on the US Management Committee. With 20+ years in consulting and executive leadership, he previously led KPMG US Consulting and as KPMG US Vice Chair of AI and Digital Innovation, he scaled AI adoption and alliances with Google and Microsoft. He holds MBA and MPA degrees from UT Austin.



Swami Chandrasekaran

Global Head of AI & Data Labs
KPMG International
Principal, KPMG in the US
swamchan@kpmg.com

Swami leads KPMG's enterprise-wide AI strategy across Tax, Audit, Advisory, and other functions, supporting 200,000 professionals. He oversees R&D in AI architecture, advanced assistants/agents, domain-tuned Small Language Models (SLMs), synthetic data, enterprise search, and hardware-optimized solutions, and chairs KPMG's AI Technology Review Board for trusted, scalable adoption. He has worked with 500+ clients in 20+ countries on large-scale digital and AI transformations. Swami holds 30+ patents, is an IBM Distinguished Engineer and Master Inventor, and earned a Master's in Electrical Engineering.



Per Edin

Global AI Sales Acceleration Leader for Advisory
KPMG in the US
pedin@kpmg.com

Per holds dual global US leadership roles for AI. He is responsible for the alignment of AI client offerings, messaging and go-to-market approach across the global advisory network and for driving AI-related growth in the US technology, media and telecom sector. His work has included leading the firm's efforts to help clients navigate the Gen AI market disruption, developing a portfolio of new AI services, AI thought leadership and key accelerators like the KPMG AI value assessment.



The views and opinions expressed herein are those of the survey respondents and do not necessarily represent the views and opinions of KPMG.

TACO framework™ is owned by KPMG LLP (US) and trademarked within the US.

The information contained herein is of a general nature and is not intended to address the circumstances of any particular individual or entity. Although we endeavor to provide accurate and timely information, there can be no guarantee that such information is accurate as of the date it is received or that it will continue to be accurate in the future. No one should act on such information without appropriate professional advice after a thorough examination of the particular situation.

© 2025 Copyright owned by one or more of the KPMG International entities. KPMG International entities provide no services to clients. All rights reserved.

KPMG refers to the global organization or to one or more of the member firms of KPMG International Limited (“KPMG International”), each of which is a separate legal entity. KPMG International Limited is a private English company limited by guarantee and does not provide services to clients. For more details about our structure please visit kpmg.com/governance.

The KPMG name and logo are trademarks used under license by the independent member firms of the KPMG global organization.

Throughout this document, “we”, “KPMG”, “us” and “our” refers to the KPMG global organization, to KPMG International Limited (“KPMG International”), and/or to one or more of the member firms of KPMG International, each of which is a separate legal entity.

Designed by Evalueserve.

Publication name: Agentic AI advantage: Unlocking next-level value | Publication number: 140142-G | Publication date: October 2025