

```
! pip install --upgrade git+https://github.com/keras-team/keras-cv -q
Installing build dependencies ... ents to build wheel ... etadata
(pyproject.toml) ...
```

```
import os
from tqdm.auto import tqdm
import xml.etree.ElementTree as ET
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib.image as mpimg
import tensorflow as tf
from tensorflow import keras
from keras_cv.losses import FocalLoss
import pandas as pd
import keras_cv
from keras_cv import bounding_box
from keras_cv import visualization

%matplotlib inline
```

Hyperparameters

```
# Try these modifications

LEARNING_RATE = 0.001    # Lower learning rate
BATCH_SIZE = 4    # Increase batch size if memory allows
WEIGHT_DECAY = 1e-5    # Adjust weight decay
GLOBAL_CLIPNORM = 1.0
SPLIT_RATIO=0.2
```

```

class_ids = ['car', 'trafficLight-YellowLeft', 'pedestrian',
'trafficLight-GreenLeft', 'trafficLight-Green', 'biker',
'trafficLight-Yellow', 'truck',
            'trafficLight-RedLeft', 'trafficLight-Red',
'trafficLight']
class_mapping = dict(zip(range(len(class_ids)), class_ids))

```

Path to images and annotations

```

path_image="/kaggle/input/self-driving-car/export/"
annot_path="/kaggle/input/self-driving-car/export/"

# Get all XML file paths in path_annot and sort them
xml_files = sorted(
    [
        os.path.join(annot_path, file_name)
        for file_name in os.listdir(annot_path)
        if file_name.endswith(".xml")
    ]
)

# Get all JPEG image file paths in path_images and sort them
jpg_files = sorted(
    [
        os.path.join(path_image, file_name)
        for file_name in os.listdir(path_image)
        if file_name.endswith(".jpg")
    ]
)

class_names = set()

# Function to parse XML annotations and extract class names
def extract_class_names(xml_file):
    tree = ET.parse(xml_file)
    root = tree.getroot()

    # Loop through each object in the XML and extract the class name
    for obj in root.iter("object"):
        class_name = obj.find("name").text
        class_names.add(class_name)

# Loop through XML files and extract class names with tqdm for
progress
for xml_file in tqdm(xml_files, desc="Extracting class names",

```

```

unit="file"):
    extract_class_names(xml_file)

# Print unique class names
print("Unique class names:")
print(class_names)

{"model_id": "ca8690c719134e72ab36742ae8c21bdd", "version_major": 2, "version_minor": 0}

Unique class names:
{'truck', 'pedestrian', 'trafficLight-YellowLeft', 'trafficLight',
'trafficLight-GreenLeft', 'biker', 'trafficLight-Green',
'trafficLight-Yellow', 'car', 'trafficLight-Red', 'trafficLight-
RedLeft'}

rows, cols = 5, 5
fig, ax = plt.subplots(rows, cols, figsize=(15, 15))
axes = ax.flatten()

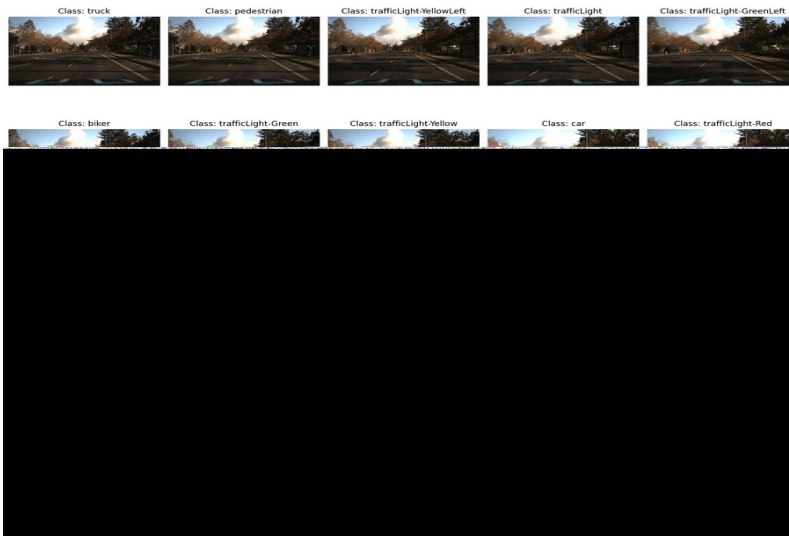
# Make sure the class_names list has enough elements for the available
images
num_images = min(len(jpg_files), rows * cols)
class_names_list = list(class_names)[:num_images]

for i, ax in enumerate(axes[:num_images]):
    img = mpimg.imread(jpg_files[i])
    ax.imshow(img)
    ax.set_title(f"Class: {class_names_list[i %
len(class_names_list)]}") # Cycle through class names if less than
images
    ax.axis('off')

# Hide any remaining axes
for j in range(num_images, len(axes)):
    axes[j].axis('off')

plt.tight_layout()
plt.show()

```



```

class_count = {class_name: 0 for class_name in class_names}

def count_classes_in_xml(xml_file):
    tree = ET.parse(xml_file)
    root = tree.getroot()

    for obj in root.iter("object"):
        class_name = obj.find("name").text
        if class_name in class_count:
            class_count[class_name] += 1

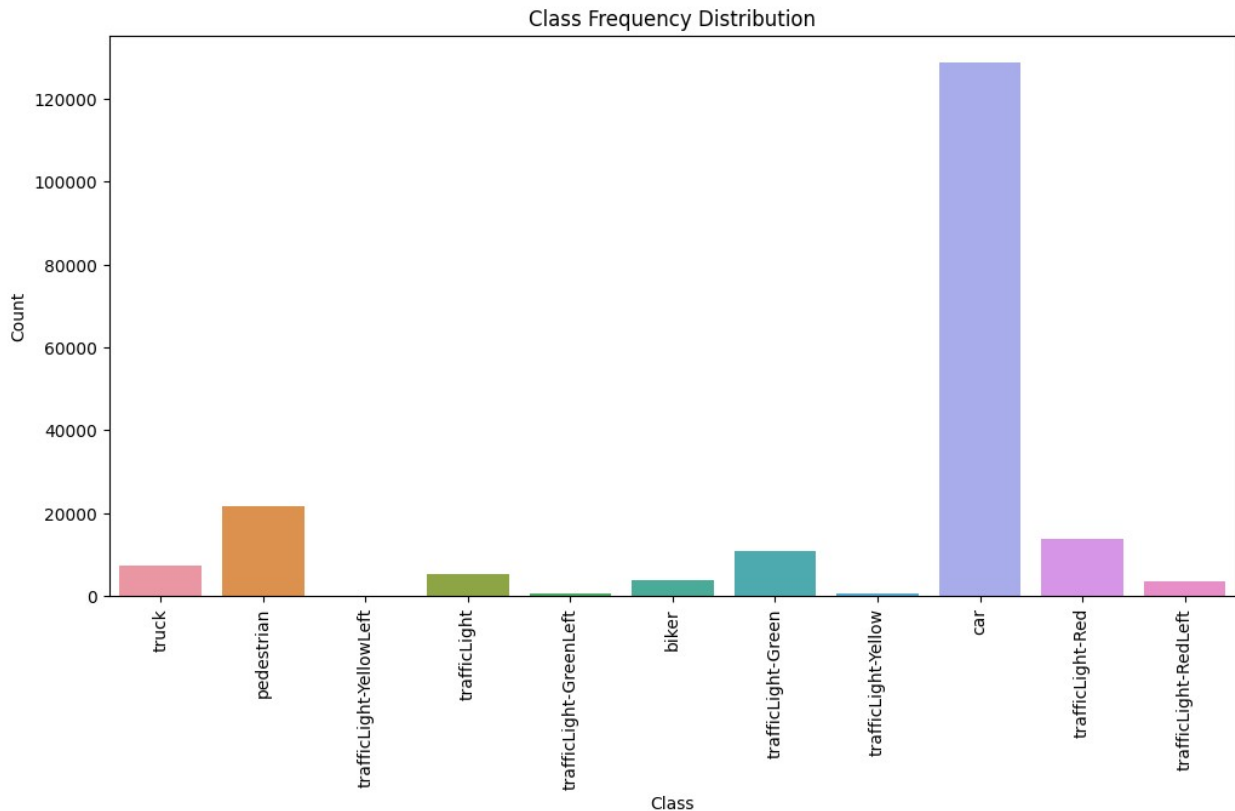
for xml_file in tqdm(xml_files, desc="Counting class occurrences",
unit="file"):
    count_classes_in_xml(xml_file)

class_count_df = pd.DataFrame(list(class_count.items()),
columns=["Class", "Count"])

plt.figure(figsize=(12, 6))
sns.barplot(x="Class", y="Count", data=class_count_df)
plt.xticks(rotation=90)
plt.title("Class Frequency Distribution")
plt.show()

{"model_id": "3aeace96ff254091ac317e5c5e26531c", "version_major": 2, "version_minor": 0}

```



- The function below reads the XML file and finds the image name and path, and then iterates over each object in the XML file to extract the bounding box coordinates and class labels for each object.
- The function returns three values: the image path, a list of bounding boxes (each represented as a list of four floats: xmin, ymin, xmax, ymax), and a list of class IDs (represented as integers) corresponding to each bounding box. The class IDs are obtained by mapping the class labels to integer values using a dictionary called `class_mapping`.

```
def parse_annotation(xml_file):
    tree = ET.parse(xml_file)
    root = tree.getroot()

    image_name = root.find("filename").text
    image_path = os.path.join(path_image, image_name)

    boxes = []
    classes = []
    for obj in root.iter("object"):
        cls = obj.find("name").text
        classes.append(cls)

        bbox = obj.find("bndbox")
```

```

        xmin = float(bbox.find("xmin").text)
        ymin = float(bbox.find("ymin").text)
        xmax = float(bbox.find("xmax").text)
        ymax = float(bbox.find("ymax").text)
        boxes.append([xmin, ymin, xmax, ymax])

    class_ids = [
        list(class_mapping.keys())
        [list(class_mapping.values()).index(cls)]
        for cls in classes
    ]
    return image_path, boxes, class_ids

image_paths = []
bbox = []
classes = []
for xml_file in tqdm(xml_files):
    image_path, boxes, class_ids = parse_annotation(xml_file)
    image_paths.append(image_path)
    bbox.append(boxes)
    classes.append(class_ids)

{"model_id": "8f489067b8db4680b04bd20fa2c79502", "version_major": 2, "version_minor": 0}

bbox = tf.ragged.constant(bbox)
classes = tf.ragged.constant(classes)
image_paths = tf.ragged.constant(image_paths)

data = tf.data.Dataset.from_tensor_slices((image_paths, classes,
bbox))

```

split data from train & valid

```

# Determine the number of validation samples
num_val = int(len(xml_files) * SPLIT_RATIO)

# Split the dataset into train and validation sets
val_data = data.take(num_val)
train_data = data.skip(num_val)

```

- The dictionary has two keys, 'boxes' and 'classes', each of which maps to a TensorFlow RaggedTensor or Tensor object. The 'boxes' Tensor has a shape of [batch, num_boxes, 4], where batch is the number of images in the batch and num_boxes is the maximum number of bounding boxes in any image. The 4 represents the four values needed to define a bounding box: xmin, ymin, xmax, ymax.

- The 'classes' Tensor has a shape of [batch, num_boxes], where each element represents the class label for the corresponding bounding box in the 'boxes' Tensor. The num_boxes dimension may be ragged, which means that the number of boxes may vary across images in the batch.

Final dict should be:

```
def load_image(image_path):
    image = tf.io.read_file(image_path)
    image = tf.image.decode_jpeg(image, channels=3)
    return image

def load_dataset(image_path, classes, bbox):
    # Read Image
    image = load_image(image_path)
    bounding_boxes = {
        "classes": tf.cast(classes, dtype=tf.float32),
        "boxes": bbox,
    }
    return {"images": tf.cast(image, tf.float32), "bounding_boxes":
bounding_boxes}

augmenter = keras.Sequential(
    layers=[
        keras_cv.layers.RandomFlip(mode="horizontal",
bounding_box_format="xyxy"),
        keras_cv.layers.RandomShear(
            x_factor=0.2, y_factor=0.2, bounding_box_format="xyxy"
        ),
        keras_cv.layers.JitteredResize(
            target_size=(640,640), scale_factor=(0.75, 1.3),
bounding_box_format="xyxy"
        ),
    ]
)
```

Creating Training Dataset





```
train_ds = train_data.map(load_dataset,
num_parallel_calls=tf.data.AUTOTUNE)
train_ds = train_ds.shuffle(BATCH_SIZE * 4)
train_ds = train_ds.ragged_batch(BATCH_SIZE, drop_remainder=True)
train_ds = train_ds.map(augmenter,
num_parallel_calls=tf.data.AUTOTUNE)
```

create valid dataset

```
resizing = keras_cv.layers.JitteredResize(  
    target_size=(640,640),  
    scale_factor=(0.75, 1.3),  
    bounding_box_format="xyxy")  
  
val_ds = val_data.map(load_dataset,  
    num_parallel_calls=tf.data.AUTOTUNE)  
val_ds = val_ds.shuffle(BATCH_SIZE * 4)  
val_ds = val_ds.ragged_batch(BATCH_SIZE, drop_remainder=True)  
val_ds = val_ds.map(resizing, num_parallel_calls=tf.data.AUTOTUNE)
```

Visualization

```
def visualize_dataset(inputs, value_range, rows, cols,  
    bounding_box_format):  
    inputs = next(iter(inputs.take(1)))  
    images, bounding_boxes = inputs["images"],  
    inputs["bounding_boxes"]  
    visualization.plot_bounding_box_gallery(  
        images,  
        value_range=value_range,  
        rows=rows,  
        cols=cols,  
        y_true=bounding_boxes,  
        scale=5,  
        font_scale=0.7,  
        bounding_box_format=bounding_box_format,  
        class_mapping=class_mapping,  
    )  
  
visualize_dataset(train_ds, bounding_box_format="xyxy",  
    value_range=(0, 255), rows=2, cols=2)
```

visualize from valid dataset

```
def visualize_dataset(inputs, value_range, rows, cols,
    bounding_box_format):
    inputs = next(iter(inputs.take(1)))
    images, bounding_boxes = inputs["images"],
    inputs["bounding_boxes"]
    visualization.plot_bounding_box_gallery(
        images,
        value_range=value_range,
        rows=rows,
```

```


    cols=cols,
    y_true=bounding_boxes,
    scale=5,
    font_scale=0.7,
    bounding_box_format=bounding_box_format,
    class_mapping=class_mapping,
)

```

```

visualize_dataset(val_ds, bounding_box_format="xyxy", value_range=(0,
255), rows=2, cols=2)

```



- We need to extract the inputs from the preprocessing dictionary and get them ready to be fed into the model.

```
def dict_to_tuple(inputs):
    return inputs["images"], bounding_box.to_dense(
        inputs["bounding_boxes"], max_boxes=32
    )

train_ds = train_ds.take(100)
train_ds = train_ds.map(dict_to_tuple,
    num_parallel_calls=tf.data.AUTOTUNE)
train_ds = train_ds.prefetch(tf.data.AUTOTUNE)

val_ds = val_ds.take(100)
val_ds = val_ds.map(dict_to_tuple,
    num_parallel_calls=tf.data.AUTOTUNE)
val_ds = val_ds.prefetch(tf.data.AUTOTUNE)
```

Creating Model

- Next, let's build a YOLOV8 model using the YOLOV8Detector, which accepts a feature extractor as the backbone argument, a num_classes argument that specifies the number of object classes to detect based on the size of the class_mapping list, a bounding_box_format argument that informs the model of the format of the bbox in the dataset, and a finally, the feature pyramid network (FPN) depth is specified by the fpn_depth argument.
- It is simple to build a YOLOV8 using any of the aforementioned backbones thanks to KerasCV

```
backbone =
keras_cv.models.YOLOV8Backbone.from_preset("yolo_v8_l_backbone_coco")

model =
keras_cv.models.YOLOV8Detector( num_classes=len(class_mapping), bounding_box_format="xyxy", backbone=backbone, fpn_depth=2,

prediction_decoder=keras_cv.layers.NonMaxSuppression(bounding_box_format="xyxy", from_logits=False,
    iou_threshold=0.5,
    confidence_threshold=0.1,
    max_detections=50))

optimizer = tf.keras.optimizers.AdamW(learning_rate=LEARNING_RATE,
weight_decay=WEIGHT_DECAY, global_clipnorm=GLOBAL_CLIPNORM)
```

```
# Reasonable starting point
```

```
lr_schedule =
```

```
tf.keras.optimizers.schedules.CosineDecay(initial_learning_rate=LEARNING_RATE,decay_steps=10000,alpha=0.01)
```

```
optimizer =
```

```
tf.keras.optimizers.AdamW(learning_rate=lr_schedule,weight_decay=WEIGHT_DECAY,global_clipnorm=GLOBAL_CLIPNORM)
```

```
model.compile(optimizer=optimizer,  
classification_loss=FocalLoss(alpha=0.25, gamma=2.0), box_loss="ciou")
```

```
model.summary()
```

```
Model: "yolov8_detector"
```

Layer (type) Connected to	Output Shape	Param #
input_layer_1 (InputLayer)	(None, None, None, 3)	0
functional (Functional) input_layer_1[0][0]	[(None, None, None, 256), (None, None, None, 512), (None, None, None, 512)]	19,831,744
repeat (Repeat) functional[0][2]	(None, None, None, 512)	0
repeat_1 (Repeat) repeat[0][0]	(None, None, None, 512)	0
concatenate_5 repeat_1[0][0],	(None, None, None, 512)	0

(Concatenate) functional[0][1]	1024)	
pa_fpn_p4p5_pre_conv concatenate_5[0][0] (Conv2D)	(None, None, None, 512)	524,288
pa_fpn_p4p5_pre_bn pa_fpn_p4p5_pre_conv[... (BatchNormalization)	(None, None, None, 512)	2,048
pa_fpn_p4p5_pre pa_fpn_p4p5_pre_bn[0]... (Activation)	(None, None, None, 512)	0
split_4 (Split) pa_fpn_p4p5_pre[0][0]	[(None, None, None, 256), (None, None, None, 256)]	0
pa_fpn_p4p5_pre_0_1_pad split_4[0][1] (ZeroPadding2D)	(None, None, None, 256)	0
pa_fpn_p4p5_pre_0_1_conv pa_fpn_p4p5_pre_0_1_p... (Conv2D)	(None, None, None, 256)	589,824
pa_fpn_p4p5_pre_0_1_bn pa_fpn_p4p5_pre_0_1_c... (BatchNormalization)	(None, None, None, 256)	1,024
pa_fpn_p4p5_pre_0_1	(None, None, None,	0

pa_fpn_p4p5_pre_0_1_bn (Activation)	256)	
pa_fpn_p4p5_pre_0_2_pad pa_fpn_p4p5_pre_0_1[0... (ZeroPadding2D)	(None, None, None, 256)	0
pa_fpn_p4p5_pre_0_2_conv pa_fpn_p4p5_pre_0_2_p... (Conv2D)	(None, None, None, 256)	589,824
pa_fpn_p4p5_pre_0_2_bn pa_fpn_p4p5_pre_0_2_c... (BatchNormalization)	(None, None, None, 256)	1,024
pa_fpn_p4p5_pre_0_2 pa_fpn_p4p5_pre_0_2_bn (Activation)	(None, None, None, 256)	0
pa_fpn_p4p5_pre_1_1_pad pa_fpn_p4p5_pre_0_2[0... (ZeroPadding2D)	(None, None, None, 256)	0
pa_fpn_p4p5_pre_1_1_conv pa_fpn_p4p5_pre_1_1_p... (Conv2D)	(None, None, None, 256)	589,824
pa_fpn_p4p5_pre_1_1_bn pa_fpn_p4p5_pre_1_1_c... (BatchNormalization)	(None, None, None, 256)	1,024
pa_fpn_p4p5_pre_1_1 pa_fpn_p4p5_pre_1_1_bn	(None, None, None,	0

(Activation)	256)	
pa_fpn_p4p5_pre_1_2_pad pa_fpn_p4p5_pre_1_1[0... (ZeroPadding2D)	(None, None, None, 256)	0
pa_fpn_p4p5_pre_1_2_conv pa_fpn_p4p5_pre_1_2_p... (Conv2D)	(None, None, None, 256)	589,824
pa_fpn_p4p5_pre_1_2_bn pa_fpn_p4p5_pre_1_2_c... (BatchNormalization)	(None, None, None, 256)	1,024
pa_fpn_p4p5_pre_1_2 pa_fpn_p4p5_pre_1_2_b... (Activation)	(None, None, None, 256)	0
concatenate_6 split_4[0][0], (Concatenate) split_4[0][1], pa_fpn_p4p5_pre_0_2[0... pa_fpn_p4p5_pre_1_2[0...	(None, None, None, 1024)	0
pa_fpn_p4p5_output_conv concatenate_6[0][0] (Conv2D)	(None, None, None, 512)	524,288
pa_fpn_p4p5_output_bn pa_fpn_p4p5_output_co... (BatchNormalization)	(None, None, None, 512)	2,048

pa_fpn_p4p5_output pa_fpn_p4p5_output_bn... (Activation)	(None, None, None, 512)	0
repeat_2 (Repeat) pa_fpn_p4p5_output[0]...	(None, None, None, 512)	0
repeat_3 (Repeat) repeat_2[0][0]	(None, None, None, 512)	0
concatenate_7 repeat_3[0][0], (Concatenate) functional[0][0]	(None, None, None, 768)	0
pa_fpn_p3p4p5_pre_conv concatenate_7[0][0] (Conv2D)	(None, None, None, 256)	196,608
pa_fpn_p3p4p5_pre_bn pa_fpn_p3p4p5_pre_con... (BatchNormalization)	(None, None, None, 256)	1,024
pa_fpn_p3p4p5_pre pa_fpn_p3p4p5_pre_bn[... (Activation)	(None, None, None, 256)	0
split_5 (Split) pa_fpn_p3p4p5_pre[0][...]	[(None, None, None, 128), (None, None, None, 128)]	0

pa_fpn_p3p4p5_pre_0_1_pad split_5[0][1] (ZeroPadding2D)	(None, None, None, 128)	0
pa_fpn_p3p4p5_pre_0_1_co... pa_fpn_p3p4p5_pre_0_1... (Conv2D)	(None, None, None, 128)	147,456
pa_fpn_p3p4p5_pre_0_1_bn pa_fpn_p3p4p5_pre_0_1... (BatchNormalization)	(None, None, None, 128)	512
pa_fpn_p3p4p5_pre_0_1 pa_fpn_p3p4p5_pre_0_1... (Activation)	(None, None, None, 128)	0
pa_fpn_p3p4p5_pre_0_2_pad pa_fpn_p3p4p5_pre_0_1... (ZeroPadding2D)	(None, None, None, 128)	0
pa_fpn_p3p4p5_pre_0_2_co... pa_fpn_p3p4p5_pre_0_2... (Conv2D)	(None, None, None, 128)	147,456
pa_fpn_p3p4p5_pre_0_2_bn pa_fpn_p3p4p5_pre_0_2... (BatchNormalization)	(None, None, None, 128)	512
pa_fpn_p3p4p5_pre_0_2 pa_fpn_p3p4p5_pre_0_2... (Activation)	(None, None, None, 128)	0

pa_fpn_p3p4p5_pre_1_1_pad pa_fpn_p3p4p5_pre_0_2... (ZeroPadding2D)	(None, None, None, 128)	0
pa_fpn_p3p4p5_pre_1_1_co... pa_fpn_p3p4p5_pre_1_1... (Conv2D)	(None, None, None, 128)	147,456
pa_fpn_p3p4p5_pre_1_1_bn pa_fpn_p3p4p5_pre_1_1... (BatchNormalization)	(None, None, None, 128)	512
pa_fpn_p3p4p5_pre_1_1 pa_fpn_p3p4p5_pre_1_1... (Activation)	(None, None, None, 128)	0
pa_fpn_p3p4p5_pre_1_2_pad pa_fpn_p3p4p5_pre_1_1... (ZeroPadding2D)	(None, None, None, 128)	0
pa_fpn_p3p4p5_pre_1_2_co... pa_fpn_p3p4p5_pre_1_2... (Conv2D)	(None, None, None, 128)	147,456
pa_fpn_p3p4p5_pre_1_2_bn pa_fpn_p3p4p5_pre_1_2... (BatchNormalization)	(None, None, None, 128)	512
pa_fpn_p3p4p5_pre_1_2 pa_fpn_p3p4p5_pre_1_2... (Activation)	(None, None, None, 128)	0

concatenate_8 split_5[0][0], (Concatenate) split_5[0][1], pa_fpn_p3p4p5_pre_0_2... pa_fpn_p3p4p5_pre_1_2...	(None, None, None, 512)	0
pa_fpn_p3p4p5_output_conv concatenate_8[0][0] (Conv2D)	(None, None, None, 256)	131,072
pa_fpn_p3p4p5_output_bn pa_fpn_p3p4p5_output_... (BatchNormalization)	(None, None, None, 256)	1,024
pa_fpn_p3p4p5_output pa_fpn_p3p4p5_output_... (Activation)	(None, None, None, 256)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_output[... (ZeroPadding2D)	(None, None, None, 256)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Conv2D)	(None, None, None, 256)	589,824
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (BatchNormalization)	(None, None, None, 256)	1,024
pa_fpn_p3p4p5_downsample1 pa_fpn_p3p4p5_downsam... (Activation)	(None, None, None, 256)	0

concatenate_9 pa_fpn_p3p4p5_downsam... (Concatenate) pa_fpn_p4p5_output[0]...	(None, None, None, 768)	0
pa_fpn_p3p4p5_downsample... concatenate_9[0][0] (Conv2D)	(None, None, None, 512)	393,216
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (BatchNormalization)	(None, None, None, 512)	2,048
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Activation)	(None, None, None, 512)	0
split_6 (Split) pa_fpn_p3p4p5_downsam...	[(None, None, None, 256), (None, None, None, 256)]	0
pa_fpn_p3p4p5_downsample... split_6[0][1] (ZeroPadding2D)	(None, None, None, 256)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Conv2D)	(None, None, None, 256)	589,824
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam...	(None, None, None, 1,024)	1,024

(BatchNormalization)	256)	
pa_fpn_p3p4p5_downsample...	(None, None, None,	0
pa_fpn_p3p4p5_downsam...	256)	
(Activation)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	0
pa_fpn_p3p4p5_downsam...	256)	
(ZeroPadding2D)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	589,824
pa_fpn_p3p4p5_downsam...	256)	
(Conv2D)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	1,024
pa_fpn_p3p4p5_downsam...	256)	
(BatchNormalization)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	0
pa_fpn_p3p4p5_downsam...	256)	
(Activation)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	0
pa_fpn_p3p4p5_downsam...	256)	
(ZeroPadding2D)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	589,824
pa_fpn_p3p4p5_downsam...	256)	
(Conv2D)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	1,024
pa_fpn_p3p4p5_downsam...	256)	
(BatchNormalization)		

pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Activation)	(None, None, None, 256)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (ZeroPadding2D)	(None, None, None, 256)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Conv2D)	(None, None, None, 256)	589,824
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (BatchNormalization)	(None, None, None, 256)	1,024
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Activation)	(None, None, None, 256)	0
concatenate_10 split_6[0][0], (Concatenate) split_6[0][1], pa_fpn_p3p4p5_downsam... pa_fpn_p3p4p5_downsam...	(None, None, None, 1024)	0
pa_fpn_p3p4p5_downsample... concatenate_10[0][0] (Conv2D)	(None, None, None, 512)	524,288
pa_fpn_p3p4p5_downsample...	(None, None, None,	2,048

pa_fpn_p3p4p5_downsam... (BatchNormalization)	512)	
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Activation)	(None, None, None, 512)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (ZeroPadding2D)	(None, None, None, 512)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Conv2D)	(None, None, None, 512)	2,359,296
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (BatchNormalization)	(None, None, None, 512)	2,048
pa_fpn_p3p4p5_downsample2 pa_fpn_p3p4p5_downsam... (Activation)	(None, None, None, 512)	0
concatenate_11 pa_fpn_p3p4p5_downsam... (Concatenate) functional[0][2]	(None, None, None, 1024)	0
pa_fpn_p3p4p5_downsample... concatenate_11[0][0] (Conv2D)	(None, None, None, 512)	524,288
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam...	(None, None, None, 	2,048

(BatchNormalization)	512)	
pa_fpn_p3p4p5_downsample...	(None, None, None,	0
pa_fpn_p3p4p5_downsam...	512)	
(Activation)		
split_7 (Split)	[(None, None, None,	0
pa_fpn_p3p4p5_downsam...	256), (None, None,	
	None, 256)]	
pa_fpn_p3p4p5_downsample...	(None, None, None,	0
split_7[0][1]	256)	
(ZeroPadding2D)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	589,824
pa_fpn_p3p4p5_downsam...	256)	
(Conv2D)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	1,024
pa_fpn_p3p4p5_downsam...	256)	
(BatchNormalization)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	0
pa_fpn_p3p4p5_downsam...	256)	
(Activation)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	0
pa_fpn_p3p4p5_downsam...	256)	
(ZeroPadding2D)		
pa_fpn_p3p4p5_downsample...	(None, None, None,	589,824

pa_fpn_p3p4p5_downsam... (Conv2D)	256)	
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (BatchNormalization)	(None, None, None, 256)	1,024
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Activation)	(None, None, None, 256)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (ZeroPadding2D)	(None, None, None, 256)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Conv2D)	(None, None, None, 256)	589,824
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (BatchNormalization)	(None, None, None, 256)	1,024
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (Activation)	(None, None, None, 256)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam... (ZeroPadding2D)	(None, None, None, 256)	0
pa_fpn_p3p4p5_downsample... pa_fpn_p3p4p5_downsam...	(None, None, None,	589,824

(Conv2D)	256)	
pa_fpn_p3p4p5_downsample...	(None, None, None,	1,024
pa_fpn_p3p4p5_downsam... (BatchNormalization)	256)	
pa_fpn_p3p4p5_downsample...	(None, None, None,	0
pa_fpn_p3p4p5_downsam... (Activation)	256)	
concatenate_12 split_7[0][0], (Concatenate) split_7[0][1],	(None, None, None, 1024)	0
pa_fpn_p3p4p5_downsam...		
pa_fpn_p3p4p5_downsam...		
pa_fpn_p3p4p5_downsample...	(None, None, None,	524,288
concatenate_12[0][0] (Conv2D)	512)	
pa_fpn_p3p4p5_downsample...	(None, None, None,	2,048
pa_fpn_p3p4p5_downsam... (BatchNormalization)	512)	
pa_fpn_p3p4p5_downsample...	(None, None, None,	0
pa_fpn_p3p4p5_downsam... (Activation)	512)	
yolo_v8_head_1_class_1_p...	(None, None, None,	0
pa_fpn_p3p4p5_output[... (ZeroPadding2D)	256)	

yolo_v8_head_2_class_1_p... pa_fpn_p3p4p5_downsam... (ZeroPadding2D)	(None, None, None, 512)	0
yolo_v8_head_3_class_1_p... pa_fpn_p3p4p5_downsam... (ZeroPadding2D)	(None, None, None, 512)	0
yolo_v8_head_1_box_1_pad pa_fpn_p3p4p5_output[... (ZeroPadding2D)	(None, None, None, 256)	0
yolo_v8_head_1_class_1_c... yolo_v8_head_1_class_... (Conv2D)	(None, None, None, 256)	589,824
yolo_v8_head_2_box_1_pad pa_fpn_p3p4p5_downsam... (ZeroPadding2D)	(None, None, None, 512)	0
yolo_v8_head_2_class_1_c... yolo_v8_head_2_class_... (Conv2D)	(None, None, None, 256)	1,179,648
yolo_v8_head_3_box_1_pad pa_fpn_p3p4p5_downsam... (ZeroPadding2D)	(None, None, None, 512)	0
yolo_v8_head_3_class_1_c... yolo_v8_head_3_class_... (Conv2D)	(None, None, None, 256)	1,179,648

yolo_v8_head_1_box_1_conv yolo_v8_head_1_box_1_... (Conv2D)	(None, None, None, 64)	147,456
yolo_v8_head_1_class_1_bn yolo_v8_head_1_class_... (BatchNormalization)	(None, None, None, 256)	1,024
yolo_v8_head_2_box_1_conv yolo_v8_head_2_box_1_... (Conv2D)	(None, None, None, 64)	294,912
yolo_v8_head_2_class_1_bn yolo_v8_head_2_class_... (BatchNormalization)	(None, None, None, 256)	1,024
yolo_v8_head_3_box_1_conv yolo_v8_head_3_box_1_... (Conv2D)	(None, None, None, 64)	294,912
yolo_v8_head_3_class_1_bn yolo_v8_head_3_class_... (BatchNormalization)	(None, None, None, 256)	1,024
yolo_v8_head_1_box_1_bn yolo_v8_head_1_box_1_... (BatchNormalization)	(None, None, None, 64)	256
yolo_v8_head_1_class_1 yolo_v8_head_1_class_... (Activation)	(None, None, None, 256)	0
yolo_v8_head_2_box_1_bn	(None, None, None, 64)	256

yolo_v8_head_2_box_1_... (BatchNormalization)		
yolo_v8_head_2_class_1 yolo_v8_head_2_class_... (Activation)	(None, None, None, 256)	0
yolo_v8_head_3_box_1_bn yolo_v8_head_3_box_1_... (BatchNormalization)	(None, None, None, 64)	256
yolo_v8_head_3_class_1 yolo_v8_head_3_class_... (Activation)	(None, None, None, 256)	0
yolo_v8_head_1_box_1 yolo_v8_head_1_box_1_... (Activation)	(None, None, None, 64)	0
yolo_v8_head_1_class_2_p... yolo_v8_head_1_class_... (ZeroPadding2D)	(None, None, None, 256)	0
yolo_v8_head_2_box_1 yolo_v8_head_2_box_1_... (Activation)	(None, None, None, 64)	0
yolo_v8_head_2_class_2_p... yolo_v8_head_2_class_... (ZeroPadding2D)	(None, None, None, 256)	0
yolo_v8_head_3_box_1 yolo_v8_head_3_box_1_...	(None, None, None, 64)	0

(Activation)		
yolo_v8_head_3_class_2_p... yolo_v8_head_3_class_... (ZeroPadding2D)	(None, None, None, 256)	0
yolo_v8_head_1_box_2_pad yolo_v8_head_1_box_1[... (ZeroPadding2D)	(None, None, None, 64)	0
yolo_v8_head_1_class_2_c... yolo_v8_head_1_class_... (Conv2D)	(None, None, None, 256)	589,824
yolo_v8_head_2_box_2_pad yolo_v8_head_2_box_1[... (ZeroPadding2D)	(None, None, None, 64)	0
yolo_v8_head_2_class_2_c... yolo_v8_head_2_class_... (Conv2D)	(None, None, None, 256)	589,824
yolo_v8_head_3_box_2_pad yolo_v8_head_3_box_1[... (ZeroPadding2D)	(None, None, None, 64)	0
yolo_v8_head_3_class_2_c... yolo_v8_head_3_class_... (Conv2D)	(None, None, None, 256)	589,824
yolo_v8_head_1_box_2_conv yolo_v8_head_1_box_2_... (Conv2D)	(None, None, None, 64)	36,864

yolo_v8_head_1_class_2_bn yolo_v8_head_1_class_... (BatchNormalization)	(None, None, None, 256)	1,024
yolo_v8_head_2_box_2_conv yolo_v8_head_2_box_2_... (Conv2D)	(None, None, None, 64)	36,864
yolo_v8_head_2_class_2_bn yolo_v8_head_2_class_... (BatchNormalization)	(None, None, None, 256)	1,024
yolo_v8_head_3_box_2_conv yolo_v8_head_3_box_2_... (Conv2D)	(None, None, None, 64)	36,864
yolo_v8_head_3_class_2_bn yolo_v8_head_3_class_... (BatchNormalization)	(None, None, None, 256)	1,024
yolo_v8_head_1_box_2_bn yolo_v8_head_1_box_2_... (BatchNormalization)	(None, None, None, 64)	256
yolo_v8_head_1_class_2 yolo_v8_head_1_class_... (Activation)	(None, None, None, 256)	0
yolo_v8_head_2_box_2_bn yolo_v8_head_2_box_2_... (BatchNormalization)	(None, None, None, 64)	256

yolo_v8_head_2_class_2 yolo_v8_head_2_class_... (Activation)	(None, None, None, 256)	0
yolo_v8_head_3_box_2_bn yolo_v8_head_3_box_2_... (BatchNormalization)	(None, None, None, 64)	256
yolo_v8_head_3_class_2 yolo_v8_head_3_class_... (Activation)	(None, None, None, 256)	0
yolo_v8_head_1_box_2 yolo_v8_head_1_box_2_... (Activation)	(None, None, None, 64)	0
yolo_v8_head_1_class_3_c... yolo_v8_head_1_class_... (Conv2D)	(None, None, None, 11)	2,827
yolo_v8_head_2_box_2 yolo_v8_head_2_box_2_... (Activation)	(None, None, None, 64)	0
yolo_v8_head_2_class_3_c... yolo_v8_head_2_class_... (Conv2D)	(None, None, None, 11)	2,827
yolo_v8_head_3_box_2 yolo_v8_head_3_box_2_... (Activation)	(None, None, None, 64)	0

yolo_v8_head_3_class_3_conv yolo_v8_head_3_class_3_conv (Conv2D)	(None, None, None, 11)	2,827
yolo_v8_head_1_box_3_conv yolo_v8_head_1_box_2_conv (Conv2D)	(None, None, None, 64)	4,160
yolo_v8_head_1_classifier yolo_v8_head_1_class_3_conv (Activation)	(None, None, None, 11)	0
yolo_v8_head_2_box_3_conv yolo_v8_head_2_box_2_conv (Conv2D)	(None, None, None, 64)	4,160
yolo_v8_head_2_classifier yolo_v8_head_2_class_3_conv (Activation)	(None, None, None, 11)	0
yolo_v8_head_3_box_3_conv yolo_v8_head_3_box_2_conv (Conv2D)	(None, None, None, 64)	4,160
yolo_v8_head_3_classifier yolo_v8_head_3_class_3_conv (Activation)	(None, None, None, 11)	0
concatenate_13 yolo_v8_head_1_box_3_conv (Concatenate) yolo_v8_head_1_classifier	(None, None, None, 75)	0
concatenate_14	(None, None, None, 75)	0

yolo_v8_head_2_box_3_...		
(Concatenate)		
yolo_v8_head_2_classi...		
concatenate_15	(None, None, None, 75)	0
yolo_v8_head_3_box_3_...		
(Concatenate)		
yolo_v8_head_3_classi...		
yolo_v8_head_1_output_re...	(None, None, 75)	0
concatenate_13[0][0]		
(Reshape)		
yolo_v8_head_2_output_re...	(None, None, 75)	0
concatenate_14[0][0]		
(Reshape)		
yolo_v8_head_3_output_re...	(None, None, 75)	0
concatenate_15[0][0]		
(Reshape)		
concatenate_16	(None, None, 75)	0
yolo_v8_head_1_output...		
(Concatenate)		
yolo_v8_head_2_output...		
yolo_v8_head_3_output...		
box_outputs (Activation)	(None, None, 75)	0
concatenate_16[0][0]		
get_item (GetItem)	(None, None, 64)	0
box_outputs[0][0]		
get_item_1 (GetItem)	(None, None, 11)	0
box_outputs[0][0]		

box (Concatenate) get_item[0][0]	(None, None, 64)	0
class (Concatenate) get_item_1[0][0]	(None, None, 11)	0

Total params: 39,417,761 (150.37 MB)

Trainable params: 39,374,881 (150.20 MB)

Non-trainable params: 42,880 (167.50 KB)

```
# keras.utils.plot_model(model, show_shapes=True)
```

COCO Metric Callback

```
class EvaluateCOCOMetricsCallback(keras.callbacks.Callback):
    def __init__(self, data, save_path):
        super().__init__()
        self.data = data
        self.metrics = keras_cv.metrics.BoxCOCOMetrics(
            bounding_box_format="xyxy",
            evaluate_freq=1e9,
        )
        self.save_path = save_path
        self.best_map = -1.0

    def on_epoch_end(self, epoch, logs):
        self.metrics.reset_state()
        for batch in self.data:
            images, y_true = batch[0], batch[1]
            y_pred = self.model.predict(images, verbose=0)
            self.metrics.update_state(y_true, y_pred)

        metrics = self.metrics.result(force=True)
        logs.update(metrics)

        current_map = metrics["MaP"]
        if current_map > self.best_map:
            self.best_map = current_map
            self.model.save(self.save_path)  # Save the model when mAP
improves

        return logs
```

```

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=3,
    callbacks=[EvaluateCOCOMetricsCallback(val_ds,
"model.h5"), tf.keras.callbacks.EarlyStopping(patience=5,
restore_best_weights=True)])

```

Epoch 1/3

```

100/100 ————— 293s 2s/step - box_loss: 2.7060 -
class_loss: 0.2422 - loss: 2.9482 - val_box_loss: 0.0492 -
val_class_loss: 0.0461 - val_loss: 0.0953 - MaP: 0.0000e+00 -
MaP@[IoU=50]: 0.0000e+00 - MaP@[IoU=75]: 0.0000e+00 -
MaP@[area=small]: 0.0000e+00 - MaP@[area=medium]: 0.0000e+00 -
MaP@[area=large]: 0.0000e+00 - Recall@[max_detections=1]: 0.0000e+00 -
Recall@[max_detections=10]: 0.0000e+00 - Recall@[max_detections=100]:
0.0000e+00 - Recall@[area=small]: 0.0000e+00 - Recall@[area=medium]:
0.0000e+00 - Recall@[area=large]: 0.0000e+00

```

Epoch 2/3

```

100/100 ————— 167s 2s/step - box_loss: 2.6770 -
class_loss: 0.0063 - loss: 2.6834 - val_box_loss: 1.0912 -
val_class_loss: 0.0011 - val_loss: 1.0924 - MaP: 0.0000e+00 -
MaP@[IoU=50]: 0.0000e+00 - MaP@[IoU=75]: 0.0000e+00 -
MaP@[area=small]: 0.0000e+00 - MaP@[area=medium]: 0.0000e+00 -
MaP@[area=large]: 0.0000e+00 - Recall@[max_detections=1]: 0.0000e+00 -
Recall@[max_detections=10]: 0.0000e+00 - Recall@[max_detections=100]:
0.0000e+00 - Recall@[area=small]: 0.0000e+00 - Recall@[area=medium]:
0.0000e+00 - Recall@[area=large]: 0.0000e+00

```

Epoch 3/3

```

100/100 ————— 165s 2s/step - box_loss: 2.9708 -
class_loss: 0.0022 - loss: 2.9730 - val_box_loss: 1.7389 -
val_class_loss: 5.6286e-04 - val_loss: 1.7395 - MaP: 1.9426e-07 -
MaP@[IoU=50]: 1.2951e-06 - MaP@[IoU=75]: 0.0000e+00 -
MaP@[area=small]: 0.0000e+00 - MaP@[area=medium]: 1.5360e-07 -
MaP@[area=large]: 1.1380e-06 - Recall@[max_detections=1]: 0.0000e+00 -
Recall@[max_detections=10]: 3.5282e-05 - Recall@[max_detections=100]:
3.5282e-05 - Recall@[area=small]: 0.0000e+00 - Recall@[area=medium]:
3.8197e-05 - Recall@[area=large]: 1.1299e-04

```

Visualize Predictions

```

import tensorflow as tf
from keras_cv import visualization, bounding_box

def visualize_detections(model, dataset, bounding_box_format,
class_mapping, conf_threshold=0.5, max_detections=5):
    images, y_true = next(iter(dataset.take(1)))
    y_pred = model.predict(images)

```

```

# Extract boxes from y_pred
if isinstance(y_pred, dict):
    y_pred_boxes = y_pred["boxes"] # Shape: (batch_size, 100, 4)
    y_pred_classes = y_pred["classes"] # Shape: (batch_size, 100)
    y_pred_confidence = y_pred["confidence"] # Shape:
(batch_size, 100)
else:
    raise ValueError("y_pred must be a dictionary with 'boxes',
'classes', and 'confidence'")

# Filter and trim predictions
y_pred_boxes_trimmed = []
y_pred_classes_trimmed = []
for i in range(len(images)):
    conf = y_pred_confidence[i]
    boxes = y_pred_boxes[i]
    classes = y_pred_classes[i]

    # Filter by confidence threshold
    mask = conf > conf_threshold
    filtered_conf = conf[mask]
    filtered_boxes = boxes[mask]
    filtered_classes = classes[mask]

    # Sort by confidence and take top max_detections
    if tf.shape(filtered_conf)[0] > 0:
        top_indices = tf.argsort(filtered_conf,
direction="DESCENDING")[:max_detections]
        filtered_boxes = tf.gather(filtered_boxes, top_indices)
        filtered_classes = tf.gather(filtered_classes,
top_indices)
    else:
        filtered_boxes = tf.zeros((0, 4), dtype=tf.float32)
        filtered_classes = tf.zeros((0,), dtype=tf.int32)

    # Pad to max_detections
    num_detections = tf.shape(filtered_boxes)[0]
    padding = [[0, max_detections - num_detections], [0, 0]]
    padded_boxes = tf.pad(filtered_boxes, padding,
constant_values=-1) # Pad with -1 for invalid boxes
    padded_classes = tf.pad(filtered_classes, [[0, max_detections
- num_detections]], constant_values=-1)

    y_pred_boxes_trimmed.append(padded_boxes)
    y_pred_classes_trimmed.append(padded_classes)

# Stack into fixed-size tensors
y_pred_boxes_padded = tf.stack(y_pred_boxes_trimmed, axis=0) #

```

```

Shape: (batch_size, max_detections, 4)
    y_pred_classes_padded = tf.stack(y_pred_classes_trimmed, axis=0)
# Shape: (batch_size, max_detections)

# Combine into a dictionary
y_pred_formatted = {
    "boxes": y_pred_boxes_padded,
    "classes": y_pred_classes_padded
}

# Visualize
visualization.plot_bounding_box_gallery(
    images,
    value_range=(0, 255),
    bounding_box_format=bounding_box_format,
    y_true=y_true,
    y_pred=y_pred_formatted,
    scale=8,
    rows=2,
    cols=2,
    show=True,
    font_scale=0.7,
    class_mapping=class_mapping,
)

# Run the function
visualize_detections(
    model,
    dataset=val_ds,
    bounding_box_format="xyxy",
    class_mapping=class_mapping, # Assuming defined elsewhere
    conf_threshold=0.5,
    max_detections=5
)

```

1/1 ————— 0s 154ms/step

