



AI & Partners

Amsterdam - London - Singapore



EU AI Act

Prohibited AI Practices:
A Primer

November 2025

Authors

Sean Musch, Michael Borrelli, Charles Kerrigan, Hr Dr. Takahisa Karita,
Saeed Al Dhaheri, Agata Wolska, Cristiana Vergara



High-algorithmic management countries face elevated EU AI Act risks under Article 5 (Prohibited AI Practices): prioritize compliance by restricting emotion-inference AI in workplaces and educational institutions

European Parliament: AI Act: Digitalisation, artificial intelligence and algorithmic management in the workplace: Shaping the future of work (October 2025)

Region	Share of Workers (%), EWCS 2024	Share of Workers (%), EB2024	Share of Companies (%), ESENER 2024	Share of Companies (%), OECD 2024
	42.3	51.3	14.5	-
	50.5	65.6	20.6	-
	58.6	58.9	16.7	-
	35.6	44.4	18.5	-
	34.9	62.7	15.2	-
	43.8	38.7	13.6	-
	38.1	47.5	16.5	78
	70.1	47	17.6	-
	43.1	53.3	20.1	-
	27.2	36.1	6	-
	41.9	51.8	12	78
	41.7	69.1	31.2	-
	37.1	53.1	19.9	81



	48.9	41	11.6	-
	33.4	54.6	13.2	-
	52.5	56.6	15.1	-
	47.3	52.8	7.7	76
	31	47.6	21.6	-
	58.4	66.2	13.9	-
	40.9	52.7	16.9	-
	39.5	63.3	18.6	-
	65.9	41.3	10.7	-
	39.4	52.2	10.8	-
	26.8	41.3	11.2	-
	41.9	68.8	14.6	-
	50.9	55	12.6	-
	40	49.9	12.1	-
	40	46.4	13	-
	-	-	-	90
	-	-	-	40

The 2024 EWCS and EB surveys show that over 50% of workers in countries like AT, BE, DK, IE, and NL are exposed to advanced management (AM) practices, while ESENER reports that 10–20% of companies deploy AI-enabled systems in workplaces. Article 5 of the EU AI Act prohibits AI systems that infer natural persons' emotions in workplaces and educational institutions.

This creates a regulatory imperative for high-AM countries where AI adoption is concentrated. A critical insight is that organizations in these markets may unintentionally breach the prohibition if emotion-inference features are embedded in performance tracking, engagement monitoring, or adaptive learning tools.

Constructively, firms should map all AI systems against the Article 5 criteria, conduct pre-deployment risk assessments, and implement mitigation protocols to maintain compliance, particularly in countries with high AM prevalence where enforcement and exposure risk are elevated.



AI & Partners

Amsterdam - London - Singapore

Who Are We

AI That You Can Trust

Why Us?

Stay on the right side of history. At AI & Partners, we believe AI should unlock potential—not cause harm. We've seen the fear and fallout when teams lose control of AI, but also the trust and innovation that follow when it's handled responsibly. That's why we exist: to help you build AI you can trust and stand behind—for the long run.

80%

of AI systems
are unknown


What Do We Do?

We enable safe AI usage—for your organization and your clients. Unknown AI adoption leads to confusion, risk, and reputational damage. We help you take control with tools to identify, monitor, and govern all AI systems—so you're not reacting to AI, you're leading it.

How Do We Do It?

Do you know what AI systems you have? Identify all known and unknown AI systems (algorithms, LLMs, prompts, and models) from all internal and external AI vendors, automated by generating your inventory. Overall, 80% of AI inventory is unknown to our clients.

How do you guarantee ongoing safe AI use? Continuously monitor deployed AI systems for performance drift, anomalies or failures, real-world impacts, and emerging risks (e.g. data poisoning). Any malfunction of an AI system has severe implications for organisations (e.g. inability to assess online misinformation that leads to widespread public mistrust), so monitoring becomes a matter of urgency.



AI Discovery & AI Inventory

Automatically detect all AI systems, including models, algorithms, and prompts, and maintain a live, always-updated register for full visibility and compliance.



Responsible AI

Embed fairness, transparency, and control into every stage of AI use—aligning with the EU AI Act and building 'Trustworthy-by-Design'.



Model Monitoring

Continuously track your AI models after deployment to detect drift, bias, or failure—so you stay in control and prevent harm before it happens.

Contact: Team | contact@ai-and-partners.com

Contents

Who Are We	3
Introduction	5
Key questions being asked about Prohibited AI Practices	6
1. What are prohibited AI practices under the EU AI Act?	7
2. How can an organisation ensure it avoids harmful manipulation through AI?	7
3. Is emotion recognition entirely banned under the AI Act?	7
4. Why is social scoring banned by the AI Act?	7
5. What is untargeted scraping, and why is it prohibited for facial recognition?	7
6. Can AI be used to predict criminal behavior under any circumstance?	8
7. What measures should be taken to avoid biometric categorisation violations?	8
8. What are the conditions for using real-time remote biometric identification (RBI)?	8
9. How can AI developers comply with the AI Act during design and development?	8
10. What penalties exist for violating the AI Act's prohibited practices?	8
Understanding Prohibited AI Practices	9
1. Social Scoring	10
2. Individual criminal offence risk assessment and prediction	11
3. Untargeted scraping to develop facial recognition databases	12
4. Emotion recognition	13
5. Biometric categorisation	14
6. Real-time remote biometric identification ("RBI")	15
7. Harmful exploitation of vulnerabilities	16
8. Harmful manipulation, and deception	17
Mapping Prohibited AI Practices to AI System Lifecycle Stages	18
Calls to Action	24
Conclusion	26
References	27

Introduction

As the EU Artificial Intelligence Act (AI Act) advances toward full enforcement, the regulation's prohibition regime has emerged as a decisive safeguard against AI applications that pose an "unacceptable risk" to fundamental rights, safety, and democratic values. These prohibited practices—ranging from social scoring and biometric categorisation to harmful manipulation and untargeted facial data scraping—are banned outright because they threaten human dignity, autonomy, equality, and privacy. Their inclusion in Article 5 of the AI Act reflects the EU's determination to establish clear red lines for AI development and use.

This report serves as a practical reference for providers, deployers, compliance officers, policymakers, and civil society actors seeking to understand, identify, and prevent prohibited AI practices in both design and operational contexts. Drawing directly from the AI Act's prohibitions and informed by real-world implementation considerations, it distils the underlying rationales, legal definitions, and risk scenarios that trigger compliance breaches. From mapping lifecycle stages to embedding technical and organisational safeguards, the report offers operational guidance to ensure AI innovation proceeds without crossing into banned territory.

The prohibition framework represents a cornerstone of the EU's strategy to build public trust in AI by removing the most harmful applications from the market entirely. It ensures that high-risk practices are not merely regulated but eliminated at the source, reducing the possibility of post-deployment harm and avoiding reliance on reactive enforcement alone. This approach protects individuals and communities from exploitative, discriminatory, or coercive uses of AI while providing legal clarity for industry actors.

In setting out the actors' roles, responsibilities, and practical mitigation measures, this document promotes a preventative, evidence-based approach to compliance. It illustrates how providers can design out prohibited capabilities, how deployers can monitor against unintended violations, and how regulators and advocacy groups can collaborate to ensure adherence.

Whether you are a public authority evaluating AI procurement, a technology provider architecting new systems, or a corporate deployer integrating AI into sensitive domains, this report provides a structured blueprint for staying on the right side of the AI Act's prohibition rules—supporting lawful, ethical, and rights-respecting AI across the EU.


Best regards,

Sean Musch

Founder/CEO

AI & Partners

A stylized, handwritten signature in white ink, likely belonging to Sean Musch, positioned to the right of his name and title.



Key questions being asked about Prohibited AI Practices

1. What are prohibited AI practices under the EU AI Act?

Prohibited AI practices are uses of AI deemed to pose unacceptable risks to individuals' fundamental rights, safety, or dignity. These include systems that manipulate or deceive users, exploit vulnerabilities, use social scoring, predict criminal behavior through profiling, scrape facial data indiscriminately, recognize emotions in certain settings, categorize by sensitive biometric traits, or conduct real-time biometric identification in public. These practices are banned because they violate principles such as autonomy, equality, and privacy. The prohibitions apply to placing systems on the market, putting them into service, or using them within the EU, with some narrowly defined exceptions.

2. How can an organisation ensure it avoids harmful manipulation through AI?

To avoid harmful manipulation, organisations must design AI systems transparently and ethically. They should avoid subliminal content, deceptive nudges, or interface tricks that impair informed decision-making. All persuasive techniques should be disclosed, optional, and not exploit cognitive biases. Regular ethics reviews, user testing, and compliance audits can help detect manipulative behaviors early. Documentation must include clear statements on intended use and risk mitigation. Training staff in ethical design and ensuring user interfaces support autonomy and clarity are critical. This ensures AI aligns with user rights and EU legal standards, preventing violations and reputational damage.

3. Is emotion recognition entirely banned under the AI Act?

No, emotion recognition is not entirely banned, but it is prohibited in workplaces and educational institutions, except for specific medical or safety purposes. This restriction exists because emotional inference is often unreliable and can lead to privacy breaches or unfair assessments. For example, using AI to monitor student focus or employee mood without consent is banned. Even where allowed, the use must be justified, documented, and accompanied by safeguards like informed consent and oversight. Organisations should implement strict controls, restrict use to approved scenarios, and regularly audit systems to ensure compliance with the AI Act's boundaries.

4. Why is social scoring banned by the AI Act?

Social scoring is banned because it enables discrimination, unjustified profiling, and social control. AI systems that rate individuals based on behavior, personality, or lifestyle—and apply these scores in unrelated contexts—can cause significant harm. Such practices may restrict access to jobs, services, or benefits based on irrelevant or biased data. They undermine fairness, dignity, and equality, especially if applied without transparency or appeal mechanisms. The ban prevents digital reputation systems from becoming tools of exclusion or coercion, ensuring people are not judged on arbitrary or cross-contextual metrics. It protects societal trust and democratic principles in AI applications.

5. What is untargeted scraping, and why is it prohibited for facial recognition?

Untargeted scraping involves indiscriminately collecting facial images from the internet or CCTV footage to build facial recognition databases. This practice is banned because it violates privacy and data protection laws. Individuals often have no knowledge of or control over their data being harvested. Such scraping facilitates mass surveillance, profiling, and misuse by both public and private actors. To comply, organisations must verify the legal provenance of any biometric data used. Data must be sourced through targeted, consent-based, and lawful means only. Prohibiting scraping protects public trust and prevents abuse of biometric surveillance technologies.

6. Can AI be used to predict criminal behavior under any circumstance?

AI cannot be used to autonomously predict criminal behavior based solely on profiling, personality traits, or personal characteristics. However, it may support human decision-making if based on objective, case-specific, and verifiable facts linked to a criminal investigation. For example, AI could analyze a known suspect's behavior within the scope of a crime scene, but not predict risk based on their demographics or psychological profile. Organisations must ensure transparency, explainability, and human oversight. This protects the presumption of innocence and prevents discriminatory policing or preemptive surveillance, aligning with both legal and ethical norms.

7. What measures should be taken to avoid biometric categorisation violations?

Organisations must prohibit AI from deducing or acting on sensitive personal characteristics—like race, religion, sexual orientation, or political beliefs—based on biometric data. This includes configuring systems to avoid such inferences, auditing training datasets for bias, and clearly documenting lawful, narrow purposes (e.g., identity verification only). Biometric systems must not be used for profiling or targeting individuals based on inferred traits. Internal policies should reflect these prohibitions, and any use in law enforcement or public services must align with legal exceptions. Regular audits, purpose limitation, and transparency are critical safeguards to prevent misuse.

8. What are the conditions for using real-time remote biometric identification (RBI)?

Real-time RBI in publicly accessible spaces is generally prohibited unless it meets narrow legal exceptions: finding missing persons, preventing imminent threats (like terrorism), or locating specific suspects. Even then, prior judicial or independent administrative authorization is required, along with strict documentation, transparency, and oversight. Organisations must implement technical restrictions that prevent unauthorized activation and log all uses. Misuse or circumvention can lead to serious legal consequences. This framework ensures RBI is only used when absolutely necessary, balancing law enforcement needs with the public's right to privacy, freedom of movement, and protection from surveillance.

9. How can AI developers comply with the AI Act during design and development?

Developers must ensure AI systems do not incorporate or enable prohibited practices. This includes designing for transparency, fairness, and safety from the outset. They should conduct fundamental rights impact assessments, eliminate manipulative or exploitative functions, and implement technical constraints to prevent misuse. All data used must be lawfully sourced and representative, especially for biometric or profiling features. Systems must be explainable and respect user autonomy. Developers must collaborate with legal, ethical, and UX experts to align their work with the AI Act. Documentation of intended use and safeguards is essential for compliance and accountability.


10. What penalties exist for violating the AI Act's prohibited practices?

Violations of prohibited AI practices carry the highest penalties under the EU AI Act: up to €35 million or 7% of global annual turnover, whichever is higher. This reflects the severity of harm these practices pose to individuals' rights and safety. Public bodies can also face fines up to €1.5 million. Fines may apply to both providers and deployers, and violations may also trigger civil liability or enforcement under other EU laws (e.g., GDPR). To mitigate risk, organisations should implement proactive compliance frameworks, regularly review their AI systems, and maintain full transparency with regulators and stakeholders.



Understanding Prohibited AI Practices

1. Social Scoring



What is this?

- Social scoring involves evaluating or ranking individuals based on social behavior or personality traits, especially when the scores lead to unfair treatment in unrelated areas.
- For example, someone's personal conduct might unjustifiably impact their access to credit, employment, or services.


How can it be prevented?

- Avoid linking personal data across contexts without clear justification.
- Disable or restrict profiling components that aggregate behavioral data.
- Conduct legal and ethical reviews before implementing scoring mechanisms.
- Ensure any risk assessment tools are limited to contextually relevant, objective criteria.

Why is it prohibited?

- It violates privacy, fairness, and equality by enabling discriminatory decisions and reputational harm.
- Individuals may be punished or rewarded for behaviors irrelevant to a specific context, leading to systemic injustice, social exclusion, or coercion to conform—akin to surveillance-based control.

2. Individual criminal offence risk assessment and prediction



What is this?

- This refers to AI systems predicting whether a person will commit a crime, based solely on profiling or psychological traits rather than objective, case-specific evidence.
- Such systems risk labeling individuals as potential criminals based on generalizations.


How can it be prevented?

- Design AI systems to support, not replace, human judgment. Use only verifiable, case-specific data and avoid reliance on inferred behavioral patterns.
- Apply strict oversight for any predictive tool used in law enforcement and prohibit automated decision-making in personal crime predictions.

Why is it prohibited?

- It infringes on the presumption of innocence and fosters discriminatory outcomes.
- Profiling can perpetuate bias and stigmatize individuals, violating rights to liberty, fair trial, and privacy.
- It risks false positives, psychological harm, and wrongful surveillance or policing.

3. Untargeted scraping to develop facial recognition databases



What is this?

- This involves collecting facial images from public sources—like social media or CCTV—without consent or targeting, to build or expand facial recognition systems.
- It constitutes mass data harvesting without lawful basis or individual awareness.


How can it be prevented?

- Implement strict data sourcing controls. Vet datasets for provenance and consent, and prohibit use of scraped images. Incorporate technical safeguards to block data ingestion from open or unverified sources.
- Adopt a “no scraping” policy and require dataset audits.

Why is it prohibited?

- It violates data protection, privacy, and consent rights. Individuals are unknowingly subjected to biometric profiling, which can lead to surveillance, discrimination, or misuse by state or private actors.
- The practice undermines trust and opens doors to mass tracking.

4. Emotion recognition



What is this?

- Emotion recognition AI attempts to infer a person's emotional state based on biometric or behavioral cues.
- It's prohibited in workplaces and education, where it could be used to monitor, assess, or influence individuals without consent or scientific reliability.


How can it be prevented?

- Ban such systems in employment and educational settings.
- For permitted contexts (e.g., medical use), require expert oversight, informed consent, and accuracy validation.
- Regularly audit the system's deployment and disable emotion inference by default in user-facing modules.

Why is it prohibited?

- Emotion recognition is often unreliable and can lead to intrusive, discriminatory, or coercive outcomes.
- In sensitive settings, it risks misuse to monitor or manipulate individuals' behavior or performance, infringing privacy, dignity, and freedom of expression.

5. Biometric categorisation



What is this?

- Biometric categorisation involves using physical traits (e.g., facial structure) to infer sensitive personal characteristics like race, religion, sexual orientation, or political views.
- The practice is banned when used to deduce or act on such attributes.


How can it be prevented?

- Disable or remove classifiers that infer sensitive traits.
- Review training data to ensure it doesn't reflect or reproduce these inferences.
- Establish policies banning categorisation beyond identification/authentication purposes.
- Document lawful basis and limits of biometric processing.

Why is it prohibited?

- It enables profiling and discrimination, often with deeply harmful consequences.
- Such systems can lead to unjust denial of services, harassment, or persecution based on identity markers, undermining equality, freedom of belief, and personal dignity.

6. Real-time remote biometric identification ("RBI")



What is this?

- This is the real-time identification of individuals in public spaces using biometric data (e.g., facial recognition), typically by law enforcement.
- Except in exceptional, narrowly defined cases, its use is prohibited due to its surveillance potential


How can it be prevented?

- Ensure any biometric system in public spaces cannot function in real time without prior legal authorization.
- Implement technical restrictions and require judicial oversight for every use.
- Maintain logs, obtain consent where applicable, and assign accountability roles.

Why is it prohibited?

- RBI enables mass surveillance, undermines privacy, and deters free movement, expression, and assembly.
- Its pervasive nature can have a chilling effect on democratic freedoms and risks abuse by authorities, leading to unwarranted targeting and social control.

7. Harmful exploitation of vulnerabilities



What is this?

- This refers to AI systems that exploit users' vulnerabilities—based on age, disability, or socio-economic status—to distort their behavior in a harmful way.
- These systems take advantage of reduced autonomy or understanding to influence decisions or actions that may result in harm.


How can it be prevented?

- Implement vulnerability impact assessments and limit functionality in high-risk contexts (e.g., children's apps or services for the elderly). Include consent prompts in accessible formats and establish safeguards for adaptive behavior that could target vulnerable traits.
- Train staff on ethical boundaries in product design.

Why is it prohibited?

- This practice endangers the most at-risk populations by using their conditions against them, violating rights to equality, safety, and human dignity.
- It can lead to undue influence, exclusion, or harm and contradicts the EU's commitment to protecting the vulnerable in digital environments.

8. Harmful manipulation, and deception



What is this?


- This refers to AI systems that use subliminal, manipulative, or deceptive techniques—either beyond a person’s awareness or designed to influence behavior—in ways that impair decision-making.
- These systems may nudge users to act against their own interests or without informed consent, leading to decisions they wouldn’t otherwise make.

How can it be prevented?

- Prevent manipulation by embedding transparency and user control into the design. Avoid dark patterns, ensure clear disclosures, and conduct ethics reviews to identify manipulative interactions.
- Test for unintended influence during development and document consent and comprehension mechanisms throughout the user experience.

Why is it prohibited?

- Such manipulation undermines autonomy and can lead to significant harm, including financial loss, coercion, or psychological distress.
- It erodes trust in AI, violates the right to free decision-making, and contradicts principles of dignity and transparency, especially when used in commercial, political, or sensitive contexts.

The background is a dark, teal-green gradient. It features several wavy, glowing lines that sweep across the frame from the bottom left towards the top right. These lines are composed of many small, bright green and yellow dots, creating a sense of motion and depth. Scattered throughout the background are numerous out-of-focus light spots, or bokeh, in shades of yellow, orange, and green, adding to the abstract, digital feel of the image.

Mapping Prohibited AI Practices to AI System Lifecycle Stages

Design

The initial stage defining AI model architecture, data needs, and goals.



Development

The creation phase transforming AI designs into trained, functional models.



Deployment

The operational stage integrating AI models into real-world user environments.



00

Prohibited AI Practice		AI System Lifecycle Stage		
Practice	Description	Design	Deeopment	Deployment
Harmful manipulation, and deception	AI systems that use subliminal, manipulative, or deceptive techniques to distort human behavior, impair decision-making, and cause significant harm. These techniques operate beyond consciousness or intentionally manipulate users, undermining autonomy and free will, often exploiting cognitive biases for unethical influence.	Design should proactively identify any manipulative, deceptive, or subliminal techniques, ensuring they are excluded from the system's functionality. Consider how user interfaces, nudges, or personalization features may impair informed decision-making. Ethical reviews and human rights impact assessments should guide architecture choices. Transparency and explainability must be built-in from the outset to avoid hidden influence on users' behavior.	During development, scrutinize content delivery methods to avoid subliminal audio/visual cues or manipulation of attention. Validate that algorithms do not deploy deceptive persuasion tactics, misrepresent functionality, or exploit cognitive biases. Conduct adversarial testing to reveal unintended manipulative behaviors. Ensure design documentation clearly defines lawful user interaction standards.	Deployers must ensure the AI is not used in contexts where it might influence vulnerable users or important decisions (e.g., financial or medical choices) through hidden or deceptive nudges. Training and usage guidelines must reinforce ethical usage. Any customization should retain transparency and safeguard users' autonomy, especially in consumer-facing applications.
Harmful exploitation of vulnerabilities	AI systems that exploit vulnerabilities due to age, disability, or socio-economic conditions, influencing behavior in harmful ways. This prohibition protects individuals whose capacity for informed decisions is compromised, especially when such exploitation leads to significant physical, psychological, or social harm.	Systems should be designed to be inclusive, ensuring they do not target or behave differently toward users based on age, disability, or socio-economic status. Designers must assess whether any functionality could exploit vulnerable individuals or impair their autonomy. Human oversight mechanisms and accessibility	Implement strict controls to prevent AI from adapting in ways that disproportionately influence or coerce vulnerable populations. Training data must be carefully reviewed for representativeness and fairness. Ensure user testing includes diverse demographic groups to identify potential exploitation paths and mitigate them before deployment.	Deployers must avoid using AI in contexts where it could take advantage of individuals' vulnerabilities, especially in public services, education, or elder care. Use should comply with ethical guidelines and be regularly audited to ensure it does not evolve into an exploitative tool. Provide clear user consent and support options.

		considerations should be embedded early.		
Social scoring	AI systems that classify individuals based on behavior or personality traits, producing scores that result in unjustified or disproportionate treatment in unrelated contexts. These systems threaten fairness, dignity, and equal treatment by institutionalizing discrimination through automated social control mechanisms.	Design systems to avoid evaluating individuals based on aggregated social behavior or personality traits across unrelated contexts. Establish strict boundaries around data use and processing logic to prevent unintended social scoring. The system's purpose should not involve creating classifications that affect rights or opportunities outside the original data context.	Ensure algorithms do not generate composite scores or profiles that rank users or groups in ways that could result in unjustified treatment. Eliminate feedback loops that adjust behavior based on assumed social value. Development documentation must detail how the system avoids cross-contextual judgment or reputational scoring.	Deployers must not repurpose systems to assess individuals' eligibility for services or benefits based on behavior or characteristics unrelated to the specific domain. Monitor usage to prevent mission creep into social evaluation. Any profiling or scoring components should be narrowly applied and fully justified with legal and ethical oversight.
Individual criminal offence risk assessment and prediction	AI systems predicting criminal behavior based solely on profiling or personality traits are banned. Such predictions lack objectivity and reinforce biases, violating fundamental rights and the presumption of innocence. Exceptions apply only if based on verifiable facts directly linked to criminal activities.	Ensure systems are not designed to predict criminal behavior based solely on personal traits or profiling. Clearly differentiate between systems supporting human analysis based on verifiable facts and those that autonomously infer future criminality. Exclude risk assessment modules that rely on subjective behavioral indicators.	Avoid using datasets or training labels that reflect bias or imply causation between personality traits and criminal activity. Clearly document exclusions of any profiling-based prediction logic. Implement audit trails to ensure all model predictions are grounded in objective data and not statistical inference about individual behavior.	Deployers must not use systems to assign crime risk to individuals unless supporting a human judgment using specific, verifiable facts. The system must not act autonomously to trigger policing, denial of services, or surveillance. Ensure output is contextual, traceable, and explainable, especially in law enforcement environments.
Untargeted scraping to develop facial recognition databases	AI systems that collect facial images indiscriminately from the internet or CCTV to create facial recognition	Design must prohibit any plan to build facial recognition datasets from untargeted	Ensure no scraping modules are implemented to indiscriminately collect facial data. Verify that any data input is	Deployers must refrain from using systems to expand or create facial recognition databases using broad scraping methods. Limit usage to

	databases are prohibited. This untargeted data collection infringes on privacy rights and data protection principles, often occurring without consent or oversight, raising serious surveillance and ethical concerns.	scraping of online images or CCTV footage. Ensure the system cannot automatically gather or ingest facial data without clear, targeted consent or legal basis. Integrate ethical constraints into data acquisition workflows.	acquired legally, with consent and purpose limitations. Implement barriers against the inclusion of scraped or public images. Conduct legal and ethical audits on all training datasets containing facial imagery.	datasets with verified provenance and lawful basis. Actively monitor data ingestion pathways and system logs to prevent unauthorized data expansion. Maintain full documentation of dataset sources and purposes.
Emotion recognition	AI systems that infer emotions in workplaces or educational settings are banned, except for specific medical or safety uses. These systems risk manipulating emotional states and intruding on personal boundaries, posing threats to mental health, dignity, and freedom of expression.	If designing emotion recognition, restrict it from being used in workplaces or education settings unless for medical or safety reasons. Validate that the system's function is clearly explained and avoids manipulating or judging users based on inferred emotional states. Default settings should minimize invasiveness.	Emotion recognition algorithms must be tested to ensure they don't misclassify emotional states or reinforce stereotypes. Implement clear limits on application contexts. Include consent mechanisms and transparency protocols. Development logs should record the purpose and anticipated use environments, with strict access control for sensitive inference data.	Deployers must not use emotion recognition systems in classrooms or workplaces unless explicitly allowed under safety or health exceptions. Train users on appropriate boundaries. Log and audit usage to ensure it aligns with medical or safety intentions and doesn't influence performance evaluation, hiring, or discipline.
Biometric categorisation	AI systems categorizing individuals based on biometric data to infer sensitive characteristics (e.g., race, religion, political views) are prohibited. These systems threaten equality, non-discrimination, and privacy by enabling profiling that can lead to systemic bias and social exclusion.	Designers must exclude any capability that infers sensitive traits such as race, religion, sexual orientation, or political beliefs from biometric data. Limit biometric processing to lawful, narrow purposes such as identity verification. Clearly separate any categorisation functionality from	Prohibit model training that maps biometric data to sensitive attributes. Strip such labels from datasets. Implement strong privacy-preserving techniques and differential privacy where applicable. Document all constraints and purpose limitations in technical documentation, ensuring that the system cannot be repurposed for profiling protected characteristics.	Deployers must not use biometric categorisation for deducing or acting on sensitive personal characteristics. Ensure that systems used in law enforcement or other sectors strictly adhere to dataset filtering or lawful categorisation exemptions. Regularly assess usage logs to prevent function creep into discriminatory profiling.

		personal profiling logic.		
Real-time remote biometric identification ("RBI")	The use of real-time RBI systems in public spaces for law enforcement is banned, except in narrowly defined cases like preventing terrorism or locating missing persons. These systems enable mass surveillance, undermining privacy, freedom of assembly, and civil liberties.	Avoid designing systems capable of real-time biometric identification in public spaces unless they meet one of the narrow legal exceptions. Design for minimal data retention, decentralised processing, and maximum user privacy. Include modular controls that allow lawful disabling of RBI features.	Implement technical safeguards that limit use to predefined, legally permissible scenarios (e.g., missing persons). Prevent continuous surveillance capabilities or expansion into general public monitoring. Ensure interoperability with authorisation protocols for law enforcement use. Testing should simulate real-world legal constraints and authorisation workflows.	Deployers must ensure RBI is only used in strict accordance with legally defined exceptions, such as imminent threats or searches for specific suspects. Obtain prior judicial or administrative authorisation. Log every activation, conduct oversight reviews, and maintain transparency with oversight bodies and the public.

The background is a dark, deep green gradient. Overlaid on this are several wavy, translucent lines in a lighter green and yellowish-gold color. These lines create a sense of depth and movement, resembling a digital or data-driven landscape. Scattered throughout the scene are numerous small, out-of-focus circular lights in shades of yellow, orange, and green, creating a bokeh effect that adds to the abstract, high-tech feel of the image.

Calls to Action



Implement AI Ethics Review Panels

Establish cross-functional ethics panels to review all AI system designs and deployments for manipulative, subliminal, or deceptive features. This panel should include legal, technical, and behavioral experts to assess risks to user autonomy and ensure any persuasive features are transparent, optional, and non-coercive



Conduct Vulnerability Impact Assessments

Before deployment, assess whether the AI system interacts with or targets potentially vulnerable groups (e.g., elderly, children, low-income users). Implement safeguards such as restricted access, enhanced transparency, and user support features to prevent exploitation based on age, disability, or socio-economic status



Prohibit Cross-Context Data Use in Policy

Adopt internal policies explicitly banning the use of AI systems to evaluate individuals based on behavior or traits across unrelated social domains. Enforce strict data compartmentalization practices and conduct periodic audits to prevent repurposing of personal data for reputational or social classification



Ban Sole Profiling- Based Risk Models

Establish internal rules prohibiting AI models that predict criminal behavior based solely on personality or profiling. Require all criminal risk assessments to be grounded in objective, case-specific evidence and validated by human reviewers with relevant expertise in legal and ethical standards

Conclusion

The prohibition framework under the EU AI Act marks a fundamental shift in how AI systems are designed, developed, and deployed—moving from reactive enforcement to proactive prevention of unacceptable risks. By defining and banning certain practices outright, the AI Act translates its rights-protection mandate into clear, non-negotiable boundaries that eliminate the most harmful AI applications before they reach the market. This approach provides legal certainty for providers and deployers while safeguarding individuals from discriminatory, coercive, or privacy-invasive uses of AI.

Throughout this report, we have outlined how prohibited AI practices intersect with different stages of the AI system lifecycle—highlighting the importance of early detection, ethical design choices, and ongoing operational oversight. When implemented with rigour, these safeguards can help organisations avoid costly violations, reputational harm, and irreversible damage to public trust, while supporting the creation of AI systems that are safe, lawful, and aligned with democratic values.

However, ensuring compliance with the prohibition regime requires more than simply knowing what is banned.

It demands operational readiness, governance structures that embed compliance into every lifecycle stage, and sustained collaboration between technical teams, legal advisors, and ethics specialists. Smaller providers and new entrants may face sharper learning curves, underscoring the need for proactive training, sector-specific guidance, and engagement with regulators and civil society.

Despite these challenges, leading organisations are already demonstrating that building within these boundaries can strengthen product quality, foster innovation in ethical AI design, and differentiate trustworthy providers in a competitive market. As AI adoption accelerates across sectors—from law enforcement and healthcare to education and consumer technology—the prohibition rules offer a powerful mechanism to ensure that innovation proceeds without compromising fundamental rights.

To conclude, consistent and transparent adherence to the AI Act's prohibition framework could not only protect individuals and communities in the EU but also establish a global benchmark for safe, rights-respecting AI governance.

References

Artificial Intelligence Act, (2025), 'Article 5: Prohibited AI Practices', accessible at <https://artificialintelligenceact.eu/article/5/> (last accessed 7th May 2025)

European Parliament, (2025), Digitalisation, artificial intelligence and algorithmic management in the workplace: Shaping the future of work', accessible at: [https://www.europarl.europa.eu/thinktank/en/document/EPRS_STU\(2025\)774670](https://www.europarl.europa.eu/thinktank/en/document/EPRS_STU(2025)774670) (last accessed 25th October 2025)

European Parliament and The Council of the European Union, (2024), 2024/1689 Regulation (EU) 2024/1689 of the European Parliament and of The Council of 13 June 2024 laying down harmonised rules on artificial intelligence and amending Regulations (EC) No 300/2008, (EU) No 167/2013, (EU) No 168/2013, (EU) 2018/858, (EU) 2018/1139 and (EU) 2019/2144 and Directives 2014/90/EU, (EU) 2016/797 and (EU) 2020/1828 (Artificial Intelligence Act), accessible at https://eur-lex.europa.eu/legal-content/EN/TXT/PDF/?uri=OJ:L_202401689 (last accessed 7th May 2025)



Important notice

This document has been prepared by AI & Partners B.V. for the sole purpose of enabling the parties to whom it is addressed to evaluate the capabilities of AI & Partners B.V. to supply the proposed services.

Other than as stated below, this document and its contents are confidential and prepared solely for your information and may not be reproduced, redistributed or passed on to any other person in whole or in part. If this document contains details of an arrangement that could result in a tax or National Insurance saving, no such conditions of confidentiality apply to the details of that arrangement (for example, for the purpose of discussion with tax authorities). No other party is entitled to rely on this document for any purpose whatsoever and we accept no liability to any other party who is shown or obtains access to this document.

This document is not an offer and is not intended to be contractually binding. Should this proposal be acceptable to you, and following the conclusion of our internal acceptance procedures, we would be pleased to discuss terms and conditions with you prior to our appointment. Images used throughout the document have either been produced in-house or sourced from publicly available sources (see **References** for details).

AI & Partners B.V. is the Dutch headquarters of AI & Partners, a global professional services firm. Please see <https://www.ai-and-partners.com/> to learn more about us.

© 2025 AI & Partners B.V. All rights reserved.

Designed and produced by AI & Partners B.V.