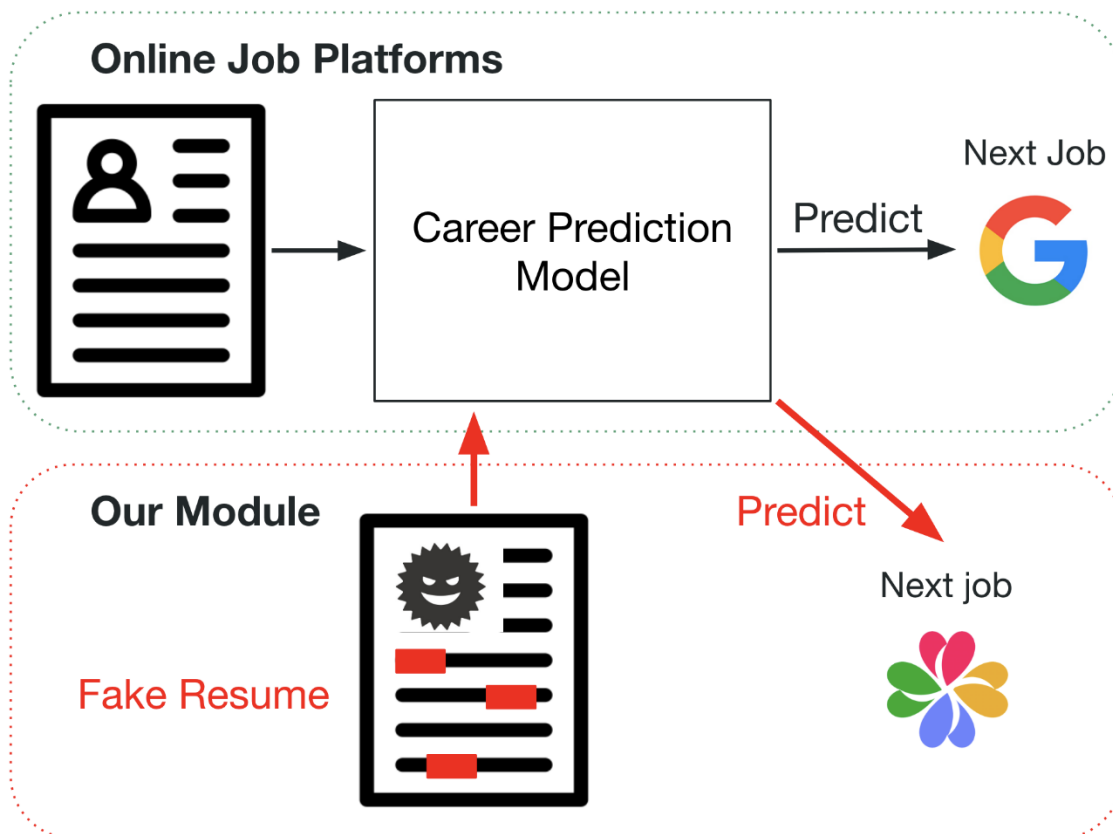


## Fake Resume Classification with Novel ABACUS based LSTM



```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Embedding

from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from datasets import load_dataset
import nltk
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import WordNetLemmatizer
```

```

import re
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
from sklearn.utils import resample
from sklearn.metrics import classification_report, roc_auc_score
from sklearn.model_selection import train_test_split
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Dropout
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from wordcloud import WordCloud

np.random.seed(42)
tf.random.set_seed(42)

sns.set(style="whitegrid")
plt.rcParams['figure.figsize'] = (12, 6)
import matplotlib

nltk.download('punkt')
nltk.download('stopwords')
nltk.download('wordnet')

def clean_text(text):
    if pd.isna(text):
        return ""
    text = text.lower()
    text = re.sub(r'\b[\w\.-]+@[ \w\.-]+\.\w+\b', '', text)
    text = re.sub(r'https?://\S+|www\.\S+', '', text)
    text = re.sub(r'^a-zA-Z\s', '', text)
    text = re.sub(r'\s+', ' ', text).strip()
    tokens = word_tokenize(text)
    stop_words = set(stopwords.words('english'))
    tokens = [token for token in tokens if token not in stop_words]
    lemmatizer = WordNetLemmatizer()
    tokens = [lemmatizer.lemmatize(token) for token in tokens]
    return ' '.join(tokens)

try:
    dataset = load_dataset("AzharAli05/Resume-Screening-Dataset")
    df = dataset['train'].to_pandas()
except Exception as e:
    print(f"Error loading dataset: {e}")
    exit()

df['Resume_clean'] = df['Resume'].apply(clean_text)
df['Job_Description_clean'] = df['Job_Description'].apply(clean_text)
df['Reason_for_decision_clean'] = df['Reason_for_decision'].apply(clean_text)

```

```

print("Original Decision distribution:",
df['Decision'].value_counts(dropna=False).to_dict())
df['Decision'] = df['Decision'].str.lower().map({'select': 1, 'reject':
0}).fillna(0).astype(int)
print("Mapped Decision distribution:",
df['Decision'].value_counts().to_dict())

def flag_fake_resume(reason_clean, resume_clean, decision):
    if not reason_clean or not resume_clean:
        return 0
    fake_keywords = [
        'inconsistent', 'exaggerated', 'unverified', 'false', 'fabricated',
'misleading',
        'overstated', 'questionable', 'unsubstantiated', 'implausible',
'discrepancy',
        'overqualified', 'unrealistic', 'suspicious', 'invalid', 'unreliable'
    ]
    if decision == 0 and any(keyword in reason_clean for keyword in
fake_keywords):
        return 1
    buzzwords = ['expert', 'worldclass', 'leading', 'unparalleled',
'visionary', 'guru']
    resume_words = word_tokenize(resume_clean)
    buzzword_count = sum(resume_words.count(bw) for bw in buzzwords)
    if buzzword_count > 5:
        return 1
    return 0

df['is_fake'] = df.apply(lambda x:
flag_fake_resume(x['Reason_for_decision_clean'], x['Resume_clean'],
x['Decision']), axis=1)

def generate_fake_resume(original_resume, role):
    buzzwords = ['expert', 'worldclass', 'leading', 'unparalleled',
'visionary']
    large_companies = ['tesla', 'google', 'microsoft', 'amazon', 'facebook']
    fake_resume = original_resume + " " + "
".join(np.random.choice(buzzwords, size=10))
    fake_resume += f" chief {role.lower()} at
{np.random.choice(large_companies)} years experience developed aidriven
blockchain solution"
    return clean_text(fake_resume)

n_fake = int(len(df) * 0.5 / (1 - 0.5)) - df['is_fake'].sum()
if n_fake > 0:
    fake_resumes = df[['Resume_clean', 'Role']].sample(n_fake,
random_state=42, replace=True)
    fake_resumes['Resume_clean'] = fake_resumes.apply(lambda x:

```

```

generate_fake_resume(x['Resume_clean'], x['Role']), axis=1)
fake_df = pd.DataFrame({
    'Role': fake_resumes['Role'],
    'Resume': fake_resumes['Resume_clean'],
    'Resume_clean': fake_resumes['Resume_clean'],
    'Decision': 0,
    'Reason_for_decision': 'Generated fake resume with exaggerated
claims',
    'Reason_for_decision_clean': 'generated fake resume exaggerated
claim',
    'Job_Description': df['Job_Description'].sample(n_fake,
random_state=42, replace=True),
    'Job_Description_clean': df['Job_Description_clean'].sample(n_fake,
random_state=42, replace=True),
    'is_fake': 1
})
df = pd.concat([df, fake_df], ignore_index=True)

tfidf = TfidfVectorizer(max_features=5000, stop_words='english')
resume_tfidf = tfidf.fit_transform(df['Resume_clean'])
job_tfidf = tfidf.transform(df['Job_Description_clean'])
similarity_scores = [cosine_similarity(resume_tfidf[i], job_tfidf[i])[0][0]
for i in range(len(df))]
df['similarity_score'] = similarity_scores
df.loc[df['similarity_score'] < 0.25, 'is_fake'] = 1

max_words = 5000
max_len = 50
beads_per_column = 5
tokenizer = Tokenizer(num_words=max_words)
tokenizer.fit_on_texts(df['Resume_clean'] + ' ' +
df['Job_Description_clean'])
X_sequences = tokenizer.texts_to_sequences(df['Resume_clean'] + ' ' +
df['Job_Description_clean'])
X_padded = pad_sequences(X_sequences, maxlen=max_len, padding='post',
truncating='post')

tfidf_abacus = TfidfVectorizer(max_features=max_words, stop_words='english')
X_tfidf = tfidf_abacus.fit_transform(df['Resume_clean'] + ' ' +
df['Job_Description_clean']).toarray()
X_abacus = np.zeros((X_padded.shape[0], max_len, beads_per_column))
for i in range(X_padded.shape[0]):
    for t in range(max_len):
        word_idx = X_padded[i, t]
        if word_idx > 0:
            word = tokenizer.index_word.get(word_idx, '')
            if word in tfidf_abacus.vocabulary_:
                tfidf_idx = tfidf_abacus.vocabulary_[word]
                bead_value = X_tfidf[i, tfidf_idx]
                bead_idx = min(int(bead_value * beads_per_column),

```

```

beads_per_column-1)
        X_abacus[i, t, bead_idx] = 1

y_fake = df['is_fake'].values
y_decision = df['Decision'].values

df_fake_0 = df[df['is_fake'] == 0]
df_fake_1 = df[df['is_fake'] == 1]
if len(df_fake_0) == 0 or len(df_fake_1) == 0:
    print("Warning: One class missing in is_fake. Skipping balancing for
is_fake.")
    X_fake_train, X_fake_test, y_fake_train, y_fake_test = train_test_split(
        X_abacus, y_fake, test_size=0.2, random_state=42
    )
else:
    n_majority = max(len(df_fake_0), len(df_fake_1))
    df_fake_0_resampled = resample(df_fake_0, replace=True,
n_samples=n_majority, random_state=42) if len(df_fake_0) < n_majority else
df_fake_0
    df_fake_1_resampled = resample(df_fake_1, replace=True,
n_samples=n_majority, random_state=42) if len(df_fake_1) < n_majority else
df_fake_1
    df_fake_balanced = pd.concat([df_fake_0_resampled, df_fake_1_resampled])
    X_fake_sequences =
tokenizer.texts_to_sequences(df_fake_balanced['Resume_clean'] + ' ' +
df_fake_balanced['Job_Description_clean'])
    X_fake_padded = pad_sequences(X_fake_sequences, maxlen=max_len,
padding='post', truncating='post')
    X_fake_tfidf = tfidf_abacus.transform(df_fake_balanced['Resume_clean'] +
' ' + df_fake_balanced['Job_Description_clean']).toarray()
    X_fake_abacus = np.zeros((X_fake_padded.shape[0], max_len,
beads_per_column))
    for i in range(X_fake_padded.shape[0]):
        for t in range(max_len):
            word_idx = X_fake_padded[i, t]
            if word_idx > 0:
                word = tokenizer.index_word.get(word_idx, '')
                if word in tfidf_abacus.vocabulary_:
                    tfidf_idx = tfidf_abacus.vocabulary_[word]
                    bead_value = X_fake_tfidf[i, tfidf_idx]
                    bead_idx = min(int(bead_value * beads_per_column),
beads_per_column-1)
                    X_fake_abacus[i, t, bead_idx] = 1
            y_fake_balanced = df_fake_balanced['is_fake'].values
            X_fake_train, X_fake_test, y_fake_train, y_fake_test = train_test_split(
                X_fake_abacus, y_fake_balanced, test_size=0.2,
stratify=y_fake_balanced, random_state=42
            )

df_dec_0 = df[df['Decision'] == 0]

```

```

df_dec_1 = df[df['Decision'] == 1]
if len(df_dec_0) == 0 or len(df_dec_1) == 0:
    print("Warning: One class missing in Decision. Skipping balancing for Decision.")
    X_dec_train, X_dec_test, y_dec_train, y_dec_test = train_test_split(
        X_abacus, y_decision, test_size=0.2, random_state=42
    )
else:
    n_majority_dec = max(len(df_dec_0), len(df_dec_1))
    df_dec_0_resampled = resample(df_dec_0, replace=True,
n_samples=n_majority_dec, random_state=42) if len(df_dec_0) < n_majority_dec
    else df_dec_0
    df_dec_1_resampled = resample(df_dec_1, replace=True,
n_samples=n_majority_dec, random_state=42) if len(df_dec_1) < n_majority_dec
    else df_dec_1
    df_dec_balanced = pd.concat([df_dec_0_resampled, df_dec_1_resampled])
    X_dec_sequences =
tokenizer.texts_to_sequences(df_dec_balanced['Resume_clean'] + ' ' +
df_dec_balanced['Job_Description_clean'])
    X_dec_padded = pad_sequences(X_dec_sequences, maxlen=max_len,
padding='post', truncating='post')
    X_dec_tfidf = tfidf_abacus.transform(df_dec_balanced['Resume_clean'] + '
' + df_dec_balanced['Job_Description_clean']).toarray()
    X_dec_abacus = np.zeros((X_dec_padded.shape[0], max_len,
beads_per_column))
    for i in range(X_dec_padded.shape[0]):
        for t in range(max_len):
            word_idx = X_dec_padded[i, t]
            if word_idx > 0:
                word = tokenizer.index_word.get(word_idx, '')
                if word in tfidf_abacus.vocabulary_:
                    tfidf_idx = tfidf_abacus.vocabulary_[word]
                    bead_value = X_dec_tfidf[i, tfidf_idx]
                    bead_idx = min(int(bead_value * beads_per_column),
beads_per_column-1)
                    X_dec_abacus[i, t, bead_idx] = 1
    y_dec_balanced = df_dec_balanced['Decision'].values
    X_dec_train, X_dec_test, y_dec_train, y_dec_test = train_test_split(
        X_dec_abacus, y_dec_balanced, test_size=0.2, stratify=y_dec_balanced,
random_state=42
    )

def build_abacus_network(max_len, beads_per_column):
    model = Sequential([
        LSTM(128, input_shape=(max_len, beads_per_column),
return_sequences=True),
        Dropout(0.3),
        LSTM(64),
        Dropout(0.3),
        Dense(1, activation='sigmoid')
    ])

```

```

    ])
    model.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
    return model

if len(np.unique(y_fake_train)) > 1:
    model_fake = build_abacus_network(max_len, beads_per_column)
    model_fake.fit(X_fake_train, y_fake_train, epochs=20, batch_size=32,
validation_split=0.2, verbose=1)
else:
    print("Skipping training for is_fake: Only one class present.")

if len(np.unique(y_dec_train)) > 1:
    model_decision = build_abacus_network(max_len, beads_per_column)
    model_decision.fit(X_dec_train, y_dec_train, epochs=20, batch_size=32,
validation_split=0.2, verbose=1)
else:
    print("Skipping training for Decision: Only one class present.")

if len(np.unique(y_fake_train)) > 1:
    y_fake_pred = (model_fake.predict(X_fake_test) > 0.5).astype(int)
    print("Fake Resume Classification Report:")
    print(classification_report(y_fake_test, y_fake_pred))
    if len(np.unique(y_fake_test)) > 1:
        print("ROC-AUC (Fake):", roc_auc_score(y_fake_test,
model_fake.predict(X_fake_test)))
    else:
        print("ROC-AUC (Fake): Skipped due to single class in test set")
else:
    print("No evaluation for is_fake: Model not trained.")

if len(np.unique(y_dec_train)) > 1:
    y_dec_pred = (model_decision.predict(X_dec_test) > 0.5).astype(int)
    print("Selection/Rejection Classification Report:")
    print(classification_report(y_dec_test, y_dec_pred))
    if len(np.unique(y_dec_test)) > 1:
        print("ROC-AUC (Decision):", roc_auc_score(y_dec_test,
model_decision.predict(X_dec_test)))
    else:
        print("ROC-AUC (Decision): Skipped due to single class in test set")
else:
    print("No evaluation for Decision: Model not trained.")

def visualize_abacus_words(features, title, tokenizer, sample_idx, X_padded):
    plt.figure(figsize=(12, 4))
    bead_positions = features
    word_labels = []
    for t in range(min(10, features.shape[0])):
        word_idx = X_padded[sample_idx, t]

```

```

word = tokenizer.index_word.get(word_idx, 'PAD') if word_idx > 0 else
'PAD'
word_labels.append(word)
for row in range(features.shape[1]):
    color = 'red' if bead_positions[t, row] else 'gray'
    plt.scatter(t, row, c=color, s=100)
plt.xticks(range(min(10, features.shape[0])), word_labels, rotation=45,
ha='right')
plt.yticks(range(features.shape[1]), [f"Bead {i+1}" for i in
range(features.shape[1])])
plt.title(title)
plt.grid(True)
plt.tight_layout()
plt.show()

sample_idx = 0
visualize_abacus_words(X_abacus[sample_idx], f"Abacus Word Activation (Fake:
{y_fake[sample_idx]})", tokenizer, sample_idx, X_padded)
visualize_abacus_words(X_abacus[sample_idx], f"Abacus Word Activation
(Decision: {y_decision[sample_idx]})", tokenizer, sample_idx, X_padded)

top_roles = df['Role'].value_counts().nlargest(10).index
df['Role_plot'] = df['Role'].where(df['Role'].isin(top_roles), 'Other')
top_reasons = df['Reason_for_decision'].value_counts().nlargest(10).index
df['Reason_plot'] =
df['Reason_for_decision'].where(df['Reason_for_decision'].isin(top_reasons),
'Other')

plt.figure()
sns.countplot(data=df, x='Role_plot', hue='is_fake',
order=top_roles.append(pd.Index(['Other'])))
plt.title('Count of Real vs. Fake Resumes by Role (Top 10 + Other)')
plt.xticks(rotation=45, ha='right')
plt.xlabel('Role')
plt.ylabel('Count')
plt.tight_layout()
plt.show()

plt.figure()
sns.countplot(data=df, x='Decision', hue='is_fake')
plt.title('Count of Real vs. Fake Resumes by Decision')
plt.xlabel('Decision')
plt.ylabel('Count')
plt.tight_layout()
plt.show()

plt.figure()
sns.countplot(data=df, x='Reason_plot', hue='is_fake',
order=top_reasons.append(pd.Index(['Other'])))

```



```
plt.title('Count of Real vs. Fake Resumes by Reason for Decision (Top 10 + Other)')
plt.xticks(rotation=45, ha='right')
plt.xlabel('Reason for Decision')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

```
plt.figure()
sns.countplot(data=df, x='is_fake')
plt.title('Count of Real vs. Fake Resumes')
plt.xlabel('Is Fake (0 = Real, 1 = Fake)')
plt.ylabel('Count')
plt.tight_layout()
plt.show()
```

```
plt.figure()
df['Decision'].value_counts().plot.pie(autopct='%1.1f%%', colors=['#4CAF50', '#F44336'], labels=['Reject', 'Select'])
plt.title('Proportion of Select vs. Reject Decisions')
plt.ylabel('')
plt.tight_layout()
plt.show()
```

```
plt.figure()
df['is_fake'].value_counts().plot.pie(autopct='%1.1f%%', labels=['Real', 'Fake'], colors=['#4CAF50', '#F44336'])
plt.title('Proportion of Real vs. Fake Resumes')
plt.ylabel('')
plt.tight_layout()
plt.show()
```

```
plt.figure()
df['Role_plot'].value_counts().nlargest(5).plot.pie(autopct='%1.1f%%')
plt.title('Proportion of Top 5 Roles')
plt.ylabel('')
plt.tight_layout()
plt.show()
```

```
plt.figure()
df['Reason_plot'].value_counts().nlargest(5).plot.pie(autopct='%1.1f%%')
plt.title('Proportion of Top 5 Reasons for Decision')
plt.ylabel('')
plt.tight_layout()
plt.show()
```

```
df['resume_word_count'] = df['Resume_clean'].apply(lambda x: len(word_tokenize(str(x))))
df['job_word_count'] = df['Job_Description_clean'].apply(lambda x:
```

```

len(word_tokenize(str(x))))

plt.figure()
sns.histplot(data=df, x='resume_word_count', hue='is_fake', bins=50)
plt.title('Resume Word Count Distribution by Real/Fake')
plt.xlabel('Word Count')
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()

plt.figure()
sns.histplot(data=df, x='job_word_count', hue='is_fake', bins=50)
plt.title('Job Description Word Count Distribution by Real/Fake')
plt.xlabel('Word Count')
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()

def generate_wordcloud(text, title):
    wordcloud = WordCloud(width=800, height=400, background_color='white',
max_words=100).generate(text)
    plt.figure()
    plt.imshow(wordcloud, interpolation='bilinear')
    plt.title(title)
    plt.axis('off')
    plt.tight_layout()
    plt.show()

real_resumes = ' '.join(df[df['is_fake'] == 0]['Resume_clean'])
fake_resumes = ' '.join(df[df['is_fake'] == 1]['Resume_clean'])
generate_wordcloud(real_resumes, 'Word Cloud for Real Resumes')
generate_wordcloud(fake_resumes, 'Word Cloud for Fake Resumes')

real_jobs = ' '.join(df[df['is_fake'] == 0]['Job_Description_clean'])
fake_jobs = ' '.join(df[df['is_fake'] == 1]['Job_Description_clean'])
generate_wordcloud(real_jobs, 'Word Cloud for Real Job Descriptions')
generate_wordcloud(fake_jobs, 'Word Cloud for Fake Job Descriptions')

real_reasons = ' '.join(df[df['is_fake'] == 0]['Reason_for_decision_clean'])
fake_reasons = ' '.join(df[df['is_fake'] == 1]['Reason_for_decision_clean'])
generate_wordcloud(real_reasons, 'Word Cloud for Real Reasons for Decision')
generate_wordcloud(fake_reasons, 'Word Cloud for Fake Reasons for Decision')

# Step 13: Print diagnostics
print("is_fake distribution:", df['is_fake'].value_counts().to_dict())
print("Decision distribution:", df['Decision'].value_counts().to_dict())
print("Top 10 Roles:", df['Role'].value_counts().head(10).to_dict())
print("Top 10 Reasons:",
df['Reason_for_decision'].value_counts().head(10).to_dict())

```

```

print("is_fake test distribution:",
pd.Series(y_fake_test).value_counts().to_dict())
if len(np.unique(y_dec_test)) > 0:
    print("Decision test distribution:",
pd.Series(y_dec_test).value_counts().to_dict())
else:
    print("Decision test distribution: Empty due to missing classes")

```

```

[nltk_data] Downloading package punkt to /usr/share/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /usr/share/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /usr/share/nltk_data...
[nltk_data] Package wordnet is already up-to-date!

```

```

Original Decision distribution: {'reject': 5114, 'select': 5060}
Mapped Decision distribution: {0: 5114, 1: 5060}
Epoch 1/20

```

```

/usr/local/lib/python3.11/dist-packages/keras/src/layers/rnn/rnn.py:200:
UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer.
When using Sequential models, prefer using an `Input(shape)` object as the
first layer in the model instead.

```

```

    super().__init__(**kwargs)

```

```

702/702 ————— 52s 66ms/step - accuracy: 0.5519 - loss: 0.6841
- val_accuracy: 0.5743 - val_loss: 0.6796

```

```

Epoch 2/20

```

```

702/702 ————— 47s 67ms/step - accuracy: 0.5776 - loss: 0.6778
- val_accuracy: 0.6177 - val_loss: 0.6603

```

```

Epoch 3/20

```

```

702/702 ————— 47s 66ms/step - accuracy: 0.6342 - loss: 0.6529
- val_accuracy: 0.6517 - val_loss: 0.6328

```

```

Epoch 4/20

```

```

702/702 ————— 47s 66ms/step - accuracy: 0.6315 - loss: 0.6488
- val_accuracy: 0.6405 - val_loss: 0.6281

```

```

Epoch 5/20

```

```

702/702 ————— 47s 67ms/step - accuracy: 0.7123 - loss: 0.5661
- val_accuracy: 0.7660 - val_loss: 0.4872

```

```

Epoch 6/20

```

```

702/702 ————— 46s 66ms/step - accuracy: 0.7649 - loss: 0.4913
- val_accuracy: 0.7876 - val_loss: 0.4657

```

```

Epoch 7/20

```

```

702/702 ————— 47s 67ms/step - accuracy: 0.7786 - loss: 0.4799
- val_accuracy: 0.7933 - val_loss: 0.4554

```

```

Epoch 8/20

```

```

702/702 ————— 47s 66ms/step - accuracy: 0.7807 - loss: 0.4681
- val_accuracy: 0.7853 - val_loss: 0.4709

```

```

Epoch 9/20

```

```

702/702 ————— 46s 66ms/step - accuracy: 0.7784 - loss: 0.4687
- val_accuracy: 0.7936 - val_loss: 0.4449

```

Epoch 10/20  
702/702 \_\_\_\_\_ 46s 65ms/step - accuracy: 0.7895 - loss: 0.4556  
- val\_accuracy: 0.8079 - val\_loss: 0.4289  
Epoch 11/20  
702/702 \_\_\_\_\_ 46s 65ms/step - accuracy: 0.7968 - loss: 0.4417  
- val\_accuracy: 0.8034 - val\_loss: 0.4281  
Epoch 12/20  
702/702 \_\_\_\_\_ 46s 66ms/step - accuracy: 0.8037 - loss: 0.4349  
- val\_accuracy: 0.8134 - val\_loss: 0.4186  
Epoch 13/20  
702/702 \_\_\_\_\_ 46s 65ms/step - accuracy: 0.8003 - loss: 0.4462  
- val\_accuracy: 0.8212 - val\_loss: 0.4068  
Epoch 14/20  
702/702 \_\_\_\_\_ 46s 65ms/step - accuracy: 0.8117 - loss: 0.4169  
- val\_accuracy: 0.8253 - val\_loss: 0.4015  
Epoch 15/20  
702/702 \_\_\_\_\_ 46s 65ms/step - accuracy: 0.8130 - loss: 0.4140  
- val\_accuracy: 0.8264 - val\_loss: 0.3991  
Epoch 16/20  
702/702 \_\_\_\_\_ 45s 65ms/step - accuracy: 0.8201 - loss: 0.4053  
- val\_accuracy: 0.8203 - val\_loss: 0.4010  
Epoch 17/20  
702/702 \_\_\_\_\_ 46s 65ms/step - accuracy: 0.8243 - loss: 0.3977  
- val\_accuracy: 0.8262 - val\_loss: 0.3906  
Epoch 18/20  
702/702 \_\_\_\_\_ 45s 65ms/step - accuracy: 0.8304 - loss: 0.3911  
- val\_accuracy: 0.8244 - val\_loss: 0.3916  
Epoch 19/20  
702/702 \_\_\_\_\_ 46s 65ms/step - accuracy: 0.8307 - loss: 0.3828  
- val\_accuracy: 0.8321 - val\_loss: 0.3760  
Epoch 20/20  
702/702 \_\_\_\_\_ 45s 65ms/step - accuracy: 0.8338 - loss: 0.3769  
- val\_accuracy: 0.8283 - val\_loss: 0.3808  
Epoch 1/20  
612/612 \_\_\_\_\_ 45s 66ms/step - accuracy: 0.5014 - loss: 0.6943  
- val\_accuracy: 0.5094 - val\_loss: 0.6932  
Epoch 2/20  
612/612 \_\_\_\_\_ 40s 65ms/step - accuracy: 0.5062 - loss: 0.6933  
- val\_accuracy: 0.5061 - val\_loss: 0.6931  
Epoch 3/20  
612/612 \_\_\_\_\_ 39s 64ms/step - accuracy: 0.5059 - loss: 0.6932  
- val\_accuracy: 0.5061 - val\_loss: 0.6930  
Epoch 4/20  
612/612 \_\_\_\_\_ 39s 64ms/step - accuracy: 0.5044 - loss: 0.6931  
- val\_accuracy: 0.5123 - val\_loss: 0.6930  
Epoch 5/20  
612/612 \_\_\_\_\_ 39s 64ms/step - accuracy: 0.5072 - loss: 0.6933  
- val\_accuracy: 0.5102 - val\_loss: 0.6925  
Epoch 6/20  
612/612 \_\_\_\_\_ 40s 65ms/step - accuracy: 0.5137 - loss: 0.6928

- val\_accuracy: 0.5117 - val\_loss: 0.6926  
Epoch 7/20  
612/612 \_\_\_\_\_ 39s 64ms/step - accuracy: 0.5802 - loss: 0.6595  
- val\_accuracy: 0.7500 - val\_loss: 0.5571  
Epoch 8/20  
612/612 \_\_\_\_\_ 39s 64ms/step - accuracy: 0.7500 - loss: 0.5597  
- val\_accuracy: 0.7537 - val\_loss: 0.5502  
Epoch 9/20  
612/612 \_\_\_\_\_ 39s 64ms/step - accuracy: 0.7550 - loss: 0.5505  
- val\_accuracy: 0.7553 - val\_loss: 0.5420  
Epoch 10/20  
612/612 \_\_\_\_\_ 39s 64ms/step - accuracy: 0.7587 - loss: 0.5434  
- val\_accuracy: 0.7535 - val\_loss: 0.5425  
Epoch 11/20  
612/612 \_\_\_\_\_ 40s 65ms/step - accuracy: 0.7575 - loss: 0.5398  
- val\_accuracy: 0.7533 - val\_loss: 0.5409  
Epoch 12/20  
612/612 \_\_\_\_\_ 39s 64ms/step - accuracy: 0.7573 - loss: 0.5392  
- val\_accuracy: 0.7545 - val\_loss: 0.5440  
Epoch 13/20  
612/612 \_\_\_\_\_ 39s 64ms/step - accuracy: 0.7581 - loss: 0.5372  
- val\_accuracy: 0.7520 - val\_loss: 0.5451  
Epoch 14/20  
612/612 \_\_\_\_\_ 39s 64ms/step - accuracy: 0.7573 - loss: 0.5358  
- val\_accuracy: 0.7565 - val\_loss: 0.5394  
Epoch 15/20  
612/612 \_\_\_\_\_ 40s 65ms/step - accuracy: 0.7587 - loss: 0.5346  
- val\_accuracy: 0.7510 - val\_loss: 0.5434  
Epoch 16/20  
612/612 \_\_\_\_\_ 40s 65ms/step - accuracy: 0.7595 - loss: 0.5333  
- val\_accuracy: 0.7531 - val\_loss: 0.5441  
Epoch 17/20  
612/612 \_\_\_\_\_ 41s 67ms/step - accuracy: 0.7593 - loss: 0.5312  
- val\_accuracy: 0.7522 - val\_loss: 0.5455  
Epoch 18/20  
612/612 \_\_\_\_\_ 41s 67ms/step - accuracy: 0.7602 - loss: 0.5294  
- val\_accuracy: 0.7543 - val\_loss: 0.5440  
Epoch 19/20  
612/612 \_\_\_\_\_ 40s 66ms/step - accuracy: 0.7598 - loss: 0.5287  
- val\_accuracy: 0.7543 - val\_loss: 0.5407  
Epoch 20/20  
612/612 \_\_\_\_\_ 40s 66ms/step - accuracy: 0.7593 - loss: 0.5276  
- val\_accuracy: 0.7555 - val\_loss: 0.5385  
220/220 \_\_\_\_\_ 6s 24ms/step

Fake Resume Classification Report:

	precision	recall	f1-score	support
0	0.81	0.87	0.84	3510
1	0.86	0.80	0.83	3510

accuracy			0.83	7020
macro avg	0.83	0.83	0.83	7020
weighted avg	0.83	0.83	0.83	7020

220/220 \_\_\_\_\_ 5s 22ms/step

ROC-AUC (Fake): 0.9111811186597512

192/192 \_\_\_\_\_ 5s 23ms/step


Selection/Rejection Classification Report:



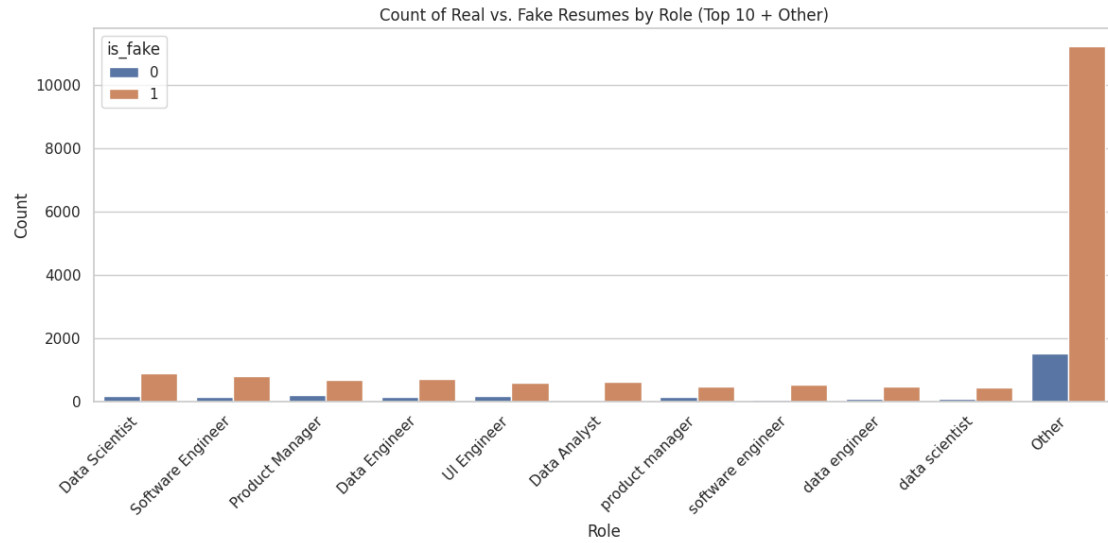
	precision	recall	f1-score	support
0	0.84	0.66	0.74	3058
1	0.72	0.88	0.79	3057

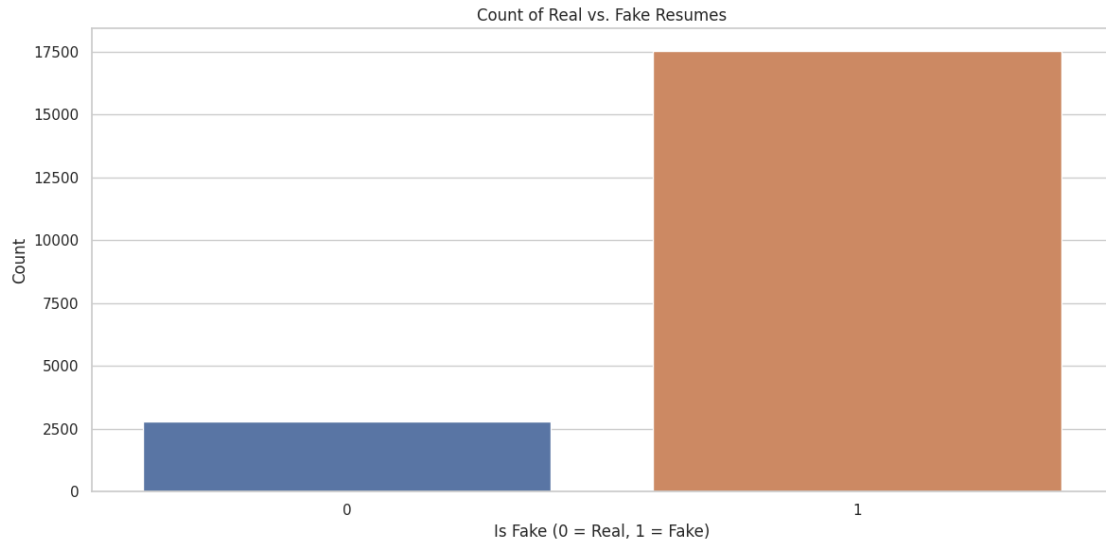
accuracy			0.77	6115
macro avg	0.78	0.77	0.76	6115
weighted avg	0.78	0.77	0.76	6115

192/192 \_\_\_\_\_ 4s 21ms/step


ROC-AUC (Decision): 0.8079580407402155








Proportion of Select vs. Reject Decisions






Proportion of Real vs. Fake Resumes




Proportion of Top 5 Roles



Proportion of Top 5 Reasons for Decision


Generated fake resume with exaggerated claims



```

/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1075:
FutureWarning: When grouping with a length-1 list-like, you will need to pass
a length-1 tuple to get_group in a future version of pandas. Pass `(name,)`
instead of `name` to silence this warning.
    data_subset = grouped_data.get_group(pd_key)
/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1075:
FutureWarning: When grouping with a length-1 list-like, you will need to pass
a length-1 tuple to get_group in a future version of pandas. Pass `(name,)`
instead of `name` to silence this warning.
    data_subset = grouped_data.get_group(pd_key)
/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1075:
FutureWarning: When grouping with a length-1 list-like, you will need to pass
a length-1 tuple to get_group in a future version of pandas. Pass `(name,)`
instead of `name` to silence this warning.
    data_subset = grouped_data.get_group(pd_key)

```



```

/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1119:
FutureWarning: use_inf_as_na option is deprecated and will be removed in a
future version. Convert inf values to NaN before operating instead.
    with pd.option_context('mode.use_inf_as_na', True):
/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1075:
FutureWarning: When grouping with a length-1 list-like, you will need to pass
a length-1 tuple to get_group in a future version of pandas. Pass `(name,)`
instead of `name` to silence this warning.
    data_subset = grouped_data.get_group(pd_key)
/usr/local/lib/python3.11/dist-packages/seaborn/_oldcore.py:1075:
FutureWarning: When grouping with a length-1 list-like, you will need to pass
a length-1 tuple to get_group in a future version of pandas. Pass `(name,)`
instead of `name` to silence this warning.
    data_subset = grouped_data.get_group(pd_key)

```



[illegible]



[illegible][illegible]

[illegible]

Decision distribution: {0: 15286, 1: 5060}

```
Top 10 Reasons: {'Generated fake resume with exaggerated claims': 10172,
'Insufficient system design expertise for senior role.': 730, 'Solid
experience in machine learning and AI.': 728, 'Strong technical skills in AI
and ML.': 718, 'No experience in back-end development.': 717, 'Needs
improvement in machine learning algorithms.': 707, 'Impressive leadership and
communication abilities.': 704, 'Lacked leadership skills for a senior
position.': 685, 'Excellent full-stack development experience.': 677,
'Perfectly aligned with data engineering needs.': 668}
```

Decision test distribution: {0: 3058, 1: 3057}