




AI & Partners

Amsterdam - London - Singapore



EU AI Act

Providers of General-
Purpose AI Models:
A Breakdown

December 2025


Authors

Sean Musch, Michael Borrelli, Charles Kerrigan, Alisar
Mustafa, Eevamaija Virtanen, Tulsi Soni, Diane Tea



Generative AI's Acceleration Curve: Text Models Dominate, Image Systems Mature, and Audio Applications Emerge — Redefining the Competitive Landscape Across AI Modalities

*OECD.AI: Generative AI Models by Modality
(October 2025)*



Between January 2023 and April 2025, the number of generative AI models expanded dramatically across all modalities—most notably in text-to-text systems, which grew from six to 616 models. This trajectory underscores an unprecedented acceleration in language-based AI innovation, driven by foundation model scaling, open-weight releases, and fine-tuning democratization.

Text-to-image models experienced early volatility but maintained steady expansion, rising from seven to 187 models, reflecting the maturing diffusion model ecosystem and commercial diversification in creative and marketing applications. By contrast, audio-to-text systems grew more modestly, suggesting that transcription and speech-to-text have reached functional maturity with fewer differentiated entrants. Text-to-audio models emerged late in the dataset, indicating the early stages of a new growth curve as generative speech, music, and multimodal interfaces gain traction.

Critically, the proliferation of models signals both opportunity and fragmentation. Rapid diffusion enhances accessibility but risks redundancy, uneven quality, and regulatory complexity—particularly as models blur jurisdictional and ethical boundaries. Constructively, this trend calls for disciplined model governance, investment in evaluation standards, and ecosystem coordination to sustain innovation without diluting reliability. In sum, generative AI's evolution is shifting from breakthrough creation to competitive scaling. Organizations that integrate governance, interoperability, and responsible deployment will capture strategic advantage in an increasingly crowded landscape.



AI & Partners

Amsterdam - London - Singapore

Who Are We

AI That You Can Trust

Why Us?

Stay on the right side of history. At AI & Partners, we believe AI should unlock potential—not cause harm. We've seen the fear and fallout when teams lose control of AI, but also the trust and innovation that follow when it's handled responsibly. That's why we exist: to help you build AI you can trust and stand behind—for the long run.

80%

of AI systems
are unknown


What Do We Do?

We enable safe AI usage—for your organization and your clients. Unknown AI adoption leads to confusion, risk, and reputational damage. We help you take control with tools to identify, monitor, and govern all AI systems—so you're not reacting to AI, you're leading it.

How Do We Do It?

Do you know what AI systems you have? Identify all known and unknown AI systems (algorithms, LLMs, prompts, and models) from all internal and external AI vendors, automated by generating your inventory. Overall, 80% of AI inventory is unknown to our clients.

How do you guarantee ongoing safe AI use? Continuously monitor deployed AI systems for performance drift, anomalies or failures, real-world impacts, and emerging risks (e.g. data poisoning). Any malfunction of an AI system has severe implications for organisations (e.g. inability to assess online misinformation that leads to widespread public mistrust), so monitoring becomes a matter of urgency.



AI Discovery & AI Inventory

Automatically detect all AI systems, including models, algorithms, and prompts, and maintain a live, always-updated register for full visibility and compliance.



Responsible AI

Embed fairness, transparency, and control into every stage of AI use—aligning with the EU AI Act and building 'Trustworthy-by-Design'.



Model Monitoring

Continuously track your AI models after deployment to detect drift, bias, or failure—so you stay in control and prevent harm before it happens.

Contact: Team | contact@ai-and-partners.com

Contents

Who Are We	2
Introduction	4
Key questions being asked about General-Purpose AI Models.....	5
1. What qualifies as a general-purpose AI (GPAI) model under the AI Act?	6
2. Who is considered the provider of a GPAI model under the AI Act?.....	6
3. What obligations apply to providers of GPAI models with systemic risk?	6
4. How is “placing a GPAI model on the market” defined?	6
5. When do open-source GPAI model exemptions apply?	6
6. How is compute calculated to determine GPAI model status?	7
7. What happens if a downstream actor modifies a GPAI model significantly?	7
8. What transitional rules exist for GPAI model providers?	7
9. How does signing the Code of Practice affect a GPAI provider’s obligations?	7
10. How will the rules for GPAI models be enforced in practice?	7
Understanding General-Purpose AI Models.....	8
1. Understanding GPAI Models	9
2. Provider Obligations and Systemic Risk.....	10
3. Open-Source and Compute Requirements.....	11
4. Downstream Modification and Legal Impact.....	12
5. Transition, Compliance, and Enforcement	13
Mapping General-Purpose AI Models to Ecosystem Actors	14
Calls to Action.....	0
Conclusion.....	2
Appendix: Additional Considerations for GPAI Compliance	3
References.....	5

Introduction

As the European Union's Artificial Intelligence Act (AI Act) advances toward full enforcement, the regulation of General-Purpose AI (GPAI) models has emerged as one of its most transformative components. These models—capable of performing a wide range of tasks and serving as the foundation for countless downstream AI systems—sit at the core of Europe's efforts to ensure safe, transparent, and accountable AI. By introducing specific obligations for GPAI providers, particularly those whose models pose systemic risk, the AI Act establishes a governance framework designed to balance innovation with responsibility.

Managed under the oversight of the EU AI Office, this framework ensures that the development, deployment, and modification of GPAI models adhere to high standards of safety, cybersecurity, and legal compliance. It provides a structured mechanism for documentation, transparency, and ongoing risk assessment—ensuring that the entities behind these models remain accountable for their real-world impact.

This report offers a comprehensive guide to understanding the obligations, roles, and responsibilities of GPAI model providers under the AI Act. It explains key definitions, thresholds, and compliance mechanisms, including the treatment of open-source models, compute-based classification, and reclassification rules for modified systems. Through clear, actionable insights, it equips AI developers, compliance officers, legal practitioners, and policymakers with the knowledge needed to navigate this evolving regulatory landscape.


Best regards,

Sean Musch

Founder/CEO

AI & Partners

A handwritten signature in white ink, appearing to read 'Sean Musch', is positioned to the right of the text 'Founder/CEO'.

A stylized, glowing green brain composed of circuitry and data lines, set against a dark background with a vertical beam of light. The brain is depicted with intricate circuit patterns and glowing nodes, suggesting a digital or artificial intelligence theme. A bright vertical beam of light passes through the center of the brain, and numerous smaller light points are scattered throughout the circuitry.

Key questions
being
asked about
General-Purpose
AI Models

1. What qualifies as a general-purpose AI (GPAI) model under the AI Act?

A GPAI model is defined as an AI model that displays significant generality and can competently perform a wide range of distinct tasks. It can be integrated into various downstream systems or applications. The release method (e.g., open weights or API) does not affect this classification unless the model is solely used for research or prototyping before being placed on the market. The AI Office proposes a compute-based threshold to classify a model as GPAI: models trained using over 10^{22} FLOP (floating point operations) are presumed to be GPAI models under this definition.

2. Who is considered the provider of a GPAI model under the AI Act?

A provider is any entity (natural or legal) that develops or commissions the development of a GPAI model and places it on the market under its own name or trademark. This applies whether the model is offered for free or for payment. Providers can be original developers or downstream actors who significantly modify existing models. If the compute used in the modification exceeds one-third (i.e., 3×10^{21} FLOP) of the GPAI model threshold, the downstream actor is presumed to become a provider, subject to the same obligations as the original creator.

3. What obligations apply to providers of GPAI models with systemic risk?

GPAI model providers with systemic risk (GPAISR) must comply with Articles 53 and 55 of the AI Act. These obligations include maintaining model documentation, implementing copyright policies, conducting model evaluations, adversarial testing, and ensuring cybersecurity protections. Providers must also track and report serious incidents. These obligations are more stringent due to the heightened societal impact of GPAISR models. The AI Office assumes currently that no downstream modifications substantially change systemic risk, but this could change based on future evidence and compute usage thresholds defined in the guidelines.

4. How is “placing a GPAI model on the market” defined?

Placing a GPAI model on the market includes making the model publicly available via software libraries, APIs, cloud services, or similar platforms. This definition applies regardless of whether the model is released commercially or for free. The guidelines emphasize that open-source releases also count as placement on the market unless they fall under specific exemptions. The goal is to ensure that distribution mechanisms do not circumvent compliance obligations, especially when models are capable of wide usage and integration across multiple downstream applications or services.

5. When do open-source GPAI model exemptions apply?

Open-source exemptions apply under narrow conditions where the model is released for unrestricted access, usage, modification, and redistribution under open-source terms. However, GPAI models with systemic risk are explicitly excluded from these exemptions regardless of how they are released. The AI Office defines terms like ‘free’, ‘open-source’, and ‘modification’ to clarify the boundary of exemptions. Only models released without significant constraints and without systemic risk may benefit from open-source relief, ensuring transparency and safety without placing undue burden on genuine open-source contributors.

6. How is compute calculated to determine GPAI model status?

Compute is calculated based on the training resources used, focusing on floating point operations (FLOP). The AI Office recommends two estimation methods: hardware-based and architecture-based approaches, detailed in Annex A.1 of the guidelines. Only compute used for pre-training or modifications directly contributing to model capabilities is counted. Activities prior to pre-training or improvements at inference time are excluded. The threshold for a GPAI model is set at 10^{22} FLOP, while modifications exceeding 3×10^{21} FLOP may reclassify an actor as a new provider, triggering compliance obligations.

7. What happens if a downstream actor modifies a GPAI model significantly?

If a downstream actor modifies a GPAI model using compute that exceeds one-third of the GPAI threshold (i.e., 3×10^{21} FLOP), they are presumed to become a provider of a new GPAI model. If the original model is already one with systemic risk, and the modification adds over 3×10^{24} FLOP, or raises total compute to 10^{25} FLOP, then the modifier is presumed to provide a GPAI model with systemic risk. This means full compliance obligations apply, not limited to the modification itself. This encourages transparency and responsibility for substantial modifications.

8. What transitional rules exist for GPAI model providers?


Transitional provisions are designed to help new providers adjust to the AI Act. The AI Office encourages early dialogue for those anticipating difficulties in compliance. For example, when launching a GPAI model with systemic risk on the EU market for the first time, deadlines and obligations may be adapted. The guidelines promote pragmatic enforcement to allow innovation while supporting safety and compliance. These measures help reduce burdens for startups or first-time entrants, ensuring that the regulatory framework scales reasonably with company size and market impact.

9. How does signing the Code of Practice affect a GPAI provider's obligations?

Providers who sign and adhere to the voluntary Code of Practice benefit from enhanced trust and potentially reduced regulatory scrutiny. The Commission may consider Code adherence a mitigating factor when determining fines or investigating compliance. Non-signatories must independently prove compliance with Articles 53 and 55, often facing more frequent information requests and audits. Signing the Code of Practice provides a structured path to compliance and can facilitate smoother market access, especially for providers handling GPAI models with significant societal or economic implications.


10. How will the rules for GPAI models be enforced in practice?

The AI Office is responsible for overseeing compliance and enforcement for GPAI model providers. The guidelines emphasize collaboration and proportionality, especially during model training and deployment. Providers of GPAI models with systemic risk are expected to report proactively and engage early with the Office. Enforcement may involve audits, requests for documentation, or evaluations. The Office aims to avoid unnecessary delays in market placement while ensuring safety and accountability.

A stylized brain composed of glowing green circuitry and data lines, symbolizing artificial intelligence. The brain is centered against a dark background with a vertical beam of light passing through its center. The circuitry is intricate, with many small glowing points and lines that resemble a complex network. The overall effect is futuristic and high-tech.

Understanding General-Purpose AI Models

1. Understanding GPAI Models



What are the key goals?

- This section aims to define what constitutes a general-purpose AI (GPAI) model and classify it within the scope of the AI Act.
- It clarifies who is considered a "provider" of such models, whether developers, deployers, or those modifying the systems.
- Establishing clear definitions is critical to ensuring appropriate allocation of legal responsibilities.


Why is it needed?

- Without a clear definition of GPAI models and their providers, regulatory uncertainty could hinder innovation or allow legal loopholes.
- Clarifying these terms helps differentiate between high-impact and general-use systems, ensuring that only the appropriate models are subject to strict regulatory obligations.
- It also provides legal certainty for organizations.

How does it work in practice?

- In practice, regulators and organizations assess whether an AI model meets the criteria for GPAI classification based on its general-purpose nature and scale of use.
- They determine provider status by examining who develops, modifies, or places the model on the market.
- This helps identify which entities are legally accountable under the AI Act.

2. Provider Obligations and Systemic Risk



What are the key goals?

- This section outlines the specific obligations of GPAI model providers, with added scrutiny for those whose models pose systemic risks.
- The aim is to promote transparency, accountability, and safety in high-impact AI systems.
- It seeks to mitigate potential harm by requiring detailed documentation, risk assessments, and safeguards.


Why is it needed?

- As GPAI models grow in influence, they can have far-reaching effects across sectors, including misinformation, discrimination, and economic disruption.
- Without regulation, providers may overlook critical safety and ethical considerations.
- This section is needed to prevent such harms, especially from large, general-purpose models that could be easily repurposed.

How does it work in practice?

- In practice, providers must conduct risk assessments, maintain detailed technical documentation, and implement safety mechanisms.
- For systemic-risk models, additional oversight applies, such as incident reporting and model evaluations.
- Regulators assess compliance through audits and engagement.

3. Open-Source and Compute Requirements



What are the key goals?

- This section aims to clarify under what conditions open-source GPAI models are exempt from regulatory obligations, and to define thresholds—especially compute-related—that trigger regulatory oversight.
- It seeks to ensure fairness and proportionality by distinguishing between small-scale, collaborative projects and high-impact deployments.


Why is it needed?

- Regulating all GPAI models equally would stifle open-source research and innovation.
- Many open models are critical to academic and low-resource development environments.
- At the same time, regulators need a consistent metric to identify high-risk models.
- Compute usage provides a quantifiable, technology-neutral measure.

How does it work in practice?

- Organizations track the compute used in training models (e.g., FLOPs or GPU hours) to assess if they exceed thresholds set by the AI Act.
- If they do, even open-source models may face obligations unless they qualify for specific exemptions.
- Developers must document compute metrics and consider whether sharing their models changes their classification.

4. Downstream Modification and Legal Impact



What are the key goals?

- This section seeks to define when a third party that significantly modifies a GPAI model becomes a “provider” under the AI Act.
- It focuses on assigning legal responsibility based on the nature and impact of modifications.
- The goal is to prevent accountability gaps when powerful models are changed post-release and to clarify legal duties for modifiers introducing new risks or use cases.


Why is it needed?

- AI models are often fine-tuned or repurposed by downstream users.
- Without clarity, it’s unclear who is responsible for new risks introduced during modification.
- This could allow harmful versions of models to be distributed without oversight.
- Establishing legal criteria for provider reclassification ensures modified models remain safe and accountable.

How does it work in practice?

- In practice, if a third party makes substantial changes—particularly if they increase risk or require significant compute—they may be reclassified as a provider.
- This imposes the same legal obligations as on original developers, such as safety checks and documentation.
- Developers must evaluate their changes and possibly consult regulators to assess regulatory thresholds.

5. Transition, Compliance, and Enforcement



What are the key goals?


- This section aims to support organizations, especially new or smaller ones, in transitioning to AI Act compliance.
- It introduces mechanisms like the Code of Practice and outlines how the AI Office will enforce the rules.
- The goals include fostering a collaborative compliance environment, encouraging best practices, and enabling smooth adoption of legal obligations.

Why is it needed?





- Compliance with the AI Act can be complex and resource-intensive.
- Without structured support and predictable enforcement, smaller players may struggle or exit the market.
- This section is needed to prevent fragmentation, support innovation, and ensure all actors—regardless of size—can align with the law.
- It also ensures that enforcement is fair, and transparent.

How does it work in practice?

- The AI Office provides guidance, monitors compliance, and may use collaborative tools like the Code of Practice to help providers meet their obligations.
- Enforcement actions prioritize high-risk or non-cooperative actors, while others receive support and time to comply.
- This approach balances rigor with flexibility, encouraging voluntary alignment while reserving legal action for violations.

A stylized brain composed of glowing green circuitry and data lines, symbolizing artificial intelligence. The brain is centered against a dark background with a vertical beam of light passing through its center. The circuitry is intricate, with many small glowing points and lines that resemble a complex network. The overall effect is one of high-tech, futuristic intelligence.

Mapping General-Purpose AI Models to Ecosystem Actors



Providers of General-Purpose AI Models		EU AI Act	
Section	Description	Article(s)	Explanation
Understanding GPAI Models	Clarifies the definition and classification of general-purpose AI models and identifies who qualifies as a provider under the AI Act, setting the foundation for understanding obligations and responsibilities.	53, 3, 55, 51, Recital 106, Recital 111, Recital 107, Recital 97, Recital 114, Recital 99	Definition and Classification - A general-purpose AI model is an AI model capable of performing a wide range of tasks, typically trained on large datasets using methods like self-supervised learning. - These models are not AI systems by themselves but can be integrated into AI systems. - A common example is large generative AI models that create text, audio, images, or video. - A general-purpose AI model is classified as having systemic risk if it: - Possesses high-impact capabilities. - Significantly impacts the internal market. - High-impact capabilities are assessed by factors such as the amount of computation used for training (e.g., measured in floating point operations). - The European Commission can designate a model as having systemic risk based on these criteria. Providers of General-Purpose AI Models - A provider is any entity that develops or places an AI system or model on the market under its name or trademark, whether for payment or free of charge. - This includes natural persons, legal persons, public authorities, agencies, or other bodies. Obligations and Responsibilities Providers of general-purpose AI models must: - Maintain up-to-date technical documentation, including details of the training and testing processes. - Provide information to downstream providers to help them understand the model's capabilities and limitations. Comply with EU

			copyright laws, especially regarding the use of copyrighted content for training. • Make publicly available a summary of the content used for training the model. For models with systemic risk, providers have additional obligations: • Conduct adversarial testing. • Ensure cybersecurity protection.
Provider Obligations and Systemic Risk	Outlines the regulatory duties for providers, especially those developing models with systemic risk, including documentation, safety measures, and legal consequences tied to market placement of GPAI models.	55, 53, 51, 101, Recital 114, Recital 110	Regulatory Duties for Providers Providers of general-purpose AI models, particularly those with systemic risk, must: • Documentation: Maintain up-to-date technical documentation, including details of training and testing processes. This documentation must be available to the AI Office and national competent authorities upon request. • Safety Measures: Conduct model evaluations using standardized protocols (including adversarial testing) to identify and mitigate systemic risks. Ensure an adequate level of cybersecurity protection for the model and its physical infrastructure. • Risk Assessment: Assess and mitigate possible systemic risks at the Union level, including risks arising from the development, market placement, or use of general-purpose AI models with systemic risk. Legal Consequences • Market Placement: Comply with EU copyright and related rights law, ensuring training content is legally sourced. Publish a summary of the content used for training. • Reporting Obligations: Report serious incidents and corrective measures to the AI Office and national competent authorities without undue delay. • Compliance and Penalties: Rely on codes of practice to demonstrate compliance until harmonized standards are published. Non-

			compliance can result in fines up to 3% of global annual turnover or €15 million, whichever is higher. Systemic Risk • Definition and Impact: A general-purpose AI model has systemic risk if it possesses high-impact capabilities or significantly affects the internal market. Risks may include negative effects on public health, safety, and fundamental rights. • Mitigation Measures: Continuously assess and mitigate systemic risks, implement risk-management policies, and cooperate with relevant actors across the AI value chain.
Open-Source and Compute Requirements	Explains when open-source models are exempt from obligations and details how training compute is measured to determine whether a model crosses thresholds for regulatory classification.	54, 51, Recital 104, Recital 112, Recital 102, Recital 111	Exemptions for Open-Source Models • Open-source models released under a free and open-source license are generally exempt from certain obligations under the EU AI Act. • Exemptions apply when: ◦ The model's parameters (weights, architecture, usage information) are made publicly available. ◦ The model does not present systemic risks. • If an open-source model is deemed to have systemic risks, it must comply with obligations regardless of its open-source status. • The rationale for these exemptions is to promote transparency and innovation while ensuring systemic-risk models are properly regulated. Obligations Despite Exemptions Even when exempt, open-source models must still: • Produce a summary of the content used for training. • Implement a policy to comply with EU copyright law, particularly regarding reservation of rights. Compute Requirements for Regulatory Classification Measuring Training Compute

			• The EU AI Act uses the amount of computation (measured in floating point operations) as a key metric for classifying systemic risk. • A model is presumed to have high-impact capabilities if cumulative training compute exceeds 10²⁵ floating point operations. • This threshold acts as a proxy for the model's capabilities and potential impact, reflecting the state of the art in AI development. Implications of Crossing the Threshold • Crossing the computation threshold presumes systemic risk. • Providers must notify the AI Office. • Providers may need to demonstrate that the model does not present systemic risks due to its specific characteristics.
Downstream Modification and Legal Impact	Covers how significant modifications by third parties may reclassify them as providers, depending on the level of compute used and the risk introduced by changes to the original model.	51, 25, 101, Recital 84	Significant Modifications by Third Parties • Third parties making significant modifications to AI systems can be reclassified as providers under the EU AI Act. • Reclassification occurs when: ◦ A third party places its name or trademark on a high-risk AI system already on the market, or makes a substantial modification while it remains high-risk. ◦ The intended purpose of an AI system (including a general-purpose AI system) is modified in a way that it becomes high-risk. Reclassification as Providers • Once reclassified as a provider, the third party assumes all relevant obligations under the EU AI Act. • This includes compliance with requirements for high-risk AI systems, such as: ◦ Conducting conformity assessments. ◦ Maintaining technical documentation. Level of Compute and Risk Assessment

			• The compute used in training is a key factor in assessing risk and possible reclassification. • A general-purpose AI model is presumed to have high-impact capabilities if cumulative training compute exceeds 10²⁵ floating point operations. • Models meeting this threshold may be classified as having systemic risk, triggering additional regulatory obligations. Legal Implications • Third parties reclassified as providers must comply with all obligations for high-risk AI systems, including documentation, risk management, and transparency. • Non-compliance can result in penalties of up to 3% of global annual turnover or €15 million, whichever is higher.
Transition, Compliance, and Enforcement	Describes support for new providers, the role of the Code of Practice, and how the AI Office will supervise compliance using a collaborative and risk-based enforcement strategy.	57, 62, 53, 55, 56, 89, 95, 75	Support for New Providers The EU AI Act includes measures to support SMEs and start-ups, such as: • Priority Access to AI Regulatory Sandboxes: ◦ SMEs and start-ups receive priority access to controlled environments where they can develop, test, and validate AI systems before market placement. ◦ These sandboxes provide regulatory guidance and help providers demonstrate compliance. • Training and Awareness: ◦ Member States must organize awareness-raising and training activities tailored to SMEs and start-ups. ◦ These initiatives help smaller providers understand and implement the AI Act effectively. Role of the Code of Practice • Voluntary Application: ◦ The AI Office promotes the development of codes of practice to encourage voluntary compliance, especially for non-high-risk AI systems.

			○ These codes help providers align with best practices and industry standards. • Harmonization and Standards: ○ Providers of general-purpose AI models may rely on codes of practice to demonstrate compliance until harmonized standards are established. ○ Adhering to these codes grants a presumption of conformity with the AI Act. Supervision and Enforcement by the AI Office • Monitoring and Evaluation: ○ The AI Office monitors implementation of the AI Act by providers, including adherence to codes of practice. ○ It evaluates whether objectives in the codes contribute to effective application of the regulation. • Risk-Based Approach: ○ Focus is placed on systemic risks and high-impact AI models. ○ Providers are expected to assess and mitigate risks at the Union level. • Collaboration with National Authorities: ○ The AI Office works with national competent authorities to ensure consistent enforcement across the EU. ○ It provides guidance and support to market surveillance authorities.
--	--	--	--



Calls to Action



Classify AI Models and Identify Provider Status

Enterprises must evaluate whether their AI systems qualify as GPAI models under the AI Act. They should also determine if they meet the criteria of a "provider," which triggers compliance obligations. Legal and technical teams must align on model classification and responsibilities early in the development lifecycle.



Implement Risk Management and Documentation Processes

Establish internal processes for documenting model development, conducting risk assessments, and applying safety controls. For models posing systemic risks, ensure readiness for audits and ongoing compliance. Legal and engineering teams must collaborate to embed regulatory safeguards and ensure models meet AI Act obligations before market release.



Establish Governance for Model Modifications

Create policies to evaluate and document any modifications to GPAI models. Assess whether changes reclassify the organization as a provider under the AI Act. This includes tracking compute used in modifications and evaluating newly introduced risks. Legal teams should guide developers on when obligations are triggered.



Engage with the AI Office and Adopt Best Practices

Appoint a compliance lead to monitor regulatory updates, liaise with the AI Office, and integrate the Code of Practice. Prioritize early alignment with AI Act requirements, especially for new models. Establish training and risk-based compliance frameworks to ensure a smooth transition and ongoing regulatory alignment.

Conclusion

The regulation of **General-Purpose AI (GPAI) models** under the EU Artificial Intelligence Act represents a pivotal evolution in the governance of advanced AI technologies. By setting out clear obligations for GPAI providers—particularly those whose models present **systemic risks**—the Act moves the European Union toward a framework of proactive oversight, transparency, and accountability. This shift ensures that the most powerful AI systems are developed and deployed responsibly, safeguarding privacy, fairness, and fundamental rights across the Union.

Throughout this report, we have outlined how the AI Act establishes a structured regime for GPAI providers—defining their duties in documentation, compute transparency, model evaluation, and risk mitigation. Compliance with these provisions not only demonstrates adherence to legal standards but also reinforces **public trust** and strengthens the integrity of the European AI ecosystem. By embedding responsible governance and transparency into model development and deployment, organizations can align innovation with ethical and regulatory expectations.

For successful implementation, providers should integrate **robust compliance** and **governance** throughout the AI lifecycle—covering data sourcing, model training, deployment, and modification. Smaller developers and open-source communities may require targeted support and clear guidance to meet these obligations effectively. Yet, by embracing these responsibilities early, organizations position themselves as leaders in ethical and compliant AI innovation.

As AI continues to expand its influence across industries and societies, the EU's approach to regulating GPAI models sets a **global benchmark** for safe, transparent, and rights-respecting AI governance. Adherence to these standards will ensure that general-purpose AI models contribute positively to Europe's digital future—fostering innovation that is both sustainable and trustworthy.

Appendix: Additional Considerations for GPAI Compliance

Compute Thresholds and Their Limitations

The AI Act establishes 10^{25} FLOP as the presumption trigger for classifying a GPAI model as having systemic risk. Separately, Commission guidance (the GPAI guidelines published in July 2025) uses a threshold around 10^{23} FLOP as a presumptive indicator for a model being in-scope as a GPAI model. While these thresholds provide useful reference points, they function as proxies rather than definitive measures of capability or risk. Advances in training efficiency, model distillation, fine-tuning techniques, and architectural innovations can produce highly capable models below these thresholds. Regulators may look beyond raw compute when assessing individual models, and these thresholds are subject to revision as the technology evolves. Organizations should monitor AI Office guidance for updates to classification criteria.

Open-Source Models: Key Limitations

The AI Act provides certain exemptions for models released under free and open-source licenses, but the exemption is limited and fact-dependent. Eligibility hinges on license terms plus public availability of weights, architecture, and usage information. Critically, the exemption does not apply to systemic-risk GPAI models regardless of how they are released. Even where the exemption applies, providers may still owe other GPAI obligations, notably around copyright policy and training data summaries. Monetization through adjacent services can also affect eligibility, Commission guidelines discuss this as a practical consideration. Organizations releasing or building on open-source GPAI models should seek legal guidance and monitor forthcoming AI Office clarifications.

Extraterritorial Application

The AI Act applies to providers placing GPAI models on the EU market regardless of where they are headquartered. Non-EU providers must appoint an authorized representative in the EU and comply with applicable obligations as if established within the Union (with a narrow exception for certain open-source models that do not pose systemic risk). Cross-jurisdictional compliance presents practical challenges, including coordinating with multiple regulatory frameworks, managing documentation requirements across legal systems, and navigating potential conflicts between EU obligations and domestic laws. Organizations operating across

jurisdictions should assess their EU market exposure and develop compliance strategies accordingly.

Enforcement: What to Expect

The AI Office is responsible for supervising GPAI model providers and has signaled a collaborative, proportionate approach to enforcement. However, full enforcement powers, including the ability to request information, require corrective actions, restrict availability or placing on the market, and impose fines, do not take effect until August 2, 2026. GPAI obligations have applied since August 2, 2025. Models placed on the market before that date have until August 2, 2027 to achieve compliance. Organizations should use this transitional period to establish compliance infrastructure and engage proactively with the AI Office where appropriate.

Model Evaluation and Adversarial Testing

Providers of GPAI models with systemic risk must conduct model evaluations and adversarial testing, but standardized methodologies are still maturing. The Code of Practice provides a framework, and the AI Office has launched a €9 million procurement effort to develop evaluation capacity for specific risk domains, organized into lots covering CBRN risks, cyber offence, harmful manipulation, loss of control, and cross-cutting support for agentic evaluations. Organizations should document their evaluation methodologies, track emerging standards, and be prepared to demonstrate that their testing approach is appropriate to the risks posed by their models.

Code of Practice: Finalized but Evolving in Practice

The GPAI Code of Practice was published on July 10, 2025, and the Commission and AI Board confirmed its adequacy on August 1, 2025. Major providers including Amazon, Anthropic, Google, Microsoft, and OpenAI have signed the full Code, though Meta declined and xAI signed only the Safety and Security chapter. While the Code provides a pathway to demonstrate compliance and signatories benefit from a presumption of conformity, questions remain about how strictly commitments will be interpreted and enforced in practice. Ultimate legal interpretation rests with the AI Office and, eventually, the Court of Justice of the European Union.

References

Artificial Intelligence Act, (2025), 'Section 1: Classification Rules', accessible at <https://artificialintelligenceact.eu/section/5-1/> (last accessed 7th May 2025)


Artificial Intelligence Act, (2025), 'Section 2: Obligations for Providers of General-Purpose AI Models', accessible at <https://artificialintelligenceact.eu/section/5-2/> (last accessed 7th May 2025)

Artificial Intelligence Act, (2025), 'Section 3: Obligations of Providers of General-Purpose AI Models with Systemic Risk', accessible at <https://artificialintelligenceact.eu/section/5-3/> (last accessed 7th May 2025)

Artificial Intelligence Act, (2025), 'Section 4: Codes of Practice', accessible at <https://artificialintelligenceact.eu/section/5-4/> (last accessed 7th May 2025)

European Parliament and The Council of the European Union, (2024), 2024/1689 Regulation (EU) 2024/1689 of the European Parliament and of The Council of 13 June 2024 laying down harmonised rules on artificial intelligence and amending Regulations (EC) No 300/2008, (EU) No 167/2013, (EU) No 168/2013, (EU) 2018/858, (EU) 2018/1139 and (EU) 2019/2144 and Directives 2014/90/EU, (EU) 2016/797 and (EU) 2020/1828 (Artificial Intelligence Act), accessible at https://eur-lex.europa.eu/legal-content/EN/TXT/PDF/?uri=OJ:L_202401689 (last accessed 7th May 2025)

OECD.AI, (2025), 'Generative AI models by modality', accessible at: <https://oecd.ai/en/data?selectedArea=ai-models-and-datasets&selectedVisualization=generative-ai-models-by-modality> (last accessed 28th October 2025)



Important notice

This document has been prepared by AI & Partners B.V. for the sole purpose of enabling the parties to whom it is addressed to evaluate the capabilities of AI & Partners B.V. to supply the proposed services.

Other than as stated below, this document and its contents are confidential and prepared solely for your information, and may not be reproduced, redistributed or passed on to any other person in whole or in part. If this document contains details of an arrangement that could result in a tax or National Insurance saving, no such conditions of confidentiality apply to the details of that arrangement (for example, for the purpose of discussion with tax authorities). No other party is entitled to rely on this document for any purpose whatsoever and we accept no liability to any other party who is shown or obtains access to this document.

This document is not an offer and is not intended to be contractually binding. Should this proposal be acceptable to you, and following the conclusion of our internal acceptance procedures, we would be pleased to discuss terms and conditions with you prior to our appointment. Images used throughout the document have either been produced in-house or sourced from publicly available sources (see **References** for details).

AI & Partners B.V. is the Dutch headquarters of AI & Partners, a global professional services firm. Please see <https://www.ai-and-partners.com/> to learn more about us.

© 2025 AI & Partners B.V. All rights reserved.

Designed and produced by AI & Partners B.V.