



World Health
Organization

Ethics and governance of artificial intelligence for health

Guidance on large multi-modal models





**World Health
Organization**

Ethics and governance of artificial intelligence for health

Guidance on large multi-modal models

Ethics and governance of artificial intelligence for health. Guidance on large multi-modal models

ISBN 978-92-4-008475-9 (electronic version)

ISBN 978-92-4-008476-6 (print version)

© **World Health Organization 2024**

Some rights reserved. This work is available under the Creative Commons Attribution-NonCommercial-ShareAlike 3.0 IGO licence (CC BY-NC-SA 3.0 IGO; <https://creativecommons.org/licenses/by-nc-sa/3.0/igo>).

Under the terms of this licence, you may copy, redistribute and adapt the work for non-commercial purposes, provided the work is appropriately cited, as indicated below. In any use of this work, there should be no suggestion that WHO endorses any specific organization, products or services. The use of the WHO logo is not permitted. If you adapt the work, then you must license your work under the same or equivalent Creative Commons licence. If you create a translation of this work, you should add the following disclaimer along with the suggested citation: “This translation was not created by the World Health Organization (WHO). WHO is not responsible for the content or accuracy of this translation. The original English edition shall be the binding and authentic edition”.

Any mediation relating to disputes arising under the licence shall be conducted in accordance with the mediation rules of the World Intellectual Property Organization. (<http://www.wipo.int/amc/en/mediation/rules/>).

Suggested citation. Ethics and governance of artificial intelligence for health. Guidance on large multi-modal models. Geneva: World Health Organization; 2024. Licence: CC BY-NC-SA 3.0 IGO.

Cataloguing-in-Publication (CIP) data. CIP data are available at <http://apps.who.int/iris>.

Sales, rights and licensing. To purchase WHO publications, see <http://apps.who.int/bookorders>. To submit requests for commercial use and queries on rights and licensing, see <http://www.who.int/about/licensing>.

Third-party materials. If you wish to reuse material from this work that is attributed to a third party, such as tables, figures or images, it is your responsibility to determine whether permission is needed for that reuse and to obtain permission from the copyright holder. The risk of claims resulting from infringement of any third-party owned component in the work rests solely with the user.

General disclaimers. The designations employed and the presentation of the material in this publication do not imply the expression of any opinion whatsoever on the part of WHO concerning the legal status of any country, territory, city or area or of its authorities, or concerning the delimitation of its frontiers or boundaries. Dotted and dashed lines on maps represent approximate border lines for which there may not yet be full agreement.

The mention of specific companies or of certain manufacturers' products does not imply that they are endorsed or recommended by WHO in preference to others of a similar nature that are not mentioned. Errors and omissions excepted, the names of proprietary products are distinguished by initial capital letters.

All reasonable precautions have been taken by WHO to verify the information contained in this publication. However, the published material is being distributed without warranty of any kind, either expressed or implied. The responsibility for the interpretation and use of the material lies with the reader. In no event shall WHO be liable for damages arising from its use.

Graphics design: Joanna Sleight (ETH Zurich, Zurich, Switzerland)
Layout: Imprimerie Centrale (Luxembourg)

Contents

Acknowledgements	v
Abbreviations.....	vii
Executive summary	viii
1 Introduction	1
1.1 Significance of LMMs.....	3
1.2 WHO guidance on ethics and governance of AI for health	4
I. Applications, challenges and risks of LMMs.....	7
2 Applications and challenges of use of LMMs in health	8
2.1 Diagnosis and clinical care.....	8
2.2 Patient-centred applications.....	12
2.3 Clerical functions and administrative tasks.....	16
2.4 Medical and nursing education	17
2.5 Scientific and medical research and drug development.....	17
3 Risks to health systems and society and ethical concerns about use of LMMs.....	20
3.1 Health systems.....	20
3.2 Compliance with regulatory and legal requirements.....	23
3.3 Societal concerns and risks	24
II. Ethics and governance of LMMs in health care and medicine.....	31
4 Design and development of general-purpose foundation models (LMMs)	34
4.1 Risks to be addressed during the development of general-purpose foundation models (LMMs).....	34
4.2 Measures developers can take to address risks with general-purpose foundation models (LMMs).....	35
4.3 Government laws, policies and public sector investments.....	39
4.4 Open-source LMMs.....	41

5	Provision with general-purpose foundation models (LMMs)	45
5.1	Risks to be addressed when providing a health-care service or application with a general-purpose foundation model (LMM).....	45
5.2	Measures that governments can introduce to address such risks and ethical principles that should be upheld	46
6	Deployment with general-purpose foundation models (LMMs)	53
6.1	Risks to be addressed when deploying a health-care service or application with a general-purpose foundation model (LMM).....	53
6.2	On-going responsibilities of developers and providers during deployment	54
6.3	Responsibilities of deployers.....	54
6.4	Government programmes and practices	55
7	Liability for LMMs	58
8	International governance of LMMs	60
	References	62
	Annex. Methods	77

Acknowledgements

Development of this World Health Organization (WHO) guidance was led by Andreas Reis (co-lead of the Health Ethics and Governance unit in the Department of Research for Health) and Sameer Pujari (Department of Digital Health and Innovation), under the overall guidance of John Reeder (Director, Research for Health), Alain Labrique (Director, Digital Health and Innovation) and Jeremy Farrar (Chief Scientist).

Rohit Malpani (consultant, France) was the lead writer. The co-chairs of the WHO expert group on ethics and governance of AI for health, Effy Vayena (ETH Zurich, Switzerland) and Partha Majumder (Indian Statistical Institute and National Institute of Biomedical Genomics, India), provided overall guidance on drafting of the report and leadership of the expert group.

WHO is grateful to the following individuals who contributed to development of this guidance.

WHO expert group on ethics and governance of AI for health

Najeeb Al Shorbaji, eHealth Development Association, Amman, Jordan; Maria Paz Canales, Global Partners Digital, Santiago de Chile, Chile; Arisa Ema, University of Tokyo, Tokyo, Japan; Amel Ghouila, Bill & Melinda Gates Foundation, Seattle (WA), USA; Jennifer Gibson, WHO Collaborating Centre for Bioethics, University of Toronto, Toronto, Canada; Kenneth Goodman, Institute of Bioethics and Health Policy, University of Miami Miller School of Medicines, Miami (FL), USA; Malavika Jayaram, Digital Asia Hub, Singapore; Daudi Jjingo, Makerere University, Kampala, Uganda; Tze Yun Leong, National University of Singapore, Singapore; Alex John London, Carnegie Mellon University, Pittsburgh (PA), USA; Partha Majumder, Indian Statistical Institute and National Institute of Biomedical Genomics, Kolkata, India; Thilidzi Marwala, University of Johannesburg, Johannesburg, South Africa; Roli Mathur, Indian Council of Medical Research, Bangalore, India; Timo Minssen, Centre for Advanced Studies in Biomedical Innovation Law, Faculty of Law, University of Copenhagen, Copenhagen, Denmark; Andrew Morris, Health Data Research UK, London, United Kingdom; Daniela Paolotti, ISI Foundation, Turin, Italy; Jerome Singh, University of Kwa-Zulu Natal, Durban, South Africa; Jeroen van den Hoven, University of Delft, Delft, Netherlands (Kingdom of the); Effy Vayena, ETH Zurich, Zurich, Switzerland; Robyn Whittaker, University of Auckland, Auckland, New Zealand; and Yi Zeng, Chinese Academy of Sciences, Beijing, China.

Observers

David Gruson, Luminess, Paris, France; Lee Hibbard, Council of Europe, Strasbourg, France

External reviewers

Oren Asman, Tel Aviv University, Tel Aviv, Israel; I. Glenn Cohen, Harvard Law School, Boston (MA), USA; Alexandrine Pirlot de Corbion, Privacy International, London, United Kingdom; Rodrigo Lins, Federal University of Pernambuco, Recife, Brazil; Doug McNair, Deputy Director, Integrated Development, Bill & Melinda Gates Foundation, Seattle (WA), USA; Keymanthri Moodley, Stellenbosch University, Cape Town, South Africa; Amir Tal, Tel Aviv University, Tel Aviv, Israel; Tom West, Privacy International, London, United Kingdom.

External contributors

Box 2 (Ethical considerations for the use of LMMs by children) of the guidance was drafted by Vijaytha Muralidharan, Alyssa Burgart, Roxana Daneshjou and Sherri Rose, Stanford University, Stanford (CA), USA. Box 3 (Ethical considerations associated with LMMs and their impact on individuals with disabilities) of the guidance was drafted by Yonah Welker, independent consultant, Geneva, Switzerland.

All external reviewers, experts and contributors declared their interests in line with WHO policies. None of the interests declared were assessed to be significant.

WHO

Shada Al-Salamah, Technical Officer, Department of Digital Health and Innovation, Geneva; Mariam Otmani Del Barrio, Scientist, Special Programme on Tropical Diseases Research, Geneva; Marcelo D'Agostino, Unit Chief, Information Systems and Digital Health, WHO Regional Office for the Americas, Washington (DC); Jeremy Farrar, Chief Scientist, Geneva; Clayton Hamilton, Technical Officer, WHO Regional Office for Europe, Copenhagen, Denmark; Kanika Kalra, Consultant, Department of Digital Health and Innovation, Geneva; Ahmed Mohamed Amin Mandil, Coordinator, Research and Innovation, WHO Regional Office for the Eastern Mediterranean, Cairo; Issa T. Matta, Legal Affairs, Geneva; Jose Eduardo Diaz Mendoza, Consultant, Department of Digital Health and Innovation, Geneva; Mohammed Hassan Nour, Technical Officer, Department of Digital Health and Innovation, WHO Regional Office for the Eastern Mediterranean, Cairo; Denise Schalet, Technical Officer, Department of Digital Health and Innovation, Geneva; Yu Zhao, Technical Officer, Department of Digital Health and Innovation, Geneva.

Abbreviations

AI artificial intelligence

LMM large multi-modal model

USA United States of America

Executive summary

Artificial Intelligence (AI) refers to the capability of algorithms integrated into systems and tools to learn from data so that they can perform automated tasks without explicit programming of every step by a human. Generative AI is a category of AI techniques in which algorithms are trained on data sets that can be used to generate new content, such as text, images or video. This guidance addresses one type of generative AI, large multi-modal models (LMMs), which can accept one or more type of data input and generate diverse outputs that are not limited to the type of data fed into the algorithm. It has been predicted that LMMs will have wide use and application in health care, scientific research, public health and drug development. LMMs are also known as “general-purpose foundation models”, although it is not yet proven whether LMMs can accomplish a wide range of tasks and purposes.

LMMs have been adopted faster than any consumer application in history. They are compelling because they facilitate human–computer interaction to mimic human communication and to generate responses to queries or data inputs that may appear human-like and authoritative. With rapid consumer adoption and uptake and in view of its potential to disrupt core social services and economic sectors, many large technology companies, start-ups and governments are investing in and competing to guide the development of generative AI.

In 2021, WHO published comprehensive guidance (1) on the ethics and governance of AI for health. WHO consulted 20 leading experts in AI, who identified both potential benefits and potential risks of use of AI in health care and issued six principles arrived at by consensus for consideration in the policies and practices of governments, developers, and providers that are using AI. The principles should guide the development and deployment of AI in health care by a wide range of stakeholders, including governments, public sector agencies, researchers, companies and implementers. The principles are: (1) protect autonomy; (2) promote human well-being, human safety and the public interest; (3) ensure transparency, “explainability” and intelligibility; (4) foster responsibility and accountability; (5) ensure inclusiveness and equity; and (6) promote AI that is responsive and sustainable (Figure 1).

WHO is issuing this guidance to assist Member States in mapping the benefits and challenges associated with use of LMMs for health and in developing policies and practices for appropriate development, provision and use. The guidance includes recommendations for governance, within companies, by governments and through international collaboration, aligned with the guiding principles. The principles and recommendations, which account for the unique ways in which humans can use generative AI for health, are the basis of this guidance.

Figure 1: WHO consensus ethical principles for use of AI for health

Applications, challenges and risks of large multi-modal models

The potential applications of LMMs in health care are similar to those of other forms of AI, yet how LMMs are accessed and used is new, with both novel benefits and risks that societies, health systems and end-users may not yet be prepared to address fully. Table 1 summarizes the main applications of LMMs and their potential benefits and risks.

The systemic risks associated with use of LMMs include risks that could affect health systems (Table 2).

Broader regulatory and systemic risks could emerge with use of LMMs. One concern (being examined by several data protection authorities) is whether LMMs comply with existing legal or regulatory regimes, including international human rights obligations, and with national data protection regulations. Algorithms might not comply with such laws because of the way in which data are collected to train LMMs, the management and processing of data that have been collected (or put into LMMs by end users), the transparency and accountability of entities that develop LMMs, and the possibility that LMMs “hallucinate”. LMMs could also be non-compliant with consumer protection laws.

Broader societal risks associated with the growing use of LMMs (including and beyond the use of such algorithms in health care) include the fact that LMMs are often developed and deployed by large technology companies, due partly to the significant computing, data, human and financial resource required for development of LMMs. This may reinforce the dominance of these companies vis-a-vis smaller enterprises and governments with respect to the development and use of AI, including the focus of AI research in the public and private sectors. Additional concerns about the potential dominance of large technology companies include insufficient corporate commitment to ethics and transparency. New voluntary

Table 1. Potential benefits and risks in various uses of LMMs in health care

Use	Potential or proposed benefits	Potential risks
Diagnosis and clinical care	Assist in managing complex cases and review of routine diagnoses Reduce the communication workload of health-care providers (“keyboard liberation”) Provide novel insights and reports from various unstructured forms of health data	Inaccurate, incomplete or false responses Poor quality training data Bias (of training data and responses) Automation bias Degradation of skills (of health-care professionals) Informed consent (of patients)
Patient-guided use	Generate information to improve understanding of a medical condition (as a patient or as a caregiver) Virtual health assistant Clinical trial enrolment	Inaccurate, incomplete or false statements Manipulation Privacy Less interaction between clinicians and patients Epistemic injustice Risk of delivery of care outside the health system
Clerical and administrative tasks	Assist with paperwork and documentation required for clinical care Assist in language translation Completion of electronic health records Draft clinical notes after a patient visit	Inaccuracies and errors Inconsistent responses depending on prompts
Medical and nursing education	Dynamic texts suited to each student’s needs Simulated conversation to improve communication and to practise in diverse situations and with diverse patients Responses to questions accompanied by chain-of-thought reasoning	Contribute to automation bias Errors or false information undermine the quality of medical education New burden of learning digital skills
Scientific research and drug development	Generate insights from scientific data and research Generate text for use in scientific articles, manuscript submission or peer-review Analyse and summarize data for research Proofreading De novo drug design	Cannot hold algorithms accountable for content Algorithms encode bias towards the perspectives of high-income countries Generate information and/or references that do not exist Undermine key tenets of scientific research, such as peer review Exacerbate differential access to scientific knowledge

Table 2. Risks to health systems associated with use of LMMs in health care


Type of risk	Description
Overestimation of the benefits of LMMs	There may be a tendency to ‘technological solutionism’, or over-estimation of the benefits of LMMs while ignoring or downplaying challenges in its use, including its safety, efficacy and utility.
Accessibility and affordability	Equitable access to LMMs may be lacking for several reasons, including the “digital divide” and subscription fees to access LMMs.
System-wide biases	Use of ever-larger data sets could increase biases encoded in LMMs, which could be automated throughout a health-care system.
Impacts on labour	Use of LMMs could lead to job losses in some countries and require health workers to retrain and adjust to use of LMMs. Data annotation and filtering can lead to low wages and to untreated psychological distress.
Dependence of health systems on ill-suited LMMs	Dependence on LMMs could make health systems vulnerable if LMMs are not maintained or (in low- and middle-income countries) are updated only for use in high-income countries. Furthermore, lack of preservation and protection of privacy and confidentiality could undermine trust in health-care systems by people who are not confident that their privacy will be protected.
Cybersecurity risks	Malicious attacks or hacking could undermine safety and trust in the use of LMMs in health care.

commitments by such companies, with one another and with governments, could mitigate several risks in the short-term but are not an alternative to governmental oversight that might eventually be enacted.

Another societal risk is the carbon and water footprints of LMMs, which like other forms of AI, require both significant energy and contribute to AI’s growing water footprint. While LMMs and other forms of AI can provide important societal benefits, the growing carbon footprint may become a major contributor to climate change, and increasing water consumption can have a further negative impact in water-stressed communities. Another societal risk associated with the emergence of LMMs is that, by providing plausible responses that are increasingly considered a source of knowledge, LMMs may eventually undermine human epistemic authority, including in the domains of health care, science and medicine.

Ethics and governance of LMMs in health care and medicine

LMMs can be considered products of a series (or chain) of decisions on programming and product development by one or more actors (Figure 2). Decisions made at each stage of an AI value chain may have both direct and indirect consequences on those who participate in the development, deployment and use of LMMs downstream. The decisions can be influenced and regulated by governments by enacting and enforcing laws and policies nationally, regionally and globally.

Figure 2: Value chain of the development, provision and deployment of LMMs

The AI value chain often begins in a large technology company, referred to as a “developer” in this guidance. The developer could also be a university, a smaller technology company, national health systems, public-private consortiums or other entities that have the resources and capacity to use several inputs, which comprise the “AI infrastructure”, such as data, computing power and AI expertise, to develop general-purpose foundation models (a term used by governments to describe LMMs in legislation and regulation). These models can be used directly to perform various, often unanticipated tasks, including those related to health care. Several general-purpose foundation models are trained specifically for use in health care and medicine.

A general-purpose foundation model can be used by a third party (a “provider”) through an active programming interface for a specific purpose or use. This involves (i) fine-tuning a new LMM, which may require additional training of the foundation model; (ii) integrating the LMM into applications or a larger software system to provide a service to users; or (iii) integrating

components known as “plug-ins” to channel, filter and organize the LMMs into formal or regulated formats to generate “digestible” results.¹

Thereafter, the provider can market a product or service based on the LMM to a customer (or “deployer”), such as a ministry of health, a health-care system, a hospital, a pharmaceutical company or even an individual, such as a health-care provider. The customer who acquires or licenses the product or application may then use it directly for patients, health-care providers, other entities in the health system, lay people or in its own business. The value chain can be “vertically integrated”, so that a company (or other entity, such as a national health system) that collects data and trains a general-purpose foundation model can modify the LMM for a particular use and provide the application directly to users.

Governance is a means for enshrining ethical principles and human rights obligations through existing laws and policies and through new or revised laws, norms, internal codes of practice and procedures for developers and international agreements and frameworks.

One way of framing the governance of LMMs is in the three stages of the AI value chain: (i) the design and development of general-purpose foundation models or LMMs; (ii) provision of a service, application or product based on a general-purpose foundation model; and (iii) deployment of a health-care service or application. In this guidance, each phase is examined with respect to three areas of enquiry:

- What risks (described above) should be addressed at each stage of the value chain, and which actors are best placed to address those risks?
- What can a relevant actor do to address the risks, and which ethical principles must be upheld?
- What is the role of government, including relevant laws, policies and regulations?

Certain risks can be addressed at each phase of the AI value chain, and certain actors are likely to play more important roles in mitigating each risk and upholding ethical values. While there is likely to be disagreement and tension about where responsibility rests between developers, providers and deployers, there are clear areas in which each actor is best placed or is the only entity with the capacity to address a potential or actual risk.

Design and development of general-purpose foundation models (LMMs)

During the design and development of general-purpose foundation models, the responsibility rests with the developers. Governments bear the responsibility to set laws and standards to require or forbid certain practices. Section 4 of this guidance provides recommendations to help address risks and maximize benefits during the development of LMMs.

1 Communication from Leong Tze-Yun, WHO expert on the ethics and governance of AI for health.

Provision with general-purpose foundation models (LMMs)

During provision of a service or application, governments are responsible for defining the requirements and obligations of both developers and providers to address specific risks associated with AI-based systems to be used in health-care settings. Section 5 of this report provides recommendations to address risks and maximize benefits during provision of services and applications for health care with LMMs.

Deployment with general-purpose foundation models (LMMs)

Even if relevant laws, policies and ethical practices are applied during development and provision of an LMM, risks will materialize during their use due partly to the unpredictability of LMMs and the responses they provide, the possibility that a user applies a general-purpose foundation model in a way that neither the developer nor the provider had anticipated, and because LMM outputs may change over time. Section 6 of this report provides recommendations with respect to risks and challenges that should be addressed during use of LMMs and applications.

Liability for LMMs

As LMMs gain broader use in health care and medicine, errors, misuse and ultimately harm to individuals are inevitable. Therefore, liability rules might ensure that users harmed by an LMM are adequately compensated or have other forms of redress to both reduce the burden of proof by a user who is harmed, ensuring that such individuals are adequately and fairly compensated.

Governments can do so by introducing a presumption of causality. They might also consider introducing a strict liability standard for any harm that results from deployment of an LMM. Although strict liability rules may ensure compensation for those who suffer harm, they may also discourage use of increasingly sophisticated LMMs. Governments might also consider no-fault, no-liability compensation funds.

International governance of LMMs

Governments must work together to build new institutional structures and rules to ensure that international governance keeps pace with globalization of these technologies. They should also ensure stronger cooperation and collaboration within the United Nations system to respond to the opportunities and challenges for deploying AI in health care, as well as its wider applications in society and the economy.

International governance is necessary to ensure that all governments are accountable for their investments and participation in the development and deployment of AI-based systems

and that governments introduce appropriate regulations that uphold ethical principles, human rights and international law. International governance can also ensure that companies develop and deploy LMMs that meet adequate international standards of safety and efficacy and are upholding ethical principles and human rights obligations. Governments should also avoid introducing regulations that provide a competitive advantage or disadvantage for either companies or themselves.

In order for international governance to be meaningful, such rules must be shaped by all countries and not only high-income countries (and technology companies that work with high-income country governments). International governance of AI may require that all stakeholders cooperate through networked multilateralism, as proposed by the United Nations Secretary-General in 2019, which would bring together the United Nations family, international financial institutions, regional organizations, trading blocs and others, including civil society, cities, businesses, local authorities and young people, to work more closely, effectively and inclusively.

Ethics and governance of artificial intelligence for health: Large multi-modal models

Risks to be addressed

What can be done, and by who

Development phase



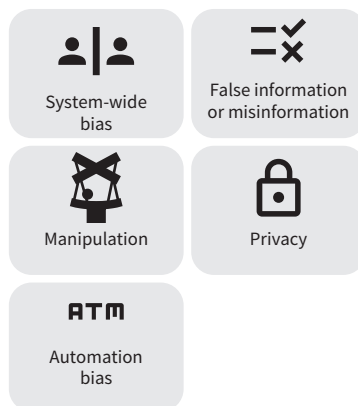
Developer actions

- ☐ Certification/training for programmers
- ☐ Data protection impact assessments
- ☐ Training data collected with 'best-practice' data protection rules
- ☐ Training data are refreshed, up-to-date, and context-appropriate
- ☐ Ensure transparency of training data
- ☐ Fair wages and support to data workers
- ☐ Involve diverse stakeholders in design
- ☐ Design for accuracy and predictability
- ☐ Design for values based on consensus principles and ethical norms
- ☐ Design to improve energy efficiency of models

Government actions

- ☐ Have and enforce strong data protection laws
- ☐ Issue target product profiles
- ☐ Mandate outcomes (predictability, interpretability, corrigibility, safety, cybersecurity)
- ☐ Introduce pre-certification programmes to identify and avoid ethical risks
- ☐ Conduct audits during early AI development
- ☐ Require developers to address carbon and water footprints
- ☐ Require developers to label AI-generated content for users
- ☐ Encourage or require early-stage registration of algorithms
- ☐ Invest in or provide public or not-for-profit infrastructure
- ☐ Promote open-source LMMs with funding

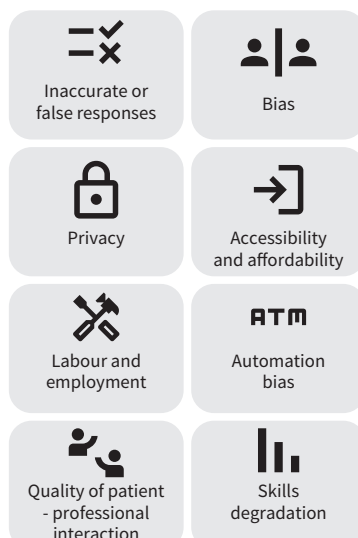
Provision phase



Government actions

- ☐ Assign regulatory agency to assess and approve LMMs for health
- ☐ Require transparency, including source code and data inputs
- ☐ Enforce data protection laws for user-inputted data
- ☐ Mandate ethical and human rights standards, irrespective of risk or benefit
- ☐ Enact laws requiring impact assessments, audited by third parties and disclosed publicly
- ☐ Require proof of performance and compliance with medical device regulations
- ☐ Prohibit non-trial experimental use; explore regulatory sandboxes for controlled testing
- ☐ Apply consumer protection laws to prevent negative impacts on end-users and patients

Deployment phase



Deployer actions

- ☐ Avoid using LMMs in inappropriate settings
- ☐ Communicate known risks, errors and harms with clear warnings and measures
- ☐ Enforce affordability and availability by ensuring pricing and languages offered are inclusive

Government actions

- ☐ Mandate independent post-release audits and impact assessments for LMM deployment
- ☐ Hold developers or providers responsible for false or toxic information
- ☐ Enforce operational disclosures, including technical documentation
- ☐ Train healthcare workers on LMM decision-making, avoiding bias, patient engagement and cybersecurity risks
- ☐ Facilitate public participation through human oversight colleges to ensure appropriate use
- ☐ Engage the public to understand data sharing, assess social/cultural acceptability, improve AI literacy, and gauge acceptable LMM uses
- ☐ Use procurement authority to encourage transparency and responsible practices by value chain actors

1 Introduction

This guidance addresses the emerging uses of large multi-modal models (LMMs) for health-related applications.² It includes the potential benefits and risks of use of LMMs in health care and medicine and approaches to the governance of LMMs that could best ensure compliance with guidelines and obligations for ethics, human rights and safety. The guidance builds on WHO guidance issued in June 2021 – Ethics and governance of artificial intelligence for health (1), which addressed the ethical challenges and risks of use of artificial intelligence (AI) in health, identified six principles for ensuring that AI is used to the public benefit of all countries, and issued recommendations to enhance the governance of AI for health in order to maximize the promise of the technology.

AI refers to the capability of algorithms integrated into systems and tools to learn from data to perform automated tasks, without programming of each step explicitly by a human. Generative AI is a category of AI techniques in which machine learning models are used to train algorithms on data sets to create new outputs, such as text, images, videos and music. Generative AI models learn patterns and structures from training data and produce new data based on predictions made from the learnt patterns. Generative AI models can be improved by reinforcement learning with human feedback, wherein human trainers rank responses provided by generative AI models to train algorithms to issue responses that maximize how much humans will value the response. Generative AI has potential applications in various fields, including design, content generation, simulation and scientific discovery.

Much attention has been focused on a specific type of generative AI, large language models, which receive one type of input – text – and provide a response that is also in text. Large language models are examples of large unimodal models, which are the basis for the operation of early versions of chatbots that integrate these models. Although large language models engage in dialogue, the models themselves have no conception of what they are producing. They merely predict the next word according to previous words, learnt patterns or combinations of words (2).

This document addresses the growing use of LMMs (including large language models), which, for use in health care and medicine, are trained with highly diverse datasets, extending beyond text, and include biosensor, genomic, epigenomic, proteomic, imaging, clinical, social and environmental data (3). Therefore, LMMs can accept more than one type of input and generate outputs that are not limited to the type of data entered. LMMs are envisioned for diverse applications in health care and drug development.

LMMs differ from previous types of AI and machine learning. While AI has already been integrated widely into many consumer applications, the outputs of most algorithms neither

² For the purposes of this guidance, the terms “large-multi-modal models” and “general-purpose foundation models” are used interchangeably, the latter term being used particularly in discussions of governance. It is not yet known, however, whether LMMs can accomplish a wide range of tasks for general purposes.

require nor invite participation of the customer or user, except for rudimentary forms of AI integrated into social media platforms that curate user-generated content to capture attention (4). Another difference between LMMs and other forms of AI is their versatility. Previous and existing AI models, including for medical uses, are designed for specific tasks and are therefore inflexible. They can execute only tasks defined in the training set and its labels (5) and cannot adapt or carry out other functions without retraining with a different dataset. Thus, even though more than 500 AI models for clinical medicine have been approved by the Food and Drug Administration in the USA (5), most are approved for only one or two narrow tasks. In contrast, LMMs are trained on various datasets and can be used in numerous tasks, including some for which they were not explicitly trained (5).

LMMs usually have an interface and format that facilitate human–computer algorithm interactions that might mimic human communication and which can therefore lead users to imbue the algorithm with human-like qualities. Thus, the way in which LMMs are used and the content they generate and provide as responses, which may appear to be “human-like”, are different from those of other forms of AI and have contributed to the unprecedented public adoption of LMMs. Furthermore, because the responses they provide appear to be authoritative, many users uncritically accept them as correct, even if an LMM cannot guarantee a correct response and cannot integrate ethical norms or moral reasoning into the responses it generates. While this guidance illustrates the different ways in which LMMs are used (or imagined for use) in health care and medicine, they are already used in numerous domains, including education, finance, communications and computer science.

LMMs can be considered products of a series (or chain) of decisions on programming and product development by one or more actors. Decisions made at each stage of an AI value chain may have both direct and indirect consequences on those that participate in the development, deployment and use of LMMs downstream. The decisions can be influenced and regulated by governments that enact and enforce laws and policies nationally, regionally and globally.

The AI value chain often begins in a large technology company. The developer could also be a university, a small technology company, a national health system, public-private consortiums or other entities that have the resources and capacity to use several inputs, which comprise “AI infrastructure”, such as data, computing power and AI expertise, to develop general-purpose foundation models. These models can be used directly by a user to perform various, often unanticipated tasks (including those related to health care). Several general-purpose foundation models are trained specifically for use in health care and medicine.

A general-purpose foundation model can be used by a third party (a “provider”) through an active programming interface for a specific purpose or use. This involves (i) fine-tuning a new LMM, which may require additional training of the foundation model; (ii) integrating the LMM into applications or a larger software system to provide a service to users; or (iii) integrating components known as “plug-ins” to channel, filter and organize the LMMs into formal or regulated formats to generate “digestible” results.³

3 Communication from Leong Tze-Yun, WHO expert on the ethics and governance of AI for health.

Thereafter, the provider can market a product or service based on the LMM to a customer (or “deployer”), such as a ministry of health, a health-care system, a hospital, a pharmaceutical company or even an individual, such as a health-care provider. The customer who acquires or licenses the product or application may then use it directly for patients, health-care providers, other entities in the health system or laypeople or in its own business. The value chain can be “vertically integrated”, so that a company (or other entity, such as a health system) that collects data and trains a general-purpose foundation model can modify the LMM for a particular use and provide the application directly to users.

WHO recognizes the tremendous benefits that AI could provide to health systems, including improving public health and achieving universal health coverage. Yet, as described in the WHO guidance on the ethics and governance of AI for health (1), it entails significant risks that could both undermine public health and imperil individual dignity, privacy and human rights. Even though LMMs are relatively new, the speed of their uptake and diffusion led WHO to provide this guidance to ensure that they could potentially be used successfully and sustainably worldwide. WHO recognizes that this guidance is being issued at a time of many competing views about the potential benefits and risks of AI, the ethical principles that should apply to its design and use and approaches to governance and regulation. As the guidance is being issued shortly after the first applications of LMMs in health care and before successively more powerful models will be released, WHO will update the guidance to keep up with the rapid evolution of the technology, how society addresses its use and the health consequences of the use of LMMs beyond health care and medicine.

1.1 Significance of LMMs

LMMs, although relatively new and untested, have had an inordinate impact on society in various domains, including health care and medicine. Chat GPT, a large-language model, successive versions of which have been released by a US technology company, was estimated to have 100 million monthly active users in January 2023, just 2 months after its launch. At the time, this made it the fastest-growing consumer application in history (6).

Many companies are now developing LMMs or integrating LMMs into consumer applications, such as Internet search engines. Large technology companies are rapidly integrating LMMs into most applications or creating new ones (7,8). New companies, backed by millions of US dollars in private investment, are developing competing LMMs (9). Open-source LMMs are also emerging more quickly and inexpensively than those developed by the largest companies, due to the availability of such platforms (10).

Even as the advent of LMMs is fuelling new investment in the technology sector, and new products are being released, some companies themselves admit that they do not fully understand why LMMs generate certain responses (11). Despite the use of reinforcement learning with human feedback, LMMs can generate outputs that are not always predictable or controlled, including LMMs taking part in “conversations”, which can be uncomfortable for users (12), or publishing content that is erroneous or otherwise faulty but highly

convincing (13). Nevertheless, much of the support for LMMs is not just enthusiasm for its functionality, but also by unqualified, uncritical claims about the performance of LMMs in publications that have not been peer-reviewed (14).

LMMs have been adopted rapidly, although the datasets used to train them have not been disclosed (15), making it difficult or impossible to know whether the data are biased, whether they were legally acquired and in accordance with data protection rules and principles and whether the performance of any task or query reflects that it has been trained on the same or a similar problem or, rather, has acquired the ability to solve problems. Other concerns about the data used to train LMMs, including consistency with data protection laws, are discussed below.

Neither individuals nor governments were prepared for the release of LMMs. Individuals have not been trained in using LMMs effectively and may not understand that the responses are not always accurate or reliable, even if an LMM-powered chatbot creates such an impression. One study found that, while one large language model, GPT-3, “in comparison to humans... can product accurate information that is easier to understand”, it can also produce “more compelling disinformation”, and humans cannot distinguish content generated by the LMM from that generated by a human (16).

Governments have also been largely unprepared. Regulations and laws written to govern the use of AI may not be fit to address either the challenges or opportunities associated with LMMs. The European Union, which has reached an agreement to enact an European Union-wide Artificial Intelligence Act, had to revise its legislative framework in the final stages of drafting to account for LMMs (17). Other governments are rapidly developing new laws or regulations (18) or have instituted temporary bans (several of which have already been rescinded) (19). Companies are expected to release successively more powerful and capable LMMs in the coming months, which may introduce new benefits but also new regulatory challenges. In this dynamic environment, this guidance, which builds on previous guidance, including on ethics, provides suggestions and recommendations for the use of LMMs in health care and medicine.

1.2 WHO guidance on ethics and governance of AI for health

The original WHO guidance on ethics and governance of AI for health (1) examined various approaches to machine learning and various applications of AI in health care but did not specifically examine generative AI or LMMs. During development of that guidance and at the time of its publication in 2021, there was no evidence that generative AI and LMMs would be widely available so soon and would be applied to clinical care, health research and public health.

Nevertheless, the underlying ethical challenges identified in the guidance and the core ethical principles and recommendations (see Box 1) remain relevant both for assessing and for effectively and safely using LMMs, even as additional gaps in governance and challenges

have and will continue to arise with respect to this new technology. The challenges, principles and recommendations were the basis for the expert group's approach to LMMs presented in this guidance.


Box 1. Brief overview of WHO consensus ethical principles for use of AI for health

- **Protect autonomy:** Humans should remain in control of health-care systems and medical decisions. Providers have the information necessary to use AI systems safely and effectively. People understand the role that AI systems play in their care. Data privacy and confidentiality are protected by valid informed consent through appropriate legal frameworks for data protection.
- **Promote human well-being, human safety and the public interest:** Designers of AI satisfy regulatory requirements for safety, accuracy and efficacy for well-defined uses or indications. Measures of quality control in practice and quality improvement in the use of AI over time should be available. AI is not used if it results in mental or physical harm that could be avoided by use of an alternative practice or approach.
- **Ensure transparency, “explainability” and intelligibility:** AI technologies should be intelligible or understandable to developers, medical professions, patients, users and regulators. Sufficient information is published or documented before the design or deployment of AI, and the information facilitates meaningful public consultation and debate on how the AI is designed and how it should or should not be used. AI is explainable according to the capacity of those to whom it is explained.
- **Foster responsibility and accountability** to ensure that AI is used under appropriate conditions and by appropriately trained people. Patients and clinicians evaluate development and deployment of AI. Regulatory principles are applied upstream and downstream of the algorithm by establishing points of human supervision. Appropriate mechanisms are available for questioning and for redress for individuals and groups that are adversely affected by decisions based on AI.
- **Ensure inclusiveness and equity:** AI is designed and shared to encourage the widest possible, appropriate, equitable use and access, irrespective of age, sex, gender identity, income, race, ethnicity, sexual orientation, ability or other characteristics. AI is available for use not only in high-income settings but also in low- and middle-income countries. AI does not encode biases to the disadvantage of identifiable groups. AI minimizes inevitable disparities in power. AI is monitored and evaluated to identify disproportionate effects on specific groups of people.
- **Promote AI that is responsive and sustainable:** AI technologies are consistent with the wider promotion of the sustainability of health systems, the environment and workplaces.





I. Applications, challenges and risks of LMMs



2 Applications and challenges of use of LMMs in health

Applications of AI for health include diagnosis, clinical care, research, drug development, health-care administration, public health and surveillance. Many applications of LMMs are not novel uses of AI; however, clinicians, patients, laypeople and health-care professionals and workers access and use LMMs differently. This section addresses the potential applications of LMMs in health care and the actual and expected challenges and risks associated with its use. Many applications and uses are still unproven and may ultimately not deliver the benefits that have been advertised.

2.1 Diagnosis and clinical care

AI is already used in diagnosis and clinical care, for instance to assist diagnosis in fields such as radiology and medical imaging, tuberculosis and oncology. It had been hoped that clinicians could use AI to integrate patient records during consultation, to identify at-risk patients and as an aid in difficult treatment decisions and to catch clinical errors (1). LMMs could make it possible to extend use of AI-based systems throughout diagnosis and clinical care – both virtual and in-person consultations, with some experts expecting that LMMs “will be more important to doctors than the stethoscope in the past” (20). Several LMMs have passed the US medical licensing examination; however, passing a written medical test by regurgitating medical knowledge is not the same as providing safe, effective clinical services (21), and LMMs have failed tests with material not previously published online or that could be easily solved by children (22). One study of the clinical knowledge of a large language model concluded that “transitioning from a large language model that is used for answering medical questions to a tool that can be used by healthcare providers, administrators, and consumers will require considerable additional research to ensure the safety, reliability, efficacy and privacy of the technology” (23).

Diagnosis is seen as a particularly promising area, because LMMs could be used to identify rare diagnoses or “unusual presentations” in complex cases (24). Doctors are already using Internet search engines, online resources and differential diagnosis generators, and LMMs would be an additional instrument for diagnosis. LMMs could also be used in routine diagnosis, to provide doctors with an additional opinion to ensure that obvious diagnoses are not ignored. All this can be done quickly, partly because an LMM can scan a patient’s full medical record much more quickly than can doctors (24).

Several popular LMMs now being used in pilot programmes to support clinicians were not, however, trained specifically on electronic health records or medical or other relevant health

data, although their datasets do include such information. For example, in several health-care systems in the USA, an LMM provided by one technology company is being pilot-tested to read messages from patients and to draft responses from doctors to reduce the time that medical staff spend in replying to patient queries. This practice is intended to decrease the burn-out of health-care workers, who field thousands of messages daily, and to enable them to focus on their clinical duties (“keyboard liberation”) (25). Thus, when a patient message is received, the LMM displays a draft reply based on both information from the patient and a version of their electronic medical history. While the AI is used only for some patient questions, the replies require heavy editing (25). Nevertheless, a study in the USA found that a Chat GPT-powered chatbot performed better than qualified doctors in responding to questions posed on an online forum. Of the 195 questions selected, independent evaluators preferred the chatbot responses to the physician responses in nearly 80% of cases (26). Chatbots should also be helpful in answering standardized “curb-side consult” questions and providing information and responses on the initial presentation of a patient or to summarize laboratory test results (27).

Companies and universities are also developing LMMs trained with medical and health data or electronic health records, which include LMMs based on small data sets. For example, one LMM was trained on a dataset of an estimated 30 000 medical case reports to learn the relations between medical conditions and symptoms to assist in diagnoses (28). Another LMM was trained on a dataset of over 100 000 chest X-rays to identify abnormalities and eventually provide insights or identify conditions (29). Several LMMs that have been publicly evaluated were trained with an algorithm on millions of electronic health records and other sources of specialized and general medical knowledge. The approach improved the ability of an algorithm to process different forms of written medical information and to issue responses (“medical question answering”) (30).

Several of the largest technology companies are adapting their general-purpose LMMs to ones that could assist in clinical diagnoses and care. One technology company is developing Med-PaLM 2, which is intended to answer questions and summarize insights from medical texts and is now evolving to synthesize images (such as X-rays and mammograms) to write reports and respond to follow-up questions, to facilitate additional queries by clinicians, a functionality that could mitigate “peer disagreement” between a health-care worker and a computer (31).

The long-term vision is to develop “generalist medical artificial intelligence”, which will allow health-care workers to dialogue flexibly with an LMM to generate responses according to customized, clinician-driven queries. Thus, a user could adapt a generalist medical AI model to a new task by describing what is required in common speech, without having to retrain the LMM or training the LMM to accept different types of unstructured data to generate a response (5).

Risks of use of LMMs in diagnosis and clinical care

The promise of LMMs in clinical care is accompanied by significant risks associated with their use, several of which predate LMMs. Five major risks have been identified of use of LMMs in diagnosis and clinical care.

- *Inaccurate, incomplete, biased or false responses:* One concern with respect to LMMs has been the propensity of chatbots to produce incorrect or wholly false responses from data or information (such as references) “invented” by the LMM (32) and responses that are biased in ways that replicate flaws encoded in training data (33). LMMs could also contribute to contextual bias, in which assumptions about where an AI technology is used result in recommendations for a different setting (1). For example, there is under-representation of training data and perspectives from low- and middle-income countries. Thus, if an LMM is asked to summarize a treatment paradigm for a disease to guide a ministry of health in a low-income country, it might reproduce an approach that is appropriate only for a high-income context (34). Furthermore, an LMM may provide an incomplete answer, no response at all or a response that does not account for changed circumstances in the setting in which it is being used.

False responses, known colloquially as “hallucinations”, are indistinguishable from factually accurate responses generated by an LMM, because even LMMs with reinforcement learning from human feedback are not trained to produce facts but to produce information that looks like facts. One study found that large language models, when provided a simple set of facts to summarize, would hallucinate at least 3 percent of the time and as high as 27 percent (35). Current LMMs also depend on human “prompt engineering”, in which an input is optimized to communicate effectively with an LMM (36). Thus, LMMs, even if trained specifically on medical data and health information, may not necessarily produce correct responses. For certain LMM-based diagnoses, there may be no confirmatory test or other means to verify its accuracy (24). In medicine and other areas of public health decision-making, use of LMMs, even if they are factually correct most of the time, may not be accurate enough to justify the cost of developing them or the cost of implementing them in health-care systems safely and effectively.

- *Data quality and data bias:* One reason that LMMs produce biased or inaccurate responses is poor data quality. Many of the LMMs currently available for public use were trained on large datasets, such as on the Internet, which may be rife with misinformation and bias. Most medical and health data are also biased, whether by race, ethnicity, ancestry, sex, gender identity or age. LMMs trained on health data often encode such biases, as most data are collected in high-income settings. For example, genetic data tend to be collected disproportionately on people of European descent (1). LMMs are also often trained on electronic health records, which are full of errors and inaccurate information (24) or rely on information obtained from physical examinations that may be inaccurate, thus affecting the output of an LMM (25). Problems of data quality and bias affect all AI models, including LMMs (1).

Box 2. Ethical considerations for the use of LMMs by children

While broad guidance has recently been published for safe, ethical use of paediatric data in AI and machine learning (ACCEPT-AI) (43), special consideration must be given to the potential impacts of use of LMMs by children.

The wide availability of open LMMs provides access by users of various ages. There is, however, limited evidence on how children engage with or use LMMs. While the potential opportunities and drawbacks of LMM use have been discussed in wider educational contexts (44), it is unclear how such engagement by children affects their mental or physical well-being. Their use of LMMs must be monitored over time to understand the benefits and potential harms.

Laws and policies on paediatric consent, assent and stipulations for legal parental involvement differ among and within countries. Thus, lack of cohesive, unified, global, child-specific regulation and oversight could result in unidentified, unmonitored harm, especially from use of LMMs.

Specifically, it is not clear how accurately LMMs generalize paediatric health. Studies have shown that generalization of adult data sets to the paediatric population may be limited. Paediatric data should therefore be kept separate in testing and training datasets (45).

Developers should include demographic information on training data that includes age and must be encouraged to provide clear descriptors of target populations, including age ranges, for appropriate, safe engagement with LMMs, as relevant. When legally possible, LMMs should be improved by including appropriate engagement and feedback from young users.

One technology company states on the system card for its LMM (GPT-4): “We found that GPT-4-early and GPT-4-launch exhibit many of the same limitations as earlier language models, such as producing biased and unreliable content” (37). LMMs can be limited by the end date of data with which the algorithm was trained, although some LMMs can now access up-to-date information from the Internet. For example, Chat GPT-4 was trained on data up to September 2021 (38) but can now search or browse the Internet for up-to-date information (39). This could, however, lead to generation of more false or inaccurate information. The previous “end-date” prevented introduction of false newly published material (39). In medicine, both up-to-date and highly accurate information is critical for meeting standards of care as well as for understanding certain diseases.

- *Automation bias*: Concern that LMMs generate false, inaccurate or biased responses is heightened by the fact that, as with other forms of AI, LMMs are likely to encourage automation bias in experts and health-care professionals (and patients, see below). In automation bias, a clinician may overlook errors that should have been spotted by a human (1). There is also concern that physicians and health-care workers

might use LMMs in making decisions for which there are competing ethical or moral considerations (20). LMMs such as Chat GPT may be very inconsistent as moral advisers, although, as recent experiments indicated, they can influence users' moral judgement, even if users know that they are being advised by a chatbot (40). Use of LMMs for moral judgments could lead to "moral de-skilling", as physicians become unable to make difficult judgements or decisions (20).

- *Skills degradation*: There is a long-term risk that increased use of AI in medical practice will degrade or erode clinicians' competence as medical professionals, as they increasingly transfer routine responsibilities and duties to computers. Loss of skills could result in physicians being unable to overrule or challenge an algorithm's decision confidently or that, in the event of a network failure or security breach, a physician would be unable to complete certain medical tasks and procedures (1).
- *Informed consent*: Increased use of LMMs, in person but especially virtually, should require that patients are made aware that an AI technology may be either assisting in a response or could eventually be responsible for generating a response that will mimic a clinician's feedback. Yet, if and when LMMs and other forms of AI are merged into regular medical practice, patients or their caregivers, even if they are uncomfortable or unwilling to rely wholly or partially on an AI technology, may be unable to withhold consent for its use. This is true especially if other options (not based on AI) are not easily available or if the clinician who has handed over responsibility for such functions to a computer cannot provide medical care without use of AI.

2.2 Patient-centred applications

AI is beginning to change how patients manage their own medical conditions. Patients already take significant responsibility for their own care, including taking medicines, improving their nutrition and diet, engaging in physical activity, caring for wounds or delivering injections. AI tools have been projected to increase self-care, including by the use of chatbots, health monitoring and risk prediction tools and systems designed for people with disabilities (1).

LMMs could accelerate the trend towards use of AI by patients and laypeople for medical purposes. Individuals have used Internet searches to obtain medical information for two decades. Therefore, LMMs could play a central role in providing information to patients and laypeople, including by integrating them into Internet searches. Large language model-powered chatbots could replace search engines for seeking information (41), including for self-diagnosis and before visiting a medical provider.

LMM-powered chatbots, with increasingly diverse forms of data, could serve as highly personalized, broadly focused virtual health assistants. According to one study, "virtual health assistants can leverage individual profiles...to promote behaviour change, answer

Box 3. Ethical considerations associated with LMMs and their impact on individuals with disabilities

In the past, individuals with disabilities have been excluded from workplaces, educational systems and appropriate medical support (56) and therefore from data sets used to train AI systems. The systems may discriminate against individuals with facial asymmetry, different gesticulation, styles of communication, behaviour and action patterns. The groups that are most severely affected are people with disabilities, cognitive or sensory impairments or autism spectrum disorder (57).

Such bias and exclusion may apply to generative AI. For example, LMMs may introduce a negative connotation or sentiment to keywords or phrases associated with “disability” in a patient’s description or biography (58). Chatbots may recognize an individual with a disability as “not alive”, “non-human” or “emotionally flat” because of a different behaviour or pattern of actions. Speech recognition systems may be less accurate for individuals with speech impairments, leading to misinterpretation.

Addressing and overcoming biases related to disability require interventions throughout the development of AI: inclusion of people with disabilities in the development and design of AI systems; audits to evaluate disability bias in a dataset and the performance of an AI system; and ensuring that legislation designed to protect and promote the rights of people with disabilities takes into account the challenges associated with AI technologies, while also ensuring laws and policies to regulate AI for probable challenges and barriers faced by people with disabilities with increased use of AI-based systems. AI-specific legislation could include “disability-specific” categorization, including how specific spectrums and conditions are affected by AI systems.

health-related questions, triage symptoms, or communicate with healthcare providers when appropriate” (3). Specific LMM-powered chatbots could provide treatment in, for example, mental health (2).

A third application of patient-centred LMMs could be for identifying clinical trials or for enrolment in such trials (28). While AI-based programmes already assist both patients and clinical trial researchers in identifying a match (42), LMMs could be used in the same way by using a patient’s relevant medical data (28). This use of AI could both lower the cost of recruitment and increase speed and efficiency, while giving individuals more opportunities to seek appropriate trials and treatment that are difficult to identify and access through other channels (42).

Risks and challenges

The ease with which LMMs can be used by individuals may portend significant risks, such as those listed below.

- *Inaccurate, incomplete or false statements:* As in the use of LMMs by clinicians and health-care professionals, use of LMMs by patients and laypeople is associated with risks of false, biased, incomplete or inaccurate statements, including from AI programmes that claim to provide medical information. The risks are heightened when used by a person without medical expertise, who will have no basis for challenging the response, have no access to another source of information or when used by children (Box 2). Although people have used Internet searches to obtain medical information for several decades, an LMM can provide answers that will appear to be “correct”, with reference only to other LMMs (which carry the same risks) for a quick comparison.
- *Manipulation:* Many LMM-powered chatbot applications have distinct approaches to chatbot dialogue, which is expected to become both more persuasive and more addictive (46), and chatbots may eventually be able to adapt conversational patterns to each user (41). Chatbots can provide responses to questions or engage in conversation to persuade individuals to undertake actions that go against their self-interest or well-being (12). Several experts have called for urgent action to manage the potential negative consequences of chatbots, noting that they could become “emotionally manipulative” (47,48). One highly publicized case involved a person in Belgium with anxiety, who committed suicide after 6 weeks of intensive conversation with a chatbot (49).
- *Privacy:* Use of LMMs by patients and laypeople may not be private and may not respect the confidentiality of personal and health information that they share. Users of LMMs for other purposes have tended to share sensitive information, such as company proprietary information (50). Data that are shared on an LMM do not necessarily disappear, as companies may use them to improve their AI models (50), even though there may be no legal basis for doing so, even though the data may eventually be removed from company servers (51). A related problem is sharing of information on an LMM with other users of the LMM, whether because the other user specifically requests the LMM to disclose such information (52) or, in the case of one LMM, mistaken disclosure of other people’s chat histories (even if not the substance of their conversations) (53). Thus, if a person’s identifiable medical information is fed into an LMM, it could be disclosed to third parties (54).
- *Degradation of interactions between clinicians, laypeople and patients:* Use of LMMs by patients or their caregivers could change the physician–patient relationship fundamentally. The increase in Internet searches by patients during the past two decades has already changed these relationships, as patients can use the information they find to challenge or seek more information from their health-care provider. While an LMM could improve such dialogue, a patient or caregiver might decide to rely wholly on an LMM for prognosis and treatment and thereby reduce or eliminate appropriate reliance on professional medical judgement and support. A related concern is that, if an AI technology reduces contact between a provider and a patient, it could reduce the opportunities for clinicians to promote health and could

undermine general supportive care, such as human–human interactions when people are often most vulnerable (1). Generally, there is concern that clinical care could be “de-humanized” by AI.

- *Epistemic injustice*: An additional potential consequence of substituting a health-care provider’s judgement with that of an LMM is introduction of epistemic injustice to the patient. “Epistemic injustice” is a “wrong done to someone specifically in their capacity as a subject of knowledge”, such as a patient in a health-care system (55). One form of epistemic injustice, hermeneutic injustice, occurs when there is a gap in shared understanding and knowledge (so-called “collective interpretive resources”) that puts some people at a disadvantage with respect to their lived experience, social experience or, in the case of health care, their own understanding of their physical or mental condition (55). LMMs, even when trained on large quantities of data, have boundaries with respect to what they can recognize and respond to and the concepts and notions that are outside their vocabulary. If a patient’s experience is not acknowledged or recognized in a clinical setting by an LMM, it can obviate appropriate care from a medical provider, which could harm the patient. This is especially likely for vulnerable groups that are already neglected and underrepresented in data (55), such as people with disabilities (Box 3).
- *Delivery of health care outside the health-care system*: AI applications in health are no longer used exclusively or accessed and used within health-care systems or in home care, as AI technologies for health can be readily acquired and used by non-health system entities or simply introduced by a company, such as those that offer LMMs for public use. This raises questions about whether such technologies should be regulated as clinical applications, which require greater regulatory scrutiny, or as “wellness applications”, which require less regulatory scrutiny. At present, such technologies arguably fall into a grey zone between the two categories.

LMMs that are “lightly” regulated could present a risk if they are used by a patient with no regulatory safeguard. This includes use of an LMM for medical advice or self-diagnosis. There is concern that patients may receive false or misleading advice (see above) and that patient safety may be compromised if individuals are not in contact with health-care services, including lack of supportive care for individuals with suicidal ideation who use an AI chatbot, even if the chatbot is not “manipulative”. Even if the information is correct, individuals without medical training who use such information for self-diagnosis could misinterpret or misuse it. As such applications, including LMMs, continue to proliferate and are not necessarily labelled as health-care applications, the overall quality of health care could be compromised. This could further exacerbate unequal access to good-quality health care, especially as people who lack other options may resort to such applications (1).

2.3 Clerical functions and administrative tasks

LMMs are beginning to be used to assist health professionals in clerical, administrative and financial aspects of practising medicine. Medicine exposes physicians and other health-care professionals in many settings to ever-growing paperwork for numerous obligations for recording patient information and data in electronic health records, billing in private, insurance or public health-care systems and other administrative tasks. While many such obligations, such as completing electronic health records, were intended to “liberate” health-care professionals, most are now a major cause of physician and health-worker burn-out (59). According to one study, documentation constituted one fourth to one half of a doctor’s time and one fifth of a nurse’s time (59).

LMMs have been singled out as a means of returning to health-care professionals their most valuable commodity – time – both to reduce burn-out, to allocate more time to providing care to each patient or to see more patients. One physician who used software that encodes an LMM to document patient visits reported that “AI has allowed me, as a physician, to be 100 percent present for my patients” and that the software had freed up to 2 h daily (60).

Examples of current and anticipated uses of LMMs include:

- better communication to assist in translation or improving clinician–patient communication by simplifying medical jargon and making communication more “patient-friendly” (34);
- to fill in missing information in electronic health records (61); and
- with other forms of AI, to draft clinical notes after each patient visit (virtual or in person) (62).
- Use of LMMs is also expected to include pre-emptive writing of automated prescriptions and appointments, billing codes, scheduling of tests, pre-authorization by insurance companies, procedure notes and discharge summaries (5). As more sophisticated LMMs are developed, they could be used for even more complex clerical note-taking, such as with radiologists, by “automatically draft(ing) radiology reports that both describe abnormalities and relevant normal findings, while taking into account the patient’s history...these models can provide further assistance to clinicians by pairing text reports with interactive visualization, such as by highlighting the region described by each phase” (5).

Risks and challenges

As with other uses of LMMs, serious errors could arise, due to inaccuracies, mistakes (for example in transcription, translation or simplification) or “hallucinations”. It is therefore important that most clerical and administrative functions not be completely automated. Even

if oversight and review claim the time of a health-care professional, an LMM is still likely to be less burdensome than the status quo. Another problem is that LMMs can be inconsistent; slight changes to a prompt or question can generate a completely different electronic health record, although inconsistencies are expected to decrease over time (63).

2.4 Medical and nursing education

LMMs are also projected to play a role in medical and nursing education. They could be used to create “dynamic texts” that, in comparison with generic texts, are tailored to the specific needs and questions of a student (63). LMMs integrated into chatbots can provide simulated conversations to improve clinician–patient communication and problem-solving, including practising medical interviewing, diagnostic reasoning and explaining treatment options. A chatbot could also be tailored to provide a student with various virtual patients, including those with disabilities or rare medical conditions. LMMs could also provide instruction, in which a medical student asks questions and receives responses accompanied by reasoning through a “chain-of-thought” including physiological and biological processes (63).

Risks and challenges

Although use of AI can improve or sharpen the training and skills of a health-care professional, it could also pose the risk that professionals suspend their judgement (or that of a human peer) in favour of that of a computer. If an LMM provides incorrect information or responses (or fabricates a response), it could affect the quality of medical education.

An additional concern is that use of LMMs in education or to simplify administrative and clerical functions could place an additional burden on health-care workers who are not yet digitally literate and will have to develop new competence in use of AI-supported technologies in everyday practice (1). The new functionalities of LMMs are expected to require health-care professionals to continue retraining and adjusting (1). Developers may eventually introduce AI-supported technologies with communication interfaces that can be used easily by laypeople, such as natural language or vision.

2.5 Scientific and medical research and drug development

AI is already used in scientific and clinical research and in drug development. With AI, electronic health records can be analysed to identify clinical practice patterns and to develop models of new clinical practice. Machine learning is also used in genomic medicine, for example to improve understanding of a disease and identify new biomarkers (1). AI is used in almost every stage of the drug development cycle, including to streamline compound screening, to predict the three-dimensional shape of a protein (the “protein folding problem”) (1), to predict the toxicity and effectiveness of compounds in preclinical

development, and to improve the recruitment, enrolment and monitoring of individuals during clinical trials.⁴

LMMs are extending the ways in which AI can support scientific and medical research as well as drug discovery. LMMs can be used in a variety of aspects of scientific research. They can generate text to be used in a scientific article, for submitting manuscripts or in writing a peer review (34). They can be used to summarize texts, including summaries for academic papers, or can generate abstracts. LMMs can also be used to analyse and summarize data to gain new insights in clinical and scientific research. They can be used to edit text, improving the grammar, readability and conciseness of written documents such as articles and grant proposals. Additionally, LMMs can be used to gain insights from the data with which it was trained (34). One LMM, which was trained on millions of academic articles, is claimed to be able to analyse scientific research to answer questions, extract information or generate relevant text (64). LMMs are also used in drug discovery and specifically in de-novo drug design to develop new compounds with specific properties (65).

Risks and challenges

Leading medical and scientific journals have already responded to the emergence of LMMs, their potential and their impact on scientific research. For example, one academic publishing company has set two rules: (1) an LMM will not be accepted as a credited author on a research paper, and (2) researchers who use LMMs should document their use in the section on methods and in the acknowledgements (66). The World Association of Medical Editors has restricted authorship to humans (67).

General concerns about use of LMMs in scientific research include the following:

- *Lack of accountability:* The authorship of a scientific or medical research paper requires accountability, which cannot be assumed by AI tools (66). Lack of accountability was the basis for the decision of a major academic publisher and the World Association of Medical Editors not to accept an LMM as a credited author.
- *High-income country bias:* Most of the scientific and medical research used to train LMMs is conducted in high-income countries. Thus, the outputs of any LMM query are likely to be biased towards a high-income country perspective (34). This can reinforce and possibly exacerbate the trend of ignoring and not citing research that originates from a low- or middle-income country (68), especially if the publication is written in non-Latin script.
- *Hallucination and/or misinformation:* An LMM may “hallucinate” by summarizing or citing academic articles or other information that does not exist (69).

4 Ethics and governance of artificial intelligence in drug development. Geneva: World Health Organization; forthcoming.

- *Undermining of trust:* Use of LMMs for activities such as generating peer reviews could undermine trust in that process (69).
- *Accessibility of LMMs and of knowledge generated by LMMs:* Like other tools, technologies and information used in scientific and medical research, LMMs are expected to “sit behind a paywall”, coming at a cost and thereby exacerbating the digital and knowledge divide and affecting scientists with less support and funding who seek to participate in scientific and medical research (34).

Although AI (including LMMs) can benefit drug development, there is also concern about the use of AI in this area, which is examined in a forthcoming WHO publication.

3 Risks to health systems and society and ethical concerns about use of LMMs

Whereas many risks and concerns associated with LMMs affect individual users (such as health-care professionals, patients, researchers or caregivers), they may also pose systemic risks. Emerging or anticipated risks associated with use of LMMs and other AI-based technologies in health care include: (i) risks that could affect a country's health system, (ii) risks for regulation and governance and (iii) international societal concerns.

3.1 Health systems

Health systems are based on six building blocks: service delivery, the health workforce, health information systems, access to essential medicines, financing, and leadership and governance (70). LMMs could directly or indirectly impact these building blocks. The risks associated with use of LMMs that could affect health systems are described below.

Overestimating the benefits of LMMs and discounting risks

Some people tend to overstate and overestimate what AI can do, which can result in the uptake of unproven products and services that have not been subjected to rigorous evaluation for safety and efficacy (1). This is due partly to the enduring appeal of “technological solutionism”, in which technologies such as AI and LMMs are considered to be “magic bullets” for removing deeper social, structural, economic and institutional barriers (1), before it is shown that such technologies are useful, safe and effective.

LMMs are novel and untested and, as noted above, do not produce facts but information that resembles facts and may be inaccurate. They have been the subject of intense consumer, political and public interest but could lead policy-makers, providers and patients to overestimate their benefits and to dismiss the challenges and problems that LMMs could introduce. For policy-makers, it may not be possible to obtain the evidence necessary to determine how widely LMMs should be used until they are already developed and used. Use of an LMM should not be prioritized over AI-based technologies that are already in use or over non-AI or non-digital solutions that may be underfunded and underused but have proven therapeutic or public health benefits. Unbalanced health care policy and misguided investments, especially in resource-constrained low- and middle-income countries, can divert attention and resources from proven interventions and intensify the pressure on health ministries to reduce public expenditure on health care (1).

Accessibility and affordability

Several factors can undermine equitable access to LMMs that could provide benefits to providers and patients. One is the digital divide, which limits use of digital tools to certain countries, regions or segments of a population. The digital divide leads to other disparities, many of which affect the use of AI, and AI itself can reinforce and exacerbate disparity. Another factor that could undermine access to LMMs is that, unlike the Internet, many LMMs are available only by paying a fee or subscription, as both developing and operating an LMM can be expensive. It has been estimated that Chat GPT costs US\$ 700 000 per day to operate (71). Some companies are introducing subscription fees for new versions of LMMs (72), which could make certain LMMs unaffordable, not just in low- and middle-income countries, but also for individuals, health systems or local governments in resource-poor settings in high-income countries (54). Conversely, poor people, in all countries, could be limited to using LMMs as a “cost-effective solution”, whereas access to “real” health-care professionals will be limited to wealthier individuals. A third factor is that most LMMs, at present, operate only in English. Thus, while they can receive inputs and provide outputs in other languages, they are more likely to generate false information or misinformation (73).

System-wide biases

As noted above, the data sets used to train AI models are biased, as many exclude girls and women, ethnic minorities, elderly people, rural communities and disadvantaged groups. In general, AI is biased towards the populations for which there are most data, so that, in unequal societies, AI may disadvantage a minority population (1). Biases are likely to increase with the scale of a model (74), which may be a particular problem with LMMs, as the data used to train successive models continues to increase, even though so-called smaller LMMs are being developed. Bias could introduce discrimination throughout a health-care system, affecting people’s access to basic goods, including health services and high-quality care (75). At the same time, LMMs are likely to include data that could “push back” against forms of bias and stereotyping. Researchers have found that prompting a model not to rely on stereotyping had a dramatic positive effect on the algorithm’s response (74).

Impact on labour and employment

An investment bank estimated that LMMs will eventually result in the loss (or “degradation”) of at least 300 million jobs (76). A report by the Organisation for Economic Co-operation and Development noted that, in its member countries, the occupations at highest risk from AI-driven automation are highly skilled jobs, due partly to use of LMMs and specifically that “occupations in finance, medicine, and legal activities...may suddenly find themselves at risk of automation from AI” (77). For many countries, however, health care is not an industry but a core government function, and health-care workers might not be replaced by technology. Furthermore, many countries continue to have shortages of health-care workers (1), including in the aftermath of the coronavirus disease 2019 pandemic (78). WHO has estimated that, by

2030, there will be a shortage of 10 million health workers (79), mainly in low- and low-middle income countries. Therefore, LMMs that are proven to be safe and effective might be used to narrow the gap between the workforce required to provide health care and that which is available.

A separate concern is the impact of introduction of LMMs on the number of current and future health-care professionals. A major technology company estimated that up to 80% of jobs will be affected by the arrival of AI (80). Accenture, a consulting firm, estimated that 40% of working hours could be affected by LMMs and noted optimistically that “the positive impact on human creativity and productivity will be massive” (81). Yet, as noted above, the introduction of LMMs could create significant challenges for many health-care workers, who will require training and adjustment to LMMs. Health systems will have to account for the challenges it presents to providers and the risks to patients and caregivers.

A third concern is the mental and psychological toll on people responsible for reviewing content, annotating data used to train LMMs and removing abusive, violent or mentally disturbing content from a data set. Those responsible for filtering out such content are often based in low- and middle-income countries, earn low wages and can suffer psychological distress from reviewing such content, without access to counselling or other forms of medical care (73).

Dependence of health systems on unsuitable LMMs

While LMMs could counter persistent shortages of health workers and extend the reach of health systems, those systems could become dependent on LMMs and specifically on LMM technologies developed by industry. Thus, if LMMs used in health care or public health are not maintained, are diminished or are designed and updated only for use in high-income contexts, the health systems that have relied on them will have to adjust and potentially deliver health care without LMMs. This may be difficult if medical professionals have been “de-skilled” and have outsourced certain responsibilities to AI or if patients expect its use. A related risk is that, if LMMs do not preserve a patient’s privacy and confidentiality, overdependence on LMMs could undermine individual and societal trust in health-care systems, as people would no longer be confident in accessing health-care services without risking their privacy.

Cybersecurity risks

As health-care systems become increasingly dependent on AI, the technologies could be targeted for malicious attacks and hacking and some systems shut down, with manipulation of the data used to train the algorithm, thereby changing its performance and recommendations, or data could be “kidnapped” for ransom (1). A particular security risk, noted above, is feeding of sensitive data into LMMs that are not protected from unauthorized disclosure or use. LMMs themselves may also be vulnerable to cyber-security risks, such as “prompt injection”, which is an attack in which data are fed to an LMM by a third party,

causing it to behave in ways that the developer did not intend (82). A prompt injection could, for example, instruct an LMM designed to answer questions about a database to delete information from the database or to change the information. There is as yet no known solution to address this flaw. Even though prompt injections are currently used by security researchers to illustrate the challenges of LMMs, they could be used by malicious actors to steal data or to defraud users (83).

3.2 Compliance with regulatory and legal requirements

Although new laws may be enacted to regulate the use of AI, certain existing laws and regulations, in particular data protection laws as well as international human rights obligations, are applicable to the development, provision and deployment of LMMs. Some LMMs, as currently developed and introduced for public use, may violate several major data protection laws, such as the European Union's General Data Protection Regulation (84), which covers various rights, such as protection from automated decision-making. Such rights, protections and requirements must guide the development of AI (85).

Some such violations have led to investigations of LMMs in European Union Member States (83) and elsewhere, such as in Canada (86). The violations have included: (i) LMMs scraped and used the personal data of individuals from the internet without their consent (and without a 'legitimate interest' for collecting such data) (87); (ii) LMMs that cannot inform people that they are using their data or give them the right to correct mistakes, to have their data erased (the "right to be forgotten") or to object to use of such data (87); (iii) LMMs that are not fully transparent in their use of sensitive data provided to a chatbot or other consumer interface, although, by law, a user must be able to delete chat log data (83); (iv) LMMs that do not have an appropriate "age-gating" system to filter out users under the age of 13 and those aged 13–18 for whom parental consent has not been given (88); (v) LMMs that cannot prevent breaches of personal information (87); and (vi) LMMs that publish inaccurate personal information, due partly to hallucinations (89). Other possible contraventions of the General Data Protection Regulation include the "right to explanation" requirement, such that an entity that uses personal data for automated processing can explain how the system, such as an LMM, makes decisions.⁵ As noted above, companies cannot yet explain how LMMs make decisions, although some are working on approaches to satisfy the "explainability" requirement (90).

Many violations are significant. They are associated with how LMMs were trained, how they are used and how they can be managed by a data controller. It is possible that LMMs may never be compliant with the General Data Protection Regulation or other data protection laws (91). A complaint filed in 2023 with the data protection authority of one European Union Member State alleged that the large language model of one company and the approach for which it was developed and is now operated systematically breached the General Data Protection Regulation (92).

⁵ See Article 15(1)(h) and Recital 71 of the General Data Protection Regulation.

Many violations of data protection may also violate consumer protection laws (93). More broadly, if these problems cannot be resolved, they are also in direct contravention of the WHO guiding principles on use of AI in health, including the principles of protecting autonomy and of ensuring transparency, “explainability” and intelligibility.

The inability of companies to abide by existing laws may be the reason for the expression of serious concern about forthcoming AI regulations by some companies. In response to the planned introduction by the European Union of an Act on AI, the head of a major company stated that it might be unable to offer its key LMM product in Europe, as it might not be able to comply with its regulations (94). This ultimatum could result in the erosion of privacy rights and other protections, or provision of health care which will depend on the willingness of a country to forego certain human rights.

3.3 Societal concerns and risks

Like other AI technologies, LMMs are expected to have broader societal impacts, beyond the health system, that cannot be addressed by one law or policy. They include the likelihood that LMMs will reinforce the power and authority of a small group of technology companies (and their executives) that are at the forefront of commercializing LMMs. LMMs may also have a negative impact on the environment and climate, due to the carbon and water used in training and using LMMs. Such technologies quickly become “entangled with the lives of billions of people at a pace faster than cultures can safely absorb them” (4), including in the domains of health care and medicine, before humans can ensure that AI technologies do not displace our epistemic authority with information, evidence and recommendations that are often inaccurate, false or biased and lack any moral or contextual reasoning. There is also serious concern that LMMs could augment technology-facilitated gender-based violence, including cyber-bullying, hate speech and non-consensual use of images and videos, such as “deep fakes”. The latter risk is not addressed in this report, but it merits broader consideration by WHO, as it has serious negative implications for the health and well-being of populations, especially adolescent girls and women, who are targeted by such uses of AI (95).

Challenges related to large technology companies

The emergence of LMMs, which continue to grow larger with more and more parameters, has reinforced the dominance and centrality of a few large technology companies that develop and deploy AI (96). Few companies and governments have the human and financial resources, expertise, data and computing power to develop increasingly sophisticated LMMs (96). Computing power and investment in LMMs have increased, and, with the demand for AI growing, recruiting “AI talent” is expensive (97,98). LMMs that contain the most powerful microchips available require many computers and thousands of chips working together to be trained, the computers working non-stop for weeks or even months (99).

As the cost of training, deploying and maintaining LMMs continues to rise, there is a risk of “industrial capture” by a few companies of what is potentially a building block of many products and services (including in health care) in a way that could crowd out universities (academics), start-ups and even governments (100). In research on AI, there is already compelling evidence that the largest companies are crowding out both universities and governments.

One indication is where doctoral graduates in AI are choosing to work. Presently, “unprecedented” numbers are choosing to work in companies. While, in 2004, only about 20% of graduates went into industry, by 2020, almost 70% were in industry (101). Faculty members who specialize in AI are being hired away from universities to work in industry, the number having risen by eight times since 2006, not only in the USA but also in other countries (101). Industry has also come to dominate both governments and academia with respect to computing power and the use of large datasets. In 2021, industry models were 29 times larger than academic models (101). Furthermore, raw spending, especially by high-income governments, lags far behind that of industry. As noted in one study: “In 2021, nondefense US government agencies allocated US\$ 1.5 billion on AI. In that same year, the European Commission planned to spend €1 billion (US\$ 1.2 billion). By contrast, globally, industry spent more than US\$ 340 billion on AI in 2021, vastly outpacing public investment” (101).

The dominance of “AI inputs” means that large technology companies also now dominate outputs and outcomes. The industry share of the largest AI models increased from 11% in 2010 to 96% in 2021, while the number of research reports with one or more industry co-authors increased by 16% between 2000 and 2020 (101).

The dominance of large technology companies defines not only the applications and uses of AI but also, increasingly, the priorities for early-stage research (101). The dominance of industry and lack of government investment also mean that alternatives for important AI technologies that are in the public interest, including health care and medicine, could become increasingly rare in the absence of appropriate public investment. This differs from the case in the pharmaceutical sector, for example, where there is substantial government, not-for-profit and philanthropic investment in research and development, especially in the critical early stages of drug development and also for late-stage development of certain therapeutics (102). Companies will therefore increasingly oversee operation of the systems that underpin our economies and social sectors, including health care, raising concern about the ability of citizens and officials to manage their lives (101).

In the absence of alternatives and of regulation (which may take several years to be fully implemented, even if laws are enacted in 2023), how large technology companies make internal decisions and how they relate to societies and governments becomes more and more relevant. Companies might start to address various concerns, in partnership, through, for example, the Frontier Model Forum (103) or with the governments of high-income countries, including several voluntary commitments with the US Government (104) and forthcoming commitments with the European Union (105).

A further concern is that companies may not maintain a corporate commitment to ethics. For example, major technology companies have been either side-lining or eliminating their ethics teams that were established to ensure that the design and development of AI models would adhere to internal ethics principles (106), and to introduce “friction” that would require the company to slow or cease certain development activities. Eliminating entire teams working on ethics-related issues for AI means that principles are not “closely tied to product design” (108) and therefore leaves a gap.

Several large technology companies, through the Frontier Model Forum, have committed themselves to ensuring “responsible and safe development of frontier AI models”, including LMMs, “identifying best practices for the responsible development and deployment of frontier models” and “collaborating with policymakers, academics, civil society, and companies to share knowledge about trust and safety risks” (103). In voluntary commitments with the US Government, technology companies have pledged to avoid harmful bias and discrimination and protect privacy (104). It is not clear, however, whether voluntary commitments or partnerships will suffice to replace a robust commitment to ethics. Ethics teams in one company, for example, that had recommended halting the release of a new LMM changed their documents and downplayed previously documented risks (106).

Large technology companies have neither a history of nor are they specialized in the development of health products and services. They may therefore not be sensitive to the requirements of health-care systems, providers and patients and may not address privacy or quality assurance, for example, which are familiar to traditional health-care companies and public health providers. Their sensitivity may improve with time, as occurred in other companies that have provided health-care goods and services for several decades.

Many companies that are developing LMMs are not transparent with either governments or regulators or with those companies that may use their models and which may require (i) evidence, data, performance and other information to assess the risks and benefits of an LMM (97,106) and (ii) the number of parameters in the model, which is an indicator of how powerful it is (8). Companies that use such models in developing their own products and services also do not disclose how they assess ethical challenges and risks, the safeguards in place, the reaction of LMMs to those safeguards and when use of technology should be limited or stopped. The Foundation Model Transparency Index, which assesses ten leading developers of large language models against one hundred indicators, found that “no major foundation model developer is close to providing adequate transparency, revealing a fundamental lack of transparency in the AI industry” (107). The voluntary agreement between the US Federal Government and several large technology companies includes two commitments to transparency. The companies commit themselves to (i) share information on managing risks with the industry and with governments, civil society and academia and (ii) to publicly report the capabilities, limitations and areas of appropriate and inappropriate use of their AI systems (104). While these commitments might be an improvement over the status quo, they are voluntary and are open to interpretation by each company, which may not result in adequate disclosure without concrete regulatory requirements.

Companies are rushing new LMMs to market as quickly as possible, because of internal commercial pressure or external competition (106), before they fully understand how the LMMs function (109) and irrespective of whether appropriate testing, safeguards and ethical risks and concerns have been identified and addressed (106,110). One company executive remarked that it was an “absolutely fatal error in this moment to worry about things that can be fixed later” (106). Companies seek a “first-mover” advantage, because the market share of LMMs in certain domains, such as Internet search, provide revenues. According to one company, every 1% of market share in a search engine equals US\$ 2 billion in additional revenue (108). An executive at a major technology company remarked that its LMM was “not perfect” but that it would be released because “the market demands it” (8). Companies that release LMMs without fully identifying, validating, accounting for and mitigating risks accumulate an “ethical debt”, the eventual consequences of which will be felt not by the companies but by those most vulnerable to the negative impacts of such technologies (109). Members of the Frontier Model Forum have committed themselves to “advancing AI safety research” and “identifying best practices” (103), and voluntary commitments to the US Government include internal and external testing of AI systems before their release (104).

Commercial pressure may not only result in companies rushing LMMs to market as quickly as possible but also to their de-prioritizing or abandoning products and services that have a significant public health benefit in order to prioritize services that could generate revenue. In 2023, one major technology company “axed” a team that had developed an LMM called ESMFold, a protein language model that can predict a full atomic-level protein structure from a single sequence and which generated a database of more than 600 million protein structures. There is concern that the company may not be willing to “absorb the costs to keep the database running, as well as another service that allows scientists to run the ESM algorithm on new protein sequences” (111).

Carbon and water footprints of LMMs

Another consequence of the increasing size of LMMs is their environmental impact. LMMs require large amounts of data, and training with data therefore consumes significant amounts of energy (112). In one large company, training of a new LMM used an estimated 3.4 GWh over 2 months, equivalent to the annual energy consumption of 300 US households (112). While some LMMs are trained at data centres that use renewable or carbon-free energy, most AI models are trained with electricity grids powered by fossil fuels (112). Electricity consumption will continue to increase as more companies introduce LMMs, which could eventually significantly affect climate change.

WHO considers climate change an urgent global health challenge that requires prioritized action, now and in the decades to come. Between 2030 and 2050, climate change is expected to cause approximately 250 000 additional deaths per year due to malnutrition, malaria, diarrhoea and heat stress. The cost of direct damage to health by 2030 is estimated to be US\$ 2–4 billion per year. Areas with weak health infrastructure, most of which are in low- and middle-income countries, will be least able to cope without assistance to prepare and respond (1).

LMMs also have a significant water-use footprint. Training of an early LMM in a large technology company consumed 700 000 L of fresh water, which would have been even greater at other data centres (113). Although many developers are increasingly aware of their carbon footprint, many are not aware of their water footprint (114). A short conversation by an LMM (20–50 questions and answers) requires the equivalent of a 500-mL bottle of water. The overall water footprint of training an LMM, comprising all the water consumed, including AI server manufacture, transport and chipmaking, may be significantly larger. (114). Data centres can stress local water supplies. For example, the data centre of one company used more than 25% of all the water in a city in Oregon, USA (114). Another large technology company is planning to construct a data centre in a country that is in the midst of severe drought, such that residents are obliged to drink salty water (115). Tracking water footprints is difficult, because, while there is greater awareness, measurement and transparency about carbon footprints, companies are not equally transparent about their water footprint or do not measure it (114).

“Dangerous” algorithms that displace the epistemic authority of humans

A more general societal risk associated with the emergence of LMMs is that, by providing plausible responses that are increasingly considered a source of knowledge, LMMs may eventually undermine the epistemic authority of humans, including in the domains of health care, science and medicine. LMMs do not in fact generate knowledge or understand what they are “saying” or have any moral or contextual reasoning in answering questions.


If the concerns persist, societies may not be prepared for the consequences of computer-generated reasoning. Earlier forms of AI that supplied information through social media algorithms spread misinformation, with negative impacts on mental health and increased polarization and division (4). Even as technology companies issue repeated warnings of the dangers of LMMs, they continue to release them directly into society without safeguards or regulatory oversight in ways that could not just displace human control of knowledge production but reduce the ability of humans to use knowledge safely, in health care, medicine and elsewhere, in systems on which societies depend. Such harm could affect people and communities in resource-poor settings in particular, as their data are unlikely to have been used to train an AI system, thereby reducing the accuracy of the responses. Such groups might, however, be likely to follow the advice of an AI system, especially if no health-care professional or medical provider is available to contextualize or correct an LMM-generated response that is false or inaccurate.

Release of ever-more imperfect information or disinformation into the public domain and knowledge bases by LMMs could eventually lead to “model collapse”, whereby LMMs trained on inaccurate or false information also pollute public sources of information, such as the Internet (116,117). To avert such a scenario while maximizing the benefits of LMMs in health care and other areas of social importance, governments, civil society and the private sector must steer these technologies towards the common good.





II. Ethics and governance of LMMs in health care and medicine



The ethical principles defined by the WHO expert group (see above) provide guidance to stakeholders on the basic ethical requirements that should direct their decisions and actions in development, deployment and assessment of the use of LMMs in health care and medicine. The principles should be the basis for how governments, public sector agencies, researchers, companies and implementers govern the use of LMMs.

Governance comprises the steering and rule-making functions of governments and other decision-makers, including international health agencies, for achieving national and global health policy conducive to universal health coverage. Governance is also a political process that involves balancing competing influences and demands (1). Current laws and policies are unlikely to suffice for effective management of the use of LMMs, as many were written before the release of earlier versions of LMMs. Governance of LMMs, as with overall governance of AI, involves applying current and new legislation and regulations, “soft law” (such as ethical principles), human rights obligations, codes of practice and internal procedures of companies, industry associations and standard-setting bodies.

Currently, LMMs are deployed more quickly than our ability to fully understand their capability and frailty. An early suggestion for addressing concern about LMMs was for a ban or a moratorium on their development (118). While some countries do restrict the use of or even ban certain LMMs, most governments now seek to ensure that their use can be directed towards socially beneficial outcomes through appropriate governance. Leading AI companies have also called for careful, deliberative development of LMMs and other forms of AI. Neither governments nor companies, however, are immune to competition. Several governments are locked in an “arms race” to technological supremacy, while even AI companies that are calling for regulation are not immune to commercial pressure (119). While optimists consider that many of the challenges and risks of AI can be addressed through design, including ever-larger data sets and more powerful algorithms, critics have pointed out that the limitations of LMMs are systemic and that increasing the size of training data and model parameters will not overcome shortcomings but will in fact amplify them (59).

Governance of LMMs must keep pace with its rapid development and increasing uses and should privilege neither governments seeking a technological advantage nor companies seeking commercial gain. The initial suggestions and recommendations below place ethical principles and human rights obligations at the centre of appropriate governance, comprising both procedures and practices that could be introduced by companies and laws and policies enacted by governments.

LMMs can be considered products of a series (or chain) of decisions on programming and product development by one or more actors. Decisions made at each stage of an AI value chain may have both direct and indirect consequences on those that participate in the development, deployment and use of LMMs downstream. The decisions can be influenced and regulated by governments that enact and enforce laws and policies nationally, regionally and globally. The AI value chain begins with integration of several inputs, which comprise the “AI infrastructure”, such as data, computing power and AI expertise, to the development of general-purpose foundation models. These models can be used directly by a user to perform

various, often unanticipated tasks (including those related to health care). Several general-purpose foundation models are trained specifically for use in health care and medicine.

Appropriate governance of LMMs used in health care and medicine should be defined at each stage of the value chain, from collection of data to deployment of applications in health care. Therefore, the three critical stages of the AI value chain discussed are:

the design and development of general-purpose foundation models (design and development phase);

- definition of a service, application or product with a general-purpose foundation model (provision phase); and
- deployment of a health-care service application or service (deployment phase).

At each stage of the AI value chain, the following questions are asked.

- Which actor (the developer, the provider and/or the deployer) is best placed to address relevant risks? What risks should be addressed in the AI value chain?
- How can the relevant actor(s) address such risks? What ethical principles must they uphold?
- What is the role of a government in addressing risks? What laws, policies or investment might a government introduce or apply to require actors in the AI value chain to uphold specific ethical principles?

During the design and development phase, the focus is on the practices that developers can introduce to uphold ethical commitments and norms and government policies and investments. During the provision phase, the focus is on the measures that governments can introduce to assess and regulate use of LMMs in health care and medicine. During the deployment phase, measures are used by governments and all actors in the value chain to ensure that any potential or actual harm to users is identified and avoided.

4 Design and development of general-purpose foundation models (LMMs)

General-purpose foundation models are usually trained on a vast amount of data, requiring tremendous computing power. Development of LMMs also requires specialized human resources, including scientific and engineering expertise. The WHO guidance on ethics and governance of AI for health (1) recommends that developers of medical AI “should invest in measures to improve the design, oversight, reliability and self-regulation of their products”.

Although most findings and recommendations below could apply to all general-purpose foundation models, the guidance is intended for such models that may be or are used in health care and medicine (either directly by a user or through an application or service). The recommendations below are also intended to guide the design and use of LMMs trained specifically for use in health care and medicine, which may be used directly by users or through an application or service.

4.1 Risks to be addressed during the development of general-purpose foundation models (LMMs)

The design and development of general-purpose foundation models can introduce serious risks that, if left uncorrected, could have either a broad societal impact or specific negative consequences on the users of an LMM. Elimination or mitigation of such risks is the responsibility of the developer, because it is only the developer who can (or could) make certain decisions during design and development, which are beyond of the control of providers and deployers that may use the algorithm (and which cannot be mitigated by correct use of the technology by a provider, deployer or user) (120). The decisions refer, for example, to the data used to train an LMM (121). Obligations to ensure data protection and quality and to mitigate bias are also outside the control of the downstream developer of an application (121), as are measures that might have to be introduced to ensure that LMMs do not issue “AI-fuelled toxicity” (122). Failure to hold the developers of LMMs accountable for such design flaws shields the companies with the most resources from, as one report noted, “responsibility to tackle problems...which their methods may be thoughtlessly baking in as they rush to dominate a new form of applied AI” (122).

At least eight risks should be addressed by the developer of a general-purpose foundation model, including through government laws and regulations:

- bias (associated with the design and training data);

- privacy (of training and other input data);
- labour concerns (outsourced filtering of data to remove offensive content);
- the carbon and water footprints;
- false information, hate speech or misinformation;
- safety and cybersecurity; and
- preserving the epistemic authority of humans
- exclusive control of LMMs

4.2 Measures developers can take to address risks with general-purpose foundation models (LMMs)

A developer could use many measures or practices to address such risks, whether as a commitment to ethical principles or policies or to meet the requirements of governments.

AI expertise (scientific and engineering personnel): A developer can ensure that its scientific and programming personnel can identify and avoid risks. The WHO ethics guidance (1) made several recommendations for the training of scientific and engineering personnel and on the inclusiveness of the design process. In particular, the WHO expert group recommended that developers consider “licensing or certification requirements for developers of ‘high-risk’ AI, including AI for health”.

Companies and other entities that develop an LMM that shall or could be used in health care, scientific research or medicine should consider certification or training to align themselves with requirements in the medical profession and also to increase trust in their products and services (1). Any standards, introduced and enforced by either developers or professional societies, should be written in collaboration with or by government regulators and should be consistent with the WHO ethical principle of promoting human well-being, safety and the public interest. Developers that do not intend but can foresee that their LMM might be used in health may wish to ensure internal expertise to anticipate and address such uses.

Data: While human resources and computing power are essential for the development of LMMs, data are probably the most critical infrastructure requirement. The quality and type of data used to train LMMs determine whether it meets core ethical principles and legal requirements (123). Although AI developers have agreed in qualitative surveys that data quality “matters” and require a significant commitment of time, work related to data is often under-valued, which can have significant negative repercussions for AI in “high-stakes” domains such as health care and medicine (123). If data are not of the appropriate quality, several WHO guiding principles could be violated, including the promotion of human well-being, safety and the public interest and the principle of ensuring inclusiveness and equity if the data are biased.

As use of data for health care will probably require strict adherence to laws for informed consent, developers who train LMMs intended for use in health care and medicine might have to rely on smaller data sets (59). Smaller data sets might also be preferable to ensure the quality of data, that the data are diversified in order to avoid bias (59) and that they reflect the composition and reality of the population(s) that will be served by the LMM. Smaller data sets might, however, increase the risk of re-identification of individuals, which would expose them to current or future harm. Reliance on smaller data sets may have further benefits, including reducing the carbon and water footprints (112) of the models and also making it possible for smaller entities to participate in or develop LMMs that require fewer data, computing, human and financial resources (59). Regardless of the size of the data set, developers should undertake “data protection impact assessments”, which would, as required under the General Data Protection Regulation, require developers to assess the risks of data-processing operations to the rights and freedom of individuals and their impact on the protection of personal data (1) before processing such data. Collection of data from low- and middle-income countries could amount to “data colonialism”, in which data are used for commercial or non-commercial purposes without due respect for consent, privacy or autonomy (1).

The assessments could extend beyond risks to privacy and include the quality of data, such as whether they are unbiased and accurate. AI researchers who examine or audit such datasets are reluctant to invest their resources, noting that, while creation of a data set for AI is simple, auditing is difficult, time-consuming and costly. As one researcher noted: “Doing the dirty work is just a lot harder” (124).

Developers can take other measures to improve data quality and adhere to data protection laws. Irrespective of the model size, developers should, in contrast to how early LMMs were developed, train LMMs on data collected according to best-practice data protection rules. Developers should thus avoid using data from third-party sources such as data brokers, as their data may be old, biased, combined incorrectly or have other flaws that may not have been corrected (125). Careful collection of data could also ensure that an LMM does not violate copyright or data protection laws, for which there may be legal repercussions, such that certain LMMs could be labelled unlawful (126).

If third-party data providers are used, they could, for example, be certified, in order to build trust and to ensure their expertise and legitimacy (127). All data used to train LMMs by developers, whether collected directly or from third parties, must be kept up to date. As noted above, some of the leading AI models were not trained with up-to-date data (38), which can jeopardize the performance of the model in health care and medicine, in which new evidence and information meaningfully affect decisions. Data sets should be updated and accurate so that LMMs are appropriate and relevant for the contexts in which they are used.

It may be difficult to ensure that data are adequately transparent. Companies that launch new LMMs have become more and more opaque about the data used to train its models. One leading AI company that released a new LMM stated that: “Given both the competitive landscape and the safety implications of large-scale models like GPT-4, this report contains

no further details about the architecture (including model size), hardware, training compute, dataset construction, training method, or similar” (128).

Unwillingness to be transparent about data is, however, inconsistent with the WHO ethical principle of ensuring transparency, “explainability” and intelligibility. Developers should be transparent about the data they used to train a model so that downstream users, including those who fine-tune the LMM, use the LMM to develop a health-care application and those who use the LMM directly, are aware of any insufficiency or incompleteness of the training data set.

When developers improve data quality by using data workers in low- and middle-income countries to screen the content for abusive, violent or offensive material and to annotate data, those workers should be paid a living wage and be provided with mental health services and other forms of counselling; and developers should introduce safeguards to protect workers from any distress. Governments should update their labour standards to extend benefits to all data workers, to promote a “level playing field” among companies and to ensure that labour standards are maintained and improved over time.

Ethical design and design for values: One approach to integrating ethics and human rights standards into the development of AI technologies is “design for values”, a paradigm for basing design on the values of human dignity, freedom, equality and solidarity and for considering them non-functional requirements (1). Several recommendations in the original WHO expert guidance for the design of AI technologies, including “design for values”, bear repeating here.

The guidance recommended that the design and development of AI technologies not be done solely by scientists and engineers and that “potential end-users and all direct and indirect stakeholders should be engaged from the early stages of AI development in structured inclusive, transparent design and given opportunities to raise ethical issues, voice concerns and provide input for the AI application under consideration” (1). Thus, in development of foundation models, people who might either use or benefit from the models could be involved in the initial development. One proposal is to introduce so-called “human oversight colleges”, which would facilitate inclusion of patient representatives in the development of an LMM that is intended to benefit a patient or caregiver, either directly or indirectly through a medical provider.⁶ Medical and health-care professionals, research scientists, patients, laypeople and vulnerable populations could also be included in the design of LMMs, labelling of data and testing. Inclusiveness in the design of the LMM could, for example, protect human autonomy, as the participation of medical providers could forestall or reduce automation bias by providers. Inclusive design could promote the WHO guiding principle of ensuring inclusiveness and equity, especially if design teams include diverse viewpoints by age, ability, race, ethnicity, sex or gender identity.

6 Communication from David Gruson, WHO expert on ethics and governance of AI for health.

The original WHO guidance also recommended that “designers and other stakeholders should ensure that AI systems are designed to perform well-defined tasks with the accuracy and reliability necessary to improve the capacity of health systems and advance patient interests. Designers and other stakeholders should also be able to predict and understand potential secondary outcomes” (1). Even before initiating development of an LMM, the developer could conduct a so-called “pre-mortem” (33) to consider “hypothetical failures”, so that the development team can reverse-engineer such unanticipated failures. This allows developers to identify known and unknown risks and to formulate alternatives (33). A second suggestion, by several developers of general-purpose foundation models, is “red teaming” (129), an evaluation of a model or system that identified vulnerability in real-world simulations that might result in undesirable behaviour, such as an LMM providing a biased opinion, so that the developer can correct the model or system to ensure its reliability and safety. One company announced that it would submit its latest LMM models to the DEFCON conference, a hacker convention, in August 2023, so that “experts can further analyse and stress test their capabilities” (130).

The original WHO guidance also recommended that “the procedures that designers use to ‘design for values’ should be informed and updated by the consensus principles, best practices (e.g. privacy preserving technologies and techniques), standards of ethics by design and evolving professional norms” (1). Appropriate design could limit unauthorized disclosure of data entered into an LMM or address environmental (carbon and water) concerns associated with the training and use of LMMs (see below). It could also ensure that users know that the content produced by an LMM is generated by an AI system and not a human, in order to avoid displacing humans from the centre of epistemic authority. Such notification can remind users, communities and societies that, while an LMM can produce useful information, it cannot be a substitute for knowledge production by humans.

Design to respect environmental considerations: As discussed above, a major concern with respect to LMMs is their carbon and water footprints. Developers should take all possible steps to reduce energy consumption, such as by improving a model’s energy efficiency, and several large technology companies are experimenting with such approaches. For example, one company developed an LMM that is combined with an external database that operates more efficiently than an LMM, which is trained with more variables and outperforms less energy-efficient LMMs (112). Another company is experimenting with an LMM that is based not on one neural network but distributes its variables among 64 smaller neural networks. It is trained to use only two neural networks to complete each task, thereby using only a small percentage of its variables for making each inference (112).

Another means for improving energy efficiency is to develop smaller LMMs that are trained on smaller data sets and which therefore do not require as much energy to train or to operate. Smaller LMMs may not only reduce energy consumption but also open opportunities for smaller companies or entities to develop LMMs and improve the accuracy of outputs (59). Smaller LMMs might be particularly useful for developing “specialized LMMs”, such as those intended specifically for use in health care, scientific research and medicine. Several such LMMs have been introduced, including some developed by large technology companies (59).

4.3 Government laws, policies and public sector investments

Several existing or potential laws or policies could be enforced or written to reduce or avoid risks during the design and development of general-purpose foundation models. Furthermore, governments could make public-sector investments to promote or support the ethical design and development of general-purpose foundation models.

Laws and policies governing use of data: WHO supports application and enforcement of standards, including data protection rules, that govern how data have and will be used to train LMMs. Data protection laws are usually based on rights-based approaches and include standards for regulation of data-processing that both protect the rights of individuals and establish obligations for public and private data controllers and processors and include sanctions and remedies in case of actions that violate statutory rights. As data protection laws have been adopted in over 150 countries, they provide a solid foundation for the development of all AI technologies, including LMMs (1). A limitation of data protection laws is that most were enacted before the emergence of generative and other types and uses of AI, and data protection authorities may be unwilling to apply them too aggressively, as the original laws may not have had the same intent (120).

One requirement of data protection that should be enforced, especially for health data used to train LMMs, is that the data are obtained and processed lawfully. This will often require provision of meaningful informed consent by a data subject for use of their data for the stated purpose. Any further processing should have its own legal basis, as further processing cannot be assumed to be compatible with the original purpose. Companies and other entities that have developed and released LMMs are already under scrutiny for potential use of data obtained without informed consent. The pursuit of ever-larger LMMs, requiring increasingly larger data sets, may lead developers to ignore legal requirements (83). This would also violate the WHO guiding principle of protecting human autonomy. Thus, the WHO expert group recommended that governments “should have clear data protection laws and regulations for the use of health data and protecting individual rights, including the right to meaningful informed consent”.

Other government measures to oversee and regulate the collection and use of data for training LMMs include regulations for generative AI enacted by the Government of China that came into effect in August 2023. The Cyberspace Administration of China imposes several obligations, including that: (i) providers shall employ effective measures to avoid discrimination and bias in selecting training data, (ii) providers use clear labelling and assess the quality of data labelling; and (iii) developers take “effective measures” to meet the goals of authenticity, accuracy, objectivity and diversity of data (131). The requirements are not expected to be applied strictly to companies, which will be required only to take effective measures to ensure appropriate data quality. The measures will apply only to those companies that offer services to the Chinese public (132).

Legislative provisions related to data could include requirements requirements to describe the data sources used to train a foundation model and to use data that are subject to data governance, including for suitability, bias and appropriate mitigation (133).

Other measures that governments could take during design and development are described below.

- *Target product profiles:* Governments and international agencies could issue target product profiles to state the preferences and characteristics of LMMs intended for use in health care and medicine, especially if governments anticipate purchasing such technologies for use in government-run health systems.
- *Design and development standards and requirements:* Governments could require developers to ensure that the design and development of a general-purpose foundation model achieve certain outcomes throughout its life cycle. They could include requirements for the predictability of the model and its interpretability, corrigibility, safety and cybersecurity (134).
- *Pre-certification programmes:* Regulatory agencies could introduce legal obligations and establish incentives to both require and encourage developers to identify and avoid ethical risks, such as bias or undermining autonomy, through measures including pre-certification programmes (1). The previous WHO AI ethics guidance recommends that “government regulators should provide incentives to developers to identify, monitor and address relevant safety- and human rights-related concerns during product design and development and should integrate relevant guidelines into precertification programs” (1).
- *Audits:* Governments could introduce audits of the initial stages of development of foundation models. One proposal is for three types of audit: a “governance audit” of an LMM provider, an audit of the LMMs and an “application audit” of downstream products and services built on LMMs, which would not apply during development of an LMM (121). Audits could be integrated into requirements for approval of LMMs intended for use in health care or medicine (see below). In order for audits to be effective, their quality should be assessed to ensure that they fulfil their intended purpose.
- *Environmental footprint:* Governments could require developers of general-purpose foundation models to address concern about their carbon and water footprints. For example, governments could require developers to measure their energy consumption, to reduce energy use during training (133) and to meet as-yet undefined environmental standards (134).
- *Notification that content from an LMM is “machine-generated”:* Governments could require developers to ensure that any deployment of a general-purpose foundation model includes notification and reminders to end-users that the content was generated by a machine and not a human being (133).

- Governments might also consider requiring or creating incentives for developers to register early-stage AI algorithms or systems to be used in health care and medicine. Early registration could encourage publication of negative results, prevent publication bias or over-optimistic interpretation of results and facilitate integration of knowledge that benefits patients.

Public infrastructure to develop LMMs in the public interest: As uses of LMMs for health proliferate, development of LMMs that adhere to ethical principles could be encouraged by the provision of not-for-profit or public infrastructure, including computing power and public datasets. Such infrastructure, which could be accessible to developers in the public, private and not-for-profit sectors, could require users to adhere to ethical principles and values in exchange for access. It could also help to avoid exclusive control of an LMM by a developer and "level the playing field" between the largest companies and developers that do not have access to such infrastructure and resources.

Governments, subject to independent oversight, could construct infrastructure that is then used by developers to construct LMMs for health care and medicine. For example, an international team of 1000 academic volunteers, the company Hugging Face, and others, with funding from the French Government, trained an LMM called BLOOM with 175 billion parameters, which also required US\$ 7 million of computing time (112).

Efforts to level the playing field are also applicable to academia and its resource disadvantage. The Canadian Government's national Advanced Research Computing Platform serves the country's academic sector, the Chinese Government has approved a national computing power network system to enable academics and others to access data and computing power, and, in the USA, the National AI Research Resource task force has "proposed the creation of a public research cloud and public datasets" (101). There have also been calls by civil society in Europe for governments to play a more assertive role in build so-called "European Large Generative Models", for which AI-specific computing, data infrastructure, science and research support would be provided by governments (135).

4.4 Open-source LMMs

The role of open-source LMMs in integrating ethical principles and addressing known risks is uncertain. Generally, transparency and participation can be increased by using open-source software for the design of an AI technology or by making the source code of software publicly available (1). Open-source software is open to both contributions and feedback, which allows users to understand how the system works, to identify potential issues and to extend and adapt the software (1). Open-source LMMs may present an opportunity to address some concerns about use of LMMs in health care. As open-source models are neither proprietary nor closed, they allow smaller firms and entities, such as not-for-profit institutions, to design LMMs at lower cost (136). LMMs built on open-source models can be scrutinized, as the code and data are available for review. Engagement and policing by a community of users helps to ensure the robustness of open-source models in the long term (136).

Open-source LMMs may not endure, however, if large technology companies that previously made their models available choose not to continue to do so (10). Development of most open-source LMMs was based on an LMM released on a limited basis by Meta (formerly Facebook) (10). Since the LMM and its weights were leaked (137), the company has stated its commitment to open-source approaches, noting that openness “leads to better products, faster innovation, and a flourishing market, which benefits [Meta] as it does many others.... ultimately, openness is the best antidote to the fears surrounding AI” (130). Independent observers have noted, however, that, while Meta has made its LMM available on a non-commercial basis, its terms of use includes restrictions, and it is thus not providing its LMM in a manner consistent with open-source principles (138,139).

Additional requirements for use of open-source models in order to monitor their performance and outcomes will be difficult for developers to address; however, the benefit of such models cannot supplant the necessity for regulation and avoidance of harm, such as security concerns associated with use of open-source models (140). Open-source models are vulnerable to misuse (141) and can be attacked to exploit such vulnerability (142). A group of researchers recently found that methods tested on open-source AI systems circumvented AI safety measures and safeguards and could also bypass the safeguards of so-called closed systems (143). Ultimately, open-source models are based on the same black-box technologies used in other LMMs.

One way of encouraging open-source LMMs would be for governments to require that foundation models built with government funding or intellectual property be widely accessible, in the same way that governments have required open access to government-funded research. Governments could also encourage open-source research and development in public facilities, including next-generation models, under controlled conditions with public oversight. Public oversight and participation might be better than the new reality in which Meta’s leaked model allows anyone to “download it and run it on a MacBook M2” (144).

Recommendations:

- Developers that design an LMM that shall or could be used in health care, scientific research or medicine shall consider ethics certification or training for programmers. This would bring AI developers in line with requirements in the medical profession and increase trust in their products and services.
- Regardless of the size of the dataset, developers should undertake “data protection impact assessments” before processing such data, which would require developers to assess the risk that data processing operations would go against the rights and freedom of individuals and its impact on the protection of personal data.
- Developers should train LMMs on data collected according to best-practice data protection rules.

- All data sets used to train LMMs, whether collected directly or via third parties by developers, should be kept up to date and appropriate for the contexts in which the system may be used.
- Developers should be transparent about the data used to train a model, so that users, including those who fine-tune the LMM, use the LMM to develop a health-care application or use the LMM directly, are aware of any insufficiency or incompleteness of the training data set.
- Developers should pay data workers a living wage and provide them with mental health services and other forms of counselling. Developers should also introduce safeguards to protect workers from any distress. Governments should update labour standards to extend such benefits to all data workers, to promote a “level playing field” among companies and to ensure that such labour standards are maintained and improved over time.
- Developers should ensure that LMMs are designed not only by scientists and engineers. Potential users and all direct and indirect stakeholders, including medical providers, scientific researchers, health-care professionals and patients, should be engaged from the early stages of AI development in structured, inclusive, transparent design and given opportunities to raise ethical issues, voice concerns and provide input for the AI application under consideration. Such input could be provided through “human oversight colleges”.
- Developers should ensure that LMMs are designed to perform well-defined tasks with the necessary accuracy and reliability to improve the capacity of health systems and advance patient interests. Developers should also be able to predict and understand potential secondary outcomes. Techniques to meet such requirements include “pre-mortems” and “red teaming”.
- Procedures used by developers to “design for values” should be informed and updated by consensus, best practices (e.g. technologies and techniques to preserve privacy), standards of ethics by design and evolving professional norms, including disclosure that content produced by an LMM is generated by an AI system.
- Developers should take all possible steps to reduce energy consumption (such as by improving the energy efficiency of a model).
- Governments should have strong, enforced data protection laws and regulations for the use of health data that apply to the development of LMMs. The laws must effectively protect people’s rights and give people the tools they need to protect their rights, including the right to meaningful informed consent. Additional tools are likely to be needed for data that are collected and processed for use of LMMs in health care.

- Governments and international agencies such as WHO should issue “target product profiles” to delineate preferences and characteristics of LMMs intended for use in health care and medicine, especially if governments anticipate eventual purchase of such tools for use in government-run health systems.
- Governments should require developers to ensure that the design and development of a general-purpose foundation model will achieve certain outcomes during the product’s life cycle. These could include requirements for the predictability of the model and its interpretability, corrigibility, safety and cybersecurity.
- Regulatory agencies should introduce legal obligations and establish incentives, such as pre-certification programmes, to require and encourage developers to identify and avoid ethical risks, including bias or undermining autonomy.
- Governments should introduce audits of the initial stages of development of foundation models.
- Governments should require developers of general-purpose foundation models to address concerns about the carbon and water footprints of general-purpose foundation models.
- Governments should require developers to ensure that, in any use of a general-purpose foundation model, users are notified and reminded that the content has been generated by a machine and not a human being.
- Governments should consider requiring or creating incentives for developers to register early-stage AI algorithms or systems that are to be used in health care and medicine. Early registration could encourage publication of negative results, prevent publication bias or over-optimistic interpretation of results and could facilitate inclusion of knowledge that benefits patients.
- Governments should invest in or provide not-for-profit or public infrastructure, including computing power and public data sets, accessible to developers in the public, private and not-for-profit sectors, that requires users to adhere to ethical principles and values in exchange for access.
- Governments should encourage the development of open-source LMMs by requiring that foundation models built with government funding or intellectual property are widely accessible, in the same way that governments have required open access to government-funded research. Governments should support open-source research and development in public facilities, including next-generation models, under controlled conditions, with public oversight.

5 Provision with general-purpose foundation models (LMMs)

The uses of general-purpose foundation models depend on whether a user prompts an LMM to generate health-care related outputs or a provider is permitted by the developer to integrate the LMM into a health-care related application, product or service. Either case introduces novel risks that must be addressed by developers, providers or both. Governments are responsible for assessing and regulating uses of such technologies before their deployment.

5.1 Risks to be addressed when providing a health-care service or application with a general-purpose foundation model (LMM)

There is likely to be disagreement about whether both general-purpose foundation models and applications should be assessed and approved when they are used in a product for health-care purposes or used directly by a user. Several of the largest technology companies have been quietly lobbying government officials (for example in the European Union) to abandon an evaluation framework for LMMs and to focus oversight instead on applications that might be used in ways that a government might consider “risky” (145). This would concern both providers that generate and market health-care applications that include the foundation model, and users, such as a provider or patient, who choose to use the LMM directly or indirectly via an AI system. The companies argue that oversight of the general-purpose foundation model would “completely shift the burden” to developers and that others in the value chain should also assume responsibility (145).

While it may not be appropriate to hold the developer of a general-purpose foundation model responsible for all uses of the LMM, it would also be inappropriate to place the burden solely on providers, deployers or users, as they will not have been involved in development of the model and may not understand the associated limitations and risks. This would allow developers of general-purpose foundation models, despite their significant power, resources, oversight and understanding of LMMs, to escape responsibility and would open a “massive hole” in attempts to govern AI technologies for health (145).

A developer may seek to avoid use of an LMM for a health-care purpose (or another use). If a developer does not wish an LMM to be used for health or medical purposes (especially in clinical medicine), it could discourage such use either by preventing entities that develop applications for health or medicine to use (license) the LMM on an application programming interface, or, if the LMM is used directly by a user (provider or patient) for a health-care

purpose, by blocking queries or attaching a clear warning to any responses that include health or medical information and to direct users to information or services that can provide appropriate assistance.

If such measures are not taken or if the developer intends that users apply its LMM for health care, either directly or indirectly through a provider, the developer will have specific responsibilities that only it can satisfy. Furthermore, both developers and providers have further obligations to address the risks associated with use of LMMs in health care.

The responsibilities, detailed below, are defined in government-mandated laws, policies and regulations, as it is governments that must ultimately determine whether an AI-based system should be permitted for use in health care. Developers and providers must also fulfil mutual responsibilities if an LMM is to be used in health care. Such responsibilities could be defined by governments or negotiated between the two parties through a contract if laws have not been written or updated to account for them.

Major risks that must be addressed before deployment include system-wide bias, false information or hallucinations for health-care uses, privacy of data entered into an LMM, manipulation and automation bias.

5.2 Measures that governments can introduce to address such risks and ethical principles that should be upheld

The speed of development of LMMs and of applications that include an LMM requires that governments rapidly develop regulations and specific criteria for using these AI algorithms in health-care systems and for other scientific and medical purposes. The approach should consist of assessment and approval of AI technologies intended for use in health care or medicine by a regulatory agency, such as a medical device or pharmaceutical agency, although governments could establish a new agency for this purpose. One challenge for low- and middle-income countries is that their regulatory agencies are already under-resourced and overwhelmed by pharmaceutical regulation.

The government of at least one high-income country has agreed with the largest technology companies that its foundation models will be assessed in a voluntary public evaluation, with disclosure of outcomes to provide information to the public and researchers about the models and to encourage the companies to correct any errors (146); however, a voluntary approach is likely to be neither sufficient nor sustainable.

Evaluation of LMMs and applications should not address only the AI systems or algorithms used in the health-care system, as there are also significant risks associated with the use of LMMs and applications in a grey zone between clinical and “wellness” applications. Given the rapid proliferation of such technologies, governments should, at least initially, identify such

applications, set common standards and regulations and prohibit applications that do not meet the standards and regulations from being deployed to the public.

Developers and providers should bear the burden of proof, when required, to demonstrate that an AI technology intended for use in health care meets the minimum requirements set out in a law or policy. It should not be assumed that, given the known risks and challenges associated with LMMs, AI algorithms and applications with an LMM are safe and effective or that they are superior to AI or non-AI based approaches that are already in wide use.

Several laws, policies and cross-cutting requirements that could apply to use of LMMs in health care and medicine are described below.

Disclosure (transparency) requirements: Appropriate regulation requires not only that governments have the capacity and discretion to decide what they can assess and approve for use but also adequate information for conducting such an evaluation. Disclosure is necessary to both adequately regulate an AI technology and to ensure that other actors in the AI value chain can use the technology safely. For example, unless a developer discloses the performance of a general-purpose foundation model (such as its propensity to hallucinate), a provider may not have the necessary information to fine-tune the model or to avoid marketing the technology. Such forms of disclosure by a provider or developer can also assist users, such as medical providers, in deciding not to use an LMM that could provide incorrect information or to scrutinize outputs more carefully.

Disclosure and transparency are WHO guiding principles, as well as measures to improve the “explainability” and intelligibility of an AI-based system, and should be required in the assessment of a general-purpose foundation model or application. The WHO guidance on the ethics and governance of AI for health (1) recommended that “government regulators should require the transparency of certain aspects of an AI technology, while accounting for proprietary rights, to improve oversight and assessment of safety and efficacy. This may include an AI technology’s source code, data inputs and analytical approach”. New forms of disclosure that are relevant for LMMs may include their performance in internal testing and their carbon and water footprints. Standards may also be required for “open weights”, which allow regulators, other developers, civil society and providers to understand the outputs of training an algorithm, or the knowledge that an LMM has obtained during its training (147,148).

Several forms of disclosure could assist a provider, user or regulator, including describing the capability and limitations of a foundation model, evaluation of the model according to public or industry standard benchmarks and reporting the results of internal and external testing of the model and its optimization (134). Disclosures, particularly of the risks that may be associated with an LMM or application, could be advertised clearly, and has been compared by one researcher to a “nutrition label” (129).

Data protection laws: The development of LMMs and how developers manage the data required to train an LMM may violate data protection laws. A separate problem is that data entered into an LMM or application to produce a specific output that may include sensitive

personal information may be disclosed either accidentally or through prompts. Potential disclosures are the reason that many large companies, including technology companies that are developing and commercializing LMMs, prohibit their own employees from using such algorithms (149).

Disclosures of data violate the developer's responsibility to protect autonomy. Developers may also violate data protection laws if sensitive data are held longer than permitted according to data minimization requirements (85). One developer allows users to opt out of any content they supply for refining its chatbot's performance (150). Governments that permit use of LMMs should ensure that they set, extend and enforce data protection rules to cover the data entered into an LMM. The Chinese Government's regulation for LMMs includes such a requirement, although the protection applies only to users in China (151).

Assessment of general-purpose foundation models and/or applications used in health care: human rights law versus risk-based frameworks: Several legislative frameworks are being developed to evaluate and regulate AI technologies. One question with respect to such frameworks is whether AI technologies must satisfy human rights obligations ("fundamental rights" according to the European Union) or whether a different approach should be used, to assess AI technologies in a risk-based framework. The European Union, under the AI Act, has adopted a risk-based framework (152). A risk-based framework, it is argued, could help to identify the requirements or burden of proof that must be provided for a technology, the burden of proof increasing with the level of risk of the technology.

All AI systems or tools used in health care and medicine should have to respect ethical obligations and human rights standards that affect, for example, a person's dignity, autonomy or privacy. These include general-purpose foundation models. Human rights and ethical principles are non-negotiable and must be upheld, irrespective of the risk associated with an AI technology or the benefit it may confer (153). The fact that an AI algorithm is considered "low risk" does not exempt it from scrutiny, and a developer or provider should ensure the algorithm respects human rights and ethical obligations. A human rights impact assessment can be conducted to determine whether an LMM or an application adheres to such commitments and can therefore be used safely.

The WHO guidance on the ethics and governance of AI for health (1) recommended that "governments should enact laws and policies that require government agencies and companies to conduct impact assessments of AI technologies, which should address ethics, human rights, safety and data protection, throughout the life cycle of an AI system". The guidance also noted that "impact assessments should be audited by an independent third party before and after introduction of an AI technology and published" (1). The results of impact assessments should be disclosed publicly, while accounting for proprietary or sensitive information, and should be available to the public and groups that may be affected. In the same way as for audits (see above), impact assessments might have to be examined closely, especially if they are conducted by third parties that offer tools or services, to ensure that they are of adequate quality and rigour.

Impact assessments can reveal, for example, whether an AI technology might introduce system-wide bias, risk the privacy of users who share personal data or lead to manipulation of a user. Risks to privacy should be addressed by collaboration between providers and developers in developing LMMs that preserve individual privacy. Work on such an LMM is under way by one company and a hospital system in the USA, although the project is considered unlikely to succeed because the data cannot be fully de-identified (154). Impact assessments can also ensure that use of a general-purpose foundation model or application maintains humans in the loop to avoid automated decision-making in which a user receives false information or misinformation or a health-care provider or patient relies uncritically on the output of an LMM, which would be a form of automation bias.

Governments may instead choose to use a risk-based framework for LMMs to be used in health care and medicine. For those functions that are considered to be higher risk, such as providing a prescription or mental health advice to a person with severe depression or use of an AI technology by a vulnerable or marginalized population, the burden of proof will be higher. There is concern that, if a government selects a risk-based approach, it will be considered sufficient or be used as a substitute for a human rights-based approach (153). A risk-based framework might exclude certain LMMs or applications from evaluation, which may appear to be low risk but could ultimately lead to harm.

Additional questions include whether an AI regulatory assessment should apply to foundation models regardless of their eventual use, whether the assessment should apply only to the largest, most widely used foundation models (“systemic foundation models”⁷) and when such assessments should apply to providers.

This guidance does not include a recommendation on whether all foundation models, regardless of how they are to be used, should be subjected to a risk-based and/or rights-based assessment. The guidance also does not recommend whether an AI regulatory assessment of general-purpose foundation models should apply only to the largest (systemic) LMMs. The expert group did note one concern about assessments designed to apply to all LMMs, irrespective of size, which is that it could “freeze in place” the dominance of the largest companies, as the standards may be such that only those companies can feasibly comply with them or that the standards best suit their business model and aims (155). This concern has gained the attention of competition authorities, which have found that first-movers may use “unfair methods of competition to entrench their current power or use that power to gain control over a new generative AI market” (141). Competition authorities are expected to exercise greater scrutiny of the use of LMMs, although they focus on the practices used by companies that develop LMMs (156).

Providers should also be subject to AI regulatory assessments, as their use of an LMM may change its purpose and function from those determined by the developer but is controlled by the provider. Thus, if a general-purpose foundation model is adapted for use in health care or medicine by a provider, to which the developer agrees, both the developer and the provider

7 Presentation by Kai Zenner, European Parliament, Head of Office for Axel Voss, Conference on AI for Good, 5 June 2023.

should comply with requirements for use of LMMs in health care and medicine. The regulatory burden on providers should be greater if their use of a product or application diverges substantially from or changes the foundation model in ways that are beyond the control of the developer.

Medical device regulation: A government may determine that a general-purpose foundation model or application qualifies as a medical device. While there is little guidance on which LMMs qualify as medical devices, one regulator said that “LMMs only directed toward general purposes and whose developers make no claim that their software can be used for a medical purpose are unlikely to qualify as medical devices” (157). The regulator also noted, however, that: “LMMs that are developed for, or adapted, modified or directed toward specifically medical purposes are likely to qualify as medical devices. Additionally, where a developer makes claims that their LMM can be used for a medical purpose, this again is likely to mean the product qualifies as a medical device” (157).

Chatbots based on LMMs that provide medical advice are likely to be characterized as medical devices under current European Union and US regulatory standards (158). The WHO guidance on the ethics and governance of AI for health (1) recommended that: “government regulators should require that an AI system’s performance be tested, and sound evidence obtained from prospective testing in randomized trials and not merely from comparison of the system to existing datasets in a laboratory”.

If an LMM or application is to be regulated as a medical device, the developer and/or provider should bear the burden of proof by providing evidence that the device performs as marketed and that it meets the requirements of current or amended national laws. This may include various requirements, such as adhering to ethical obligations related to bias and privacy. Newly proposed regulations on AI technologies for medical devices in the European Union and the USA will probably integrate ethical principles related to the use of AI in health, including “explainability”, control of bias and transparency. It is unlikely that current chatbots that include LMMs could meet such standards (158).

LMMs for clinical decision support are already being used experimentally. Although these LMMs include disclaimers, they do not obviate application of medical device laws, “which dictate that such experiments should take place only in an authorized clinical trial setting under appropriate controls to protect patients and to produce clinically relevant outcomes” (158). Governments could examine controlled experimental uses of such LMMs in regulatory “sandboxes”, which would allow testing in a live environment in actual clinical settings with safeguards and oversight to protect health systems from risks or unintended consequences. Such use may, however, be appropriate only in countries in which new health-care products and services and their specifications are subject to formal regulation and data protection regulations (1).

Consumer protection law: Governments should develop and use consumer protection laws to ensure that any negative consequences of LMMs and applications do not reach users and patients. Consumer protection laws could be applied, for example, to prevent practices that

would be tantamount to manipulation (159). In the USA, several government departments and agencies are applying consumer protection laws and other regulations to prevent discrimination and bias in automated systems (159). Such laws can enable governments to require entities that seek to commercialize such technologies to address the causes of any negative consequences and to protect patients and their families from any current or future harm (93). Consumer protection laws, or other regulations, could be used to require that LMMs and applications are restricted in use of language that could misdirect or mislead an end-user into ascribing human-like qualities to an LMM. Such laws could therefore restrict use of or prevent LMMs or applications from using words such as “I think”, “I suppose” or “I suggest”.

Recommendations:

- Governments should, as resources permit, assign an existing or new regulatory agency to assess and approve LMMs and applications intended for use in health care or medicine.
- Certain aspects of an LMM and its applications should be transparent to allow oversight and assessment of its safety and efficacy by regulators. This may include the source code, data inputs, model weights and analytical approach. Additional forms of disclosure to be considered by a government are the performance of an LMM or application in internal testing and its carbon and water footprints.
- Governments should ensure that data protection rules apply to data entered into an LMM or application by a user.
- Government laws, policies and regulations should ensure that LMMs and applications used in health care and medicine, irrespective of the risk or benefit associated with the AI technology, meet ethical obligations and human rights standards that affect, for example, a person’s dignity, autonomy or privacy.
- Governments should enact laws and policies that require providers and developers to conduct impact assessments of LMMs and applications, which should address ethics, human rights, safety and data protection, throughout the life cycle of an AI system. The impact assessments should be audited by an independent third party before and after introduction of an AI technology and should be in the public domain.
- The regulatory burden on a provider should increase if the product or application substantially diverges from or changes the foundation model in ways that are out of the control of the developer of the model.
- Governments should ensure that, for an LMM or application that is regulated as a medical device, the developer and/or provider is responsible for the burden of proof that the device performs as marketed and that it meets the requirements of the country’s laws or amended laws.

- Governments should ensure that LMMs or applications for supporting clinical decisions that are not yet approved for use not be used on an experimental basis outside an authorized clinical trial setting. Governments may facilitate controlled experimental uses of LMMs through regulatory “sandboxes”, which allow testing in a live environment in actual clinical settings, with safeguards and oversight to protect the health system from risks or unintended consequences.
- Governments should use consumer protection laws to ensure that any negative consequences of use of LMMs and applications do not affect users, including patients. Consumer protection laws could be applied, for example, to prevent practices that would be tantamount to manipulation or to address the causes of other negative consequences of LMMs or applications in order to protect patients and their families from any current or future harm.

6 Deployment with general-purpose foundation models (LMMs)

Even when an LMM or an application with an LMM has been ethically designed and undergone appropriate regulatory scrutiny, it may still carry risks when commercialized. The deployer of an AI health-care application or tool could be either the developer or the provider of an LMM or application or, for example, a ministry of health, a hospital, a health-care company or a pharmaceutical company.

6.1 Risks to be addressed when deploying a health-care service or application with a general-purpose foundation model (LMM)

Risks during deployment may be due to the unpredictability of LMMs and the responses they provide, the possibility of use of a general-purpose foundation model in a manner that was not anticipated by either the developer or the provider, and because responses generated by an LMM may change over time.

Major risks that must be addressed when deploying an LMM are:

- inaccurate or false responses,
- bias,
- privacy of data entered into and put out by an LMM,
- accessibility and affordability of an LMM,
- impacts on labour and employment,
- automation bias and skills degradation, and
- the quality of interactions between health-care providers and patients.

This section describes how the actors in the AI value chain, including users, can mitigate or prevent risks and the role of governments in regulating the use of AI tools once an LMM has been deployed, while equipping and training health-care workers and other actors in health systems to maximize appropriate use of an LMM.

6.2 On-going responsibilities of developers and providers during deployment

Developers and providers have responsibilities and obligations even after an LMM or application is approved for use, either because the developer or provider deploys the LMM or because certain risks can be addressed after deployment only by a developer or provider. Such obligations might have to be required by regulations or laws to ensure that developers and providers allocate adequate resources and attention.

First, governments should introduce mandatory post-release auditing and impact assessments, including for data protection and human rights, by independent third parties when an LMM is deployed on a large scale (155,160). Post-release auditing and impact assessments should be published and should address outcomes and impacts disaggregated by type of user, for example by age, race or disability.

Secondly, governments could hold providers or developers responsible for inaccurate, false or toxic content from an LMM after its release that neither the provider nor the developer took steps to correct or avoid. The Chinese Government's regulation on generative AI, for example, stipulates that it must not produce information that is "false and harmful" (151), with possible enforcement by the Government. In the European Union, addition of an LMM to a product or service could create additional responsibilities for the developer and the provider of the LMM. For example, if an LMM is integrated into a service that is within the scope of regulation of digital services, such as the European Union Digital Services Act, the LMM would be indirectly subject to regulatory scrutiny, which could require regulatory oversight because of the tendency of LMMs to hallucinate (120).

Thirdly, a developer and provider might be required to provide ongoing operational disclosures in order for governments and users to use an LMM safely. These could include sufficient technical documentation (133,134).

6.3 Responsibilities of deployers

Deployers are also responsible for avoiding or mitigating risks associated with use of an LMM or application.

First, a deployer should use information from developers or providers to decide not to use an LMM or application in an inappropriate setting, because of biases in the training data, contextual bias that renders the LMM inappropriate for the setting or other avoidable errors or potential risks known to the deployer. If a deployer receives clear, adequate warning of such risks and still offers the LMM for use in inappropriate settings, the deployer should be held accountable for any resulting harm.

Secondly, deployers should communicate any risks that they should reasonably know could result from use of an LMM and any errors or mistakes that have harmed users. Such warnings should not be in fine print or easy to miss. In some circumstances, a deployer may be responsible, even if not required by a law or regulation, for suspending use or removing an LMM or application from the market to avoid harm.

Thirdly, deployers can take steps to improve the affordability and accessibility of an LMM. A deployer can ensure that pricing or subscription fees for use of an LMM correspond to the capacity of a government or other user to pay and should ensure that appropriate LMMs are trained and provided in languages and scripts that can be used by people who are otherwise ignored or excluded from the benefits of technology. Deployers should also request providers and developers to ensure that current and future LMMs are available in several languages.

6.4 Government programmes and practices

Introduction of LMMs into health-care systems and for other health-care-associated uses will require significant adjustment by health-care professionals. Neither developers nor providers have the interest, resources or expertise to ensure appropriate use of an LMM by health-care professionals or for other uses that involve individuals with specialized training and/or expertise.

As in the design of a general-purpose foundation model (see above), governments could enlist both health-care professionals and patients in “human oversight colleges” to ensure that new LMMs and applications used in clinical decision-making are used appropriately and do not undermine the rights of patients.⁸

Governments, universities (health science faculties) or health-care providers such as in hospitals can also ensure that health workers use an LMM to deliver clinical care effectively and are appropriately trained in other uses. Health-care professionals and clinicians should be trained: (i) to understand how LMMs make decisions and the limits of understanding how such decisions are made, (ii) to identify concern about appropriate use, (iii) in methods for avoiding automation bias, (iv) engaging with and educating patients who may be or are considering use of LMMs, and (v) cybersecurity risks associated with the use of LMMs (161).

Training and continuing education of health workers is of particular importance for informing patients, laypeople and other third parties when advice is generated by an LMM or that information provided by an LMM has been used by the provider in making a medical decision or for another medical function. In such notifications, the patient or layperson should be fully apprised of the risks associated with use of LMMs to preserve his or her right to informed consent.

⁸ Communication from David Gruson, WHO expert on ethics and governance of AI for health.

Health-worker training is also critical to ensure that, when they use an LMM professionally, their duties do not unknowingly violate laws, especially those related to the protection of health data and information. For example, medical providers who introduce “protected health information” into an LMM chatbot may be violating laws such as the Health Insurance Portability and Accountability Act in the USA (150). As popular LMMs become “trusted” by health-care workers, for example, they may disclose more patient data than they realize (154).

Other stakeholders in a health-care system should be educated on the benefits, risks, uses and challenges of LMMs in health care and how LMMs differ from other technologies for generating information or advice and how they been used for other purposes in health care. Broader public awareness of the use of LMMs in health care and other domains should be improved. WHO guidance on the ethics and governance of AI for health (1) recommends that: “The public should be engaged in the development of AI for health to understand forms of data sharing and use, to comment on the forms of AI that are socially and culturally acceptable and to fully express their concerns and expectations. Further, the general public’s literacy in AI technology should be improved to enable them to determine which AI technologies are acceptable”.

Governments that supply an LMM or application to a health system could use their procurement authority to foster certain practices among developers, providers and deployers. Procurement of a critical LMM or application for use in a health-care system can eliminate barriers to access and affordability if the AI technology does not displace other health-care investments that may be more effective, equitable and affordable. Public procurement can establish requirements for transparency with respect to data training, quality assurance, risk assessment, mitigation and external audits. Such requirements may be critical if a country has neither relevant legislation nor a regulatory agency with the resources to regulate LMMs effectively.

Recommendations:

- Governments should introduce mandatory post-release auditing and impact assessments, including for data protection and human rights, by independent third parties when an LMM is deployed on a large scale. The auditing and impact assessments should be published and should include outcomes and impacts disaggregated by the type of user, including for example by age, race or disability.
- Governments could hold providers or developers responsible for inaccurate, false or toxic content issued by an LMM after its release, which has not been corrected or avoided by either the provider or the developer.
- Governments should require ongoing operational disclosures by both developers and providers to ensure that LMMs and applications can be used safely. These could include sufficient technical documentation.

- In accordance with information obtained from either developers or providers, deployers should not use an LMM or application in a setting that is inappropriate because of biases in the training data, contextual bias that renders the LMM inappropriate for a particular setting, or other potential errors or risks, such as inaccurate, false or toxic content published by an LMM, that are known to the deployer and can be avoided.
- Deployers should communicate any risks that they should reasonably know could result from use of an LMM, as well as errors that have caused harm to users; such warnings should not be in fine print (or easy to miss). In some circumstances, a deployer may be responsible, even if not required by a law or regulation, to suspend use of or to remove the LMM or application from the market to avoid future harm.
- Deployers should improve the affordability and accessibility of an LMM, by ensuring that pricing or subscription fees for use of an LMM are in line with the capability of a government or other user to pay, and should ensure that appropriate LMMs are trained and offered in languages and scripts that reach people who are otherwise ignored or excluded from the benefits of technology. Deployers should request providers and developers to ensure that current and future LMMs are developed with various languages.
- Governments should facilitate the participation of health-care professionals and patients, in “human oversight colleges” to ensure that new LMMs and applications used in making clinical decision are used appropriately and do not undermine the rights of patients.
- Ministries of health and universities (health science faculties) should train health-care professionals and clinicians: (i) to understand how LMMs make decisions (and the limits of understanding how such decisions are made), (ii) to identify and understand concerns about appropriate use, (iii) in methods to avoid automation bias, (iv) to engage with and educate patients who may be or are considering use of LMMs, and (v) on the cybersecurity risks associated with use of LMMs.
- Governments, health service providers, health researchers and funders should engage the public so that they understand different forms of data-sharing and use, can comment on whether and how LMMs are socially and culturally acceptable and can fully express their concerns and expectations. Further, the literacy of the public in AI technology should be improved to enable them to identify acceptable uses and types of LMMs.
- Governments that supply an LMM or application through a health system should ensure that their procurement authority fosters certain practices by developers, providers and deployers, including transparency.

7 Liability for LMMs

As LMMs gain broader use in health care and medicine, errors, misuse and ultimately harm to individuals are inevitable. Liability rules will have to be used to compensate individuals for such harm, with establishment of new forms of redress when current approaches are insufficient or out of date.

The design, development, quality assurance and deployment of AI technologies involve various entities, each of which plays a distinct role. This can complicate assignment of liability. Developers may demand that downstream entities, such as providers and deployers, be liable for any harm resulting from use of an LMM, while downstream entities may claim that previous actions, such as choice of the data used to train an algorithm, are the cause. Developers and providers may also claim that, once a medical AI technology has been approved for use by a regulator, they should no longer be held liable for harm (regulatory pre-emption) (1). Establishment of liability along the value chain is a challenge for lawmakers and policy-makers.

A critical function of rules for civil liability is to ensure that a victim of damage can claim compensation and redress, no matter how difficult it may be to assign blame and responsibility among the entities involved in development and deployment of an AI technology. If victims find it too difficult to obtain compensation, there can be no justice and no incentive for parties in the AI value chain to avoid such harm in the future. The rules should also ensure that the compensation is adequate for the harm suffered.

The European Union in its proposed AI Liability Directive simplifies the burden of proof by a victim by introducing a “presumption of causality” (162). Thus, if a victim can demonstrate that one or more entities did not comply with an obligation relevant to the harm and that a causal link with AI performance is likely, the court can presume that non-compliance was the cause of the damage (162). The onus is thus placed on the liable party to rebut the presumption, for example by indicating another party as the cause of the damage. The scope of the legislation is not limited to the original maker of an AI system but includes any participant in the AI value chain (162). When all the actors in the AI value chain are held jointly liable, they can demonstrate their effectiveness in assessing and mitigating risks in order to reduce their liability.

It is still possible, however, that a liability regime does not provide full clarity and redress for injuries caused by AI-driven products and services, especially if an individual does not know that a LMM was used in making a medical decision. New rules may leave gaps in liability for injuries caused by AI-driven medical technologies (163). As LMMs are highly speculative, poorly understood and being rushed to market, governments may wish to consider LMMs used in health care as products for which developers, providers and deployers will be held to a strict liability standard. Holding these actors accountable for any error might ensure that a

patient will be compensated if the error affects them (1), although this depends on whether the patient knew that an LMM was used. While such continuing liability might discourage the use of increasingly sophisticated LMMs, it might also temper a willingness to take unnecessary risks and to deploy new LMMs into health care or public health settings before their many risks and potential harms have been fully identified and addressed (1).

A liability regime for AI might not, however, be adequate to assign fault, as algorithms are evolving in ways that neither developers, providers nor deployers can fully control. Furthermore, there may be situations or jurisdictions in which a person who is harmed is unable to recover damages. For example, in the USA, a patient who is injured when using an LMM directly to seek advice may not be able to recover damages because AI systems themselves are not included in professional liability rules, and exceptions or limitations to product or consumer liability laws may preclude recovery (163). In other areas of health care, compensation is occasionally provided without assignment of fault or liability, such as for medical injuries resulting from adverse effects of vaccines. The original WHO guidance recommended determination of “whether no-fault, no-liability compensation funds are an appropriate mechanism for providing payments to individuals who suffer medical injuries due to the use of AI technologies, including how to mobilize resources to pay any claims” (1). That recommendation is also valid today and could be a means for determining compensation for injuries caused by LMMs or applications with LMMs.

Recommendation:

- Governments should establish liability along the value chain of the development, provision and deployment of LMMs and applications to ensure that a victim of damage can claim compensation, irrespective of the difficulty of assigning blame and of the responsibilities of the different entities involved in the development and deployment of the technology.

8 International governance of LMMs

Governments should support collective development of international rules for the governance of LMMs and other forms of AI used in health care, as such uses are proliferating globally. One example is the WHO global strategy on digital health 2020-2025. The process should include greater cooperation and collaboration within the United Nations system to respond to the opportunities and challenges of deploying AI in health care and of its wider application in society and the economy. Unless governments work together to set appropriate, enforceable standards, the number of LMMs and other forms of AI that do not meet appropriate legal, ethical and safety standards will increase, potentially causing harm if regulations and other types of protection are not introduced or are not properly enforced, whether willingly or because there are inadequate resources. WHO recently issued a new publication, in consultation with regulatory agencies worldwide, that outlines key principles that governments and regulatory authorities can follow to develop new guidance or adapt existing guidance on AI (164).

International governance can avoid a “race to the bottom” among companies seeking a first-mover advantage in which standards of safety and efficacy are ignored, and among governments seeking advantage in the geopolitical race for technological supremacy. Thus, international governance can ensure that all companies meet minimum standards of safety and efficacy and also avoid introduction of regulations that provide a competitive advantage or disadvantage for either companies or governments. International governance can make governments accountable for their investments and participation in the development and deployment of AI-based systems and ensure that they introduce appropriate regulations that respect ethical principles, human rights and international law. The absence of globally enforceable standards may also have a negative impact on product adoption.

International governance could take several forms. One suggestion is to establish a public research agency, funded by several governments, such as the European Organization for Nuclear Research, an international collaboration, with funding and human resources to pursue large, transformative projects the results of which are shared openly (165,166). In a separate proposal, it was suggested that such an entity could be charged with developing the most advanced, and most risky, forms of AI in a highly secure facility, making other attempts to build such forms of AI illegal (167). At present, such large-scale projects are not in the domain of publicly funded projects to generate public goods but are the purview of large technology companies in commercial competition with one another. Other leaders, including world leaders and technology executives, have called for AI to be treated similarly to nuclear weapons, with a global regulatory framework similar to treaties for the use of nuclear arms (109).

Whatever the form of international governance that is taken forward, it is imperative that it not be shaped solely by high-income countries or by high-income countries that work mostly or solely with the world's largest technology companies (168). Standards developed by and for high-income countries and technology firms, whether for all applications of AI or for specific use of LMMs in health care and medicine, will leave most of humanity, in low- and middle-income countries, with no role or voice in shaping the standards. This would render future AI technologies potentially dangerous or ineffective in the very countries that might ultimately benefit the most.

International governance of AI could require that all stakeholders cooperate through networked multilateralism, as proposed by the United Nations Secretary-General in 2019 (169), which would bring together the United Nations family, international financial institutions, regional organizations, trading blocs and others, including civil society, cities, businesses, local authorities and young people, to work more closely, effectively and inclusively. Placing ethics and human rights at the centre of the development and deployment of LMMs could make a substantial contribution towards achievement of universal health coverage.

Recommendation:

- Governments should support collective development of international rules for the governance of AI. Whatever the form of governance, it must not be shaped solely by high-income countries or by high-income countries working mostly or solely with the world's largest technology companies, as that approach would leave most of humanity, in low- and middle-income countries, without a role or voice in shaping international governance of AI.

References

1. Ethics and governance of artificial intelligence for health. Geneva: World Health Organization; 2021 (<https://www.who.int/publications/i/item/9789240029200>, accessed 26 May 2023).
2. Khullar D. Can A.I. treat mental illness? The New Yorker, 27 February 2023 (<https://www.newyorker.com/magazine/2023/03/06/can-ai-treat-mental-illness>, accessed 29 May 2023).
3. Acosta JN, Falcone GJ, Rajpurkar P, Topol EJ. Multimodal biomedical AI. Nat Med. 2022;28(9):1773–84. doi:10.1038/s41591-022-01981-2.
4. Hariri Y, Harris T, Raskin A. You can have the blue pill or the red pill, and we're out of blue pills. The New York Times, 24 March 2023 (<https://www.nytimes.com/2023/03/24/opinion/yuval-harari-ai-chatgpt.html>, accessed 26 May 2023).
5. Moor M, Banerjee O, Abad ZSH, Krumholz HM, Leskovec J, Topol EJ et al. Foundation models for generalist medical artificial intelligence, Nature. 2023;616(7956):259–65. doi:10.1038/s41586-023-05881-4.
6. Hu K. Chat GPT sets record for fastest growing user-base – analyst note. Reuters, 2 February 2023 (<https://www.reuters.com/technology/chatgpt-sets-record-fastest-growing-user-base-analyst-note-2023-02-01/>, accessed 26 May 2023).
7. Weise K, Grant N. Microsoft and Google unveil A.I. tools for businesses. The New York Times, 16 March 2023 (<https://www.nytimes.com/2023/03/16/technology/microsoft-google-ai-tools-businesses.html>, accessed 26 May 2023).
8. Yang Z. Chinese tech giant Baidu just released its answer to ChatGPT. MIT Technology Review, 16 March 2023 (<https://www.technologyreview.com/2023/03/16/1069919/baidu-ernie-bot-chatgpt-launch/>, accessed 26 May 2023).
9. Murgia M, Bradshaw T. Musk to launch AI start-up to rival ChatGPT. Financial Times, 15 April 2023 (<https://www.ft.com/content/2a96995b-c799-4281-8b60-b235e84aefe4>, accessed 26 May 2023).
10. Heaven WD. The open-source AI boom is built on Big Tech's handouts. How long will it last? MIT Technology Review, 12 May 2023 (<https://www.technologyreview.com/2023/05/12/1072950/open-source-ai-google-openai-eleuther-meta/>, accessed 26 May 2023).

11. Martin A. Google CEO Sunder Pichai admits people don't fully understand how chatbot AI works, Evening Standard, 17 April 2023 (<https://www.standard.co.uk/tech/google-ceo-sundar-pichai-understand-ai-chatbot-bard-b1074589.html>, accessed 26 May 2023).
12. Roose K. A conversation with Bing's chatbot left me deeply unsettled. The New York Times, 16 February 2023 (<https://www.nytimes.com/2023/02/16/technology/bing-chatbot-microsoft-chatgpt.html>, accessed 26 May 2023).
13. Marcus G. AI platforms like ChatGPT are easy to use but potentially dangerous, Scientific American, 19 December 2022 (<https://www.scientificamerican.com/article/ai-platforms-like-chatgpt-are-easy-to-use-but-also-potentially-dangerous/>, accessed 26 May 2023).
14. Bubeck S, Chandrasekaran V, Eldan R, Gehrke J, Horvitz E, Kamar E et al. Sparks of artificial general intelligence: early experiments with GPT-4. ArXiv:2302.12712.
15. McGowran L. OpenAI criticised for lack of transparency around ChatGPT-4. Silicon Republic, 16 March 2023 (<https://www.siliconrepublic.com/machines/openai-gpt4-transparency-ai-concerns-stripe-chatgpt>, accessed 26 May 2023).
16. Spitale G, Biller-Andorno N, Germani F. AI model GPT-3 (dis)informs us better than humans. Sci Adv. 2023;9(26):eadh1850. doi:10.1126/sciadv.adh1850.
17. Volpicelli G. ChatGPT broke the EU plan to regulate AI Politico, 3 March 2023 (<https://www.politico.eu/article/eu-plan-regulate-chatgpt-openai-artificial-intelligence-act/>, accessed 26 May 2023).
18. Arcesati R, Chang W. China is blazing a trail in regulating Generative AI – on the CCP's terms. The Diplomat, 28 April 2023 (<https://thediplomat.com/2023/04/china-is-blazing-a-trail-in-regulating-generative-ai-on-the-ccps-terms/>, accessed 26 May 2023).
19. Martindale J. These are the countries where ChatGPT is currently banned. Digital Trends, 12 April 2023 (<https://www.digitaltrends.com/computing/these-countries-chatgpt-banned/>, accessed 26 May 2023).
20. Johnson K. ChatGPT can help doctors – and hurt patients. Wired, 24 April 2023 (<https://www.wired.com/story/chatgpt-can-help-doctors-and-hurt-patients/>, accessed 28 May 2023).
21. Topol E. Multimodal AI for medicine, simplified. Ground Truths, 14 March 2023 (<https://erictopol.substack.com/p/multimodal-ai-for-medicine-simplified>, accessed 28 May 2023).
22. Heaven WD. AI hype is built on high test scores. Those tests are flawed. MIT Technology Review, 30 August 2023 (<https://www.technologyreview.com/2023/08/30/1078670/large-language-models-arent-people-lets-stop-testing-them-like-they-were/>, accessed 1 October 2023).

23. Singhal K, Azizi S, Tu T, Mahdavi SS, Wei J, Chung HW et al. Large language models encode clinical knowledge. *Nature*. 2023;620:172–80. doi:10.1038/s41586-023-06291-2.
24. Kulkarni PA, Singh H. Artificial intelligence in clinical diagnosis: opportunities, challenges, and hype. *JAMA*. 2023;330(4):317–8. doi:10.1001/jama.2023.11440.
25. Subbamaran N. ChatGPT will see you now: Doctors using AI to answer patient questions. *Wall Street Journal*, 28 April 2023 (<https://www.wsj.com/articles/dr-chatgpt-physicians-are-sending-patients-advice-using-ai-945cf60b>, accessed 28 May 2023).
26. Ayers JW, Poliak A, Dredze M, Leas EC, Zhu Z, Kelley JB et al. Comparing physician and artificial intelligence chatbot responses to patient questions posted to a public social media forum. *JAMA Intern Med*. 2023; 183(6):589–96. doi: 10.1001/jamainternmed.2023.1838.
27. Lee P, Bubeck S, Petro J. Benefits, limits, and risks of GPT-4 as an AI chatbot for medicine. *New Engl J Med*. 2023;388(13):1233–9. doi: 10.1056/NEJMSr2214184.
28. The potential of large language models in healthcare: Improving quality of care and patient outcomes, Medium, 7 December 2022. (<https://medium.com/@BuildGP/the-potential-of-large-language-models-in-healthcare-improving-quality-of-care-and-patient-6e8b6262d5ca>, accessed 28 May 2023).
29. Rajpurkar P, Irvin J, Zhu K, Yang B, Mehta H, Duan T et al. CheXNet: radiologist-level pneumonia detection on chest X-rays with deep learning. *arXiv*. 2017;1711.05225v3. doi:10.48550/arXiv.1711.05225.
30. Yang X, Chen A, PourNejatian N, Shin HC, Smith KE, Parisien C et al. A large language model for electronic health records. *npj Digit Med*. 2022;5:194. doi:10.1038/s41746-022-00742-2.
31. Ghahramani Z. Introducing PaLM 2. The Keyword, 10 May 2023 (<https://blog.google/technology/ai/google-palm-2-ai-large-language-model/>, accessed 28 May 2023).
32. Weise K, Metz C. When A.I. chatbots hallucinate. *The New York Times*, 9 May 2023 (<https://www.nytimes.com/2023/05/01/business/ai-chatbots-hallucination.html>, accessed 1 June 2023).
33. Bender EM, Gebru T, McMillan-Major A, Mitchell M. On the dangers of stochastic parrots: Can language models be too big? In: *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, March 2021, pp. 610–23. doi: 10.1145/3442188.3445922.
34. Li H, Moon JT, Purkayastha S, Celi LA, Trivedi H, Gichoya JW. Ethics of large language models in medicine and medical research. *Lancet Digit Health*. 2023;5(6):e333–5. doi:10.1016/s2589-7500(23)00083-3.

35. Metz, Cade. Chatbots may 'hallucinate' more often than many realize, *The New York Times*, 6 November 2023 (<https://www.nytimes.com/2023/11/06/technology/chatbots-hallucination-rates.html>, accessed 7 November 2023).
36. Acar OA. AI prompt engineering isn't the future, *Harvard Business Review*, 6 June 2023 (<https://hbr.org/2023/06/ai-prompt-engineering-isnt-the-future?registration=success>, accessed 26 June 2023).
37. GPT-4 system card. Open AI, 23 March 2023 (<https://cdn.openai.com/papers/gpt-4-system-card.pdf>, accessed 28 May 2023).
38. GPT-4. OpenAI, 14 March 2023 (<https://openai.com/research/gpt-4>, accessed 28 May 2023).
39. Radford A, Kleinman Z. ChatGPT can now access up-to-date information. *BBC News*, 27 September 2022 (<https://www.bbc.com/news/technology-66940771>, accessed 1 October 2023).
40. Krug S, Ostermaier A, Uhl M. The moral authority of ChatGPT. *ArXiv:23101.07098*. doi:10.48550/arXiv.23101.07098.
41. Mickle T, Metz C, Grant N. The chatbots are here, and the internet industry is in a tizzy, *The New York Times*, 8 March 2023 (<https://www.nytimes.com/2023/03/08/technology/chatbots-disrupt-internet-industry.html>, accessed 29 May 2023).
42. Woo M. Trial by artificial intelligence. *Nature*. 2019;573:S100–2 (<https://media.nature.com/original/magazine-assets/d41586-019-02871-3/d41586-019-02871-3.pdf>, accessed 29 May 2023).
43. Muralidharan V, Burgart A, Daneshjou D, Rose S. Recommendations for the use of pediatric data in artificial intelligence and machine learning ACCEPT-AI. *npj Dig Med*. 2023;6:166. doi:10.1038/s41746-023-00898-5.
44. Kasneci E, Sessler K, Küchemann S, Bannert M, Dementieva D, Fischer F et al. ChatGPT for good? On opportunities and challenges of large language models for education. *Learning Individual Differences*. 2023;103:102274. doi:10.1016/j.lindif.2023.102274.
45. Reddy CD, Lopez L, Ouyang D, Zou JY, He B. Video-based deep learning for automated assessment of left ventricular ejection fraction in pediatric patients. *J Am Soc Echocardiogr*. 2023;36(5):482–9. doi:10.1016/j.echo.2023.01.015.
46. Knight W. These ChatGPT rivals are designed to play with your emotions. *Wired*, 4 May 2023 (<https://www.wired.com/story/fast-forward-chatgpt-rivals-emotions/#:~:text=12%3A00%20PM-,These%20ChatGPT%20Rivals%20Are%20Designed%20to%20Play%20With%20Your%20Emotions,%2C%20-companionship%E2%80%94and%20even%20romance.>, accessed 29 May 2023).

47. Smuha NA, De Ketalaere M, Coeckelbergh M, Dewitte P, Pouillet Y. Open letter: We are not ready for manipulative AI – urgent need for action. KU Leuven, 31 March 2023 (<https://www.law.kuleuven.be/ai-summer-school/open-brief/open-letter-manipulative-ai>, accessed 29 May 2023).
48. Cuthbertson A. “No, I’m not a robot”: ChatGPT successor tricks worker into thinking it is human. Independent, 15 March 2023 (<https://www.independent.co.uk/tech/chatgpt-gpt4-ai-openai-b2301523.html>, accessed 26 June 2023).
49. Walker L. Belgian man dies by suicide following exchanges with chatbot, The Brussels Times, 28 March 2023 (<https://www.brusselstimes.com/430098/belgian-man-commits-suicide-following-exchanges-with-chatgpt>, accessed 29 May 2023).
50. DeGuerin M. Oops: Samsung employees leaked confidential data to ChatGPT. Gizmodo, 6 April 2023 (<https://gizmodo.com/chatgpt-ai-samsung-employees-leak-data-1850307376>, accessed 29 May 2023).
51. Privacy policy. OpenAI, 27 April 2023 (<https://openai.com/policies/privacy-policy>, accessed 29 March 2023).
52. Coles C. 11% of data employees paste into ChatGPT is confidential. Cyberhaven, 19 April 2023 (<https://www.cyberhaven.com/blog/4-2-of-workers-have-pasted-company-data-into-chatgpt/>, accessed 29 May 2023).
53. Mihalcik C. ChatGPT bug exposed some subscribers’ payment info. CNET, 24 March 2023. (<https://www.cnet.com/tech/services-and-software/chatgpt-bug-exposed-some-subscribers-payment-info/>, accessed 29 May 2023).
54. Moodley K, Rennie S. ChatGPT has many uses. Experts explore what this means for healthcare and medical research. The Conversation, 22 February 2023 (<https://theconversation.com/chatgpt-has-many-uses-experts-explore-what-this-means-for-healthcare-and-medical-research-200283>, accessed 2 June 2023).
55. De Proost M, Pozzi G. Conversational artificial intelligence and the potential for epistemic injustice. *Am J Bioethics*. 2023;23(5):51–3. doi:10.1080/15265161.2023.2191020.
56. Disability and employment. New York: United Nations, Department of Economic and Social Affairs (Disability); undates (<https://www.un.org/development/desa/disabilities/resources/factsheet-on-persons-with-disabilities/disability-and-employment.html>, accessed 11 September 2023).
57. Whittaker M, Alper M, Bennett CL, Hendren S, Kaziunas L, Mills Met al. Disability, bias and AI. New York: AI Now Institute; 2019 (<https://ainowinstitute.org/wp-content/uploads/2023/04/disabilitybiasai-2019.pdf>, accessed 11 September 2023).

58. Hallman J. AI language models show bias against people with disabilities, study finds. University Park (PA): Penn State University; 2022 (<https://www.psu.edu/news/information-sciences-and-technology/story/ai-language-models-show-bias-against-people-disabilities/>, accessed 11 September 2023).
59. Harrer S. Attention is not all you need: the complicated case of ethically using large language models in healthcare and medicine. *eBioMedicine*. 2023;90:194512. doi:10.1016/j.ebiom.2023.104512.
60. Lohr S. AI may someday work medical miracles. For now, it helps do paperwork. *The New York Times*, 26 June 2023 (<https://www.nytimes.com/2023/06/26/technology/ai-health-care-documentation.html>, accessed 10 July 2023).
61. Eddy N. Epic, Microsoft partner to use generative AI for better EHRs. *Healthcare IT News*, 18 April 2023 (<https://www.healthcareitnews.com/news/epic-microsoft-partner-use-generative-ai-better-ehrs>, accessed 31 May 2023).
62. Nuance and Microsoft announce the first fully AI-automated clinical documentation application for healthcare. Burlington (MA):Nuance; 2023 (<https://news.nuance.com/2023-03-20-Nuance-and-Microsoft-Announce-the-First-Fully-AI-Automated-Clinical-Documentation-Application-for-Healthcare>, accessed 31 May 2023).
63. Ahn S. The impending impacts of large language models on medical education. *Korean J Med Educ*. 2023;35(1):103–7. doi:10.3946/kjme.2023.253.
64. Luo R, Sun L, Xia Y, Qin T, Zhang S, Poon H et al. BioGPT: generative pre-trained transformer for biomedical text generation and mining. *Briefi Bioinformatics*. 2022; 23(6):bbac409. doi:10.1093/bib/bbac409.
65. Paul D, Sanap G, Shenoy S, Kalyane D, Kalia K, Tekade RK. Artificial intelligence in drug discovery and development. *Drug Discov Today*. 2021;26(1):80–93. doi:10.1016/j.drudis.2020.10.010.
66. Tools such as ChatGPT threaten transparent science; here are our ground rules for their use. *Nature*. 2023;613:612. doi:10.1038/d41586-023-00191-1.
67. Zielinski C, Winker MA, Aggarwal R, Ferris LA, Heinemann M, Lapeña JF Jr et al. Chatbots, generative AI, and scholarly manuscripts. Overijssel: World Association of Medical Editors; 2023. (<https://wame.org/page3.php?id=106>, accessed 26 June 2023).
68. Gibbs W. Lost science in the Third World. *Sci Am*. 1995;273(2):92–9 doi:10.1038/scientificamerican0895-92.
69. Birhane A, Kasirzadeh A, Leslie D, Wachter S. Science in the age of large language models. *Nat Rev Phys*. 2023;5:277–80. doi:10.1038/s42254-023-00581-4.

70. Monitoring the building blocks of health systems: a handbook of indicators and their measurement strategies. Geneva: World Health Organization; 2010 (<https://apps.who.int/iris/bitstream/handle/10665/258734/9789241564052-eng.pdf>, accessed 26 June 2023).
71. Morozov E. The true threat of artificial intelligence. The New York Times, 30 June 2023 (<https://www.nytimes.com/2023/06/30/opinion/artificial-intelligence-danger.html>, accessed 2 July 2023).
72. Introducing ChatGPTPlus. San Francisco (CA): Open AI; 2023 (<https://openai.com/blog/chatgpt-plus>, accessed 1 June 2023).
73. The hidden workforce that helped filter violence and abuse out of ChatGPT. Wall Street Journal, 11 July 2023 (<https://www.wsj.com/podcasts/the-journal/the-hidden-workforce-that-helped-filter-violence-and-abuse-out-of-chatgpt/ffc2427f-bdd8-47b7-9a4b-27e7267cf413>, accessed 13 July 2023).
74. Firth N. Language models may be able to self-correct biases – if you ask them. MIT Technology Review, 20 March 2023 (<https://www.technologyreview.com/2023/03/20/1070067/language-models-may-be-able-to-self-correct-biases-if-you-ask-them-to/>, accessed 1 June 2023).
75. Khan L. We must regulate A.I. Here's how. The New York Times, 3 May 2023 (<https://www.nytimes.com/2023/05/03/opinion/ai-lina-khan-ftc-technology.html>, accessed 2 June 2023).
76. Hatzius J, Briggs J, Kodnani D, Pierdomenico G. The potentially large effects of artificial intelligence on economic growth (Briggs/Kodnani). Goldman Sachs Economics Research, 26 May 2023 (https://www.key4biz.it/wp-content/uploads/2023/03/Global-Economics-Analyst_-The-Potentially-Large-Effects-of-Artificial-Intelligence-on-Economic-Growth-Briggs_Kodnani.pdf, accessed 1 June 2023).
77. Milmo D. AI revolution puts skilled jobs at highest risk, says OECD. The Guardian, 11 July 2023 (<https://www.theguardian.com/technology/2023/jul/11/ai-revolution-puts-skilled-jobs-at-highest-risk-oecd-says>, accessed 12 July 2023).
78. Health and care workforce in Europe: time to act. Geneva: World Health Organization; 2022 (<https://iris.who.int/handle/10665/362379>, accessed 1 June 2023).
79. Health workforce. Geneva: World Health Organization; 2023 (https://www.who.int/health-topics/health-workforce#tab=tab_1, accessed 1 June 2023).
80. Hurst L. OpenAI says 80% of workers could see their jobs impacted by AI. These are the jobs most impacted, Euronews.next, 30 March 2023. (<https://www.euronews.com/next/2023/03/23/openai-says-80-of-workers-could-see-their-jobs-impacted-by-ai-these-are-the-jobs-most-affe>, accessed 1 June 2023).

81. A new era of generative AI for everyone. Dublin: Accenture; 2023 (<https://www.accenture.com/content/dam/accenture/final/accenture-com/document/Accenture-A-New-Era-of-Generative-AI-for-Everyone.pdf>, accessed 1 June 2023).
82. Burgess M. The security hole at the heart of ChatGPT and Bing. Wired, 25 May 2023 (<https://www.wired.co.uk/article/chatgpt-prompt-injection-attack-security>, accessed 1 June 2023).
83. Heikkila M. Open AI's hunger for data is coming back to bite it. MIT Technology Review, 19 April 2023 (<https://www.technologyreview.com/2023/04/19/1071789/openais-hunger-for-data-is-coming-back-to-bite-it/>, accessed 1 June 2023).
84. General Data Protection Regulation, Regulation 2016//679 of the European Parliament and of the Council, 27 April 2016. Strasbourg: European Parliament; 2016 (<https://eur-lex.europa.eu/eli/reg/2016/679/oj>, accessed 27 September 2023).
85. The impact of the General Data Protection Regulation on artificial intelligence (STOA Options Brief). Strasbourg: European Parliament; 2020 ([https://www.europarl.europa.eu/RegData/etudes/STUD/2020/641530/EPRS_STU\(2020\)641530\(ANN1\)_EN.pdf](https://www.europarl.europa.eu/RegData/etudes/STUD/2020/641530/EPRS_STU(2020)641530(ANN1)_EN.pdf), accessed 26 June 2023).
86. OPC launches investigation into ChatGPT. Ottawa: Office of the Privacy Commissioner of Canada; 4 April 2023 (https://www.priv.gc.ca/en/opc-news/news-and-announcements/2023/an_230404/, accessed 1 June 2023).
87. Lomas N. Italy orders Chat GPT blocked citing data protection concerns. Tech Crunch, 31 March 2023 (<https://techcrunch.com/2023/03/31/chatgpt-blocked-italy/>, accessed 1 June 2023).
88. ChatGPT: Italian SA to lift temporary limitation if OpenAI implements measures. Rome: Italian Data Protection Authority; 2023 (<https://www.garanteprivacy.it/home/docweb/-/docweb-display/docweb/9874751#english>, accessed 1 June 2023).
89. Weatherbed J. OpenAI's regulatory troubles are only just beginning. The Verge, 5 May 2023 (<https://www.theverge.com/2023/5/5/23709833/openai-chatgpt-gdpr-ai-regulation-europe-eu-italy>, accessed 1 June 2023).
90. Wiggers K. Open AI's new tool attempts to explain language models' behaviours. Tech Crunch, 9 May 2023 (<https://techcrunch.com/2023/05/09/openais-new-tool-attempts-to-explain-language-models-behaviors/>, accessed 1 June 2023).
91. Libeau D. ChatGPT will probably never comply with GDPR. 10 April 2023. (<https://blog.davidlibeau.fr/chatgpt-will-probably-never-comply-with-gdpr/>, accessed 1 June 2023).

92. Lomas N. ChatGPT maker OpenAI accused of string of data protection breaches in GDPR complaint filed by privacy researcher. TechCrunch, 30 August 2023. (https://consent.yahoo.com/v2/collectConsent?sessionId=3_cc-session_6bdecae4-d7b6-448f-8e26-e7805c03b964, accessed 11 September 2023).
93. Fung B. The FTC should investigate Open AI and block GPT over “deceptive” behaviour, AI policy group claims. CNN, 30 March 2023. (<https://edition.cnn.com/2023/03/30/tech/ftc-openai-gpt-ai-think-tank/index.html>, accessed 2 June 2023).
94. Waters R, Murgia M, Espinoza J. Open AI warns over split with Europe as AI regulation advances. Financial Times, 25 May 2023 (<https://www.ft.com/content/5814b408-8111-49a9-8885-8a8434022352>, accessed 1 June 2023).
95. Technology-facilitated gender-based violence: Making all spaces safe. New York: United Nations Population Fund; 2021 (<https://www.unfpa.org/publications/technology-facilitated-gender-based-violence-making-all-spaces-safe>, accessed 1 October 2023).
96. Murgia M. DeepMind reinvents itself for AI counterattack. Financial Times, 2 May 2023 (<https://ft.pressreader.com/v99c/20230502/281724093873699>, accessed 2 June 2023).
97. Schaaake M. Regulating AI will put companies and governments at loggerheads, Financial Times, 2 May 2023 (<https://www.ft.com/content/7ef4811d-79bb-4b4f-b28f-b46430f0c9ff>, accessed 2 June 2023).
98. Metz, Cade. Tech giants are paying huge salaries for scarce A.I. talent, The New York Times, 22 October 2017 (<https://www.nytimes.com/2017/10/22/technology/artificial-intelligence-experts-salaries.html>).
99. Leswing K. Google reveals its newest AI supercomputer, says it beats Nvidia. CNBC, 5 April 2023. (<https://www.cnbc.com/2023/04/05/google-reveals-its-newest-ai-supercomputer-claims-it-beats-nvidia-.html>, accessed 2 June 2023).
100. Ahuja K. Antitrust has role in policing AI landscape. Financial Times, 10 April 2023 (<https://www.ft.com/content/953817f5-5bc4-49e1-b583-977cc4780eca>, accessed 2 June 2023).
101. Ahmed N, Wahed M, Thompson NC. The growing influence of industry in AI research. *Science*. 2023;379(6635):884–6. doi:10.1126/science.ade2420.
102. Røttingen JA, Regmi S, Eide M, Young AJ, Viergever RF, Ardal C et al. Mapping of available health research and development data: What’s there, what’s missing, and what role is there for a global observatory? *Lancet*. 2013;382(9900):1286–307. doi:10.1016/S0140-6736(13)61046-6.

103. A new partnership to promote responsible AI. Google Blogs, 26 July 2023 (<https://blog.google/outreach-initiatives/public-policy/google-microsoft-openai-anthropic-frontier-model-forum/#:~:text=Anthropic%2C%20Google%2C%20Microsoft%20and%20OpenAI%20are%20launching%20the%20Frontier%20Model,development%20of%20frontier%20AI%20models>, accessed 29 July 2023).
104. Fact sheet: Biden–Harris Administration secures voluntary commitments from leading artificial intelligence companies to manage the risks posed by AI. Washington DC: The White House, 21 July 2023 (<https://www.whitehouse.gov/briefing-room/statements-releases/2023/07/21/fact-sheet-biden-harris-administration-secures-voluntary-commitments-from-leading-artificial-intelligence-companies-to-manage-the-risks-posed-by-ai/>, accessed 29 July 2023).
105. Volpicelli G. Europe pitches AI pact to curtail the booming tech’s risk. Politico, 26 May 2023 (<https://www.politico.eu/article/big-tech-rumble-europe-global-artificial-intelligence-debate-ai-pact/>, accessed 29 July 2023).
106. Grant N, Weise K. In AI race, Microsoft and Google choose speed over caution. The New York Times, 7 April 2023 (<https://www.nytimes.com/2023/04/07/technology/ai-chatbots-google-microsoft.html>, accessed 2 June 2023).
107. Center for Research on Foundation Models. The Foundation Model Transparency Index, 2023. (<https://crfm.stanford.edu/fmti/>, accessed 21 October 2023).
108. Schiffer Z, Newton C. Microsoft lays off team that taught employees how to make AI tools responsibly. The Verge, 14 March 2023. (<https://www.theverge.com/2023/3/13/23638823/microsoft-ethics-society-team-responsible-ai-layoffs>, accessed 2 June 2023).
109. Milmo D. Google chief warns AI could be harmful if deployed wrongly. The Guardian, 17 April 2023 (<https://www.theguardian.com/technology/2023/apr/17/google-chief-ai-harmful-sundar-pichai>, accessed 2 June 2023).
110. Fiesler C. AI has social consequences, but who pays the price? Tech companies’ problem with ethical debt. The Conversation, 19 April 2023 (<https://theconversation.com/ai-has-social-consequences-but-who-pays-the-price-tech-companies-problem-with-ethical-debt-203375>, accessed 2 June 2023).
111. Criddle, Cristina and Murphy, Hannah, Meta disbands protein-folding team in shift towards commercial AI, Financial Times, 7 August 2023. (https://www.ft.com/content/919c05d2-b894-4812-aa1a-dd2ab6de794a?accessToken=zwAGBZu-oWVwkdORnAXSuJRIEtOqGt0qtt55Sg.MEQCIA1QQ1iG8KPAAnuDAuPvt-Ngds3OzxL1lt-0FnaVbAQFtAiAZvHnmKD_fABj8ZzLTNXRp1v7V38nTcUf_pPxAPdx16A&sharetype=gift&toKen=3ac5a132-e08e-412e-bc3c-08edea8a7417, accessed 18 September 2023).
112. Ananthaswamy A. In AI, is bigger always better? Nature, 8 March 2023 (<https://www.nature.com/articles/d41586-023-00641-w>, accessed 2 June 2023).

113. Li P. Making AI less “thirsty”: Uncovering and addressing the secret water footprint of AI models. *ArXiv*. 2023;2304.03271v. doi:10.48550/arXiv.2304.03271.
114. Syed N. The secret water footprint of AI technology. *The Markup*, 15 April 2023 (<https://themarkup.org/hello-world/2023/04/15/the-secret-water-footprint-of-ai-technology>, accessed 2 June 2023).
115. Livingstone G. It’s pillage: Thirsty Uruguayans blast Google’s plan to exploit water supply. *The Guardian*, 11 July 2023 (<https://www.theguardian.com/world/2023/jul/11/uruguay-drought-water-google-data-center>, accessed 12 July 2023).
116. Thornhill J. The sceptical case on generative AI. *Financial Times*, 17 August 2023. (<https://www.ft.com/content/ed323f48-fe86-4d22-8151-eed15581c337>, accessed 11 September 2023).
117. Marcus G. The imminent enshittification of the Internet. *Substack*, 16 August 2023 (<https://garymarcus.substack.com/p/the-imminent-enshittification-of>, accessed 11 September 2023).
118. Pause giant AI experiments: An open letter. Narberth (PA): Future of Life Institute; 2023 (<https://futureoflife.org/open-letter/pause-giant-ai-experiments/>, accessed 13 June 2023).
119. Perrigo B. DeepMind’s CEO helped take AI mainstream. Now he’s urging caution. *Time*, 12 January 2023 (<https://time.com/6246119/demis-hassabis-deepmind-interview/>, accessed 13 June 2023).
120. Lomas N. Unpacking the rules shaping generative AI. *Tech Crunch*, 13 April 2023 (<https://techcrunch.com/2023/04/13/generative-ai-gdpr-enforcement/>, accessed 13 June 2023).
121. Mökander J, Schuett J, Kirk HR, Floridi L. Auditing large language models: a three-layered approach *Soc Sci Res Netw*. 2023. doi:10.2139/ssrn.4361607.
122. Lomas N. Report details how Big Tech is leaning on EU not to regulate general purpose AIs. *Tech Crunch*, 23 February 2023 (<https://techcrunch.com/2023/02/23/eu-ai-act-lobbying-report/>, accessed 20 June 2023).
123. Sambasivan N, Kapania S, Highfill H, Akrong D, Paritosh P, Aroyo LM. “Everyone wants to do model work, not the data work.”: Data cascades in high-stakes AI. In: *Proceedings of the 2021 CHI Conference on Human Factors in Computing Systems*, May 2021. doi:10.1145/3411764.3445518.
124. Browne G. AI is steeped in Big Tech’s digital colonialism. *Wired*, 25 May 2023 (<https://www.wired.co.uk/article/abeba-birhane-ai-datasets>, accessed 17 June 2023).

125. Baxter K, Schelsinger, N. Managing the risks of generative AI. *Harvard Business Review*, 6 June 2023 (<https://hbr.org/2023/06/managing-the-risks-of-generative-ai>, accessed 17 June 2023).
126. Samuelson P. Generative AI meets copyright. *Science*. 2023;381(6654):158–61. doi:10.1126/science.adi0656.
127. El-Mhamdi E, Farhadkhani S, Guerraoui R, Gupta N, Hoang L, Pinot R et al. On the impossible safety of large AI models. *arXiv*. 2209.15259v2. doi:10.48550/arXiv.2209.15259.
128. Open AI. GPT-4 technical report. *arXiv*:2303.08774v3. doi:10.48550/arXiv.2302.08774.
129. Murgia M. Open AI's red team: experts hired to "break" ChatGPT. *Financial Times*, 14 April 2023 (<https://www.ft.com/content/0876687a-f8b7-4b39-b513-5fee942831e8>, accessed 10 July 2023).
130. Clegg N. Openness on AI is the way forward for tech. *Financial Times*, 11 July 2023 (<https://www.ft.com/content/ac3b585a-ce50-43d1-b71d-14dfe6dce999>, accessed 11 July 2023).
131. Huang S, Toner H, Haluza Z, Creemers R, Webster G. Measures for the management of generative artificial intelligence services (draft for comment) (translation). DigiChina. Palo Alto (CA): Stanford University, Program on Geopolitics; 2023 (<https://digichina.stanford.edu/work/translation-measures-for-the-management-of-generative-artificial-intelligence-services-draft-for-comment-april-2023/>, accessed 17 June 2023).
132. Ye J. China says generative AI rules to apply only to products for the public. *Reuters*, 13 July 2023 (<https://www.reuters.com/technology/china-issues-temporary-rules-generative-ai-services-2023-07-13/>, accessed 13 July 2023).
133. Bommasani R, Klyman K, Zhang D, Liang P. Do foundation model providers comply with the draft EU AI Act? Palo Alto (CA): Stanford University, Human-centered Artificial Intelligence; 2021 (<https://crfm.stanford.edu/2023/06/15/eu-ai-act.html>, accessed 17 June 2023).
134. Amendments adopted by the European Parliament on 14 June 2023 on the proposal for a regulation of the European Parliament and of the Council on laying down harmonised rules on artificial intelligence (Artificial Intelligence Act) and amending certain Union legislative acts. Strasbourg: European Parliament; 2023 (https://www.europarl.europa.eu/doceo/document/TA-9-2023-0236_EN.html, accessed 10 July 2023).
135. Beyond ChatGPT: How can Europe become a leader in generative AI? Kaiserslautern: German Research Centre for Artificial Intelligence; 2023. (<https://www.dfki.de/en/web/news/jenseits-von-chatgpt-wie-kann-europa-bei-der-generativen-ki-eine-fuehrungsposition-uebernehmen>, accessed 17 June 2023).

136. Spirling A. Why open-source generative AI models are an ethical way forward for science. *Nature*. 2023;616(7957):413. doi:10.1038/d41586-023-01295-4.
137. Vincent J. Meta's powerful AI language model has leaked online – What happens now? *The Verge*, 8 March 2023 (<https://www.theverge.com/2023/3/8/23629362/meta-ai-language-model-llama-leak-online-misuse>, accessed 29 July 2023).
138. Maffuli S. Meta's Llama 2 license is not open source. Open Source Initiative, 20 July 2023 (<https://blog.opensource.org/metals-llama-2-license-is-not-open-source/>, accessed 29 July 2023).
139. Marble A. Software licenses masquerading as open source. *marble.onl*, 1 June 2023 (<http://marble.onl/posts/software-licenses-masquerading-as-open-source.html>, accessed 29 July 2023).
140. Keary T. Report finds 82% of open-source software components “inherently risky”. *Venture Beat*, 17 April 2023 (<https://venturebeat.com/security/report-finds-82-of-open-source-software-components-inherently-risky/>, accessed 8 July 2023).
141. Generative AI raises competition concerns. *Technology blog*, 29 June 2023. Washington DC: Federal Trade Commission; 2023 (<https://www.ftc.gov/policy/advocacy-research/tech-at-ftc/2023/06/generative-ai-raises-competition-concerns>, accessed 29 July 2023).
142. Wishart-Smith H. Generative AI: cybersecurity friend and foe, *Forbes*, 6 June 2023 (<https://www.forbes.com/sites/heatherwishartsmith/2023/06/06/generative-ai-cybersecurity-friend-and-foe/?sh=4407e0884bd2>, accessed 29 July 2023).
143. Metz C. Researchers poke holes in safety controls of ChatGPT and other chatbots. *The New York Times*, 27 July 2023 (<https://www.nytimes.com/2023/07/27/business/ai-chatgpt-safety-research.html>, accessed 11 September 2023).
144. Harris T, Freuh S. The complexity of technology's consequences is going up exponentially, but our wisdom and awareness are not. *Issues in Science and Technology*, 16 May 2023 (<https://issues.org/tristan-harris-humane-technology-misinformation-ai-democracy/>, accessed 19 June 2023).
145. Schyns C. The lobbying ghost in the machine: Big Tech's covert defanging of Europe's AI Act. Brussels: Corporate Europe Observatory; 2023 (<https://corporateeurope.org/en/2023/02/lobbying-ghost-machine>, accessed 17 June 2023).
146. Fact sheet: Biden–Harris Administration announces new actions to promote responsible AI innovation that protects Americans' rights and safety. Washington DC: White House, 4 May 2023 (<https://www.whitehouse.gov/briefing-room/statements-releases/2023/05/04/fact-sheet-biden-harris-administration-announces-new-actions-to-promote-responsible-ai-innovation-that-protects-americans-rights-and-safety/>, accessed 19 June 2023).

147. Sijbrandij Sid. AI weights are not “open source”. Open Core Ventures, 27 June 2023. (<https://opencoreventures.com/blog/2023-06-27-ai-weights-are-not-open-source/>, accessed 29 July 2023).
148. Meeker H. Towards an open weights definition. Copyleft Currents, 8 June 2023 (<https://heathermeeker.com/2023/06/08/toward-an-open-weights-definition/>, accessed 29 July 2023).
149. Dastin J, Tong A. Google, one of AI’s biggest backers, warns its own staff about chatbots. Reuters, 15 June 2023 (<https://www.reuters.com/technology/google-one-ais-biggest-backers-warns-own-staff-about-chatbots-2023-06-15/>, accessed 9 July 2023).
150. Kanter GP, Packel EA. Health care privacy risks of AI chatbots. *JAMA*. 2023;330(4):311–2. doi:10.1001/jama.2023.9618.
151. Interim measures for the management of generative artificial intelligence services. Beijing: Cyberspace Administration of China; 13 2023. (http://www.cac.gov.cn/2023-07/13/c_1690898327029107.htm, accessed 29 July 2023).
152. Satariano A. E.U. agrees on landmark artificial intelligence rules. The New York Times, 8 December 2023 (<https://www.nytimes.com/2023/12/08/technology/eu-ai-act-regulation.html>, accessed 15 December 2023).
153. The EU should regulate on the basis of rights, not risks. Access Now, 17 February 2021 (<https://www.accessnow.org/eu-regulation-ai-risk-based-approach/>, accessed 21 June 2023).
154. Marks M, Haupt CE. AI chatbots, health privacy, and challenges to HIPAA compliance *JAMA*. 2023;330(4):309–10. doi: 10.1001/jama.2023.9458.
155. Marcus G. Two models of AI oversight – and how things could go deeply wrong. Substack, 8 June 2023 (<https://garymarcus.substack.com/p/two-models-of-ai-oversight-and-how>, accessed 17 June 2023).
156. Kang C, Metz C. FTC opens investigation into Chat GPT maker over technology’s potential harms. The New York Times, 13 July 2023 (<https://www.nytimes.com/2023/07/13/technology/chatgpt-investigation-ftc-openai.html>, accessed 29 July 2023).
157. Ordish J. Large language models and software as a medical device. MedRegs blogs, 3 March 2023 (<https://medregs.blog.gov.uk/2023/03/03/large-language-models-and-software-as-a-medical-device/>, accessed 19 June 2023).
158. Gilbert S, Harvey H, Melvin T, Vollebregt E, Wicks P. Large language model AI chatbots require approval as medical devices. *Nat Med*. 2023. doi:10.1038/s41591-023-02412-6.

159. Ghost in the machine: Addressing the harm of generative AI. Forbrukerradet. Oslo: Norwegian Consumer Council; 2023 (<https://storage02.forbrukerradet.no/media/2023/06/generative-ai-rapport-2023.pdf>, accessed 9 July 2023).
160. Mökander J, Floridi L. Ethics-based auditing to develop trustworthy AI. *Minds & Machines*, 2021. doi:10.1007/s11023-021-09557-8.
161. Minssen T, Vayena E, Cohen IG. The challenges for regulating medical use of ChatGPT and other large language models. *JAMA*. 2023;330(4):315–6. doi: 10.1001/jama.2023.9651.
162. Questions and Answers: AI Liability Directive. Brussels: European Commission; 2022 (https://ec.europa.eu/commission/presscorner/detail/en/QANDA_22_5793, accessed 20 June 2023).
163. Duffourc MN, Gerke S. The proposed EU directives for AI liability leave worrying gaps likely to impact medical AI. *NPJ Digit Med*. 2023;6(1):77. doi:10.1038/s41746-023-00823-w.
164. Regulatory considerations on artificial intelligence for health. Geneva: World Health Organization; 2023 (<https://iris.who.int/bitstream/handle/10665/373421/9789240078871-eng.pdf?sequence=1&isAllowed=y>, accessed 16 November 2023).
165. Marcus G. Artificial Intelligence is stuck. Here's how to move it forward. *The New York Times*, 29 July 2017 (<https://www.nytimes.com/2017/07/29/opinion/sunday/artificial-intelligence-is-stuck-heres-how-to-move-it-forward.html>, accessed 20 June 2023).
166. Parker G. Rishi Sunak to lobby Joe Biden for UK “leadership” role in AI development. *Financial Times*, 5 June 2023 (<https://www.ft.com/content/7c30ea28-2895-44c2-9a2d-c31ea7fa27e7>, accessed 19 June 2023).
167. Hogarth I. We must slow down the race to god-like AI. *Financial Times*, 13 April 2023 (<https://www.ft.com/content/03895dc4-a3b7-481e-95cc-336a524f2ac2>, accessed 10 July 2023).
168. Blinken A, Raimondo G. To shape the future of AI, we must act quickly. *Financial Times*, 24 July 2023 (<https://www.ft.com/content/eea999db-3441-45e1-a567-19dfa958dc8f>, accessed 30 July 2023).
169. Guterres, Antonio, Networked, inclusive multilateralism can help overcome challenges of era, says Secretary General, opening general assembly session, United Nations, 17 September 2019. (<https://press.un.org/en/2019/sgsm19746.doc.htm>, accessed 18 September 2023).

Annex. Methods

This guidance was developed by consensus. WHO relied upon its Expert Group on the Ethics and Governance of Artificial Intelligence for Health, consisting of 20 experts from all WHO regions who met fortnightly for approximately four months. The Expert Group applied the Consensus Principles and Recommendations from previously issued guidance on the ethics and governance of artificial intelligence for health to the emerging use of LMMs in healthcare and medicine.

The Expert Group first compiled a preliminary mapping of the potential uses and benefits of large multi-modal models, as well as the risks, at the level of the end-user. The Expert Group also identified risks that health systems and societies could encounter with the use of such AI systems. This was supplemented by a comprehensive literature search of: (a) existing and proposed uses of LMMs in healthcare that have materialized over several years of its progressive evolution and use, (b) anticipated uses of LMMs, and (c) critiques and analyses of LMMs that have been published prior to the issuance of this guidance.

Based on a shared understanding of the known and potential benefits and risks, the expert group identified an appropriate framework to address the various ethical challenges and opportunities associated with the use of LMMs. The Expert Group agreed that a “value chain” approach was well-suited to depict where and how to organize appropriate governance, and which actor or actors could be held responsible to carry out relevant measures.

While existing or proposed legislation and regulatory measures in several jurisdictions were referred to and utilized as a means of framing areas for recommendation, the Expert Group crafted each recommendation to be applicable across multiple countries and legal systems. The Expert Group also debated recommendations that should be applied by companies, purchasers of such AI systems, and end-users of LMMs, especially healthcare providers and patients.

WHO acknowledges that this guidance may need to be revised and updated to ensure that the Expert Group’s findings and recommendations remain relevant and useful.

World Health Organization
20 Avenue Appia
CH-1211 Geneva 27
Switzerland
Website: <https://www.who.int>

